Forecasting Demand for National Research Supercomputing

Australian Academy of Science Internship Project

Yusuf Kurnia Romadhon

Table of contents

1 Forecasting Demand for National Research Supercomputing 7

2 Introduction 8
- 2.1 Project Background . . . 8
- 2.2 Project Purpose . . . 8
- 2.3 The Forecasting Problem . . . 9
- 2.4 Project Objectives . . . 9
- 2.5 Analysis Framework . . . 10

3 Data Collection 11
- 3.1 Purpose . . . 11
- 3.2 Data Sources . . . 11
  - 3.2.1 Year-specific HPC source summary . . . 12
- 3.3 HPC Demand and Allocation Data . . . 14
- 3.4 Field of Research Data . . . 15
- 3.5 Potential Demand Drivers . . . 16
- 3.6 Data Provenance . . . 16
- 3.7 Data Collection Limitations . . . 17

4 Variable Definition 18
- 4.1 Purpose . . . 18
- 4.2 HPC Demand . . . 18
  - 4.2.1 Requested Compute . . . 19
  - 4.2.2 Allocated Compute . . . 19
  - 4.2.3 Latent Demand . . . 19
- 4.3 Demand and Allocation Indicators . . . 19
- 4.4 Field of Research . . . 20
- 4.5 Potential Structural Drivers . . . 21
  - 4.5.1 Research Funding . . . 21
  - 4.5.2 Research Workforce and Activity . . . 21
  - 4.5.3 Compute Intensity . . . 21
  - 4.5.4 HPC Infrastructure . . . 22
- 4.6 Variable Selection for Modelling . . . 22

5 Data Integration **23**
5.1 Purpose . . . 23
5.2 Main Datasets . . . 23
5.3 Yearly Source Integration . . . 24
5.3.1 Load Libraries and Define Paths . . . 24
5.4 Annual Demand and Allocation Summary Construction . . . 31
5.4.1 Reconstruct Annual Demand and Allocation Summary . . . 31
5.4.2 Construct Reconstructed Annual Summary . . . 33
5.4.3 Compare Reconstructed and Existing Annual Summaries 38
5.4.4 Resolve Unreconstructed Annual Values . . . 39
5.4.5 Derive Missing Annual Measures . . . 40
5.4.6 Final Reconciliation Check . . . 43
5.4.7 Export Reconciled Annual Demand and Allocation Summary 45
5.5 Field of Research Summary Construction . . . 46
5.5.1 Assess Historical FoR Comparability . . . 48
5.5.2 Construct FoR-Level Compute Measures . . . 50
5.5.3 Export Reconciled Field of Research Summary . . . 53

6 Data Standardisation **55**
6.1 Purpose . . . 55
6.2 Reporting Period . . . 55
6.3 Institution Names . . . 56
6.4 Field of Research Classification . . . 56
6.4.1 Historical FoR Reporting Groups . . . 58
6.5 HPC Facilities and Systems . . . 61
6.6 Resource Type . . . 62
6.7 Units . . . 63
6.8 Data Quality and Comparability . . . 63
6.8.1 Data Provenance and Estimation . . . 63
6.8.2 Comparability Across Reporting Years . . . 64
6.8.3 Treatment for Subsequent Analysis . . . 64
6.9 Output . . . 64

7 Data Conversion **66**
7.1 Purpose . . . 66
7.2 Conversion Approach . . . 67
7.3 CPU Resources . . . 75
7.4 GPU Resources . . . 86
7.4.1 Historical GPU Estimation Scenarios . . . 96
7.5 Field of Research Conversion . . . 104
7.6 Output . . . 111

8 Initial Analysis **114**
8.1 Purpose . . . 114
8.2 Historical Demand and Allocation . . . 115
8.2.1 Annual Data Validation . . . 115

8.2.2 Annual Analytical Indicators . . . 116
8.2.3 Annual Demand and Allocation Trends . . . 117
8.2.4 Initial Findings . . . 119
8.3 CPU and GPU Resources . . . 120
8.3.1 Validate the converted CPU and GPU series . . . 120
8.3.2 Create CPU and GPU trend charts . . . 121
8.3.3 Initial Findings . . . 127
8.4 Field of Research . . . 128
8.4.1 Validate the Field of Research series . . . 129
8.4.2 Research Field Allocation Composition . . . 132
8.4.3 Harmonise Research Fields for Longitudinal Comparison . 135
8.4.4 Annual Research Field Allocation Composition, 2014–2026 140
8.4.5 Field of Research Reporting Periods . . . 143
8.4.6 Field of Research: Initial Findings . . . 152

9 Forecasting Models 158
9.1 Purpose . . . 158
9.2 Forecast Target and Data Suitability . . . 159
9.2.1 Forecast Target . . . 159
9.2.2 Historical Data Suitability . . . 159
9.2.3 Forecasting Horizon . . . 159
9.3 Data Preparation and Validation . . . 160
9.4 Candidate Forecasting Models . . . 161
9.4.1 Overall HPC Demand: Linear and Log-Linear Trends . . 161
9.4.2 Initial Model Findings . . . 163
9.5 Sensitivity Analysis: Influence of 2026 . . . 163
9.5.1 Trend Sensitivity Comparison . . . 163
9.5.2 Sensitivity Analysis Findings . . . 165
9.6 CPU Demand Forecasting Models . . . 165
9.6.1 CPU Linear and Log-Linear Trends . . . 166
9.6.2 CPU Trend Findings . . . 167
9.6.3 CPU Demand Sensitivity Analysis . . . 168
9.6.4 CPU Sensitivity Findings . . . 169
9.7 GPU Demand Forecasting Models . . . 170
9.7.1 GPU Data Provenance and Modelling Period . . . 170
9.7.2 GPU Linear and Log-Linear Trends . . . 171
9.7.3 GPU Trend Findings . . . 173
9.7.4 GPU Demand Sensitivity Analysis . . . 174
9.7.5 GPU Sensitivity Findings . . . 175
9.8 Historical GPU Conversion Scenario Assessment . . . 176
9.8.1 Historical Scenario Consistency . . . 177
9.9 Demand Projections for 2027–2030 . . . 178
9.9.1 Overall HPC Demand Projections . . . 178
9.9.2 Overall HPC Projection Findings . . . 180
9.9.3 CPU Demand Projections . . . 180
9.9.4 CPU Projection Findings . . . 182

9.9.5 GPU Demand Projections . . . 182
9.9.6 GPU Projection Findings . . . 184
9.10 Field of Research Demand Modelling Framework . . . 185
9.10.1 Purpose and Modelling Approach . . . 185
9.10.2 Historical FoR Data Suitability . . . 186
9.10.3 FoR Allocation-Share Scenarios . . . 186
9.10.4 Interpretation and Limitations . . . 186
9.11 Modelling Summary and Limitations . . . 187
9.11.1 Summary of Candidate Demand Projections . . . 187
9.11.2 Key Modelling Findings . . . 187
9.11.3 Limitations and Implications for Forecast Evaluation . . . 188

10 Model Evaluation **190**
10.1 Purpose . . . 190
10.2 Evaluation Framework . . . 190
10.2.1 Candidate Models and Benchmarks . . . 190
10.2.2 Time-Series Cross-Validation . . . 191
10.2.3 Forecast Accuracy Measures . . . 191
10.2.4 Model Selection Principles . . . 192
10.3 Rolling-Origin Cross-Validation . . . 192
10.3.1 Overall HPC Demand: One-Step-Ahead Evaluation . . . 192
10.3.2 Overall HPC Forecast Accuracy . . . 195
10.3.3 Overall HPC Accuracy Findings . . . 197
10.3.4 CPU Demand: One-Step-Ahead Evaluation . . . 197
10.3.5 CPU Forecast Accuracy . . . 200
10.3.6 CPU Accuracy Findings . . . 201
10.3.7 GPU Demand: One-Step-Ahead Evaluation . . . 202
10.3.8 GPU Forecast Accuracy . . . 204
10.3.9 GPU Accuracy Findings . . . 205
10.4 Field of Research Allocation-Share Evaluation . . . 206
10.4.1 Purpose and Evaluation Approach . . . 206
10.4.2 Harmonised FoR Data Preparation . . . 207
10.4.3 Historical FoR Allocation-Share Stability . . . 208
10.4.4 Long-Term FoR Allocation-Share Variation, 2014–2026 . . . 213
10.4.5 Earlier FoR Allocation-Share Variation, 2014–2020 . . . 215
10.4.6 Comparison of FoR Reporting Periods, 2014–2020 and 2021–2026 . . . 216
10.4.7 FoR Evaluation Findings and Limitations . . . 219

11 Model Selection **222**
11.1 Purpose and Approach . . . 222
11.2 Model Selection Criteria . . . 222
11.3 Overall HPC Demand Model Selection . . . 223
11.3.1 Comparison of Candidate Models . . . 223
11.3.2 Sensitivity to the 2026 Observation . . . 224
11.3.3 Overall HPC Model Selection Decision . . . 225

11.4 CPU Demand Model Selection . . . 227
11.4.1 Comparison of Candidate Models . . . 227
11.4.2 Assessment of Growth Assumptions . . . 227
11.4.3 CPU Model Selection Decision . . . 228
11.5 GPU Demand Model Selection . . . 229
11.5.1 Comparison of Candidate Models . . . 229
11.5.2 Sensitivity to the 2026 Observation . . . 230
11.5.3 GPU Model Selection Decision . . . 231
11.6 Field of Research Scenario Assumptions . . . 232
11.6.1 Purpose and Analytical Approach . . . 232
11.6.2 Historical Evidence for Scenario Assumptions . . . 233
11.6.3 Selection of FoR Scenario Assumptions . . . 233
11.6.4 Interpretation and Limitations . . . 234
11.7 Selected Models and Limitations . . . 235
11.7.1 Summary of Model Selection Decisions . . . 235
11.7.2 Forecasting Uncertainty and Data Limitations . . . 236
11.7.3 Implications for Scenario Development . . . 237
11.8 Chapter Summary . . . 237

12 Demand Scenarios: 2027–2030 239
12.1 Purpose and Scope . . . 239
12.2 Scenario Framework . . . 240
12.3 Interpretation of Scenario Uncertainty . . . 240
12.4 Overall HPC Demand Scenarios . . . 241
12.4.1 Forecast Construction . . . 241
12.4.2 Annual Demand Projections . . . 241
12.4.3 Statistical Prediction Intervals for the Central Forecast . 245
12.4.4 Scenario Results and Interpretation . . . 249
12.5 CPU Demand Scenarios . . . 251
12.5.1 Forecast Construction . . . 251
12.5.2 Annual CPU Demand Projections . . . 251
12.5.3 Statistical Prediction Intervals for the Central CPU Forecast 255
12.5.4 Assessment of Prediction Interval Assumptions . . . 264
12.6 GPU Demand Scenarios . . . 265
12.6.1 Forecast Construction and Limitations . . . 265
12.6.2 Annual GPU Demand Projections . . . 265
12.6.3 GPU Scenario Results and Interpretation . . . 270
12.7 Field of Research Allocation-Share Scenarios . . . 271
12.7.1 Purpose and Methodology . . . 271
12.7.2 Validate Historical Allocation Shares . . . 272
12.7.3 Historical FoR Share Assumptions . . . 274
12.7.4 Earlier-Period FoR Share Assumptions, 2014–2020 . . . 277
12.7.5 Recent-Period FoR Allocation Scenario, 2027–2030 . . . 281
12.7.6 Longer-Term FoR Allocation Scenario, 2027–2030 . . . 283
12.7.7 Earlier-Period FoR Allocation Scenario, 2027–2030 . . . 286
12.7.8 Comparison and Interpretation of FoR Scenarios . . . 291

12.8 Overall Scenario Synthesis and Planning Implications 292
12.8.1 Central Forecast and Alternative Growth Assumptions . . 292
12.8.2 Forecast Uncertainty and Model Adequacy 292
12.8.3 GPU Demand and Emerging Infrastructure Requirements 293
12.8.4 Disciplinary Composition and FoR Scenarios 293
12.8.5 Implications for National HPC Planning 294
12.8.6 Chapter Conclusion 294

13 References and AI Acknowledgement **295**
13.1 References 295
13.2 AI Acknowledgement 297

Chapter 1

Forecasting Demand for National Research Supercomputing

Australian Academy of Science Internship Project

Chapter 2

Introduction

2.1 Project Background

High-performance computing (HPC) is an important component of Australia's research infrastructure, supporting computationally intensive research across a wide range of scientific fields (Australian Academy of Science 2024).

Australia has two national Tier-1 HPC facilities accessible to researchers: the National Computational Infrastructure (NCI) and the Pawsey Supercomputing Research Centre (National Computational Infrastructure, n.d.; Pawsey Supercomputing Research Centre, n.d.).

As research becomes increasingly data- and compute-intensive, understanding how demand for HPC resources is changing is important for planning Australia's future national research infrastructure (Australian Academy of Science 2024).

2.2 Project Purpose

The purpose of this project is to analyse historical Australian research computing demand and develop an evidence-based projection of future HPC demand to 2030 (Australian Academy of Science 2026).

The analysis is intended to support national policy discussions relating to public compute investment and future research infrastructure planning (Australian Academy of Science 2026).

The project combines data engineering, descriptive analysis and forecasting using publicly available NCI and Pawsey data, particularly information from the National Computational Merit Allocation Scheme (NCMAS), together with additional indicators where they are sufficiently consistent and relevant (Australian Academy of Science 2026).

Because historical reporting practices, hardware systems and compute accounting methods have changed over time, the analysis explicitly documents assumptions, derived values and limitations rather than treating all historical observations as directly comparable.

2.3 The Forecasting Problem

One of the main challenges in forecasting national research computing demand is that observed HPC allocation does not necessarily represent the full level of research demand.

National allocation schemes operate under finite capacity. As a result, allocated compute represents the amount of resource awarded rather than necessarily the amount researchers would use if additional capacity were available (Australian Academy of Science 2026).

Requested compute provides additional evidence of demand, but it may also understate underlying need where researchers constrain their requests based on expected availability or where potential users do not apply (Australian Academy of Science 2026).

The project therefore distinguishes between:

- requested compute,
- allocated compute,
- the gap between requested and allocated compute, and
- broader underlying or latent demand, which cannot be observed directly from the available allocation data.

A further challenge is that the historical series is not fully comparable across years. Reporting practices, hardware systems, compute units and service-unit accounting have changed over time, and CPU and GPU resources are not directly interchangeable (Australian Academy of Science 2026).

For this reason, the analysis uses the longest available historical series while explicitly documenting where values are directly reported, derived, estimated or based on assumptions. Forecasting results are therefore interpreted together with these data limitations rather than as precise measurements of unconstrained demand.

2.4 Project Objectives

The project aims to:

1. build a documented longitudinal dataset of Australian national HPC activity using publicly available NCI and Pawsey sources, including NCMAS allocation information,

2. document differences in reporting practices, compute units, systems and accounting methods across years,
3. improve comparability of historical compute measures where a defensible conversion or normalisation approach is available,
4. distinguish CPU and GPU resources where the available source data permits,
5. examine historical patterns in requested compute, allocation and oversubscription,
6. analyse changes in Field of Research composition using the most comparable available data,
7. assess potential indicators or structural drivers of research computing demand where suitable data are available,
8. develop and compare forecasting or projection approaches appropriate to the length and quality of the historical series, and
9. develop demand scenarios to 2030 and assess their implications for future national research computing requirements.

Throughout the analysis, assumptions, derived values, source limitations and uncertainty are documented explicitly.

The project objectives are informed by the internship brief, which emphasises longitudinal data assembly, descriptive analysis, structural demand modelling and scenario development (Australian Academy of Science 2026).

2.5 Analysis Framework

The analysis follows ten main steps:

1. Data collection
2. Variable definition
3. Data integration
4. Data standardisation
5. Data conversion and comparability assessment
6. Exploratory and descriptive analysis
7. Model development and comparison
8. Model evaluation
9. Model selection
10. Scenario forecasting and interpretation

Chapter 3

Data Collection

3.1 Purpose

The purpose of the data collection stage is to assemble a documented longitudinal evidence base for analysing Australian national research computing demand and allocation.

The project draws primarily on publicly available NCI, Pawsey and NCMAS reporting to identify historical requested compute, allocated compute, resource type, Field of Research and other available characteristics of national HPC allocations and usage (Australian Academy of Science 2026)

Additional datasets, including research funding, workforce and research activity indicators, are collected as potential structural drivers of future HPC demand where sufficiently consistent time-series data are available (Australian Academy of Science 2026).

Because the availability, scope and definition of variables differ substantially across years and sources, the purpose of this stage is not to assume that all collected observations are directly comparable. Instead, original values, units, reporting periods, source definitions and data limitations are retained so that their suitability for integration, standardisation and modelling can be assessed in subsequent stages.

3.2 Data Sources

The project uses multiple public data sources from NCI, Pawsey and NCMAS, together with research funding, workforce and Field of Research datasets.

The specific HPC source varies by analytical year. Earlier observations are primarily drawn from annual reports, allocation documents and other historical

NCI and Pawsey publications, while more recent observations increasingly use direct NCMAS outcome pages, applicant guidance and facility-specific allocation information.

In several years, different documents are used for different analytical components. Allocation outcomes, requested-demand information, resource availability and Field of Research composition may therefore originate from separate official publications.

For transparency and reproducibility, the exact source document, source location and URL are retained in the underlying yearly datasets. The principal sources used for each analytical year are summarised below.

3.2.1 Year-specific HPC source summary

```
library(tibble)
library(knitr)

source_table <- tribble(
  ~Year, ~`Demand / Context Source`, ~`Allocation Source`, ~`FoR Source`,

  2014,
  "NCI Annual Report 2013/14; NCMAS Information for Applicants 2014",
  "NCI Annual Report 2013/14",
  "NCI Annual Report 2013/14",

  2015,
  "NCI Annual Report 2014/15",
  "NCI Annual Report 2014/15; Pawsey Annual Report 2014/15",
  "NCI Annual Report 2014/15",

  2016,
  "NCI Annual Report 2015/16; CSIRO Annual Report 2015/16",
  "NCI Annual Report 2015/16",
  "NCI Annual Report 2015/16",

  2017,
  "NCI Annual Report 2016/17; Pawsey Annual Report 2016/17",
  "NCI Annual Report 2016/17; Pawsey Annual Report 2016/17",
  "Pawsey Annual Report 2016/17",

  2018,
  "NCI Annual Report 2017/18; NCMAS 2018 results",
  "NCI Annual Report 2017/18",
  "NCI Annual Report 2017/18",
```

```
  2019,
  "NCI Annual Report 2018/19; Pawsey 2019 NCMAS information",
  "NCI Annual Report 2018/19",
  "NCI Annual Report 2018/19",

  2020,
  "NCI Annual Report 2019/20; NCMAS 2020 Information for Applicants",
  "NCI 2019/20 Compute Allocations Detail",
  "NCI 2019/20 Compute Allocations Detail",

  2021,
  "NCMAS 2021 Outcomes; NCI and Pawsey annual reports",
  "NCMAS 2021 Outcomes; NCI Compute Allocations 2020/21",
  "NCMAS 2021 Outcomes",

  2022,
  "NCMAS 2023 Information for Applicants; NCI/Pawsey reports",
  "NCMAS 2022 Outcomes",
  "NCMAS 2022 Outcomes",

  2023,
  "NCMAS 2023 Outcomes; NCI/Pawsey supporting publications",
  "NCMAS 2023 Outcomes",
  "NCMAS 2023 Outcomes",

  2024,
  "NCMAS 2024 Outcomes; NCMAS 2024 Information for Applicants",
  "NCMAS 2024 Outcomes",
  "NCMAS 2024 Outcomes",

  2025,
  "NCMAS 2025 Outcomes; Pawsey 2025 Allocations",
  "NCMAS 2025 Outcomes",
  "NCMAS 2025 Outcomes",

  2026,
  "NCMAS 2026 Outcomes; NCMAS 2026 Information for Applicants",
  "NCMAS 2026 Outcomes",
  "NCMAS 2026 Outcomes"
)

knitr::kable(
  source_table,
  format = "html",
  escape = TRUE,
```

```
  align = c("c", "l", "l", "l")
)
```

Table 3.1: Principal year-specific sources used to constr
dataset

Year	Demand / Context Source	Allocation Source
2014	NCI Annual Report 2013/14; NCMAS Information for Applicants 2014	NCI Annual Report 2013/14
2015	NCI Annual Report 2014/15	NCI Annual Report 2014/15
2016	NCI Annual Report 2015/16; CSIRO Annual Report 2015/16	NCI Annual Report 2015/16
2017	NCI Annual Report 2016/17; Pawsey Annual Report 2016/17	NCI Annual Report 2016/17
2018	NCI Annual Report 2017/18; NCMAS 2018 results	NCI Annual Report 2017/18
2019	NCI Annual Report 2018/19; Pawsey 2019 NCMAS information	NCI Annual Report 2018/19
2020	NCI Annual Report 2019/20; NCMAS 2020 Information for Applicants	NCI 2019/20 Compute Alloc
2021	NCMAS 2021 Outcomes; NCI and Pawsey annual reports	NCMAS 2021 Outcomes; NC
2022	NCMAS 2023 Information for Applicants; NCI/Pawsey reports	NCMAS 2022 Outcomes
2023	NCMAS 2023 Outcomes; NCI/Pawsey supporting publications	NCMAS 2023 Outcomes
2024	NCMAS 2024 Outcomes; NCMAS 2024 Information for Applicants	NCMAS 2024 Outcomes
2025	NCMAS 2025 Outcomes; Pawsey 2025 Allocations	NCMAS 2025 Outcomes
2026	NCMAS 2026 Outcomes; NCMAS 2026 Information for Applicants	NCMAS 2026 Outcomes

Key source families include NCI annual and allocation reports (National Computational Infrastructure 2014a, 2015, 2016, 2017, 2018a, 2019, 2020a, 2020b, 2021), Pawsey annual and allocation publications (Pawsey Supercomputing Research Centre 2015, 2017, 2019), and NCMAS outcomes and applicant guidance (National Computational Infrastructure 2014b, 2018b, 2020c, 2022, 2023, 2025; National Computational Merit Allocation Scheme 2021, 2022, 2023, 2024, 2025, 2026).

3.3 HPC Demand and Allocation Data

The main HPC dataset covers source periods from 2013/14 to 2025/26, corresponding to analytical years 2014 to 2026.

The collected information includes, where available:

- requested compute,
- allocated compute,
- facility and computing system,
- resource type,
- original compute unit,
- Field of Research,
- institution,
- allocation or reporting scope,

- source document and URL, and
- source reporting period.

The availability and level of detail of these variables vary substantially across years. Earlier observations are often available only as aggregate values reported in annual reports, while more recent NCMAS outcomes provide more detailed project-, facility- or resource-level information.

Requested and allocated compute are retained as distinct measures. Allocated compute represents resources awarded through the relevant allocation process, while requested compute provides additional evidence of demand where it is directly reported or can be defensibly derived.

Historical values are also retained in their original reported units before any conversion or standardisation is applied. This is important because service units, CPU hours, GPU hours and other compute measures are not directly interchangeable across facilities, systems or hardware generations (Australian Academy of Science 2026).

Where a value is not directly reported, its derivation or assumption is recorded separately so that directly reported, derived and proxy values can be distinguished during subsequent analysis.

3.4 Field of Research Data

Field of Research (FoR) information was collected to examine how the disciplinary composition of national HPC allocations has changed over time.

The source, granularity and reliability of FoR information vary across years. Earlier observations are derived from annual reports, allocation summaries or project-level information with varying levels of detail, while more recent NCMAS outcomes provide more direct and consistent FoR information.

Where possible, reported FoR classifications are mapped to a common classification to support comparison across years. The original classification and source-reported FoR information are retained so that the standardisation process remains traceable.

The available FoR data primarily represents the composition of allocated compute rather than the composition of requested or latent demand. FoR allocation shares are therefore not interpreted directly as FoR-specific research demand.

Where allocation composition is subsequently combined with annual requested-compute totals to estimate demand by FoR, the resulting values are treated as analytical proxies rather than directly observed demand.

Because source scope and data quality differ across years, longer-term FoR trends are interpreted with caution, with the most confidence placed in years where direct and more comparable NCMAS FoR information is available.

3.5 Potential Demand Drivers

Additional datasets are collected to investigate potential structural drivers of national research computing demand.

These include research funding, research workforce, research activity and related indicators that may help explain changes in HPC demand over time. The internship brief identifies competitive grant flows, research workforce size and changes in compute intensity as potential structural drivers of future demand (Australian Academy of Science 2026).

These variables are treated as candidate explanatory indicators rather than assumed drivers. Their suitability for modelling is assessed later based on historical coverage, consistency, temporal alignment and their observed relationship with the available HPC demand measures.

Variables with insufficient coverage or weak comparability are retained for contextual analysis where useful but are not necessarily included in the final forecasting model.

3.6 Data Provenance

Data provenance is retained throughout the dataset so that each observation can be traced to its originating source information and, where relevant, the derivation used to construct it.

Where available, the retained provenance information includes:

- source document or webpage,
- source URL,
- source reporting period,
- source location or relevant section,
- originally reported value and unit,
- reporting or allocation scope, and
- notes describing extraction, derivation or interpretation.

This information provides an audit trail for subsequent data integration, standardisation and conversion. It also makes it possible to distinguish values that are directly reported from those that are derived, estimated or based on analytical assumptions.

Maintaining provenance is particularly important because the historical HPC series is assembled from publications that differ in reporting format, scope, compute units and level of detail across years. Provenance information is therefore used to assess whether observations can be compared directly or require additional qualification before analysis.

3.7 Data Collection Limitations

Several limitations were identified during the data collection stage.

The main limitations include:

- limited availability of sufficiently consistent historical data before 2013,
- requested compute may understate underlying demand where researchers constrain applications in response to expected resource availability,
- changes in reporting formats, terminology and level of detail across years,
- inconsistent availability of requested and allocated compute,
- incomplete separation of CPU and GPU resources in some historical records,
- changes in HPC systems, hardware generations and accounting methods,
- differences in reporting scope between NCI, Pawsey and NCMAS sources,
- changes in Field of Research classifications and reporting granularity, and
- the use of derived or proxy values where directly reported measures are not available.

These limitations mean that the historical series cannot be treated as perfectly homogeneous or directly comparable across all years.

The impact of these issues is assessed in later stages of the project through data integration, standardisation, conversion and comparability assessment. Observations that require assumptions or additional qualification are retained with supporting provenance information rather than being treated as equivalent to directly reported values.

Chapter 4

Variable Definition

4.1 Purpose

The purpose of the variable definition stage is to establish clear and consistent meanings for the measures used throughout the analysis.

The variables are organised into three main groups:

- HPC demand variables, including requested compute and related demand indicators,
- allocation and infrastructure variables, including awarded compute and available capacity, and
- potential structural drivers that may help explain changes in research computing demand over time.

Not all defined variables are expected to enter the final forecasting model. Some are used primarily for descriptive analysis, comparability assessment or interpretation.

Variables considered for modelling are evaluated later based on their theoretical relevance, historical coverage, consistency of measurement, temporal alignment and suitability for the relatively short annual time series available.

4.2 HPC Demand

The main outcome of interest is Australian research demand for national HPC resources.

Demand is not directly observed through a single consistent measure. The internship brief notes that national HPC allocations are supply-constrained and that researchers may also limit requests based on expected availability,

meaning that observed requests may not capture the full level of underlying need (Australian Academy of Science 2026).

The analysis therefore distinguishes between requested compute, allocated compute and underlying or latent demand.

4.2.1 Requested Compute

Requested compute represents the amount of HPC resource requested by researchers through an allocation process where this information is available.

It is treated as the primary observable indicator of research demand. However, requested compute may understate underlying need if researchers limit their requests based on expected resource availability or if potential users do not submit applications (Australian Academy of Science 2026).

4.2.2 Allocated Compute

Allocated compute represents the amount of HPC resource awarded to researchers through the relevant allocation process, as reported in NCMAS outcomes and related NCI and Pawsey allocation documentation (National Computational Merit Allocation Scheme 2024, 2025, 2026).

Allocation is analysed separately from demand because the amount awarded is constrained by available resources and the allocation process.

4.2.3 Latent Demand

Latent demand refers in this project to the broader level of HPC resources that researchers might require beyond what is observed in submitted requests or awarded allocations.

This is an analytical concept rather than a directly observed variable. The internship brief highlights that observed demand may be censored by available capacity and by researchers adapting their requests to expected resource availability (Australian Academy of Science 2026).

4.3 Demand and Allocation Indicators

Several indicators are derived from the historical demand and allocation data to describe the relationship between requested and awarded compute.

The observed request-allocation gap is defined as:

$$\text{Request-Allocation Gap}_t = \text{Requested Compute}_t - \text{Allocated Compute}_t$$

A positive value indicates that reported requests exceeded the amount of compute allocated in year (t).

The fulfilment rate is defined as:

$$\text{Fulfilment Rate}_t = \frac{\text{Allocated Compute}_t}{\text{Requested Compute}_t}$$

A fulfilment rate below 1 indicates that allocated compute was lower than reported requested compute, while a value closer to 1 indicates that a larger proportion of reported requests was satisfied.

These indicators are used to describe allocation pressure over time. However, they are not interpreted as direct measures of total unmet research demand because requested compute may itself be constrained by expected resource availability and participation in the allocation process (Australian Academy of Science 2026).

The indicators are calculated only where requested and allocated compute refer to sufficiently comparable scope and resource definitions.

4.4 Field of Research

Field of Research (FoR) variables are used to describe the disciplinary composition of national HPC allocation.

The principal FoR variables are:

- `for_classification`: the standardised research field assigned to an observation,
- `allocation_share_for`: the proportion of annual allocated compute associated with each FoR,
- `allocated_compute_for`: allocated compute associated with each FoR, and
- `estimated_demand_by_for_proxy`: an analytical estimate of requested compute by FoR where allocation composition is applied to annual requested compute.

The FoR allocation share for field (i) in year (t) is defined as:

$$\text{Allocation Share}_{i,t} = \frac{\text{Allocated Compute}_{i,t}}{\text{Total Allocated Compute}_t}$$

Where allocation composition is applied to annual requested compute, the resulting FoR-level demand measure is explicitly treated as a proxy rather than directly observed demand.

This distinction is maintained because the available FoR information primarily describes allocation composition rather than the composition of submitted requests.

4.5 Potential Structural Drivers

Potential structural drivers are candidate explanatory variables that may help describe changes in national HPC demand over time (Australian Academy of Science 2026).

They are not assumed to have a causal relationship with HPC demand and are considered for modelling only where historical coverage and measurement consistency are sufficient.

4.5.1 Research Funding

Research funding variables include available annual measures from:

- Australian Research Council (ARC),
- National Health and Medical Research Council (NHMRC), and
- Medical Research Future Fund (MRFF).

These variables are considered as indicators of changes in the scale of funded research activity that may be associated with research computing requirements.

4.5.2 Research Workforce and Activity

Research workforce and activity variables include available measures of:

- research workforce size, and
- research activity.

These variables are considered as indicators of changes in the scale of the Australian research system.

4.5.3 Compute Intensity

Compute-intensity indicators describe changes in the amount or type of computing associated with research activity.

Potential indicators include:

- compute allocation per project,
- compute allocation by Field of Research,
- changes in CPU allocation, and
- changes in GPU allocation.

These measures are used primarily to investigate whether the composition and intensity of research computing are changing over time.

4.5.4 HPC Infrastructure

HPC infrastructure variables describe changes in the supply environment within which observed demand and allocation occur.

Relevant variables may include:

- computing system,
- resource type,
- available compute capacity, and
- major system commissioning or retirement.

Infrastructure variables are interpreted separately from demand because an increase in allocated compute may reflect additional available capacity rather than an equivalent increase in underlying research demand.

4.6 Variable Selection for Modelling

Not all variables defined in this chapter will be included in the final forecasting model.

Candidate variables will be assessed based on:

- theoretical relevance to HPC demand,
- historical data coverage,
- consistency of measurement,
- temporal alignment with the demand series,
- observed relationship with the selected demand measure,
- availability or defensible assumptions over the forecast horizon, and
- the limited number of annual observations available for modelling.

Given the short annual time series, the final modelling approach will prioritise a small number of interpretable variables rather than a large set of predictors. Variables that are useful for descriptive analysis or interpretation may therefore be retained without being included in the final forecasting model.

Chapter 5

Data Integration

5.1 Purpose

The purpose of the data integration stage is to combine the different datasets collected for the project into a structured analytical dataset while preserving their original values, units, scope and provenance.

The source datasets cover different reporting periods, levels of granularity and data structures. Integration is therefore required to align common identifiers such as analytical year, facility, computing system, resource type and Field of Research before the data can be analysed together.

At this stage, integration focuses on organising and linking the available information rather than forcing observations into common units or classifications. Differences in naming, classification and compute units are addressed subsequently through data standardisation and conversion.

Where datasets cannot be combined directly because of differences in scope or granularity, those distinctions are retained explicitly rather than assuming equivalence.

5.2 Main Datasets

The integration process combines several datasets covering 2014 - 2026, with different analytical roles.

The core HPC datasets include:

- yearly HPC demand and allocation source workbooks,
- Field of Research composition data, and
- HPC system and capacity information.

Supporting datasets include:

- research funding data, and
- research workforce and research activity data.

The core datasets provide the historical evidence used to construct measures of requested compute, allocated compute and disciplinary composition. Supporting research funding, workforce and research activity datasets are retained separately for later integration into the modelling stage as candidate explanatory variables.

Because these datasets differ in temporal coverage, granularity and reporting scope, they are not merged indiscriminately. Integration is performed only where suitable common identifiers and compatible levels of aggregation are available.

Where annual demand or allocation measures are not directly available, derived or proxy values are constructed later from the integrated source evidence using explicitly documented assumptions. These selected annual measures are subsequently assembled into a separate summary dataset for modelling.

5.3 Yearly Source Integration

The detailed yearly source files are integrated before constructing the annual modelling summary.

For the HPC demand and allocation data, each yearly workbook contains multiple types of evidence. The integration therefore preserves the different analytical levels rather than reducing each workbook immediately to a single annual value.

Three recurring components are retained:

- `Project_Allocations`, containing project or allocation-level records,
- `Context_Metrics`, containing reported or derived annual and contextual measures, and
- `Resource_Envelope`, containing information on available resources and system-level capacity.

Field of Research data are integrated separately from the yearly `FoR_Composition` sheets.

The source filename and source period are retained during integration so that each observation remains traceable to its original yearly workbook. The resulting integrated tables form the evidence base from which annual summary measures are selected or derived later.

5.3.1 Load Libraries and Define Paths

```
library(tidyverse)
library(readxl)
```

```
demand_path <- "assets/data/raw data/demand and allocation"
for_path <- "assets/data/raw data/field of research"

clean_path <- "assets/data/data clean"
```

The yearly demand and allocation workbooks are identified from the raw data directory.

```
demand_files <- list.files(
  demand_path,
  pattern = "\\.xlsx$",
  full.names = TRUE
)
```

For the demand and allocation workbooks, the reporting period is derived from the source filename. A helper function is also defined to convert financial-year-style source periods to their ending analytical year.

```
get_source_period <- function(path) {

  filename <- basename(path)

  years <- str_match(
    filename,
    "(\\d{4})_(\\d{2})"
  )

  paste0(
    years[, 2],
    "/",
    years[, 3]
  )
}

get_end_year <- function(source_period) {

  end_year <- str_extract(
    source_period,
    "(?<=/)\\d{2}$"
  )

  2000 + as.numeric(end_year)
}
```

Each yearly demand and allocation workbook contains several recurring worksheets. The same set of yearly workbooks is therefore used to extract each analytical component separately.

5.3.1.1 Integrate Context Metrics

Demand and allocation context metrics are combined across the yearly workbooks.

These records contain annual or supporting measures used to interpret requested compute, allocated compute and related demand-allocation relationships.

```
context_metrics <- map_dfr(
  demand_files,
  function(file) {

    read_excel(
      file,
      sheet = "Context_Metrics",
      col_types = "text"
    ) |>
      mutate(
        source_file = basename(file),
        source_period = get_source_period(file)
      )
  }
)
```

5.3.1.2 Integrate Project Allocations

Project-level allocation records are combined across the yearly workbooks while preserving their original source-level structure.

```
project_allocations <- map_dfr(
  demand_files,
  function(file) {

    read_excel(
      file,
      sheet = "Project_Allocations",
      col_types = "text"
    ) |>
      mutate(
        source_file = basename(file),
        source_period = get_source_period(file)
      )
  }
)
```

5.3.1.3 Integrate Resource Envelope

Resource-envelope records are combined across the yearly workbooks.

These records describe available HPC systems, resource availability and capacity-related information used to interpret the supply environment associated with observed demand and allocation.

```
resource_envelope <- map_dfr(
  demand_files,
  function(file) {

    read_excel(
      file,
      sheet = "Resource_Envelope",
      col_types = "text"
    ) |>
      mutate(
        source_file = basename(file),
        source_period = get_source_period(file)
      )
  }
)
```

5.3.1.4 Align Analytical Year

A common analytical year is created to align observations across the different yearly source tables. The ending year of each source period is used as the common reference year where a consistent NCMAS or allocation year is not available.

```
context_metrics <- context_metrics |>
  mutate(
    year = get_end_year(source_period)
  )

project_allocations <- project_allocations |>
  mutate(
    year = coalesce(
      suppressWarnings(as.numeric(ncmas_year)),
      suppressWarnings(as.numeric(allocation_year)),
      get_end_year(source_period)
    )
  )

resource_envelope <- resource_envelope |>
  mutate(
    year = coalesce(
      suppressWarnings(as.numeric(ncmas_year)),
      suppressWarnings(as.numeric(allocation_year)),
      get_end_year(source_period)
```

```
    )
  )
```

5.3.1.5 Validate Demand and Allocation Integration

The integrated demand and allocation tables are checked to confirm that all yearly source workbooks have been retained and that the common analytical year is aligned consistently across the historical series.

```
context_metrics |>
  summarise(
    rows = n(),
    source_periods = n_distinct(source_period),
    source_files = n_distinct(source_file),
    years = n_distinct(year)
  )
```

```
# A tibble: 1 x 4
   rows source_periods source_files years
  <int>          <int>        <int> <int>
1   350             13           13    13
```

```
project_allocations |>
  summarise(
    rows = n(),
    source_periods = n_distinct(source_period),
    source_files = n_distinct(source_file),
    years = n_distinct(year)
  )
```

```
# A tibble: 1 x 4
   rows source_periods source_files years
  <int>          <int>        <int> <int>
1  5495             13           13    13
```

```
resource_envelope |>
  summarise(
    rows = n(),
    source_periods = n_distinct(source_period),
    source_files = n_distinct(source_file),
    years = n_distinct(year)
  )
```

```
# A tibble: 1 x 4
   rows source_periods source_files years
  <int>          <int>        <int> <int>
1    47             13           13    13
```

```
bind_rows(
  Context_Metrics = context_metrics |>
    distinct(source_period, year),
  Project_Allocations = project_allocations |>
    distinct(source_period, year),
  Resource_Envelope = resource_envelope |>
    distinct(source_period, year),
  .id = "table"
) |>
  count(
    table,
    source_period,
    name = "analytical_years"
  ) |>
  filter(analytical_years > 1)
```

```
# A tibble: 1 x 3
  table               source_period analytical_years
  <chr>               <chr>                    <int>
1 Project_Allocations 2014/15                      2
```

A further check is used to identify any source period that maps to more than one analytical year within an integrated table.

5.3.1.6 Integrate Field of Research Data

Field of Research composition data are combined across the yearly FoR workbooks while retaining the source filename and reporting period.

These records describe the disciplinary composition of HPC allocation and are integrated separately from the demand and allocation evidence because they have a different analytical structure and level of aggregation.

```
for_files <- list.files(
  for_path,
  pattern = "\\.xlsx$",
  full.names = TRUE
)
```

```
for_data <- map_dfr(
  for_files,
  function(file) {

    read_excel(
      file,
      sheet = "FoR_Composition",
      col_types = "text"
```

```
    ) |>
      mutate(
        source_file = basename(file)
      )
  }
)
```

```
for_data <- for_data |>
  mutate(
    year = get_end_year(source_period)
  )
```

5.3.1.7 Validate Field of Research Integration

The integrated Field of Research data are checked to confirm that all yearly FoR workbooks have been retained and that the common analytical year is aligned consistently with the demand and allocation series.

```
for_data |>
  summarise(
    rows = n(),
    source_periods = n_distinct(source_period),
    source_files = n_distinct(source_file),
    years = n_distinct(year)
  )
```

```
# A tibble: 1 x 4
   rows source_periods source_files years
  <int>          <int>        <int> <int>
1   133             13           13    13
```

5.3.1.8 Export Integrated Source Data

The validated integrated tables are exported as a consolidated workbook to provide a reproducible intermediate dataset between the original yearly source files and the annual modelling summary.

The workbook retains the different analytical components on separate worksheets so that differences in granularity and scope are preserved.

```
writexl::write_xlsx(
  list(
    Context_Metrics = context_metrics,
    Project_Allocations = project_allocations,
    Resource_Envelope = resource_envelope,
    FoR_Composition = for_data
  ),
  path = file.path(
```

```
    clean_path,
    "integrated_hpc_source_data.xlsx"
  )
)
```

5.4 Annual Demand and Allocation Summary Construction

The annual demand and allocation summary is constructed after integration of the detailed yearly source evidence.

The existing manually assembled summary is retained temporarily as a reference for reconciliation rather than treated as an independent source of data. Each annual requested and allocated compute value is checked against the integrated `Context_Metrics`, `Project_Allocations` and `Resource_Envelope` evidence.

Where a directly reported comparable annual value is available, that value is preferred. Where it is not available, a derived or proxy value may be constructed from the integrated evidence using an explicitly documented assumption.

5.4.1 Reconstruct Annual Demand and Allocation Summary

Annual demand and allocation measures are reconstructed from the integrated source tables before comparison with the previously assembled summary.

Context metrics are used to identify reported annual totals, ratios and derived demand information. Project-level allocations and resource envelopes are retained as supporting evidence and are aggregated only where their scope and units are sufficiently consistent.

```
context_candidates <- context_metrics |>
  mutate(
    value_num = suppressWarnings(as.numeric(value))
  ) |>
  filter(
    str_detect(
      metric,
      regex(
        "request|demand|allocat|fulfil|oversub|capacity|available",
        ignore_case = TRUE
      )
    )
  ) |>
  select(
    year,
```

```
    source_period,
    metric,
    value_num,
    unit,
    scope,
    source,
    notes
  ) |>
  arrange(year, metric)
```

```
project_candidates <- project_allocations |>
  mutate(
    allocated_compute_num =
      suppressWarnings(as.numeric(allocated_compute))
  ) |>
  filter(!is.na(allocated_compute_num)) |>
  group_by(
    year,
    source_period,
    unit,
    allocation_type
  ) |>
  summarise(
    allocated_total = sum(allocated_compute_num),
    records = n(),
    .groups = "drop"
  ) |>
  arrange(year, unit, allocation_type)
```

```
resource_candidates <- resource_envelope |>
  mutate(
    resource_available_num =
      suppressWarnings(as.numeric(resource_available))
  ) |>
  filter(!is.na(resource_available_num)) |>
  group_by(
    year,
    source_period,
    unit
  ) |>
  summarise(
    resource_available_total =
      sum(resource_available_num),
    records = n(),
    .groups = "drop"
```

```
  ) |>
  arrange(year, unit)
```

5.4.2 Construct Reconstructed Annual Summary

The annual series is reconstructed from the integrated evidence without using the previously assembled summary as an input.

Directly reported annual values are used where available. Derived values are calculated only where the required source values and relationships are contained in the integrated evidence. Where the integrated evidence is insufficient to reproduce an earlier summary value, the value is left missing for subsequent reconciliation rather than being inferred from the existing summary.

```
get_metric_value <- function(target_year, metric_name) {

  values <- context_candidates |>
    filter(
      year == target_year,
      metric == metric_name
    ) |>
    pull(value_num)

  if (length(values) != 1) {
    stop(
      paste(
        "Expected one value for",
        target_year,
        metric_name,
        "but found",
        length(values)
      )
    )
  }

  values
}
```

```
reconstructed_summary <- tibble(
  year = 2014:2026,
  source_period = paste0(
    year - 1,
    "/",
    str_sub(as.character(year), -2)
  )
) |>
```

```
mutate(

  requested_compute = case_when(

    year == 2014 ~
      get_metric_value(
        2014,
        "ncmas_total_ksu_request"
      ),

    year == 2015 ~
      get_metric_value(
        2015,
        "nci_ncmas_implied_demand_2015"
      ),

    year == 2016 ~
      get_metric_value(
        2016,
        "nci_ncmas_implied_requested_compute_2016"
      ),

    year == 2017 ~
      get_metric_value(
        2017,
        "pawsey_ncmas_requested_compute_2017"
      ),

    year == 2018 ~
      get_metric_value(
        2018,
        "national_implied_requested_compute_2018"
      ),

    # Current integrated evidence does not contain the
    # 220 × 3 national proxy used previously.
    year == 2019 ~ NA_real_,

    year == 2020 ~
      resource_candidates |>
        filter(
          year == 2020,
          unit == "kSU"
        ) |>
        pull(resource_available_total) |>
```

```
        first() * 2,

    year == 2021 ~
      get_metric_value(
        2021,
        "direct_ncmas_2021_total_allocation"
      ) *
      get_metric_value(
        2021,
        "ncmas_demand_supply_factor_high"
      ),

    year == 2022 ~
      get_metric_value(
        2022,
        "ncmas_allocated_compute_total"
      ) /
      (
        get_metric_value(
          2022,
          "ncmas_allocated_share_of_requested_pct"
        ) / 100
      ),

    year == 2023 ~
      get_metric_value(
        2023,
        "ncmas_requested_compute"
      ),

    # Retrospective 57% and 56% demand relationships are
    # not currently retained in the integrated source tables.
    year == 2024 ~ NA_real_,
    year == 2025 ~ NA_real_,

    year == 2026 ~
      get_metric_value(
        2026,
        "ncmas_requested_compute"
      )
  ),

  allocated_compute = case_when(

    # 83,505 kSU is not currently retained as an annual
```

```
    # metric in the integrated evidence.
    year == 2014 ~ NA_real_,

    year == 2015 ~
      get_metric_value(
        2015,
        "nci_ncmas_allocation_2014_15"
      ),

    year == 2016 ~
      get_metric_value(
        2016,
        "nci_ncmas_allocation_2015_16"
      ),

    year == 2017 ~
      get_metric_value(
        2017,
        "pawsey_ncmas_available_compute_2017"
      ),

    year == 2018 ~
      get_metric_value(
        2018,
        "national_allocated_compute_2018"
      ),

    # The 220m national capacity proxy is not currently
    # retained in the integrated evidence.
    year == 2019 ~ NA_real_,

    year == 2020 ~
      get_metric_value(
        2020,
        "nci_supplement_ncmas_allocation_total_verified"
      ),

    year == 2021 ~
      get_metric_value(
        2021,
        "direct_ncmas_2021_total_allocation"
      ),

    year == 2022 ~
      get_metric_value(
```

```
          2022,
          "ncmas_allocated_compute_total"
        ),

      year == 2023 ~
        get_metric_value(
          2023,
          "ncmas_allocated_compute_total"
        ),

      year == 2024 ~
        get_metric_value(
          2024,
          "ncmas_total_allocation"
        ),

      year == 2025 ~
        get_metric_value(
          2025,
          "ncmas_total_allocation"
        ),

      year == 2026 ~
        get_metric_value(2026, "ncmas_gadi_allocation") +
        get_metric_value(2026, "ncmas_setonix_cpu_allocation") +
        get_metric_value(2026, "ncmas_setonix_gpu_allocation")
    )
  )
```

```
reconstructed_summary
```

```
# A tibble: 13 x 4
    year source_period requested_compute allocated_compute
   <int> <chr>                     <dbl>             <dbl>
 1  2014 2013/14                131356                NA
 2  2015 2014/15                   202.            88731
 3  2016 2015/16                   225             97112
 4  2017 2016/17                   200               100
 5  2018 2017/18                   529.              222
 6  2019 2018/19                    NA                NA
 7  2020 2019/20                607980            255148
 8  2021 2020/21               1193811            397937
 9  2022 2021/22               1162230            581115
10  2023 2022/23               1333240            817956
11  2024 2023/24                    NA            812569
12  2025 2024/25                    NA            876200
```

13 2026 2025/26 2242817 872288

5.4.3 Compare Reconstructed and Existing Annual Summaries

The reconstructed annual series is compared with the previously assembled summary to identify values that are reproduced directly from the integrated evidence and values that require additional source evidence or documented derivation.

```
current_summary <- read_excel(
  file.path(clean_path, "demand_allocation_data.xlsx")
) |>
  mutate(
    year = get_end_year(source_period)
  )
```

```
summary_comparison <- current_summary |>
  select(
    year,
    source_period,
    old_requested = requested_compute,
    old_allocated = allocated_compute
  ) |>
  full_join(
    reconstructed_summary |>
      rename(
        reconstructed_requested = requested_compute,
        reconstructed_allocated = allocated_compute
      ),
    by = c("year", "source_period")
  ) |>
  mutate(
    requested_match = case_when(
      is.na(reconstructed_requested) ~ FALSE,
      TRUE ~ near(old_requested, reconstructed_requested)
    ),
    allocated_match = case_when(
      is.na(reconstructed_allocated) ~ FALSE,
      TRUE ~ near(old_allocated, reconstructed_allocated)
    )
  ) |>
  arrange(year)
```

```
summary_comparison |>
  select(
    year,
```

```
    source_period,
    old_requested,
    reconstructed_requested,
    requested_match,
    old_allocated,
    reconstructed_allocated,
    allocated_match
  )
```

```
# A tibble: 13 x 8
    year source_period old_requested reconstructed_requested requested_match
   <dbl> <chr>                 <dbl>                   <dbl> <lgl>
 1  2014 2013/14              131356                  131356  TRUE
 2  2015 2014/15                 202.                    202. TRUE
 3  2016 2015/16                 225                     225  TRUE
 4  2017 2016/17                 200                     200  TRUE
 5  2018 2017/18                 529.                    529. FALSE
 6  2019 2018/19                 660                      NA  FALSE
 7  2020 2019/20              607980                  607980  TRUE
 8  2021 2020/21             1193811                 1193811  TRUE
 9  2022 2021/22             1162230                 1162230  TRUE
10  2023 2022/23             1333240                 1333240  TRUE
11  2024 2023/24             1425560.                     NA  FALSE
12  2025 2024/25             1564643.                     NA  FALSE
13  2026 2025/26             2242817                 2242817  TRUE
# i 3 more variables: old_allocated <dbl>, reconstructed_allocated <dbl>,
#   allocated_match <lgl>
```

5.4.4 Resolve Unreconstructed Annual Values

The initial reconstruction identified several annual values that could not be reproduced from the integrated source tables because the supporting metric or derivation had not yet been retained explicitly in the integrated evidence.

These cases are reviewed against the original yearly source evidence before the final annual summary is constructed.

```
unreconstructed_values <- summary_comparison |>
  filter(
    is.na(reconstructed_requested) |
    is.na(reconstructed_allocated)
  ) |>
  transmute(
    year,
    source_period,
    missing_requested = is.na(reconstructed_requested),
    missing_allocated = is.na(reconstructed_allocated),
```

```
    old_requested,
    old_allocated
  )

unreconstructed_values
```

```
# A tibble: 4 x 6
   year source_period missing_requested missing_allocated old_requested
  <dbl> <chr>         <lgl>             <lgl>                      <dbl>
1  2014 2013/14       FALSE             TRUE                      131356
2  2019 2018/19       TRUE              TRUE                         660
3  2024 2023/24       TRUE              FALSE                    1425560.
4  2025 2024/25       TRUE              FALSE                    1564643.
# i 1 more variable: old_allocated <dbl>
```

The unresolved cases do not necessarily represent missing research activity. Rather, they identify years where no directly usable annual measure is available from the integrated source tables. These values therefore require documented derivation or proxy treatment based on the original source evidence.

5.4.5 Derive Missing Annual Measures

Where a directly reported annual demand or allocation measure is not available, a derived or proxy value is constructed from other published evidence retained in the integrated source data.

These derivations are applied only where the relationship used is explicitly supported by the source material. Derived and proxy values are retained separately from directly reported measures so that their status remains transparent.

```
reconstructed_summary <- reconstructed_summary |>
  mutate(
    requested_compute = case_when(
      year == 2019 ~ 220 * 3,
      year == 2024 ~ 812569 / 0.57,
      year == 2025 ~ 876200 / 0.56,
      TRUE ~ requested_compute
    ),
    allocated_compute = case_when(
      year == 2014 ~ 83505,
      year == 2019 ~ 220,
      TRUE ~ allocated_compute
    )
  )
```

```
reconstructed_summary <- reconstructed_summary |>
  mutate(
```

```
    requested_basis = case_when(
      year == 2014 ~ "Direct",
      year == 2015 ~ "Derived",
      year == 2016 ~ "Derived",
      year == 2017 ~ "Proxy",
      year == 2018 ~ "Derived",
      year == 2019 ~ "Derived proxy",
      year == 2020 ~ "Derived proxy",
      year == 2021 ~ "Derived proxy",
      year == 2022 ~ "Derived proxy",
      year == 2023 ~ "Direct",
      year == 2024 ~ "Derived proxy",
      year == 2025 ~ "Derived proxy",
      year == 2026 ~ "Direct"
    ),

    allocated_basis = case_when(
      year == 2014 ~ "Direct",
      year == 2015 ~ "Direct",
      year == 2016 ~ "Direct",
      year == 2017 ~ "Capacity proxy",
      year == 2018 ~ "Direct",
      year == 2019 ~ "Capacity proxy",
      year == 2020 ~ "Direct",
      year == 2021 ~ "Direct",
      year == 2022 ~ "Direct",
      year == 2023 ~ "Direct",
      year == 2024 ~ "Direct",
      year == 2025 ~ "Direct",
      year == 2026 ~ "Direct comparable subset"
    ),

    derivation_note = case_when(
      year == 2015 ~
        "Requested compute derived from reported allocation and demand relationship.",

      year == 2016 ~
        "Requested compute derived from reported one-third fulfilment relationship.",

      year == 2017 ~
        "Requested compute retained as a lower-bound proxy; available capacity used as the a

      year == 2018 ~
        "Requested compute derived from 222 million hours and a 42% fulfilment relationship.
```

```
      year == 2019 ~
        "Requested compute derived as 220 × 3; available capacity used as the allocation pro

      year == 2020 ~
        "Requested compute derived from the resource envelope and reported demand-supply rel

      year == 2021 ~
        "Requested compute derived using the approximate three-times demand-supply relations

      year == 2022 ~
        "Requested compute derived from allocated compute and the reported 50% fulfilment re

      year == 2024 ~
        "Requested compute derived from allocated compute and the retrospective 57% fulfilme

      year == 2025 ~
        "Requested compute derived from allocated compute and the retrospective 56% fulfilme

      year == 2026 ~
        "Setonix-Q excluded from allocation to match the scope of the published requested-co

      TRUE ~ NA_character_
    )
  )
```

```
reconstructed_summary
```

```
# A tibble: 13 x 7
    year source_period requested_compute allocated_compute requested_basis
   <int> <chr>                    <dbl>             <dbl> <chr>
 1  2014 2013/14                131356              83505 Direct
 2  2015 2014/15                   202.             88731 Derived
 3  2016 2015/16                   225              97112 Derived
 4  2017 2016/17                   200                100 Proxy
 5  2018 2017/18                   529.               222 Derived
 6  2019 2018/19                   660                220 Derived proxy
 7  2020 2019/20                607980             255148 Derived proxy
 8  2021 2020/21               1193811             397937 Derived proxy
 9  2022 2021/22               1162230             581115 Derived proxy
10  2023 2022/23               1333240             817956 Direct
11  2024 2023/24               1425560.            812569 Derived proxy
12  2025 2024/25               1564643.            876200 Derived proxy
13  2026 2025/26               2242817             872288 Direct
# i 2 more variables: allocated_basis <chr>, derivation_note <chr>
```

5.4.6 Final Reconciliation Check

After the missing annual measures are reconstructed and their basis is documented, the annual series is compared again with the previously assembled summary.

This final check identifies whether the reconstructed values reproduce the existing annual summary and highlights any remaining differences due to rounding, scope or methodological treatment.

```
final_summary_comparison <- current_summary |>
  select(
    year,
    source_period,
    old_requested = requested_compute,
    old_allocated = allocated_compute
  ) |>
  full_join(
    reconstructed_summary |>
      select(
        year,
        source_period,
        reconstructed_requested = requested_compute,
        reconstructed_allocated = allocated_compute,
        requested_basis,
        allocated_basis,
        derivation_note
      ),
    by = c("year", "source_period")
  ) |>
  mutate(
    requested_difference =
      reconstructed_requested - old_requested,

    allocated_difference =
      reconstructed_allocated - old_allocated,

    requested_match =
      near(
        old_requested,
        reconstructed_requested,
        tol = 1e-3
      ),

    allocated_match =
      near(
        old_allocated,
```

```
        reconstructed_allocated,
        tol = 1e-3
      )
  ) |>
  arrange(year)
```

```
final_summary_comparison |>
  select(
    year,
    source_period,
    old_requested,
    reconstructed_requested,
    requested_match,
    old_allocated,
    reconstructed_allocated,
    allocated_match
  )
```

```
# A tibble: 13 x 8
    year source_period old_requested reconstructed_requested requested_match
   <dbl> <chr>                 <dbl>                   <dbl> <lgl>
 1  2014 2013/14              131356                  131356  TRUE
 2  2015 2014/15                 202.                    202. TRUE
 3  2016 2015/16                 225                     225  TRUE
 4  2017 2016/17                 200                     200  TRUE
 5  2018 2017/18                 529.                    529. TRUE
 6  2019 2018/19                 660                     660  TRUE
 7  2020 2019/20              607980                  607980  TRUE
 8  2021 2020/21             1193811                 1193811  TRUE
 9  2022 2021/22             1162230                 1162230  TRUE
10  2023 2022/23             1333240                 1333240  TRUE
11  2024 2023/24             1425560.                1425560. TRUE
12  2025 2024/25             1564643.                1564643. TRUE
13  2026 2025/26             2242817                 2242817  TRUE
# i 3 more variables: old_allocated <dbl>, reconstructed_allocated <dbl>,
#   allocated_match <lgl>
```

```
final_summary_comparison |>
  summarise(
    years = n(),
    requested_matches = sum(requested_match),
    allocated_matches = sum(allocated_match),
    requested_mismatches = sum(!requested_match),
    allocated_mismatches = sum(!allocated_match)
  )
```

```
# A tibble: 1 x 5
```

```
  years requested_matches allocated_matches requested_mismatches
  <int>             <int>             <int>                 <int>
1    13                13                13                     0
# i 1 more variable: allocated_mismatches <int>
```

5.4.7 Export Reconciled Annual Demand and Allocation Summary

Following reconciliation, the reconstructed annual series is retained as the reproducible demand and allocation summary for subsequent standardisation, conversion and modelling.

The original reporting unit is preserved alongside each annual observation because the historical series is not yet expressed on a common compute scale. In years where requested and allocated compute were reported in different units, the mixed-unit description is retained explicitly.

The export also preserves whether each measure is directly reported, derived or proxy-based so that these distinctions remain available during later conversion and modelling.

```
final_demand_allocation_summary <- reconstructed_summary |>
  left_join(
    current_summary |>
      select(
        year,
        source_period,
        unit,
        requested_source,
        allocated_source,
        source_url
      ),
    by = c("year", "source_period")
  ) |>
  select(
    year,
    source_period,
    requested_compute,
    allocated_compute,
    unit,
    requested_basis,
    allocated_basis,
    requested_source,
    allocated_source,
    derivation_note,
    source_url
  )
```

```
final_demand_allocation_summary
```

```
# A tibble: 13 x 11
    year source_period requested_compute allocated_compute unit  requested_basis
   <dbl> <chr>                    <dbl>             <dbl> <chr> <chr>
 1  2014 2013/14                 131356             83505 kSU   Direct
 2  2015 2014/15                    202.            88731 mixe~ Derived
 3  2016 2015/16                    225             97112 mixe~ Derived
 4  2017 2016/17                    200               100 mill~ Proxy
 5  2018 2017/18                    529.              222 mill~ Derived
 6  2019 2018/19                    660               220 mill~ Derived proxy
 7  2020 2019/20                 607980            255148 kSU   Derived proxy
 8  2021 2020/21                1193811            397937 kSU   Derived proxy
 9  2022 2021/22                1162230            581115 kSU   Derived proxy
10  2023 2022/23                1333240            817956 kSU   Direct
11  2024 2023/24                1425560.           812569 kSU   Derived proxy
12  2025 2024/25                1564643.           876200 kSU   Derived proxy
13  2026 2025/26                2242817            872288 kSU   Direct
# i 5 more variables: allocated_basis <chr>, requested_source <chr>,
#   allocated_source <chr>, derivation_note <chr>, source_url <chr>
```

```
writexl::write_xlsx(
  final_demand_allocation_summary,
  path = file.path(
    clean_path,
    "reconciled_demand_allocation_summary.xlsx"
  )
)
```

5.5 Field of Research Summary Construction

The integrated Field of Research records are next assembled into a consistent annual summary.

Unlike the national demand and allocation series, historical FoR evidence varies in classification, reporting scope and completeness. The purpose of this stage is therefore to preserve the reported disciplinary composition while identifying differences in coverage and evidence quality across years.

FoR percentages are treated primarily as allocation composition rather than directly observed demand composition. Where these shares are later applied to annual requested compute, the resulting values are treated as analytical demand proxies rather than observed FoR-specific demand.

```
for_annual_check <- for_data |>
  mutate(
```

```
    composition_pct_num =
      suppressWarnings(as.numeric(composition_pct))
  ) |>
  group_by(
    year,
    source_period
  ) |>
  summarise(
    records = n(),
    composition_total =
      sum(composition_pct_num, na.rm = TRUE),
    .groups = "drop"
  ) |>
  arrange(year)

for_annual_check
```

```
# A tibble: 13 x 4
    year source_period records composition_total
   <dbl> <chr>           <int>             <dbl>
 1  2014 2013/14             7               100
 2  2015 2014/15             5               100
 3  2016 2015/16             7               100
 4  2017 2016/17            17               100
 5  2018 2017/18             8               100
 6  2019 2018/19             9               100
 7  2020 2019/20             8               100
 8  2021 2020/21             9               100
 9  2022 2021/22            13               100.
10  2023 2022/23            12               100.0
11  2024 2023/24            17               100
12  2025 2024/25            12               100.0
13  2026 2025/26             9               100.0
```

```
for_quality_summary <- for_data |>
  mutate(
    composition_pct_num =
      suppressWarnings(as.numeric(composition_pct))
  ) |>
  group_by(
    year,
    source_period,
    evidence_quality,
    scope,
    value_basis
  ) |>
```

```
  summarise(
    records = n(),
    composition_total =
      sum(composition_pct_num, na.rm = TRUE),
    .groups = "drop"
  ) |>
  arrange(year)

for_quality_summary
```

```
# A tibble: 19 x 7
    year source_period evidence_quality                scope value_basis records
   <dbl> <chr>         <chr>                           <chr> <chr>         <int>
 1  2014 2013/14       Derived residual from publishe~ NCI ~ Residual o~       1
 2  2014 2013/14       Direct percentage / scope unce~ NCI ~ Allocation~       6
 3  2015 2014/15       Direct percentage / NCMAS scop~ NCMA~ Allocation~       5
 4  2016 2015/16       Derived residual from publishe~ Meri~ Residual o~       1
 5  2016 2015/16       Direct percentage / scheme sco~ Meri~ Allocation~       6
 6  2017 2016/17       Derived residual from publishe~ Paws~ Residual o~       1
 7  2017 2016/17       Direct percentage / Pawsey all~ Paws~ Usage comp~      16
 8  2018 2017/18       Estimated - project-title clas~ NCI ~ Estimated ~       8
 9  2019 2018/19       Derived residual from publishe~ NCI ~ Residual o~       1
10  2019 2018/19       Estimated / incomplete extract~ NCI ~ Estimated ~       8
11  2020 2019/20       Estimated from near-complete N~ NCI ~ Estimated ~       8
12  2021 2020/21       Derived residual from publishe~ Nati~ Residual o~       1
13  2021 2020/21       Estimated from complete direct~ Nati~ Allocation~       8
14  2022 2021/22       Strong composition evidence: s~ Nati~ Allocation~      13
15  2023 2022/23       Strong composition evidence     Nati~ Allocation~      12
16  2024 2023/24       Source-preserving residual      Nati~ Allocation~       1
17  2024 2023/24       Strong composition evidence     Nati~ Allocation~      16
18  2025 2024/25       Strong composition evidence     Nati~ Allocation~      12
19  2026 2025/26       Strong composition evidence     Nati~ Allocation~       9
# i 1 more variable: composition_total <dbl>
```

5.5.1 Assess Historical FoR Comparability

The yearly FoR compositions generally sum to approximately 100%, but completeness alone does not imply that all years are directly comparable.

Historical observations differ in reporting scope, classification method and evidence quality. Some years are based on direct NCMAS allocation composition, while others rely on estimated project classification, residual categories, or broader facility-level usage.

For this reason, the FoR series is classified by analytical use rather than treated as a uniformly consistent time series. The strongest and most comparable

national NCMAS composition evidence is available from 2021 onward. Earlier observations are retained primarily as historical context, with 2017 treated separately because it represents Pawsey usage across all allocation schemes rather than national NCMAS allocation composition.

```
for_year_quality <- for_annual_check |>
  mutate(
    analytical_use = case_when(
      year == 2017 ~
        "Context only",

      year %in% c(
        2014,
        2016,
        2018,
        2019,
        2020
      ) ~
        "Historical / lower confidence",

      year == 2015 ~
        "Historical / moderate confidence",

      year >= 2021 ~
        "Main FoR trend"
    )
  )

for_year_quality
```

```
# A tibble: 13 x 5
    year source_period records composition_total analytical_use
   <dbl> <chr>           <int>             <dbl> <chr>
 1  2014 2013/14             7               100   Historical / lower confidence
 2  2015 2014/15             5               100   Historical / moderate confiden~
 3  2016 2015/16             7               100   Historical / lower confidence
 4  2017 2016/17            17               100   Context only
 5  2018 2017/18             8               100   Historical / lower confidence
 6  2019 2018/19             9               100   Historical / lower confidence
 7  2020 2019/20             8               100   Historical / lower confidence
 8  2021 2020/21             9               100   Main FoR trend
 9  2022 2021/22            13               100.  Main FoR trend
10  2023 2022/23            12               100.0 Main FoR trend
11  2024 2023/24            17               100   Main FoR trend
12  2025 2024/25            12               100.0 Main FoR trend
13  2026 2025/26             9               100.0 Main FoR trend
```

5.5.2 Construct FoR-Level Compute Measures

The annual FoR composition is linked to the reconciled national demand and allocation series using analytical year and source period.

Where FoR scope is sufficiently compatible with the annual national series, allocation shares are applied to annual compute totals to construct FoR-level estimates.

Years classified as contextual only are retained for composition analysis but are not used to distribute national annual compute totals because their reporting scope is not sufficiently comparable. The same shares are also applied to annual requested compute to construct an indicative FoR-level demand proxy.

The resulting demand measure is not treated as directly observed disciplinary demand because the available FoR percentages primarily describe allocation composition rather than request composition.

```
for_summary <- for_data |>
  mutate(
    composition_pct_num =
      suppressWarnings(as.numeric(composition_pct))
  ) |>
  left_join(
    for_year_quality |>
      select(
        year,
        source_period,
        analytical_use
      ),
    by = c("year", "source_period")
  )
```

```
for_summary <- for_summary |>
  left_join(
    final_demand_allocation_summary |>
      select(
        year,
        source_period,
        requested_compute,
        allocated_compute,
        unit,
        requested_basis,
        allocated_basis
      ),
    by = c("year", "source_period")
  )
```

```
for_summary <- for_summary |>
  mutate(
    allocation_share_for =
      composition_pct_num / 100,

    scope_alignment = case_when(
      year == 2017 ~
        "Not comparable",

      year == 2026 ~
        "Review required: FoR includes Setonix-Q while annual comparable allocation excludes

      analytical_use == "Main FoR trend" ~
        "Broadly comparable",

      TRUE ~
        "Historical / approximate"
    ),

    estimated_allocated_compute_by_for = case_when(
      analytical_use == "Context only" ~ NA_real_,
      TRUE ~ allocated_compute * allocation_share_for
    ),

    estimated_demand_by_for_proxy = case_when(
      analytical_use == "Context only" ~ NA_real_,
      TRUE ~ requested_compute * allocation_share_for
    )
  )
```

```
for_validation <- for_summary |>
  group_by(
    year,
    source_period,
    analytical_use
  ) |>
  summarise(
    composition_total =
      sum(composition_pct_num, na.rm = TRUE),

    estimated_allocated_for_total =
      if (
        all(is.na(estimated_allocated_compute_by_for))
      ) {
        NA_real_
```

```
      } else {
        sum(
          estimated_allocated_compute_by_for,
          na.rm = TRUE
        )
      },

    demand_proxy_total =
      if (
        all(is.na(estimated_demand_by_for_proxy))
      ) {
        NA_real_
      } else {
        sum(
          estimated_demand_by_for_proxy,
          na.rm = TRUE
        )
      },

    annual_allocated =
      first(allocated_compute),

    annual_requested =
      first(requested_compute),

    .groups = "drop"
  ) |>
  arrange(year)

for_validation
```

```
# A tibble: 13 x 8
    year source_period analytical_use   composition_total estimated_allocated_~1
   <dbl> <chr>         <chr>                        <dbl>                   <dbl>
 1  2014 2013/14       Historical / lo~               100                   83505
 2  2015 2014/15       Historical / mo~               100                   88731
 3  2016 2015/16       Historical / lo~               100                   97112
 4  2017 2016/17       Context only                   100                      NA
 5  2018 2017/18       Historical / lo~               100                     222
 6  2019 2018/19       Historical / lo~               100                     220
 7  2020 2019/20       Historical / lo~               100                  255148
 8  2021 2020/21       Main FoR trend                 100                  397937
 9  2022 2021/22       Main FoR trend                 100.                 581115.
10  2023 2022/23       Main FoR trend                 100.0                817956.
11  2024 2023/24       Main FoR trend                 100                  812569
```

```
12  2025 2024/25       Main FoR trend               100.0                 876200.
13  2026 2025/26       Main FoR trend               100.0                 872288.
# i abbreviated name: 1: estimated_allocated_for_total
# i 3 more variables: demand_proxy_total <dbl>, annual_allocated <dbl>,
#   annual_requested <dbl>
```

The validation confirms that the FoR composition shares sum to 100% for each analytical year. As a result, the sum of FoR-level allocated compute reproduces the corresponding annual allocated-compute total, while the sum of the FoR-level demand proxies reproduces the annual requested-compute total.

This confirms that the FoR allocation shares have been applied consistently to the reconciled national series. The demand values remain analytical proxies rather than directly observed FoR-specific requests.

5.5.3 Export Reconciled Field of Research Summary

The validated FoR-level measures are retained as the integrated disciplinary summary for subsequent standardisation, analysis and modelling.

The exported table preserves the original FoR classification information, allocation composition, analytical-use classification and evidence-quality metadata. FoR-level requested compute remains explicitly identified as a proxy derived from allocation composition.

```
final_for_summary <- for_summary |>
  select(
    year,
    source_period,

    research_field_reported,
    for_2008_division_code,
    for_2020_division_code,
    for_standard_2020,

    composition_pct = composition_pct_num,
    allocation_share_for,

    requested_compute_annual = requested_compute,
    allocated_compute_annual = allocated_compute,
    annual_compute_unit = unit,

    estimated_demand_by_for_proxy,
    estimated_allocated_compute_by_for,

    requested_basis,
    allocated_basis,
```

```
    analytical_use,
    scope_alignment,
    evidence_quality,
    scope,
    value_basis,
    classification_status,
    standardisation_note,
    notes,
    source,
    source_url
  ) |>
  arrange(
    year,
    desc(composition_pct)
  )
```

```
final_for_summary |>
  summarise(
    rows = n(),
    years = n_distinct(year),
    source_periods = n_distinct(source_period),
    missing_composition = sum(is.na(composition_pct)),
    missing_for_name = sum(
      is.na(for_standard_2020) |
        str_trim(for_standard_2020) == ""
    )
  )
```

```
# A tibble: 1 x 5
   rows years source_periods missing_composition missing_for_name
  <int> <int>          <int>               <int>            <int>
1   133    13             13                   0                0
```

```
writexl::write_xlsx(
  final_for_summary,
  path = file.path(
    clean_path,
    "reconciled_field_of_research_summary.xlsx"
  )
)
```

Chapter 6

Data Standardisation

6.1 Purpose

The purpose of this step is to standardise the integrated historical High-Performance Computing (HPC) data to support consistent analysis across reporting years.

This stage uses the reconciled annual demand and allocation summary and the Field of Research (FoR) summary produced during Data Integration.

Standardisation focuses on aligning reporting periods, variable names, units, research field classifications, and data quality labels where necessary. Original values, source information, and derivation assumptions are retained to maintain traceability.

Numerical demand and allocation values are not converted at this stage. Compute-unit conversion is addressed separately in the Data Conversion chapter.

6.2 Reporting Period

Reporting periods are standardised using the ending year of each annual reporting period. For example, 2013/14 is represented as 2014, and 2025/26 as 2026.

The analysis covers 13 annual observations from 2014 to 2026. The year column is retained as the analytical year, while source_period preserves the original reporting period.

Both variables are carried forward from Data Integration. This step validates their consistency without modifying the existing year assignments.

```r
library(dplyr)
library(readxl)

annual_data <- read_excel(
  "assets/data/data clean/reconciled_demand_allocation_summary.xlsx"
)

year_check <- annual_data |>
  mutate(
    expected_year = as.integer(
      sub("^.*[/\\-]", "", source_period)
    ) + 2000L
  ) |>
  filter(is.na(year) | year != expected_year)

year_check
```

```
# A tibble: 0 x 12
# i 12 variables: year <dbl>, source_period <chr>, requested_compute <dbl>,
#   allocated_compute <dbl>, unit <chr>, requested_basis <chr>,
#   allocated_basis <chr>, requested_source <chr>, allocated_source <chr>,
#   derivation_note <chr>, source_url <chr>, expected_year <int>
```

```r
stopifnot(nrow(year_check) == 0)
```

The validation identified no inconsistencies between the analytical years and their corresponding reporting periods.

6.3 Institution Names

Institution names may differ across historical sources because of abbreviations, naming conventions, or organisational changes.

The reconciled annual demand and allocation summary and Field of Research (FoR) summary do not contain institution-level observations. Therefore, no institution-name mapping is applied to these datasets.

Institution-level standardisation may be undertaken separately if institutional allocation data are included in subsequent analysis.

6.4 Field of Research Classification

Historical Field of Research (FoR) data use different classification systems, including ANZSRC 2008 and ANZSRC 2020.

The reconciled FoR dataset from Data Integration retains the originally reported

research fields and classification codes, alongside a standardised FoR category for longitudinal analysis.

Existing mappings are retained, including combined historical categories where a direct one-to-one mapping is not supported. Classification status and standardisation notes document estimated classifications, unresolved mappings, and differences between reporting systems.

This step validates the existing classifications without changing allocation shares, demand estimates, or other numerical values.

```
for_data <- read_excel(
  "assets/data/data clean/reconciled_field_of_research_summary.xlsx"
)

# Review existing classifications
for_classification_check <- for_data |>
  count(
    classification_status,
    standardisation_note,
    name = "records"
  ) |>
  arrange(desc(records))

for_classification_check
```

```
# A tibble: 32 x 3
   classification_status                              standardisation_note records
   <chr>                                              <chr>                  <int>
 1 Derived from source-reported 6-digit FoR percen~ Mapped from reporte~      41
 2 ANZSRC 2020 source codes; allocation-weighted f~ FoR 2020 division r~      21
 3 As reported; FoR 2008/2020 alignment pending     Mapped from FoR 200~      11
 4 As reported; Pawsey area-of-science category; F~ Mapped from FoR 200~      10
 5 Added residual to complete published compositio~ No FoR code               5
 6 Estimated broad research-field classification f~ Mapped from FoR 200~       5
 7 Estimated from 148 supplied positive-NCMAS NCI ~ Mapped from FoR 200~       4
 8 Estimated from direct NCMAS 2021 project titles~ Mapped from FoR 200~       4
 9 Estimated - project-title classification; alloc~ Mapped from FoR 200~       4
10 As reported; FoR 2008/2020 alignment pending     Historical combined~       3
# i 22 more rows
```

```
# Identify records requiring mapping review
for_mapping_review <- for_data |>
  filter(
    is.na(for_standard_2020) |
      grepl(
        "pending|estimated|cannot split",
        paste(classification_status,
```

```
              standardisation_note),
        ignore.case = TRUE
      )
  ) |>
  select(
    year,
    research_field_reported,
    for_standard_2020,
    classification_status,
    standardisation_note
  )

for_mapping_review
```

```
# A tibble: 64 x 5
    year research_field_reported        for_standard_2020  classification_status
   <dbl> <chr>                         <chr>              <chr>
 1  2014 Earth & Environmental Sciences Earth sciences + ~ As reported; FoR 200~
 2  2014 Physical Sciences              Physical sciences  As reported; FoR 200~
 3  2014 Chemical Sciences              Chemical sciences  As reported; FoR 200~
 4  2014 Engineering & Technology       Engineering + Tec~ As reported; FoR 200~
 5  2014 Biological Sciences            Biological scienc~ As reported; FoR 200~
 6  2014 Mathematical Sciences          Mathematical scie~ As reported; FoR 200~
 7  2015 Earth & Environmental Sciences Earth sciences + ~ As reported; FoR 200~
 8  2015 Physical Sciences              Physical sciences  As reported; FoR 200~
 9  2015 Chemical Sciences              Chemical sciences  As reported; FoR 200~
10  2015 Engineering & Technology       Engineering + Tec~ As reported; FoR 200~
# i 54 more rows
# i 1 more variable: standardisation_note <chr>
```

6.4.1 Historical FoR Reporting Groups

For visualisation, the FoR data are divided into two periods: 2014–2020 and 2021–2026.

For 2014 - 2020, originally reported research fields are grouped into consistent historical reporting categories. In particular, the separately reported Geosciences, Earth Sciences, and Environmental Sciences categories are combined into Earth & Environmental Sciences. Radioastronomy is combined with Physical Sciences, while Engineering and Technology are combined into Engineering & Technology.

For 2021 - 2026, the existing standardised ANZSRC 2020 FoR categories are retained.

These groupings are used for reporting and visualisation only. Original research field names, allocation shares, estimated compute values, and source information remain unchanged. Differences in data scope, particularly Pawsey-wide

supercomputer usage in 2016/17, are retained as comparability limitations.

```
# Create FoR reporting categories without modifying source values
for_data <- for_data |>
  mutate(
    for_analysis_group = case_when(
      year <= 2020 &
        research_field_reported %in% c(
          "Geosciences",
          "Earth Sciences (exc. Geosciences)",
          "Environmental Sciences",
          "Earth & Environmental Sciences"
        ) ~ "Earth & Environmental Sciences",

      year <= 2020 &
        research_field_reported %in% c(
          "Radioastronomy",
          "Physical Sciences (exc. radio astronomy)",
          "Physical Sciences"
        ) ~ "Physical Sciences",

      year <= 2020 &
        research_field_reported %in% c(
          "Engineering",
          "Technology",
          "Engineering & Technology"
        ) ~ "Engineering & Technology",

      year <= 2020 &
        research_field_reported %in% c(
          "Other / Unreported",
          "Other / Unclassified"
        ) ~ "Other / Unclassified",

      year <= 2020 ~ research_field_reported,

      TRUE ~ for_standard_2020
    )
  )

# Review the harmonised historical categories
historical_for_check <- for_data |>
  filter(year <= 2020) |>
  group_by(year, for_analysis_group) |>
  summarise(
    allocation_share_pct = sum(
```

```
      allocation_share_for, na.rm = TRUE
    ) * 100,
    .groups = "drop"
  ) |>
  arrange(year, desc(allocation_share_pct))

historical_for_check
```

```
# A tibble: 56 x 3
    year for_analysis_group             allocation_share_pct
   <dbl> <chr>                                         <dbl>
 1  2014 Earth & Environmental Sciences                   42
 2  2014 Physical Sciences                                22
 3  2014 Chemical Sciences                                17
 4  2014 Engineering & Technology                         11
 5  2014 Biological Sciences                               4
 6  2014 Other / Unclassified                              3
 7  2014 Mathematical Sciences                             1
 8  2015 Earth & Environmental Sciences                   39
 9  2015 Physical Sciences                                23
10  2015 Chemical Sciences                                21
# i 46 more rows
```

```
# Confirm the new grouping does not alter annual shares
share_validation <- for_data |>
  group_by(year) |>
  summarise(
    original_share = sum(allocation_share_for, na.rm = TRUE),
    .groups = "drop"
  ) |>
  left_join(
    for_data |>
      group_by(year, for_analysis_group) |>
      summarise(
        grouped_share = sum(allocation_share_for, na.rm = TRUE),
        .groups = "drop"
      ) |>
      group_by(year) |>
      summarise(
        grouped_share = sum(grouped_share),
        .groups = "drop"
      ),
    by = "year"
  ) |>
  mutate(
    difference = original_share - grouped_share
```

```
  )

share_validation
```

```
# A tibble: 13 x 4
    year original_share grouped_share difference
   <dbl>          <dbl>         <dbl>      <dbl>
 1  2014          1             1              0
 2  2015          1             1              0
 3  2016          1             1              0
 4  2017          1             1              0
 5  2018          1             1              0
 6  2019          1             1              0
 7  2020          1             1              0
 8  2021          1             1              0
 9  2022          1.00          1.00           0
10  2023          1.000         1.000          0
11  2024          1             1              0
12  2025          1.000         1.000          0
13  2026          1.000         1.000          0
```

```
stopifnot(
  all(abs(share_validation$difference) < 1e-10)
)

# Confirm every record has a reporting category
missing_groups <- for_data |>
  filter(is.na(for_analysis_group) | for_analysis_group == "")

missing_groups
```

```
# A tibble: 0 x 26
# i 26 variables: year <dbl>, source_period <chr>,
#   research_field_reported <chr>, for_2008_division_code <chr>,
#   for_2020_division_code <chr>, for_standard_2020 <chr>,
#   composition_pct <dbl>, allocation_share_for <dbl>,
#   requested_compute_annual <dbl>, allocated_compute_annual <dbl>,
#   annual_compute_unit <chr>, estimated_demand_by_for_proxy <dbl>,
#   estimated_allocated_compute_by_for <dbl>, requested_basis <chr>, ...
```

```
stopifnot(nrow(missing_groups) == 0)
```

6.5 HPC Facilities and Systems

Historical HPC demand and allocation data originate from Australia's national computing facilities, including the National Computational Infrastructure (NCI)

and Pawsey Supercomputing Research Centre.

The reconciled annual summary retains facility and system information within the original source descriptions. These include NCI allocation records, Pawsey historical reporting, and more recent NCMAS outcomes identifying Gadi and Setonix CPU and GPU resources separately.

Facility and system information is therefore available for some reporting periods, although it is not consistently represented in dedicated structured variables across the integrated annual dataset.

At this stage, the existing facility and system descriptions are retained without modifying the original records. Where system-level information is available, it will inform resource classification and conversion assumptions in the Data Conversion chapter.

Differences between NCI and Pawsey accounting conventions, system configurations, and CPU/GPU resource measures are preserved as comparability considerations. No assumption is made that allocations from different systems are directly equivalent.

6.6 Resource Type

Historical HPC demand and allocation records include different computing resources, primarily CPU and GPU systems. The availability of resource-level information varies across reporting periods.

Earlier reporting years generally provide aggregate allocation or demand figures without a consistent CPU/GPU breakdown. Some historical systems, including Raijin and Magnus, are identified in the source records, but their system-level allocation shares are not always available.

More detailed resource-level information is available for later years. The 2020/21 allocation records identify Magnus, FlashLite, MASSIVE, and Gadi, while the 2021/22 records distinguish MASSIVE, Gadi, Setonix CPU, and Setonix GPU. From 2022/23 onwards, several reporting years provide separate Gadi CPU, Setonix CPU, and Setonix GPU allocations.

However, the availability of resource-level demand data is more limited. In particular, the 2023/24 and 2024/25 demand estimates are derived from aggregate allocation and fulfilment rates, without a directly reported CPU/GPU demand breakdown.

For standardisation, resource types are classified as CPU, GPU, or Unspecified only where supported by the original source information. Records without sufficient resource-level detail retain an unspecified classification rather than being assigned an estimated CPU/GPU split.

No numerical allocation or demand values are redistributed between resource

types at this stage. Resource-specific conversion factors and comparability adjustments are addressed in the Data Conversion chapter.

6.7 Units

The reconciled annual demand and allocation summary contains different computational accounting units across reporting periods.

The 2013/14 records are expressed in kSU. The 2014/15 and 2015/16 records contain mixed units, with demand expressed in million hours and allocation expressed in kSU. The 2016/17 records use million core-hours, while 2017/18 uses million hours and 2018/19 uses million CPU-hours.

From 2019/20 to 2025/26, annual demand and allocation values are recorded in kSU. However, the use of the same unit label does not necessarily establish equivalence between different HPC systems or CPU and GPU resources.

At the standardisation stage, original unit labels are retained without numerical conversion. Mixed-unit records are explicitly identified to prevent inappropriate direct comparisons.

The Data Conversion chapter will assess system-specific conversion factors, distinguish CPU and GPU accounting where supported, and document any assumptions required to establish comparable analytical measures.

6.8 Data Quality and Comparability

The integrated historical HPC dataset combines directly reported demand and allocation figures with values derived from reported demand-to-supply relationships, fulfilment rates, and available computing capacity. Consequently, the quality and comparability of observations vary across reporting years.

6.8.1 Data Provenance and Estimation

Directly reported observations are retained as published in the original sources. Where demand figures are unavailable, existing estimates derived during Data Integration are preserved together with their calculation assumptions and source references.

For example, the 2014/15 and 2015/16 demand estimates are derived from reported demand-to-supply or fulfilment relationships. The 2017/18, 2020/21, 2023/24, and 2024/25 demand estimates similarly rely on reported allocation totals and demand-to-supply or fulfilment information.

The 2016/17 and 2018/19 records include capacity-based information or conservative demand proxies. These are distinguished from directly reported allocation and demand figures.

No demand estimates or allocation values are recalculated during Data Standardisation.

6.8.2 Comparability Across Reporting Years

Historical observations differ in their reporting scope, computational accounting units, and availability of system-level information.

In particular, the 2016/17 data include Pawsey-wide supercomputer usage and capacity information, which are not directly equivalent to NCMAS allocation outcomes in other years.

Differences between CPU and GPU resources, system-specific accounting conventions, and incomplete resource breakdowns also limit direct comparisons across reporting periods.

These limitations are retained in the analysis rather than resolved through unsupported assumptions.

6.8.3 Treatment for Subsequent Analysis

The standardisation process preserves original numerical values, source references, estimation assumptions, and existing data quality information.

Observations with different reporting scopes, derived demand values, or uncertain conversion assumptions will be explicitly considered during Data Conversion, exploratory analysis, and forecasting model evaluation.

Where appropriate, sensitivity analysis will be used to assess how historical proxies and comparability limitations affect the forecasting results.

6.9 Output

The Data Standardisation stage prepares the reconciled historical HPC datasets for subsequent conversion, analysis, and forecasting.

Two datasets are carried forward:

- **Annual Demand and Allocation Summary:** Retains the 13 annual observations from 2014 to 2026, including original demand and allocation values, reporting periods, accounting units, source references, and existing estimation information. No numerical conversions are performed at this stage.
- **Standardised Field of Research Summary:** Retains the original FoR records and classification information while adding `for_analysis_group` for consistent historical reporting and visualisation.

The standardisation process preserves original numerical values, research field labels, source references, and estimation assumptions. Historical FoR categories

are harmonised for reporting purposes without redistributing allocation shares.

Validation checks confirm reporting-period consistency, preservation of annual FoR allocation-share totals, and the absence of missing FoR reporting groups.

The resulting datasets provide the input for Data Conversion, where computational accounting units and system-specific conversion assumptions will be evaluated before numerical comparisons and forecasting are undertaken.

```
# Create a separate folder for standardisation outputs
output_dir <- "assets/data/data standardised"

dir.create(
  output_dir,
  recursive = TRUE,
  showWarnings = FALSE
)

# Save the annual summary without changing its values
saveRDS(
  annual_data,
  file.path(output_dir, "annual_demand_allocation.rds")
)

# Save the FoR dataset with its new reporting categories
saveRDS(
  for_data,
  file.path(output_dir, "field_of_research.rds")
)

# Verify that both output files were created
output_files <- c(
  file.path(output_dir, "annual_demand_allocation.rds"),
  file.path(output_dir, "field_of_research.rds")
)

stopifnot(all(file.exists(output_files)))

output_files
```

```
[1] "assets/data/data standardised/annual_demand_allocation.rds"
[2] "assets/data/data standardised/field_of_research.rds"
```

Chapter 7

Data Conversion

7.1 Purpose

The purpose of this step is to convert historical High-Performance Computing (HPC) demand and allocation values into more comparable computational resource measures across reporting years, facilities, and systems.

The integrated dataset covers 13 financial years, from 2013/14 to 2025/26, and contains observations reported in different accounting units, including kilo service units (kSU), million compute hours, CPU core-hours, and GPU-related service units.

These observations originate from different HPC facilities and systems, including the National Computational Infrastructure (NCI) and Pawsey Supercomputing Research Centre. Historical system information includes Raijin, Magnus, FlashLite, MASSIVE, Gadi, and Setonix CPU and GPU resources.

Service-unit accounting conventions can vary across facilities, systems, resource types, and reporting periods. Consequently, values expressed in the same unit, such as kSU, are not necessarily computationally equivalent.

This stage therefore evaluates conversion rules at the system and resource level rather than applying a single conversion factor across the historical dataset.

Where sufficient documentation is available, original demand and allocation values are converted into CPU core-hours or GPU hours using the applicable accounting rules. Where conversion cannot be supported reliably, original values and their limitations are retained without introducing unsupported assumptions.

The conversion process preserves the original reported values, units, source references, derivation methods, and conversion assumptions to maintain transparency and reproducibility for subsequent exploratory analysis and forecasting.

7.2 Conversion Approach

The conversion process uses the standardised annual demand and allocation dataset produced in Chapter 04, together with the system-level resource information collected during Data Integration.

Each historical observation is assessed according to its reporting period, facility, HPC system, resource type, original accounting unit, and available conversion documentation.

The conversion process follows four principles:

1. **Preserve original observations.** Original demand and allocation values, units, source references, and derivation assumptions are retained without modification.
2. **Identify resource-specific information.** CPU and GPU resources are distinguished where supported by the original source records. Aggregate observations without a documented resource split remain unallocated to a specific resource type.
3. **Apply documented conversion rules.** Conversion factors are selected according to the relevant system, resource type, and reporting period. No universal service-unit-to-compute-hour factor is assumed.
4. **Record conversion uncertainty.** Where direct conversion rules or resource-level observations are unavailable, transparent assumptions are applied to produce estimated values. Each assumption is documented with its rationale, supporting evidence, confidence level, and limitations. Reported, derived, and assumed values remain distinguishable throughout the analysis.

The general relationship between kilo service units (kSU) and service units (SU) is:

$$\text{SU} = \text{kSU} \times 1000$$

However, conversion from service units to physical computational hours depends on the accounting rules of the relevant HPC system.

For example, CPU core-hours and GPU hours require separate conversion factors and cannot be combined directly into a single physical compute-hour measure.

The conversion framework therefore maintains separate demand and allocation measures and records conversion status, supporting evidence, and limitations for each observation.

Where system-level breakdowns are available, they will be assessed separately before aggregation. Historical observations lacking sufficient detail will not be assigned unsupported CPU/GPU proportions.

```r
library(dplyr)

# Load the standardised annual dataset from Chapter 04
annual_standardised <- readRDS(
  "assets/data/data standardised/annual_demand_allocation.rds"
)

# Inspect available variables before conversion
names(annual_standardised)
```

```
 [1] "year"              "source_period"     "requested_compute"
 [4] "allocated_compute" "unit"              "requested_basis"
 [7] "allocated_basis"   "requested_source"  "allocated_source"
[10] "derivation_note"   "source_url"
```

```r
# Confirm the expected reporting years
annual_standardised |>
  select(year, source_period, unit) |>
  arrange(year)
```

```
# A tibble: 13 x 3
    year source_period unit
   <dbl> <chr>         <chr>
 1  2014 2013/14       kSU
 2  2015 2014/15       mixed: million hours requested / kSU allocated
 3  2016 2015/16       mixed: million hours requested / kSU allocated
 4  2017 2016/17       million core hours
 5  2018 2017/18       million hours
 6  2019 2018/19       million CPU-hours
 7  2020 2019/20       kSU
 8  2021 2020/21       kSU
 9  2022 2021/22       kSU
10  2023 2022/23       kSU
11  2024 2023/24       kSU
12  2025 2024/25       kSU
13  2026 2025/26       kSU
```

```r
stopifnot(
  nrow(annual_standardised) == 13,
  all(sort(annual_standardised$year) == 2014:2026)
)
```

```r
# Create the conversion framework
# Original annual values remain unchanged

# Create conversion framework and retain existing kSU values
conversion_data <- annual_standardised |>
```

```
  mutate(
    demand_ksu = if_else(
      unit == "kSU",
      as.numeric(requested_compute),
      NA_real_
    ),

    allocation_ksu = if_else(
      unit == "kSU" |
        grepl("kSU allocated", unit, fixed = TRUE),
      as.numeric(allocated_compute),
      NA_real_
    ),

    demand_cpu_core_hours = NA_real_,
    allocation_cpu_core_hours = NA_real_,
    demand_gpu_hours = NA_real_,
    allocation_gpu_hours = NA_real_,

    conversion_status = case_when(
      unit == "kSU" ~ "Original kSU retained",
      grepl("kSU allocated", unit, fixed = TRUE) ~
        "Allocation kSU retained; demand conversion required",
      TRUE ~ "Historical unit conversion required"
    ),

    conversion_note = NA_character_
  )

conversion_data |>
  select(
    year,
    demand_ksu,
    allocation_ksu,
    conversion_status
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 4
    year demand_ksu allocation_ksu conversion_status
   <dbl>      <dbl>          <dbl> <chr>
 1  2014     131356          83505 Original kSU retained
 2  2015         NA          88731 Allocation kSU retained; demand conversion r~
 3  2016         NA          97112 Allocation kSU retained; demand conversion r~
 4  2017         NA             NA Historical unit conversion required
```

```
 5  2018        NA              NA Historical unit conversion required
 6  2019        NA              NA Historical unit conversion required
 7  2020    607980          255148 Original kSU retained
 8  2021   1193811          397937 Original kSU retained
 9  2022   1162230          581115 Original kSU retained
10  2023   1333240          817956 Original kSU retained
11  2024   1425560.         812569 Original kSU retained
12  2025   1564643.         876200 Original kSU retained
13  2026   2242817          872288 Original kSU retained
```

```
# Check the framework
conversion_data |>
  select(
    year,
    source_period,
    demand_ksu,
    allocation_ksu,
    demand_cpu_core_hours,
    allocation_cpu_core_hours,
    demand_gpu_hours,
    allocation_gpu_hours,
    conversion_status
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 13 x 9
    year source_period demand_ksu allocation_ksu demand_cpu_core_hours
   <dbl> <chr>              <dbl>          <dbl>                  <dbl>
 1  2014 2013/14          131356          83505                     NA
 2  2015 2014/15              NA          88731                     NA
 3  2016 2015/16              NA          97112                     NA
 4  2017 2016/17              NA             NA                     NA
 5  2018 2017/18              NA             NA                     NA
 6  2019 2018/19              NA             NA                     NA
 7  2020 2019/20          607980          255148                     NA
 8  2021 2020/21         1193811          397937                     NA
 9  2022 2021/22         1162230          581115                     NA
10  2023 2022/23         1333240          817956                     NA
11  2024 2023/24         1425560.         812569                     NA
12  2025 2024/25         1564643.         876200                     NA
13  2026 2025/26         2242817          872288                     NA
   allocation_cpu_core_hours demand_gpu_hours allocation_gpu_hours
                       <dbl>            <dbl>                <dbl>
 1                        NA               NA                   NA
 2                        NA               NA                   NA
 3                        NA               NA                   NA
```

```
 4                       NA                NA                    NA
 5                       NA                NA                    NA
 6                       NA                NA                    NA
 7                       NA                NA                    NA
 8                       NA                NA                    NA
 9                       NA                NA                    NA
10                       NA                NA                    NA
11                       NA                NA                    NA
12                       NA                NA                    NA
13                       NA                NA                    NA
   conversion_status
   <chr>
 1 Original kSU retained
 2 Allocation kSU retained; demand conversion required
 3 Allocation kSU retained; demand conversion required
 4 Historical unit conversion required
 5 Historical unit conversion required
 6 Historical unit conversion required
 7 Original kSU retained
 8 Original kSU retained
 9 Original kSU retained
10 Original kSU retained
11 Original kSU retained
12 Original kSU retained
13 Original kSU retained
```

```r
# Confirm no historical observations were lost
stopifnot(
  nrow(conversion_data) == nrow(annual_standardised),
  identical(
    conversion_data$year,
    annual_standardised$year
  )
)
```

```r
# Review original historical values and supporting evidence
historical_conversion_check <- conversion_data |>
  filter(year >= 2014, year <= 2019) |>
  select(
    year,
    source_period,
    requested_compute,
    allocated_compute,
    unit,
    requested_basis,
    allocated_basis,
```

```
    derivation_note
  )

historical_conversion_check |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 6 x 8
   year source_period requested_compute allocated_compute
  <dbl> <chr>                    <dbl>             <dbl>
1  2014 2013/14                131356              83505
2  2015 2014/15                   202.             88731
3  2016 2015/16                   225              97112
4  2017 2016/17                   200                100
5  2018 2017/18                   529.               222
6  2019 2018/19                   660                220
  unit                                          requested_basis allocated_basis
  <chr>                                         <chr>           <chr>
1 kSU                                           Direct          Direct
2 mixed: million hours requested / kSU allocated Derived        Direct
3 mixed: million hours requested / kSU allocated Derived        Direct
4 million core hours                            Proxy           Capacity proxy
5 million hours                                 Derived         Direct
6 million CPU-hours                             Derived proxy   Capacity proxy
  derivation_note
  <chr>
1 <NA>
2 Requested compute derived from reported allocation and demand relationship.
3 Requested compute derived from reported one-third fulfilment relationship.
4 Requested compute retained as a lower-bound proxy; available capacity used as~
5 Requested compute derived from 222 million hours and a 42% fulfilment relatio~
6 Requested compute derived as 220 × 3; available capacity used as the allocati~
```

```
# Historical conversion assumptions
# Baseline: 1 service unit per CPU core-hour
# 1 kSU = 1,000 SU
# 1 million hours = 1,000,000 hours

historical_su_per_cpu_hour <- 1

conversion_data <- conversion_data |>
  mutate(
    demand_ksu = case_when(
      year %in% 2015:2019 ~
        requested_compute * 1000 * historical_su_per_cpu_hour,
      TRUE ~ demand_ksu
    ),
```

```
    allocation_ksu = case_when(
      year %in% 2017:2019 ~
        allocated_compute * 1000 * historical_su_per_cpu_hour,
      TRUE ~ allocation_ksu
    ),

    conversion_status = case_when(
      year == 2014 ~ "Reported kSU; resource conversion assumed",
      year %in% 2015:2016 ~
        "Demand kSU estimated; allocation kSU reported",
      year %in% c(2017, 2019) ~
        "Both kSU estimated; capacity proxy used",
      year == 2018 ~
        "Both kSU estimated; allocation reported",
      TRUE ~ conversion_status
    ),

    conversion_note = case_when(
      year == 2014 ~
        "Original reported kSU retained; CPU/GPU resource split not established.",
      year %in% 2015:2016 ~
        "Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU retain
      year == 2017 ~
        "Million core-hours converted using assumed 1 SU/core-hour; capacity is an allocatio
      year == 2018 ~
        "Million hours converted using assumed 1 SU/hour; original rounded values retained."
      year == 2019 ~
        "Million CPU-hours converted using assumed 1 SU/CPU-hour; available capacity is an a
      TRUE ~ conversion_note
    )
  )

# Display converted historical totals
conversion_data |>
  filter(year <= 2019) |>
  select(
    source_period,
    demand_ksu,
    allocation_ksu,
    conversion_status,
    conversion_note
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 6 x 5
```

```
  source_period demand_ksu allocation_ksu
  <chr>               <dbl>          <dbl>
1 2013/14          131356           83505
2 2014/15          202500           88731
3 2015/16          225000           97112
4 2016/17          200000          100000
5 2017/18          528571.         222000
6 2018/19          660000          220000
  conversion_status
  <chr>
1 Reported kSU; resource conversion assumed
2 Demand kSU estimated; allocation kSU reported
3 Demand kSU estimated; allocation kSU reported
4 Both kSU estimated; capacity proxy used
5 Both kSU estimated; allocation reported
6 Both kSU estimated; capacity proxy used
  conversion_note
  <chr>
1 Original reported kSU retained; CPU/GPU resource split not established.
2 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU ~
3 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU ~
4 Million core-hours converted using assumed 1 SU/core-hour; capacity is an all~
5 Million hours converted using assumed 1 SU/hour; original rounded values reta~
6 Million CPU-hours converted using assumed 1 SU/CPU-hour; available capacity i~
```

```r
conversion_data |>
  filter(year <= 2019) |>
  select(
    source_period,
    demand_ksu,
    allocation_ksu,
    requested_basis,
    allocated_basis,
    source_url,
    conversion_note
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 6 x 7
  source_period demand_ksu allocation_ksu requested_basis allocated_basis
  <chr>               <dbl>          <dbl> <chr>           <chr>
1 2013/14          131356           83505 Direct          Direct
2 2014/15          202500           88731 Derived         Direct
3 2015/16          225000           97112 Derived         Direct
4 2016/17          200000          100000 Proxy           Capacity proxy
5 2017/18          528571.         222000 Derived         Direct
```

```
6 2018/19          660000          220000 Derived proxy   Capacity proxy
  source_url
  <chr>
1 https://nci.org.au/sites/default/files/annual-reports/2019-
07/About%20Us_Annu~
2 https://nci.org.au/sites/default/files/annual-reports/2019-
07/About%20Us_Annu~
3 https://nci.org.au/sites/default/files/annual-reports/2019-
07/About%20Us_Annu~
4 https://pawsey.org.au/wp-content/uploads/2018/10/Pawsey-Annual-
Report-2016-20~
5 https://nci.org.au/news-events/news/ncmas-allocations-now-
available-online
6 https://nci.org.au/sites/default/files/documents/2019-10/NCMAS-
2020-Informati~
  conversion_note
  <chr>
1 Original reported kSU retained; CPU/GPU resource split not established.
2 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU ~
3 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU ~
4 Million core-hours converted using assumed 1 SU/core-hour; capacity is an all~
5 Million hours converted using assumed 1 SU/hour; original rounded values reta~
6 Million CPU-hours converted using assumed 1 SU/CPU-hour; available capacity i~
```

7.3 CPU Resources

CPU resource conversion translates historical service-unit allocations and reported compute hours into CPU core-hours.

The conversion is performed at the HPC system level where sufficient information is available. This is necessary because the accounting relationship between service units and CPU core-hours may differ across facilities and reporting periods.

For observations reported in kSU, CPU core-hours are calculated as:

$$\text{CPU core-hours} = \frac{\text{CPU kSU} \times 1000}{\text{SU per CPU core-hour}}$$

For observations already reported in CPU core-hours, the original values are retained without applying an additional service-unit conversion.

Where aggregate kSU values include both CPU and GPU resources, the CPU component must first be identified or estimated before the conversion formula is applied.

The conversion framework considers the historical systems identified during Data Integration, including Raijin, Magnus, FlashLite, MASSIVE, Gadi, and Setonix

CPU.

Documented accounting factors are prioritised. Where system-specific factors or resource splits are unavailable, assumptions are applied with explicit justification, source references, and confidence classifications.

All calculations preserve the original observations and distinguish directly reported CPU values from derived and assumed estimates.

```
# Historical CPU core-hour conversion: 2013/14-2018/19
# Assumption: 1 SU per CPU core-hour for historical CPU kSU
# Years already reported in millions of hours are scaled by 1,000,000

historical_cpu_share <- 1.0  # Baseline assumption: 100% CPU-equivalent

conversion_data <- conversion_data |>
  mutate(
    demand_cpu_core_hours = case_when(
      year == 2014 ~ demand_ksu * 1000 * historical_cpu_share,
      year %in% 2015:2019 ~ requested_compute * 1e6,
      TRUE ~ demand_cpu_core_hours
    ),

    allocation_cpu_core_hours = case_when(
      year %in% c(2014, 2015, 2016) ~
        allocation_ksu * 1000 * historical_cpu_share,
      year %in% 2017:2019 ~ allocated_compute * 1e6,
      TRUE ~ allocation_cpu_core_hours
    ),

    cpu_conversion_assumption = case_when(
      year == 2014 ~
        "Assumed 100% CPU-equivalent; 1 SU per CPU core-hour. Historical Fornax GPU may be i
      year %in% 2015:2016 ~
        "Assumed reported million hours represent CPU core-hours; aggregate kSU allocation t
      year == 2017 ~
        "Reported million core-hours converted directly; allocation represents available cap
      year == 2018 ~
        "Assumed reported million hours represent CPU core-hours; original system mix not fu
      year == 2019 ~
        "Reported million CPU-hours converted directly; allocation represents available capa
      TRUE ~ NA_character_
    )
  )

conversion_data |>
  filter(year <= 2019) |>
```

```
  select(
    source_period,
    demand_cpu_core_hours,
    allocation_cpu_core_hours,
    cpu_conversion_assumption
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 6 x 4
  source_period demand_cpu_core_hours allocation_cpu_core_hours
  <chr>                         <dbl>                     <dbl>
1 2013/14                   131356000                  83505000
2 2014/15                   202500000                  88731000
3 2015/16                   225000000                  97112000
4 2016/17                   200000000                 100000000
5 2017/18                   528571429.                222000000
6 2018/19                   660000000                 220000000
  cpu_conversion_assumption
  <chr>
1 Assumed 100% CPU-equivalent; 1 SU per CPU core-hour. Historical Fornax GPU ma~
2 Assumed reported million hours represent CPU core-hours; aggregate kSU alloca~
3 Assumed reported million hours represent CPU core-hours; aggregate kSU alloca~
4 Reported million core-hours converted directly; allocation represents availab~
5 Assumed reported million hours represent CPU core-hours; original system mix ~
6 Reported million CPU-hours converted directly; allocation represents availabl~
```

```
# System-level CPU allocation data
# Analytical year 2021 = FY2020/21
# Values are in kSU

cpu_allocation_systems <- tibble::tribble(
  ~year, ~system,       ~cpu_ksu, ~su_per_core_hour,
  2021,  "Magnus",         107730, 1,
  2021,  "FlashLite",         992, 1,
  2021,  "MASSIVE",          2135, 1,
  2021,  "Gadi",           287080, 2,
  2022,  "MASSIVE",          1845, 1,
  2022,  "Gadi",           300120, 2,
  2022,  "Setonix CPU",    108700, 1,
  2023,  "Gadi",           328156, 2,
  2023,  "Setonix CPU",    327400, 1,
  2024,  "Gadi",           305784, 2,
  2024,  "Setonix CPU",    339155, 1,
  2025,  "Gadi",           355200, 2,
  2025,  "Setonix CPU",    349950, 1,
  2026,  "Gadi",           387288, 2,
```

```
  2026,  "Setonix CPU",   325000, 1
)

# Calculate CPU core-hours for each system
cpu_allocation_systems <- cpu_allocation_systems |>
  mutate(
    cpu_core_hours =
      cpu_ksu * 1000 / su_per_core_hour,

    factor_status = case_when(
      system %in% c("FlashLite", "MASSIVE") ~
        "Provisional 1 SU/core-hour assumption",
      TRUE ~
        "Working system-specific factor from conversion workbook"
    )
  )

# Aggregate CPU core-hours by year
cpu_allocation_annual <- cpu_allocation_systems |>
  group_by(year) |>
  summarise(
    allocation_cpu_core_hours =
      sum(cpu_core_hours),
    allocation_cpu_ksu =
      sum(cpu_ksu),
    .groups = "drop"
  )

cpu_allocation_annual |>
  print(n = Inf)
```

```
# A tibble: 6 x 3
   year allocation_cpu_core_hours allocation_cpu_ksu
  <dbl>                     <dbl>              <dbl>
1  2021                 254397000             397937
2  2022                 260605000             410665
3  2023                 491478000             655556
4  2024                 492047000             644939
5  2025                 527550000             705150
6  2026                 518644000             712288
```

```
# Compare system-level CPU allocation with annual totals

cpu_allocation_check <- cpu_allocation_annual |>
  left_join(
    conversion_data |>
```

```
      select(year, source_period, allocation_ksu),
    by = "year"
  ) |>
  mutate(
    non_cpu_ksu =
      allocation_ksu - allocation_cpu_ksu
  )

cpu_allocation_check |>
  select(
    source_period,
    allocation_ksu,
    allocation_cpu_ksu,
    non_cpu_ksu
  ) |>
  print(n = Inf)
```

```
# A tibble: 6 x 4
  source_period allocation_ksu allocation_cpu_ksu non_cpu_ksu
  <chr>                  <dbl>              <dbl>       <dbl>
1 2020/21               397937             397937           0
2 2021/22               581115             410665      170450
3 2022/23               817956             655556      162400
4 2023/24               812569             644939      167630
5 2024/25               876200             705150      171050
6 2025/26               872288             712288      160000
```

```
stopifnot(
  all(cpu_allocation_check$non_cpu_ksu >= 0)
)
```

```
# Integrate system-level CPU allocations
# Preserve existing historical CPU values

conversion_data <- conversion_data |>
  left_join(
    cpu_allocation_annual |>
      select(
        year,
        allocation_cpu_core_hours
      ) |>
      rename(
        calculated_allocation_cpu_core_hours =
          allocation_cpu_core_hours
      ),
    by = "year"
```

```
  ) |>
  mutate(
    allocation_cpu_core_hours = coalesce(
      calculated_allocation_cpu_core_hours,
      allocation_cpu_core_hours
    )
  ) |>
  select(-calculated_allocation_cpu_core_hours)

# Complete FY2019/20 CPU allocation
# Assumption: 100% Gadi CPU-equivalent
# Working factor: 2 SU per CPU core-hour

conversion_data <- conversion_data |>
  mutate(
    allocation_cpu_core_hours = case_when(
      year == 2020 ~ allocation_ksu * 1000 / 2,
      TRUE ~ allocation_cpu_core_hours
    ),

    cpu_conversion_assumption = case_when(
      year == 2020 ~ paste(
        "FY2019/20: aggregate allocation assumed to be",
        "100% Gadi CPU-equivalent at 2 SU/core-hour.",
        "System split unavailable; sensitivity analysis required."
      ),
      TRUE ~ cpu_conversion_assumption
    )
  )

# Display results AFTER correction
conversion_data |>
  select(
    source_period,
    allocation_cpu_core_hours
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 2
   source_period allocation_cpu_core_hours
   <chr>                             <dbl>
 1 2013/14                        83505000
 2 2014/15                        88731000
 3 2015/16                        97112000
 4 2016/17                       100000000
```

```
 5 2017/18                       222000000
 6 2018/19                       220000000
 7 2019/20                       127574000
 8 2020/21                       254397000
 9 2021/22                       260605000
10 2022/23                       491478000
11 2023/24                       492047000
12 2024/25                       527550000
13 2025/26                       518644000
```

```r
# Validate all 13 years
stopifnot(
  nrow(conversion_data) == 13,
  all(!is.na(conversion_data$allocation_cpu_core_hours))
)
```

```r
# CPU demand conversion: FY2019/20-FY2025/26
# Preserve original annual demand and allocation values

# 1. Direct system-level CPU demand (kSU)
cpu_demand_systems <- tibble::tribble(
  ~year, ~system,       ~cpu_ksu, ~su_per_core_hour,
  2023,  "Gadi",          770866, 2,
  2023,  "Setonix CPU",   390033, 1,
  2026,  "Gadi",         1246722, 2,
  2026,  "Setonix CPU",   611675, 1
)

cpu_demand_direct <- cpu_demand_systems |>
  mutate(
    cpu_core_hours = cpu_ksu * 1000 / su_per_core_hour
  ) |>
  group_by(year) |>
  summarise(
    demand_cpu_ksu = sum(cpu_ksu),
    demand_cpu_core_hours = sum(cpu_core_hours),
    .groups = "drop"
  )

# 2. Estimate CPU demand using allocation composition
# FY2019/20: assume 100% Gadi CPU-equivalent at 2 SU/core-hour
# Other years: assume demand follows allocation resource/system mix

cpu_demand_estimated <- conversion_data |>
  filter(year >= 2020) |>
  left_join(
```

```
    cpu_allocation_annual |>
      select(year, allocation_cpu_ksu),
    by = "year"
  ) |>
  mutate(
    allocation_cpu_ksu = case_when(
      year == 2020 ~ allocation_ksu,
      TRUE ~ allocation_cpu_ksu
    ),

    cpu_allocation_share =
      allocation_cpu_ksu / allocation_ksu,

    estimated_demand_cpu_ksu =
      demand_ksu * cpu_allocation_share,

    estimated_demand_cpu_core_hours = case_when(
      year == 2020 ~
        estimated_demand_cpu_ksu * 1000 / 2,

      TRUE ~
        estimated_demand_cpu_ksu *
        (allocation_cpu_core_hours / allocation_cpu_ksu)
    )
  ) |>
  select(
    year,
    estimated_demand_cpu_ksu,
    estimated_demand_cpu_core_hours,
    cpu_allocation_share
  )

# 3. Prioritise direct system-level demand over estimates

cpu_demand_annual <- cpu_demand_estimated |>
  left_join(
    cpu_demand_direct,
    by = "year"
  ) |>
  mutate(
    demand_cpu_ksu = coalesce(
      demand_cpu_ksu,
      estimated_demand_cpu_ksu
    ),
```

```
    demand_cpu_core_hours = coalesce(
      demand_cpu_core_hours,
      estimated_demand_cpu_core_hours
    ),

    cpu_demand_method = case_when(
      year %in% c(2023, 2026) ~
        "Reported CPU system split; working SU conversion factors",

      year == 2020 ~
        "Assumed 100% Gadi CPU-equivalent; 2 SU/core-hour",

      TRUE ~
        "Estimated using allocation CPU share and weighted CPU conversion"
    ),

    cpu_demand_confidence = case_when(
      year %in% c(2023, 2026) ~ "Higher: reported resource split",
      year == 2020 ~ "Low: assumed CPU-only system mix",
      TRUE ~ "Moderate-low: allocation mix proxy"
    ),

    cpu_demand_assumption = case_when(
      year %in% c(2023, 2026) ~
        "CPU demand split is reported; physical hours use working system factors.",

      year == 2020 ~
        "Aggregate demand assumed entirely Gadi CPU at 2 SU/core-hour; GPU demand not indepe

      TRUE ~
        "CPU demand share and system mix assumed equal to allocation composition; actual rec
    )
  )

# 4. Validate all seven years

stopifnot(
  nrow(cpu_demand_annual) == 7,
  all(!is.na(cpu_demand_annual$demand_cpu_ksu)),
  all(!is.na(cpu_demand_annual$demand_cpu_core_hours))
)

# 5. Display results

cpu_demand_annual |>
```

```
  select(
    year,
    demand_cpu_ksu,
    demand_cpu_core_hours,
    cpu_demand_method,
    cpu_demand_confidence
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 7 x 5
   year demand_cpu_ksu demand_cpu_core_hours
  <dbl>          <dbl>                <dbl>
1  2020         607980            303990000
2  2021        1193811            763191000
3  2022         821330            521210000
4  2023        1160899            775466000
5  2024        1131472.           863240351.
6  2025        1259196.           942053571.
7  2026        1858397           1235036000
  cpu_demand_method
  <chr>
1 Assumed 100% Gadi CPU-equivalent; 2 SU/core-hour
2 Estimated using allocation CPU share and weighted CPU conversion
3 Estimated using allocation CPU share and weighted CPU conversion
4 Reported CPU system split; working SU conversion factors
5 Estimated using allocation CPU share and weighted CPU conversion
6 Estimated using allocation CPU share and weighted CPU conversion
7 Reported CPU system split; working SU conversion factors
  cpu_demand_confidence
  <chr>
1 Low: assumed CPU-only system mix
2 Moderate-low: allocation mix proxy
3 Moderate-low: allocation mix proxy
4 Higher: reported resource split
5 Moderate-low: allocation mix proxy
6 Moderate-low: allocation mix proxy
7 Higher: reported resource split
```

```
# Integrate CPU demand into the main conversion dataset
# Preserve historical CPU values and original source observations

conversion_data <- conversion_data |>
  left_join(
    cpu_demand_annual |>
      select(
        year,
```

```
        demand_cpu_core_hours,
        cpu_demand_method,
        cpu_demand_confidence,
        cpu_demand_assumption
      ) |>
      rename(
        calculated_demand_cpu_core_hours =
          demand_cpu_core_hours
      ),
    by = "year"
  ) |>
  mutate(
    demand_cpu_core_hours = coalesce(
      calculated_demand_cpu_core_hours,
      demand_cpu_core_hours
    )
  ) |>
  select(-calculated_demand_cpu_core_hours)

# Validate CPU conversion for all 13 years

stopifnot(
  nrow(conversion_data) == 13,
  all(!is.na(conversion_data$demand_cpu_core_hours)),
  all(!is.na(conversion_data$allocation_cpu_core_hours))
)

# Display consolidated CPU results

conversion_data |>
  select(
    year,
    source_period,
    demand_cpu_core_hours,
    allocation_cpu_core_hours
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 4
    year source_period demand_cpu_core_hours allocation_cpu_core_hours
   <dbl> <chr>                         <dbl>                     <dbl>
 1  2014 2013/14                   131356000                  83505000
 2  2015 2014/15                   202500000                  88731000
 3  2016 2015/16                   225000000                  97112000
 4  2017 2016/17                   200000000                 100000000
```

5	2018	2017/18	528571429.	222000000
6	2019	2018/19	660000000	220000000
7	2020	2019/20	303990000	127574000
8	2021	2020/21	763191000	254397000
9	2022	2021/22	521210000	260605000
10	2023	2022/23	775466000	491478000
11	2024	2023/24	863240351.	492047000
12	2025	2024/25	942053571.	527550000
13	2026	2025/26	1235036000	518644000

7.4 GPU Resources

GPU resource conversion translates GPU-specific service-unit allocations and demand into GPU-hours.

Unlike CPU core-hours, GPU-hours represent the use of GPU resources and depend on the accounting conventions of the relevant HPC system.

For GPU observations reported in kSU, the conversion is:

$$\text{GPU hours} = \frac{\text{GPU kSU} \times 1000}{\text{SU per GPU hour}}$$

The analysis prioritises directly reported GPU demand and allocation figures. Where GPU demand is unavailable, estimates are developed using documented assumptions and available resource-level evidence.

The conversion distinguishes reported GPU observations from estimated values. Historical years without a documented GPU breakdown require separate assessment rather than being automatically classified as having zero GPU demand.

Original source values, conversion factors, assumptions and limitations are preserved for reproducibility.

```
# GPU allocation: FY2021/22-FY2025/26
# Source: reconciled system-level allocation evidence
# Working conversion factor: 128 SU per GPU-hour

gpu_allocation_systems <- tibble::tribble(
  ~year, ~system,       ~gpu_ksu, ~su_per_gpu_hour,
  2022,  "Setonix GPU",    170450, 128,
  2023,  "Setonix GPU",    162400, 128,
  2024,  "Setonix GPU",    167630, 128,
  2025,  "Setonix GPU",    171050, 128,
  2026,  "Setonix GPU",    160000, 128
)
```

```r
gpu_allocation_annual <- gpu_allocation_systems |>
  mutate(
    allocation_gpu_hours =
      gpu_ksu * 1000 / su_per_gpu_hour
  ) |>
  group_by(year) |>
  summarise(
    allocation_gpu_ksu = sum(gpu_ksu),
    allocation_gpu_hours = sum(allocation_gpu_hours),
    .groups = "drop"
  )

# Check GPU allocation against annual non-CPU residuals
gpu_allocation_check <- cpu_allocation_check |>
  left_join(
    gpu_allocation_annual,
    by = "year"
  ) |>
  filter(year >= 2022) |>
  mutate(
    reconciliation_difference =
      non_cpu_ksu - allocation_gpu_ksu
  )

stopifnot(
  all(gpu_allocation_check$reconciliation_difference == 0)
)

# Display GPU allocation results
gpu_allocation_annual |>
  print(n = Inf)
```

```
# A tibble: 5 x 3
   year allocation_gpu_ksu allocation_gpu_hours
  <dbl>              <dbl>                <dbl>
1  2022             170450             1331641.
2  2023             162400             1268750
3  2024             167630             1309609.
4  2025             171050             1336328.
5  2026             160000             1250000
```

```r
# GPU demand: FY2021/22-FY2025/26
# Working factor: 128 SU per Setonix GPU-hour

gpu_demand_direct <- tibble::tribble(
  ~year, ~demand_gpu_ksu,
```

```
  2023, 172341,
  2026, 384420
)

gpu_demand_annual <- conversion_data |>
  filter(year %in% 2022:2026) |>
  left_join(
    gpu_allocation_annual |>
      select(year, allocation_gpu_ksu),
    by = "year"
  ) |>
  left_join(
    gpu_demand_direct,
    by = "year"
  ) |>
  mutate(
    gpu_allocation_share =
      allocation_gpu_ksu / allocation_ksu,

    estimated_demand_gpu_ksu =
      demand_ksu * gpu_allocation_share,

    demand_gpu_ksu = coalesce(
      demand_gpu_ksu,
      estimated_demand_gpu_ksu
    ),

    demand_gpu_hours =
      demand_gpu_ksu * 1000 / 128,

    gpu_demand_method = if_else(
      year %in% c(2023, 2026),
      "Reported Setonix GPU demand; working conversion factor",
      "Estimated using GPU allocation share"
    ),

    gpu_demand_confidence = if_else(
      year %in% c(2023, 2026),
      "Higher: reported GPU demand",
      "Moderate-low: allocation share assumption"
    )
  ) |>
  select(
    year,
    demand_gpu_ksu,
```

```
    demand_gpu_hours,
    gpu_demand_method,
    gpu_demand_confidence
  )

# Validate GPU demand
stopifnot(
  nrow(gpu_demand_annual) == 5,
  all(!is.na(gpu_demand_annual$demand_gpu_hours))
)

gpu_demand_annual |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 5 x 5
   year demand_gpu_ksu demand_gpu_hours
  <dbl>          <dbl>            <dbl>
1  2022         340900          2663281.
2  2023         172341          1346414.
3  2024         294088.         2297560.
4  2025         305446.         2386300.
5  2026         384420          3003281.
  gpu_demand_method
  <chr>
1 Estimated using GPU allocation share
2 Reported Setonix GPU demand; working conversion factor
3 Estimated using GPU allocation share
4 Estimated using GPU allocation share
5 Reported Setonix GPU demand; working conversion factor
  gpu_demand_confidence
  <chr>
1 Moderate-low: allocation share assumption
2 Higher: reported GPU demand
3 Moderate-low: allocation share assumption
4 Moderate-low: allocation share assumption
5 Higher: reported GPU demand
```

```
# Reconcile CPU and GPU demand against annual demand

resource_demand_check <- cpu_demand_annual |>
  select(year, demand_cpu_ksu) |>
  inner_join(
    gpu_demand_annual |>
      select(year, demand_gpu_ksu),
    by = "year"
  ) |>
```

```
  left_join(
    conversion_data |>
      select(year, demand_ksu),
    by = "year"
  ) |>
  mutate(
    combined_demand_ksu =
      demand_cpu_ksu + demand_gpu_ksu,

    reconciliation_difference =
      demand_ksu - combined_demand_ksu
  )

resource_demand_check |>
  print(n = Inf)
```

```
# A tibble: 5 x 6
   year demand_cpu_ksu demand_gpu_ksu demand_ksu combined_demand_ksu
  <dbl>          <dbl>          <dbl>      <dbl>               <dbl>
1  2022         821330         340900    1162230             1162230
2  2023        1160899         172341    1333240             1333240
3  2024        1131472.        294088.   1425560.            1425560.
4  2025        1259196.        305446.   1564643.            1564643.
5  2026        1858397         384420    2242817             2242817
# i 1 more variable: reconciliation_difference <dbl>
```

```
stopifnot(
  all(abs(resource_demand_check$reconciliation_difference) < 1e-6)
)
```

```
# Integrate GPU demand and allocation into the annual conversion dataset
# Preserve original values and existing CPU calculations

gpu_conversion_annual <- gpu_allocation_annual |>
  full_join(
    gpu_demand_annual |>
      select(
        year,
        demand_gpu_ksu,
        demand_gpu_hours,
        gpu_demand_method,
        gpu_demand_confidence
      ),
    by = "year"
  )
```

```
conversion_data <- conversion_data |>
  left_join(
    gpu_conversion_annual |>
      rename(
        calculated_demand_gpu_hours = demand_gpu_hours,
        calculated_allocation_gpu_hours = allocation_gpu_hours
      ),
    by = "year"
  ) |>
  mutate(
    demand_gpu_hours = coalesce(
      calculated_demand_gpu_hours,
      demand_gpu_hours
    ),
    allocation_gpu_hours = coalesce(
      calculated_allocation_gpu_hours,
      allocation_gpu_hours
    )
  ) |>
  select(
    -calculated_demand_gpu_hours,
    -calculated_allocation_gpu_hours
  )

# Check converted GPU values
conversion_data |>
  select(
    year,
    source_period,
    demand_gpu_hours,
    allocation_gpu_hours
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 4
    year source_period demand_gpu_hours allocation_gpu_hours
   <dbl> <chr>                    <dbl>                <dbl>
 1  2014 2013/14                     NA                   NA
 2  2015 2014/15                     NA                   NA
 3  2016 2015/16                     NA                   NA
 4  2017 2016/17                     NA                   NA
 5  2018 2017/18                     NA                   NA
 6  2019 2018/19                     NA                   NA
 7  2020 2019/20                     NA                   NA
 8  2021 2020/21                     NA                   NA
```

```
 9  2022 2021/22              2663281.            1331641.
10  2023 2022/23              1346414.            1268750
11  2024 2023/24              2297560.            1309609.
12  2025 2024/25              2386300.            1336328.
13  2026 2025/26              3003281.            1250000
```

```
# Validate years with reported Setonix GPU allocations
stopifnot(
  nrow(conversion_data) == 13,
  all(!is.na(
    conversion_data$demand_gpu_hours[
      conversion_data$year %in% 2022:2026
    ]
  )),
  all(!is.na(
    conversion_data$allocation_gpu_hours[
      conversion_data$year %in% 2022:2026
    ]
  ))
)
```

```
# Historical GPU estimation framework: FY2013/14-FY2020/21
# Assumptions must be documented and supported by historical evidence

historical_gpu_assumptions <- tibble::tibble(
  year = 2014:2021,

  gpu_share_assumed = c(
    NA_real_,  # 2013/14 - Fornax evidence requires assessment
    NA_real_,  # 2014/15
    NA_real_,  # 2015/16
    NA_real_,  # 2016/17
    NA_real_,  # 2017/18
    NA_real_,  # 2018/19
    NA_real_,  # 2019/20
    NA_real_   # 2020/21 - FlashLite evidence requires assessment
  ),

  su_per_gpu_hour = NA_real_,

  assumption_status = "Requires source-informed estimate"
)

historical_gpu_assumptions
```

```
# A tibble: 8 x 4
   year gpu_share_assumed su_per_gpu_hour assumption_status
```

```
  <int>             <dbl>           <dbl> <chr>
1  2014                NA              NA Requires source-informed estimate
2  2015                NA              NA Requires source-informed estimate
3  2016                NA              NA Requires source-informed estimate
4  2017                NA              NA Requires source-informed estimate
5  2018                NA              NA Requires source-informed estimate
6  2019                NA              NA Requires source-informed estimate
7  2020                NA              NA Requires source-informed estimate
8  2021                NA              NA Requires source-informed estimate
```

```
historical_gpu_evidence <- tibble::tribble(
  ~year, ~system, ~evidence_type, ~reported_ksu,
  ~source_url,

  2014, "Fornax",
  "2014 NCMAS GPU capacity available",
  1440,
  "https://nci.org.au/news-events/news/ncmas-information-applicants",

  2021, "FlashLite",
  "System included in 2021 NCMAS; GPU split not reported",
  NA_real_,
  "https://pawsey.org.au/ncmas2021/"
)

historical_gpu_evidence |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 2 x 5
   year system    evidence_type
  <dbl> <chr>     <chr>
1  2014 Fornax    2014 NCMAS GPU capacity available
2  2021 FlashLite System included in 2021 NCMAS; GPU split not reported
  reported_ksu source_url
         <dbl> <chr>
1         1440 https://nci.org.au/news-events/news/ncmas-information-
applicants
2           NA https://pawsey.org.au/ncmas2021/
```

```
# Historical system classification based on UQ documentation

flashlite_classification <- tibble::tibble(
  year = 2021,
  system = "FlashLite",
  allocation_ksu = 992,
  resource_class = "CPU",
  gpu_count = 0L,
```

```
  gpu_allocation_ksu = 0,
  evidence_status = "Verified system specification",
  source_url = paste0(
    "https://rcc.uq.edu.au/files/2009/",
    "FlashLite_User_Guide_20200122.pdf"
  )
)

flashlite_classification
```

```
# A tibble: 1 x 8
   year system    allocation_ksu resource_class gpu_count gpu_allocation_ksu
  <dbl> <chr>              <dbl> <chr>              <int>               <dbl>
1  2021 FlashLite            992 CPU                    0                   0
# i 2 more variables: evidence_status <chr>, source_url <chr>
```

```
# Gadi historical GPU conversion reference
# Charging factor from Pawsey's Gadi/Setonix comparison

gadi_gpu_reference <- tibble::tibble(
  year = 2020:2021,
  system = "Gadi GPU (NVIDIA V100)",
  su_per_gpu_hour = 36,
  conversion_formula = "GPU kSU * 1000 / 36",
  allocation_gpu_ksu = NA_real_,
  allocation_gpu_hours = NA_real_,
  evidence_status = paste(
    "GPU charging factor documented;",
    "historical GPU allocation split unavailable"
  ),
  source_url = paste0(
    "https://pawsey.atlassian.net/wiki/spaces/US/pages/",
    "51927586/Changes+in+Supercomputing+Services+for+2023"
  )
)

gadi_gpu_reference |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 2 x 8
   year system                 su_per_gpu_hour conversion_formula
  <int> <chr>                            <dbl> <chr>
1  2020 Gadi GPU (NVIDIA V100)              36 GPU kSU * 1000 / 36
2  2021 Gadi GPU (NVIDIA V100)              36 GPU kSU * 1000 / 36
  allocation_gpu_ksu allocation_gpu_hours
               <dbl>                <dbl>
```

```
1                 NA                    NA
2                 NA                    NA
  evidence_status
  <chr>
1 GPU charging factor documented; historical GPU allocation split unavailable
2 GPU charging factor documented; historical GPU allocation split unavailable
  source_url
  <chr>
1 https://pawsey.atlassian.net/wiki/spaces/US/pages/51927586/Changes+in+Superco~
2 https://pawsey.atlassian.net/wiki/spaces/US/pages/51927586/Changes+in+Superco~
```

```r
# Historical GPU conversion status
# Missing GPU-hours are unknown, not zero.

historical_gpu_status <- tibble::tibble(
  year = 2014:2021,
  gpu_system_evidence = c(
    "Fornax",
    "Historical GPU split not established",
    "Historical GPU split not established",
    "Historical GPU split not established",
    "Historical GPU split not established",
    "Historical GPU split not established",
    "Gadi",
    "Gadi; FlashLite classified as CPU"
  ),
  gpu_allocation_status = c(
    "Fornax capacity reported; allocation unconfirmed",
    rep("GPU allocation split unavailable", 5),
    "Gadi GPU allocation split unavailable",
    "Gadi GPU allocation split unavailable"
  ),
  gpu_conversion_status = "Not estimable from available allocation data",
  gpu_estimation_method = "No unsupported GPU share assigned"
)

# Check historical coverage
stopifnot(
  nrow(historical_gpu_status) == 8,
  identical(historical_gpu_status$year, 2014:2021)
)

historical_gpu_status |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 8 x 5
```

```
   year gpu_system_evidence
  <int> <chr>
1  2014 Fornax
2  2015 Historical GPU split not established
3  2016 Historical GPU split not established
4  2017 Historical GPU split not established
5  2018 Historical GPU split not established
6  2019 Historical GPU split not established
7  2020 Gadi
8  2021 Gadi; FlashLite classified as CPU
  gpu_allocation_status
  <chr>
1 Fornax capacity reported; allocation unconfirmed
2 GPU allocation split unavailable
3 GPU allocation split unavailable
4 GPU allocation split unavailable
5 GPU allocation split unavailable
6 GPU allocation split unavailable
7 Gadi GPU allocation split unavailable
8 Gadi GPU allocation split unavailable
  gpu_conversion_status                          gpu_estimation_method
  <chr>                                          <chr>
1 Not estimable from available allocation data No unsupported GPU share assigned
2 Not estimable from available allocation data No unsupported GPU share assigned
3 Not estimable from available allocation data No unsupported GPU share assigned
4 Not estimable from available allocation data No unsupported GPU share assigned
5 Not estimable from available allocation data No unsupported GPU share assigned
6 Not estimable from available allocation data No unsupported GPU share assigned
7 Not estimable from available allocation data No unsupported GPU share assigned
8 Not estimable from available allocation data No unsupported GPU share assigned
```

7.4.1 Historical GPU Estimation Scenarios

Historical GPU allocation and demand are estimated using low, baseline, and high GPU-share scenarios for 2013/14–2020/21. These percentages are modelling assumptions rather than observed historical shares.

The provisional conversion factor of 36 SU per GPU-hour is used as a common scenario proxy. It is not a verified conversion factor for Fornax or Topaz.

```
# Step 1: Historical GPU assumptions

historical_gpu_scenarios <- tibble::tribble(
  ~year, ~gpu_system_proxy, ~gpu_factor_proxy,
  ~gpu_share_low, ~gpu_share_baseline, ~gpu_share_high,

  2014, "Fornax",         36, 0.005, 0.010, 0.020,
```

```
  2015, "Fornax",          36, 0.005, 0.012, 0.025,
  2016, "Fornax",          36, 0.007, 0.015, 0.030,
  2017, "Fornax",          36, 0.008, 0.018, 0.035,
  2018, "Fornax",          36, 0.010, 0.020, 0.040,
  2019, "Fornax / Topaz", 36, 0.010, 0.025, 0.045,
  2020, "Gadi",            36, 0.015, 0.030, 0.050,
  2021, "Gadi",            36, 0.020, 0.035, 0.060
)

# Step 2: Calculate historical GPU estimates

historical_gpu_estimates <- conversion_data |>
  dplyr::filter(year %in% 2014:2021) |>
  dplyr::select(
    year,
    source_period,
    demand_ksu,
    allocation_ksu
  ) |>
  dplyr::left_join(
    historical_gpu_scenarios,
    by = "year"
  ) |>
  dplyr::mutate(
    allocation_gpu_ksu_low =
      allocation_ksu * gpu_share_low,

    allocation_gpu_ksu_baseline =
      allocation_ksu * gpu_share_baseline,

    allocation_gpu_ksu_high =
      allocation_ksu * gpu_share_high,

    demand_gpu_ksu_low =
      demand_ksu * gpu_share_low,

    demand_gpu_ksu_baseline =
      demand_ksu * gpu_share_baseline,

    demand_gpu_ksu_high =
      demand_ksu * gpu_share_high,

    allocation_gpu_hours_low =
      allocation_gpu_ksu_low * 1000 / gpu_factor_proxy,
```

```
    allocation_gpu_hours_baseline =
      allocation_gpu_ksu_baseline * 1000 / gpu_factor_proxy,

    allocation_gpu_hours_high =
      allocation_gpu_ksu_high * 1000 / gpu_factor_proxy,

    demand_gpu_hours_low =
      demand_gpu_ksu_low * 1000 / gpu_factor_proxy,

    demand_gpu_hours_baseline =
      demand_gpu_ksu_baseline * 1000 / gpu_factor_proxy,

    demand_gpu_hours_high =
      demand_gpu_ksu_high * 1000 / gpu_factor_proxy
  )

# Step 3: Validate historical estimates

stopifnot(
  nrow(historical_gpu_estimates) == 8,

  all(
    historical_gpu_estimates$gpu_share_low <=
      historical_gpu_estimates$gpu_share_baseline
  ),

  all(
    historical_gpu_estimates$gpu_share_baseline <=
      historical_gpu_estimates$gpu_share_high
  ),

  all(
    historical_gpu_estimates$allocation_gpu_ksu_high <=
      historical_gpu_estimates$allocation_ksu
  ),

  all(
    historical_gpu_estimates$demand_gpu_ksu_high <=
      historical_gpu_estimates$demand_ksu
  )
)

# Step 4: Display baseline results

historical_gpu_estimates |>
```

```
  dplyr::select(
    year,
    source_period,
    gpu_system_proxy,
    gpu_share_baseline,
    allocation_gpu_hours_baseline,
    demand_gpu_hours_baseline
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 8 x 6
   year source_period gpu_system_proxy gpu_share_baseline
  <dbl> <chr>         <chr>                         <dbl>
1  2014 2013/14       Fornax                        0.01
2  2015 2014/15       Fornax                        0.012
3  2016 2015/16       Fornax                        0.015
4  2017 2016/17       Fornax                        0.018
5  2018 2017/18       Fornax                        0.02
6  2019 2018/19       Fornax / Topaz                0.025
7  2020 2019/20       Gadi                          0.03
8  2021 2020/21       Gadi                          0.035
  allocation_gpu_hours_baseline demand_gpu_hours_baseline
                          <dbl>                     <dbl>
1                        23196.                    36488.
2                        29577                     67500
3                        40463.                    93750
4                        50000                    100000
5                       123333.                   293651.
6                       152778.                   458333.
7                       212623.                   506650
8                       386883.                  1160650.
```

```
# Reconcile historical CPU and GPU service units
# Baseline scenario only

historical_resource_reconciliation <- historical_gpu_estimates |>
  dplyr::mutate(
    allocation_cpu_ksu_baseline =
      allocation_ksu - allocation_gpu_ksu_baseline,

    demand_cpu_ksu_baseline =
      demand_ksu - demand_gpu_ksu_baseline,

    allocation_reconciliation =
      allocation_cpu_ksu_baseline +
      allocation_gpu_ksu_baseline -
```

```
      allocation_ksu,

    demand_reconciliation =
      demand_cpu_ksu_baseline +
      demand_gpu_ksu_baseline -
      demand_ksu
  )

# Validate reconciliation
stopifnot(
  all(abs(
    historical_resource_reconciliation$allocation_reconciliation
  ) < 1e-6),

  all(abs(
    historical_resource_reconciliation$demand_reconciliation
  ) < 1e-6)
)

# Display results
historical_resource_reconciliation |>
  dplyr::select(
    year,
    source_period,
    allocation_cpu_ksu_baseline,
    allocation_gpu_ksu_baseline,
    demand_cpu_ksu_baseline,
    demand_gpu_ksu_baseline,
    allocation_reconciliation,
    demand_reconciliation
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 8 x 8
   year source_period allocation_cpu_ksu_baseline allocation_gpu_ksu_baseline
  <dbl> <chr>                              <dbl>                        <dbl>
1  2014 2013/14                           82670.                        835.
2  2015 2014/15                           87666.                       1065.
3  2016 2015/16                           95655.                       1457.
4  2017 2016/17                           98200                        1800
5  2018 2017/18                          217560                        4440
6  2019 2018/19                          214500                        5500
7  2020 2019/20                          247494.                       7654.
8  2021 2020/21                          384009.                      13928.
  demand_cpu_ksu_baseline demand_gpu_ksu_baseline allocation_reconciliation
```

```
                  <dbl>                  <dbl>                  <dbl>
1                130042.                 1314.                      0
2                200070                  2430                       0
3                221625                  3375                       0
4                196400                  3600                       0
5                518000.                10571.                      0
6                643500                 16500                       0
7                589741.                18239.                      0
8               1152028.                41783.                      0
  demand_reconciliation
                  <dbl>
1                     0
2                     0
3                     0
4                     0
5                     0
6                     0
7                     0
8                     0
```

```r
# Adjust historical CPU core-hours after separating GPU resources
# Preserve original CPU-equivalent estimates

historical_cpu_adjusted <- historical_resource_reconciliation |>
  dplyr::left_join(
    conversion_data |>
      dplyr::select(
        year,
        original_demand_cpu_core_hours = demand_cpu_core_hours,
        original_allocation_cpu_core_hours = allocation_cpu_core_hours
      ),
    by = "year"
  ) |>
  dplyr::mutate(
    # Preserve each year's original implied CPU conversion rate
    demand_cpu_core_hours_baseline =
      original_demand_cpu_core_hours *
      (demand_cpu_ksu_baseline / demand_ksu),

    allocation_cpu_core_hours_baseline =
      original_allocation_cpu_core_hours *
      (allocation_cpu_ksu_baseline / allocation_ksu)
  )

# Check results
historical_cpu_adjusted |>
```

```
  dplyr::select(
    year,
    source_period,
    original_demand_cpu_core_hours,
    demand_cpu_core_hours_baseline,
    original_allocation_cpu_core_hours,
    allocation_cpu_core_hours_baseline
  ) |>
  print(n = Inf)
```

```
# A tibble: 8 x 6
   year source_period original_demand_cpu_core_hours demand_cpu_core_hours_bas~1
  <dbl> <chr>                                  <dbl>                        <dbl>
1  2014 2013/14                           131356000                    130042440
2  2015 2014/15                           202500000                    200070000
3  2016 2015/16                           225000000                    221625000
4  2017 2016/17                           200000000                    196400000
5  2018 2017/18                           528571429.                   518000000.
6  2019 2018/19                           660000000                    643500000
7  2020 2019/20                           303990000                    294870300
8  2021 2020/21                           763191000                    736479315
# i abbreviated name: 1: demand_cpu_core_hours_baseline
# i 2 more variables: original_allocation_cpu_core_hours <dbl>,
#   allocation_cpu_core_hours_baseline <dbl>
```

```
# Validate adjusted CPU values
stopifnot(
  nrow(historical_cpu_adjusted) == 8,
  all(
    historical_cpu_adjusted$demand_cpu_core_hours_baseline <=
      historical_cpu_adjusted$original_demand_cpu_core_hours
  ),
  all(
    historical_cpu_adjusted$allocation_cpu_core_hours_baseline <=
      historical_cpu_adjusted$original_allocation_cpu_core_hours
  )
)
```

```
# Integrate historical CPU/GPU baseline estimates
# Preserve original CPU-equivalent values

historical_resource_baseline <- historical_cpu_adjusted |>
  dplyr::transmute(
    year,
    historical_demand_cpu_core_hours =
      demand_cpu_core_hours_baseline,
```

```
    historical_allocation_cpu_core_hours =
      allocation_cpu_core_hours_baseline,
    historical_demand_gpu_hours =
      demand_gpu_hours_baseline,
    historical_allocation_gpu_hours =
      allocation_gpu_hours_baseline,
    historical_demand_gpu_ksu =
      demand_gpu_ksu_baseline,
    historical_allocation_gpu_ksu =
      allocation_gpu_ksu_baseline,
    historical_demand_cpu_ksu =
      demand_cpu_ksu_baseline,
    historical_allocation_cpu_ksu =
      allocation_cpu_ksu_baseline
  )

# Preserve original CPU-equivalent values
conversion_data <- conversion_data |>
  dplyr::mutate(
    original_demand_cpu_core_hours = demand_cpu_core_hours,
    original_allocation_cpu_core_hours = allocation_cpu_core_hours
  ) |>
  dplyr::left_join(
    historical_resource_baseline,
    by = "year"
  ) |>
  dplyr::mutate(
    demand_cpu_core_hours = dplyr::coalesce(
      historical_demand_cpu_core_hours,
      demand_cpu_core_hours
    ),
    allocation_cpu_core_hours = dplyr::coalesce(
      historical_allocation_cpu_core_hours,
      allocation_cpu_core_hours
    ),
    demand_gpu_hours = dplyr::coalesce(
      historical_demand_gpu_hours,
      demand_gpu_hours
    ),
    allocation_gpu_hours = dplyr::coalesce(
      historical_allocation_gpu_hours,
      allocation_gpu_hours
    )
  ) |>
  dplyr::select(
```

```
    -dplyr::starts_with("historical_")
  )

# Validate complete annual CPU/GPU baseline
stopifnot(
  nrow(conversion_data) == 13,
  all(!is.na(conversion_data$demand_cpu_core_hours)),
  all(!is.na(conversion_data$allocation_cpu_core_hours)),
  all(!is.na(conversion_data$demand_gpu_hours)),
  all(!is.na(conversion_data$allocation_gpu_hours))
)

conversion_data |>
  dplyr::select(
    year,
    source_period,
    demand_cpu_core_hours,
    allocation_cpu_core_hours,
    demand_gpu_hours,
    allocation_gpu_hours
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 6
    year source_period demand_cpu_core_hours allocation_cpu_core_hours
   <dbl> <chr>                         <dbl>                     <dbl>
 1  2014 2013/14                   130042440                  82669950
 2  2015 2014/15                   200070000                  87666228
 3  2016 2015/16                   221625000                  95655320
 4  2017 2016/17                   196400000                  98200000
 5  2018 2017/18                   518000000.                217560000
 6  2019 2018/19                   643500000                 214500000
 7  2020 2019/20                   294870300                 123746780
 8  2021 2020/21                   736479315                 245493105
 9  2022 2021/22                   521210000                 260605000
10  2023 2022/23                   775466000                 491478000
11  2024 2023/24                   863240351.                492047000
12  2025 2024/25                   942053571.                527550000
13  2026 2025/26                  1235036000                 518644000
# i 2 more variables: demand_gpu_hours <dbl>, allocation_gpu_hours <dbl>
```

7.5 Field of Research Conversion

The Field of Research (FoR) conversion estimates the distribution of computational resources across research disciplines using the standardised FoR allocation

dataset from Chapter 04.

The conversion uses each field's annual allocation share to distribute the corresponding annual CPU core-hours and GPU-hours. This assumes that the resource composition within each field is proportional to the overall annual allocation composition.

The resulting FoR-level CPU and GPU values are estimates rather than directly observed resource-specific allocations. Original FoR observations and classification variables are retained to support traceability and subsequent analysis.

The 2016/17 reporting period requires particular caution because the available FoR information reflects Pawsey-wide reporting rather than the same NCMAS allocation scope used in other years. This difference will be explicitly flagged.

Before conversion, the FoR allocation shares are checked for completeness and consistency within each reporting year.

```
# Load standardised Field of Research dataset from Chapter 04

for_standardised <- readRDS(
  "assets/data/data standardised/field_of_research.rds"
)

# Inspect available variables
names(for_standardised)
```

```
 [1] "year"                                "source_period"
 [3] "research_field_reported"             "for_2008_division_code"
 [5] "for_2020_division_code"              "for_standard_2020"
 [7] "composition_pct"                     "allocation_share_for"
 [9] "requested_compute_annual"            "allocated_compute_annual"
[11] "annual_compute_unit"                 "estimated_demand_by_for_proxy"
[13] "estimated_allocated_compute_by_for" "requested_basis"
[15] "allocated_basis"                     "analytical_use"
[17] "scope_alignment"                     "evidence_quality"
[19] "scope"                               "value_basis"
[21] "classification_status"               "standardisation_note"
[23] "notes"                               "source"
[25] "source_url"                          "for_analysis_group"
```

```
# Inspect the structure
dplyr::glimpse(for_standardised)
```

```
Rows: 133
Columns: 26
$ year                              <dbl> 2014, 2014, 2014, 2014, 2014, 2014,~
$ source_period                     <chr> "2013/14", "2013/14", "2013/14", "2~
$ research_field_reported           <chr> "Earth & Environmental Sciences", "~
```

```
$ for_2008_division_code             <chr> "04 + 05", "02", "03", "09 + 10", "~
$ for_2020_division_code             <chr> "37 + 41", "51", "34", "40 + unmapp~
$ for_standard_2020                  <chr> "Earth sciences + Environmental sci~
$ composition_pct                    <dbl> 42.00, 22.00, 17.00, 11.00, 4.00, 3~
$ allocation_share_for               <dbl> 0.4200, 0.2200, 0.1700, 0.1100, 0.0~
$ requested_compute_annual           <dbl> 131356.0, 131356.0, 131356.0, 13135~
$ allocated_compute_annual           <dbl> 83505, 83505, 83505, 83505, 83505, ~
$ annual_compute_unit                <chr> "kSU", "kSU", "kSU", "kSU", "kSU", ~
$ estimated_demand_by_for_proxy      <dbl> 55169.520, 28898.320, 22330.520, 14~
$ estimated_allocated_compute_by_for <dbl> 35072.100, 18371.100, 14195.850, 91~
$ requested_basis                    <chr> "Direct", "Direct", "Direct", "Dire~
$ allocated_basis                    <chr> "Direct", "Direct", "Direct", "Dire~
$ analytical_use                     <chr> "Historical / lower confidence", "H~
$ scope_alignment                    <chr> "Historical / approximate", "Histor~
$ evidence_quality                   <chr> "Direct percentage / scope uncertai~
$ scope                              <chr> "NCI compute allocation; exact NCMA~
$ value_basis                        <chr> "Allocation composition", "Allocati~
$ classification_status              <chr> "As reported; FoR 2008/2020 alignme~
$ standardisation_note               <chr> "Historical combined category; cann~
$ notes                              <chr> "Percentage retained exactly as pub~
$ source                             <chr> "NCI Annual Report 2013/14", "NCI A~
$ source_url                         <chr> "https://nci.org.au/sites/default/f~
$ for_analysis_group                 <chr> "Earth & Environmental Sciences", "~
```

```r
# Check allocation shares by financial year
for_share_check <- for_standardised |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    n_fields = dplyr::n(),

    total_allocation_share = sum(
      allocation_share_for,
      na.rm = TRUE
    ),

    missing_shares = sum(
      is.na(allocation_share_for)
    ),

    .groups = "drop"
  )

# Display all reporting years
for_share_check |>
  print(n = Inf)
```

```
# A tibble: 13 x 4
    year n_fields total_allocation_share missing_shares
   <dbl>    <int>                  <dbl>          <int>
 1  2014        7                  1                  0
 2  2015        5                  1                  0
 3  2016        7                  1                  0
 4  2017       17                  1                  0
 5  2018        8                  1                  0
 6  2019        9                  1                  0
 7  2020        8                  1                  0
 8  2021        9                  1                  0
 9  2022       13                  1.00               0
10  2023       12                  1.000              0
11  2024       17                  1                  0
12  2025       12                  1.000              0
13  2026        9                  1.000              0
```

```
# Validate that the annual dataset and FoR dataset
# use consistent analytical years
stopifnot(
  all(for_standardised$year %in% conversion_data$year),
  nrow(conversion_data) == 13
)
```

```
# Convert Field of Research allocations
# Preserve all original FoR observations and variables

for_converted <- for_standardised |>
  dplyr::left_join(
    conversion_data |>
      dplyr::select(
        year,
        annual_allocation_ksu = allocation_ksu,
        annual_allocation_cpu_core_hours =
          allocation_cpu_core_hours,
        annual_allocation_gpu_hours =
          allocation_gpu_hours
      ),
    by = "year"
  ) |>
  dplyr::mutate(
    # FoR allocation expressed in comparable annual kSU
    allocation_for_ksu =
      allocation_share_for * annual_allocation_ksu,

    # Estimated CPU allocation by research field
```

```
    allocation_for_cpu_core_hours =
      allocation_share_for *
      annual_allocation_cpu_core_hours,

    # Estimated GPU allocation by research field
    allocation_for_gpu_hours =
      allocation_share_for *
      annual_allocation_gpu_hours,

    # Record estimation method
    for_resource_conversion_method =
      "Proportional allocation using annual FoR share",

    # Flag reporting scope differences
    for_resource_scope_flag = dplyr::case_when(
      year == 2017 ~
        "Pawsey-wide FoR scope; not directly comparable to NCMAS",
      TRUE ~
        "Annual resource allocation distributed using FoR shares"
    ),

    # Distinguish historical scenario estimates
    for_resource_confidence = dplyr::case_when(
      year <= 2021 ~
        "Low: historical GPU scenario and proportional allocation",
      TRUE ~
        "Estimated: proportional CPU/GPU resource distribution"
    )
  )

# Check converted FoR results
for_converted |>
  dplyr::select(
    year,
    for_analysis_group,
    allocation_share_for,
    allocation_for_ksu,
    allocation_for_cpu_core_hours,
    allocation_for_gpu_hours
  ) |>
  print(n = 20, width = Inf)
```

```
# A tibble: 133 x 6
    year for_analysis_group             allocation_share_for allocation_for_ksu
   <dbl> <chr>                                         <dbl>              <dbl>
```

```
 1  2014 Earth & Environmental Sciences               0.42              35072.
 2  2014 Physical Sciences                            0.22              18371.
 3  2014 Chemical Sciences                            0.17              14196.
 4  2014 Engineering & Technology                     0.11               9186.
 5  2014 Biological Sciences                          0.04               3340.
 6  2014 Other / Unclassified                         0.03               2505.
 7  2014 Mathematical Sciences                        0.01                835.
 8  2015 Earth & Environmental Sciences               0.39              34605.
 9  2015 Physical Sciences                            0.23              20408.
10  2015 Chemical Sciences                            0.21              18634.
11  2015 Engineering & Technology                     0.13              11535.
12  2015 Biological Sciences                          0.04               3549.
13  2016 Earth & Environmental Sciences               0.302             29328.
14  2016 Mathematical Sciences                        0.235             22821.
15  2016 Chemical Sciences                            0.197             19131.
16  2016 Engineering & Technology                     0.165             16023.
17  2016 Biological Sciences                          0.039              3787.
18  2016 Physical Sciences                            0.032              3108.
19  2016 Other / Unclassified                         0.03               2913.
20  2017 Engineering & Technology                     0.305             30540
   allocation_for_cpu_core_hours allocation_for_gpu_hours
                           <dbl>                   <dbl>
 1                      34721379                    9742.
 2                      18187389                    5103.
 3                      14053892.                   3943.
 4                       9093694.                   2552.
 5                       3306798                     928.
 6                       2480098.                    696.
 7                        826700.                    232.
 8                      34189829.                  11535.
 9                      20163232.                   6803.
10                      18409908.                   6211.
11                      11396610.                   3845.
12                       3506649.                   1183.
13                      28887907.                  12220.
14                      22479000.                   9509.
15                      18844098.                   7971.
16                      15783128.                   6676.
17                       3730557.                   1578.
18                       3060970.                   1295.
19                       2869660.                   1214.
20                      29990280                   15270
# i 113 more rows
```

```
# Validate original FoR row count
stopifnot(
```

```
  nrow(for_converted) == nrow(for_standardised),
  nrow(for_converted) == 133
)
```

```
# Reconcile FoR allocations with annual resource totals

for_reconciliation <- for_converted |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    for_total_ksu = sum(allocation_for_ksu),
    for_total_cpu_core_hours =
      sum(allocation_for_cpu_core_hours),
    for_total_gpu_hours =
      sum(allocation_for_gpu_hours),
    .groups = "drop"
  ) |>
  dplyr::left_join(
    conversion_data |>
      dplyr::select(
        year,
        allocation_ksu,
        allocation_cpu_core_hours,
        allocation_gpu_hours
      ),
    by = "year"
  ) |>
  dplyr::mutate(
    ksu_difference =
      for_total_ksu - allocation_ksu,

    cpu_difference =
      for_total_cpu_core_hours -
      allocation_cpu_core_hours,

    gpu_difference =
      for_total_gpu_hours -
      allocation_gpu_hours
  )

# Display reconciliation results
for_reconciliation |>
  dplyr::select(
    year,
    ksu_difference,
    cpu_difference,
```

```
    gpu_difference
  ) |>
  print(n = Inf)
```

```
# A tibble: 13 x 4
    year ksu_difference cpu_difference gpu_difference
   <dbl>          <dbl>          <dbl>          <dbl>
 1  2014     0                0              0
 2  2015     0                0              0
 3  2016     0                0              0
 4  2017     0                0              0
 5  2018     0                0              0
 6  2019     0                0              0
 7  2020     0                0              0
 8  2021     0                0              0
 9  2022     0.00000581       0.00261        0.0000133
10  2023    -0.0000245       -0.0147        -0.0000381
11  2024     0                0              0
12  2025    -0.0000526       -0.0317        -0.0000802
13  2026    -0.0000785       -0.0467        -0.000113
```

```
# Validate differences using a relative tolerance
stopifnot(
  all(abs(for_reconciliation$ksu_difference) <
        pmax(1, abs(for_reconciliation$allocation_ksu)) * 1e-6),

  all(abs(for_reconciliation$cpu_difference) <
        pmax(1, abs(for_reconciliation$allocation_cpu_core_hours)) * 1e-6),

  all(abs(for_reconciliation$gpu_difference) <
        pmax(1, abs(for_reconciliation$allocation_gpu_hours)) * 1e-6)
)
```

7.6 Output

The Data Conversion stage produces three reproducible datasets for subsequent analysis and forecasting.

The annual converted dataset contains 13 financial years of HPC demand and allocation observations, including converted CPU core-hours and GPU-hours. Original reported values, source references, and conversion assumptions are retained.

The Field of Research converted dataset contains 133 observations with estimated allocations in kSU, CPU core-hours, and GPU-hours, distributed according to the standardised annual FoR allocation shares. These resource-specific field

estimates are based on proportional allocation assumptions rather than directly observed CPU or GPU usage.

A separate historical GPU scenario dataset preserves low, baseline, and high estimates for 2013/14 – 2020/21 to support sensitivity analysis.

All converted datasets are stored separately from the standardised source datasets. Reconciliation checks confirm that the FoR allocations reproduce annual resource totals within numerical tolerance.

Historical conversion factors, estimated resource shares, differences in reporting scope, and assumptions regarding CPU/GPU composition remain important limitations for interpretation and forecasting.

```
# Create output directory
output_dir <- "assets/data/data converted"

dir.create(
  output_dir,
  recursive = TRUE,
  showWarnings = FALSE
)

# Save converted annual demand and allocation dataset
saveRDS(
  conversion_data,
  file.path(
    output_dir,
    "annual_demand_allocation_converted.rds"
  )
)

# Save converted Field of Research dataset
saveRDS(
  for_converted,
  file.path(
    output_dir,
    "field_of_research_converted.rds"
  )
)

# Preserve historical GPU uncertainty scenarios separately
saveRDS(
  historical_gpu_estimates,
  file.path(
    output_dir,
    "historical_gpu_scenarios.rds"
```

```
  )
)

# Verify that all files were created
output_files <- file.path(
  output_dir,
  c(
    "annual_demand_allocation_converted.rds",
    "field_of_research_converted.rds",
    "historical_gpu_scenarios.rds"
  )
)

stopifnot(all(file.exists(output_files)))

# Verify exported dataset dimensions
stopifnot(
  nrow(readRDS(output_files[1])) == 13,
  nrow(readRDS(output_files[2])) == 133,
  nrow(readRDS(output_files[3])) == 8
)

cat("All three converted datasets exported successfully.\n")
```

```
All three converted datasets exported successfully.
```

Chapter 8

Initial Analysis

8.1 Purpose

The purpose of this stage is to examine historical demand and allocation patterns for Australian national High-Performance Computing (HPC) resources from 2014 to 2026, using the converted datasets produced in Chapter 05.

The analysis focuses on identifying trends in computational resource demand, allocation levels, and potential differences between requested and allocated resources. CPU core-hours and GPU-hours are examined separately because they represent distinct computational resource types and cannot be directly combined.

The analysis also examines the distribution of HPC allocations across Fields of Research (FoR) to identify how allocated computational resources are distributed among research disciplines and how these patterns have changed over time.

Particular attention is given to historical reporting differences, conversion assumptions, and estimated resource values. These limitations are considered when interpreting trends and assessing the suitability of the data for forecasting.

The findings from this stage will inform subsequent forecasting model development, evaluation, and scenario analysis through 2030.

```
annual_converted <- readRDS(
  "assets/data/data converted/annual_demand_allocation_converted.rds"
)

for_converted <- readRDS(
  "assets/data/data converted/field_of_research_converted.rds"
)
```

```
historical_gpu_scenarios <- readRDS(
  "assets/data/data converted/historical_gpu_scenarios.rds"
)
```

8.2 Historical Demand and Allocation

Historical HPC demand and allocation are examined across the 2014 - 2026 analytical period to identify changes in requested computational resources, awarded allocations, and the difference between them.

Requested compute is used as an indicator of researcher demand, while allocated compute reflects the resources awarded within available infrastructure and programme constraints. These measures are therefore interpreted separately.

The analysis uses the converted annual kSU series to examine demand and allocation trends, annual growth rates, request allocation gaps, and fulfilment rates. However, historical values include directly reported observations, derived estimates, and capacity-based proxies. Consequently, observed changes may partly reflect reporting differences or changes in infrastructure rather than underlying demand alone.

8.2.1 Annual Data Validation

```
library(dplyr)
library(ggplot2)

# Inspect the annual demand and allocation series
annual_converted |>
  select(
    year,
    demand_ksu,
    allocation_ksu,
    conversion_status,
    conversion_note
  ) |>
  arrange(year) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 13 x 5
    year demand_ksu allocation_ksu conversion_status
   <dbl>      <dbl>          <dbl> <chr>
 1  2014     131356          83505 Reported kSU; resource conversion assumed
 2  2015     202500          88731 Demand kSU estimated; allocation kSU reported
 3  2016     225000          97112 Demand kSU estimated; allocation kSU reported
 4  2017     200000         100000 Both kSU estimated; capacity proxy used
```

```
 5  2018    528571.         222000 Both kSU estimated; allocation reported
 6  2019    660000          220000 Both kSU estimated; capacity proxy used
 7  2020    607980          255148 Original kSU retained
 8  2021   1193811          397937 Original kSU retained
 9  2022   1162230          581115 Original kSU retained
10  2023   1333240          817956 Original kSU retained
11  2024   1425560.         812569 Original kSU retained
12  2025   1564643.         876200 Original kSU retained
13  2026   2242817          872288 Original kSU retained
   conversion_note
   <chr>
 1 Original reported kSU retained; CPU/GPU resource split not established.
 2 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU~
 3 Demand million hours x 1,000,000 x assumed 1 SU/hour / 1,000; allocation kSU~
 4 Million core-hours converted using assumed 1 SU/core-hour; capacity is an al~
 5 Million hours converted using assumed 1 SU/hour; original rounded values ret~
 6 Million CPU-hours converted using assumed 1 SU/CPU-hour; available capacity ~
 7 <NA>
 8 <NA>
 9 <NA>
10 <NA>
11 <NA>
12 <NA>
13 <NA>
```

```r
# Check coverage and completeness
stopifnot(
  nrow(annual_converted) == 13,
  all(sort(annual_converted$year) == 2014:2026),
  all(!is.na(annual_converted$demand_ksu)),
  all(!is.na(annual_converted$allocation_ksu)),
  all(annual_converted$demand_ksu > 0),
  all(annual_converted$allocation_ksu >= 0)
)
```

8.2.2 Annual Analytical Indicators

```r
annual_analysis <- annual_converted |>
  arrange(year) |>
  mutate(
    demand_allocation_gap_ksu =
      demand_ksu - allocation_ksu,

    fulfilment_rate_pct =
      100 * allocation_ksu / demand_ksu,
```

```
    demand_growth_pct =
      100 * (demand_ksu / lag(demand_ksu) - 1),

    allocation_growth_pct =
      100 * (allocation_ksu / lag(allocation_ksu) - 1)
  )

annual_analysis |>
  select(
    year,
    demand_ksu,
    allocation_ksu,
    demand_allocation_gap_ksu,
    fulfilment_rate_pct,
    demand_growth_pct,
    allocation_growth_pct
  ) |>
  print(n = Inf)
```

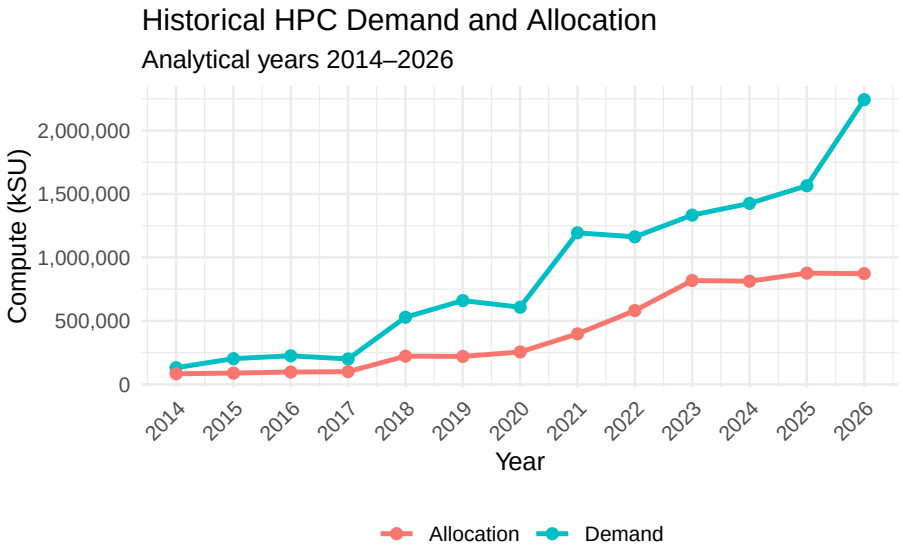
```
# A tibble: 13 x 7
    year demand_ksu allocation_ksu demand_allocation_gap_ksu fulfilment_rate_pct
   <dbl>      <dbl>          <dbl>                     <dbl>               <dbl>
 1  2014     131356          83505                     47851                63.6
 2  2015     202500          88731                    113769                43.8
 3  2016     225000          97112                    127888                43.2
 4  2017     200000         100000                    100000                50
 5  2018     528571.        222000                    306571.               42.0
 6  2019     660000         220000                    440000                33.3
 7  2020     607980         255148                    352832                42.0
 8  2021    1193811         397937                    795874                33.3
 9  2022    1162230         581115                    581115                50
10  2023    1333240         817956                    515284                61.4
11  2024    1425560.        812569                    612991.               57
12  2025    1564643.        876200                    688443.               56
13  2026    2242817         872288                   1370529                38.9
# i 2 more variables: demand_growth_pct <dbl>, allocation_growth_pct <dbl>
```

8.2.3 Annual Demand and Allocation Trends

```
annual_plot_data <- annual_analysis |>
  select(
    year,
    Demand = demand_ksu,
    Allocation = allocation_ksu
  ) |>
```

```
  tidyr::pivot_longer(
    cols = c(Demand, Allocation),
    names_to = "measure",
    values_to = "value_ksu"
  )

ggplot(
  annual_plot_data,
  aes(
    x = year,
    y = value_ksu,
    colour = measure
  )
) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  scale_x_continuous(breaks = 2014:2026) +
  scale_y_continuous(
    labels = scales::label_comma()
  ) +
  labs(
    title = "Historical HPC Demand and Allocation",
    subtitle = "Analytical years 2014-2026",
    x = "Year",
    y = "Compute (kSU)",
    colour = NULL
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1),
    legend.position = "bottom"
  )
```



8.2.4 Initial Findings

The annual HPC demand and allocation series show an overall upward trend between 2014 and 2026. Reported or estimated demand increases from approximately 131,356 kSU in 2014 to 2,242,817 kSU in 2026, while allocation increases from approximately 83,505 kSU to 872,288 kSU.

Demand exceeds allocation in every analytical year. The largest absolute gap occurs in 2026, when demand exceeds allocation by approximately 1.37 million kSU. The corresponding fulfilment rate is approximately 38.9%, compared with 56.0% in 2025.

The series also exhibits periods of substantial change, particularly around 2018, 2021, and 2026. However, these changes should not automatically be interpreted as equivalent increases in underlying researcher demand. Historical conversion assumptions, differences in reporting scope, capacity-based allocation proxies, and changes in HPC systems may contribute to the observed variation.

These findings suggest a persistent difference between requested and allocated computational resources. Further analysis of CPU and GPU resources is required to understand whether similar patterns are present across distinct computational resource types.

8.2.4.1 Interpretation of the 2026 Demand Increase

The marked increase in reported demand in 2026 requires particular caution. It is not yet established whether this represents a sustained increase in underlying researcher demand, a temporary change, or an effect of reporting scope

and conversion assumptions. Following expert statistical feedback, subsequent forecasting will examine the sensitivity of trend estimates to the inclusion of 2026 and compare simple linear and log-linear models.

8.3 CPU and GPU Resources

CPU and GPU resources are examined separately to identify historical patterns in computational demand and allocation across the 2014–2026 analytical period.

The analysis uses the converted CPU core-hours and GPU-hours produced in Chapter 05. These measures provide resource-specific perspectives on HPC demand and allocation, although their interpretation remains subject to system-level accounting conventions and historical conversion assumptions.

Historical CPU estimates are influenced by the assumed resource composition and conversion factors applied to earlier observations. GPU estimates for 2014 - 2021 are based on low, baseline, and high scenarios rather than directly observed resource-level demand and allocation. GPU values for 2022 - 2026 use reported system-level allocations and a combination of reported and estimated demand.

CPU core-hours and GPU-hours are therefore analysed independently and are not added together. The objective is to examine trends within each resource category while retaining the distinction between reported observations and modelled estimates.

8.3.1 Validate the converted CPU and GPU series

```
# Inspect converted CPU and GPU resources
annual_converted |>
  select(
    year,
    demand_cpu_core_hours,
    allocation_cpu_core_hours,
    demand_gpu_hours,
    allocation_gpu_hours
  ) |>
  arrange(year) |>
  print(n = Inf)
```

```
# A tibble: 13 x 5
    year demand_cpu_core_hours allocation_cpu_core_hours demand_gpu_hours
   <dbl>                 <dbl>                     <dbl>            <dbl>
 1  2014             130042440                  82669950           36488.
 2  2015             200070000                  87666228           67500
 3  2016             221625000                  95655320           93750
 4  2017             196400000                  98200000          100000
```

```
 5  2018            518000000.                217560000          293651.
 6  2019            643500000                 214500000          458333.
 7  2020            294870300                 123746780          506650
 8  2021            736479315                 245493105         1160650.
 9  2022            521210000                 260605000         2663281.
10  2023            775466000                 491478000         1346414.
11  2024            863240351.                492047000         2297560.
12  2025            942053571.                527550000         2386300.
13  2026           1235036000                 518644000         3003281.
# i 1 more variable: allocation_gpu_hours <dbl>
```

```
# Check completeness and non-negative values
resource_columns <- c(
  "demand_cpu_core_hours",
  "allocation_cpu_core_hours",
  "demand_gpu_hours",
  "allocation_gpu_hours"
)

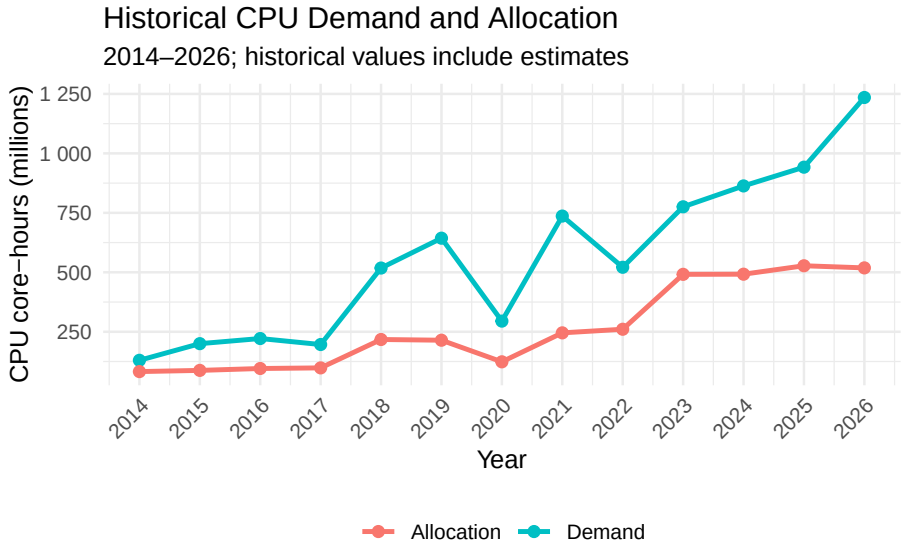
stopifnot(
  nrow(annual_converted) == 13,
  all(sort(annual_converted$year) == 2014:2026),
  all(!is.na(annual_converted[resource_columns])),
  all(as.matrix(annual_converted[resource_columns]) >= 0)
)
```

8.3.2 Create CPU and GPU trend charts

```
# CPU demand and allocation trends
cpu_plot_data <- annual_converted |>
  select(
    year,
    Demand = demand_cpu_core_hours,
    Allocation = allocation_cpu_core_hours
  ) |>
  tidyr::pivot_longer(
    cols = c(Demand, Allocation),
    names_to = "measure",
    values_to = "cpu_core_hours"
  )

ggplot(
  cpu_plot_data,
  aes(
    x = year,
    y = cpu_core_hours,
```

```
    colour = measure
  )
) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  scale_x_continuous(breaks = 2014:2026) +
  scale_y_continuous(
    labels = scales::label_number(
      scale = 1e-6,
      accuracy = 1
    )
  ) +
  labs(
    title = "Historical CPU Demand and Allocation",
    subtitle = "2014-2026; historical values include estimates",
    x = "Year",
    y = "CPU core-hours (millions)",
    colour = NULL
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1),
    legend.position = "bottom"
  )
```



The CPU demand and allocation trends represent converted baseline estimates. Historical CPU values are subject to uncertainty arising from system-specific

accounting conventions, assumed CPU/GPU resource composition, and the use of capacity proxies in selected years. Unlike the historical GPU analysis, separate low and high CPU core-hour scenarios have not been established. Therefore, the CPU chart presents the baseline series without uncertainty bands, and historical fluctuations should be interpreted cautiously.

```
# Validate historical GPU scenario coverage
stopifnot(
  nrow(historical_gpu_scenarios) == 8,

  isTRUE(all.equal(
    sort(as.numeric(historical_gpu_scenarios$year)),
    as.numeric(2014:2021)
  )),

  all(
    historical_gpu_scenarios$demand_gpu_hours_low <=
      historical_gpu_scenarios$demand_gpu_hours_baseline
  ),

  all(
    historical_gpu_scenarios$demand_gpu_hours_baseline <=
      historical_gpu_scenarios$demand_gpu_hours_high
  ),

  all(
    historical_gpu_scenarios$allocation_gpu_hours_low <=
      historical_gpu_scenarios$allocation_gpu_hours_baseline
  ),

  all(
    historical_gpu_scenarios$allocation_gpu_hours_baseline <=
      historical_gpu_scenarios$allocation_gpu_hours_high
  ),
  all(
    abs(
      annual_converted$demand_gpu_hours[
        match(
          historical_gpu_scenarios$year,
          annual_converted$year
        )
      ] -
        historical_gpu_scenarios$demand_gpu_hours_baseline
    ) < 1e-6
  ),
  all(
```

```
    abs(
      annual_converted$allocation_gpu_hours[
        match(
          historical_gpu_scenarios$year,
          annual_converted$year
        )
      ] -
        historical_gpu_scenarios$allocation_gpu_hours_baseline
    ) < 1e-6
  )
)
```

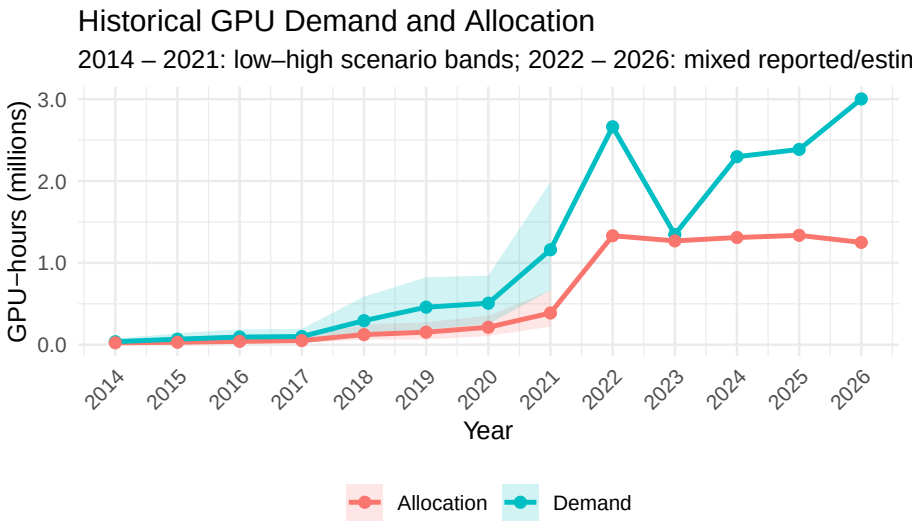
```
# Prepare GPU demand and allocation data
gpu_plot_data <- annual_converted |>
  dplyr::select(
    year,
    Demand = demand_gpu_hours,
    Allocation = allocation_gpu_hours
  ) |>
  tidyr::pivot_longer(
    cols = c(Demand, Allocation),
    names_to = "measure",
    values_to = "gpu_hours"
  )

# Use the existing historical low/high GPU scenarios
gpu_uncertainty_data <- historical_gpu_scenarios |>
  dplyr::select(
    year,
    demand_gpu_hours_low,
    demand_gpu_hours_high,
    allocation_gpu_hours_low,
    allocation_gpu_hours_high
  ) |>
  tidyr::pivot_longer(
    cols = -year,
    names_to = c("measure", ".value"),
    names_pattern = "(demand|allocation)_(gpu_hours_low|gpu_hours_high)"
  ) |>
  dplyr::mutate(
    measure = dplyr::recode(
      measure,
      demand = "Demand",
      allocation = "Allocation"
    )
```

```
  )

# Plot baseline trends with historical uncertainty bands
ggplot() +
  geom_ribbon(
    data = gpu_uncertainty_data,
    aes(
      x = year,
      ymin = gpu_hours_low,
      ymax = gpu_hours_high,
      fill = measure,
      group = measure
    ),
    alpha = 0.18,
    colour = NA
  ) +
  geom_line(
    data = gpu_plot_data,
    aes(
      x = year,
      y = gpu_hours,
      colour = measure
    ),
    linewidth = 1
  ) +
  geom_point(
    data = gpu_plot_data,
    aes(
      x = year,
      y = gpu_hours,
      colour = measure
    ),
    size = 2
  ) +
  scale_x_continuous(
    breaks = 2014:2026
  ) +
  scale_y_continuous(
    labels = scales::label_number(
      scale = 1e-6,
      accuracy = 0.1
    )
  ) +
  labs(
    title = "Historical GPU Demand and Allocation",
```

```
    subtitle = "2014 - 2021: low-high scenario bands; 2022 - 2026: mixed reported/estimated'
    x = "Year",
    y = "GPU-hours (millions)",
    colour = NULL,
    fill = NULL,
    caption = paste(
      "Shaded bands represent assumed historical GPU scenarios,",
      "not statistical confidence intervals."
    )
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(
      angle = 45,
      hjust = 1
    ),
    legend.position = "bottom"
  )
```



Shaded bands represent assumed historical GPU scenarios, not statistical confidence intervals.

The shaded bands represent the range between the low and high historical GPU estimation scenarios for 2014 - 2021, while the solid lines represent the baseline demand and allocation estimates. These bands illustrate sensitivity to assumed historical GPU resource shares rather than statistical confidence intervals. No equivalent uncertainty bands are shown for 2022 - 2026, although those observations remain subject to conversion assumptions and, for some years, estimated GPU demand.

```
historical_gpu_scenarios |>
  select(
    year,
    allocation_gpu_hours_low,
    allocation_gpu_hours_baseline,
    allocation_gpu_hours_high,
    demand_gpu_hours_low,
    demand_gpu_hours_baseline,
    demand_gpu_hours_high
  ) |>
  arrange(year) |>
  print(n = Inf)
```

```
# A tibble: 8 x 7
   year allocation_gpu_hours_low allocation_gpu_hours_b~1 allocation_gpu_hours~2
  <dbl>                    <dbl>                    <dbl>                   <dbl>
1  2014                   11598.                   23196.                  46392.
2  2015                   12324.                   29577                   61619.
3  2016                   18883.                   40463.                  80927.
4  2017                   22222.                   50000                   97222.
5  2018                   61667.                  123333.                 246667.
6  2019                   61111.                  152778.                 275000
7  2020                  106312.                  212623.                 354372.
8  2021                  221076.                  386883.                 663228.
# i abbreviated names: 1: allocation_gpu_hours_baseline,
#   2: allocation_gpu_hours_high
# i 3 more variables: demand_gpu_hours_low <dbl>,
#   demand_gpu_hours_baseline <dbl>, demand_gpu_hours_high <dbl>
```

8.3.3 Initial Findings

The resource-specific analysis indicates an overall increase in estimated CPU and GPU demand between 2014 and 2026, although the two series exhibit different patterns and levels of uncertainty.

CPU Resources

Estimated CPU demand increases from approximately 130.0 million CPU core-hours in 2014 to 1,235.0 million CPU core-hours in 2026. Over the same period, CPU allocation increases from approximately 82.7 million to 518.6 million CPU core-hours.

CPU demand remains above allocation throughout the analytical period. Substantial fluctuations are visible around 2018 – 2022, followed by a sustained increase in estimated demand from 2022 to 2026. In contrast, CPU allocation remains relatively stable between 2023 and 2026, at approximately 491 – 528 million CPU core-hours.

The widening difference between CPU demand and allocation in recent years suggests increasing pressure on allocated CPU resources. However, the magnitude of historical changes should be interpreted cautiously because earlier observations include assumed conversion factors, capacity proxies, and estimated CPU/GPU resource shares.

GPU Resources

Estimated GPU demand increases from approximately 0.04 million GPU-hours in 2014 to 3.00 million GPU-hours in 2026. GPU allocation increases from approximately 0.02 million to 1.25 million GPU-hours over the same period.

The GPU series shows a pronounced increase between 2021 and 2022, followed by fluctuations in estimated demand and comparatively stable allocations during 2022–2026. By 2026, estimated GPU demand is approximately 2.4 times the allocated GPU-hours.

However, the historical GPU series requires particular caution. Values for 2014–2021 are based on assumed GPU resource shares and a provisional conversion factor rather than directly observed GPU demand and allocation. The low, baseline, and high scenarios demonstrate that historical GPU estimates are sensitive to these assumptions.

Furthermore, GPU allocation from 2022 onward is derived from reported system-level service units using a working conversion factor, while GPU demand includes both directly reported and estimated values. The apparent increase around 2022 may therefore partly reflect changes in reporting coverage, infrastructure, and conversion methodology rather than an equivalent increase in underlying GPU demand.

Implications for Forecasting

The CPU and GPU series suggest different historical demand patterns and should be considered separately during forecasting model development. The observed differences between demand and allocation also reinforce the importance of distinguishing researcher demand from computational resources awarded under capacity constraints.

Historical uncertainty, system transitions, and the limited availability of directly observed GPU demand should be considered when selecting modelling periods and evaluating forecast reliability. Scenario analysis will be particularly important for assessing the sensitivity of GPU forecasts to historical conversion assumptions.

8.4 Field of Research

The Field of Research (FoR) analysis examines the distribution of HPC allocations across research disciplines during the 2014–2026 analytical period.

For the overall longitudinal analysis, the original `research_field_reported`

variable is harmonised into broad analytical groups using consistent mapping rules across all years. This approach avoids relying on the period-dependent `for_analysis_group` variable.

The broad groups are intended for exploratory comparison and do not represent official ANZSRC divisions. Original research-field labels and classification codes remain unchanged in the underlying datasets.

The analysis examines allocation composition rather than directly observed researcher demand. FoR-specific kSU, CPU core-hours and GPU-hours are derived using allocation shares and should not be interpreted as independently measured resource consumption by each discipline.

The 2017 observation is interpreted cautiously because its Pawsey-wide reporting scope differs from the other observations. Other source and classification limitations also affect longitudinal comparability.

8.4.1 Validate the Field of Research series

```
# Inspect the original research-field labels and converted measures
for_converted |>
  dplyr::select(
    year,
    research_field_reported,
    allocation_share_for,
    allocation_for_ksu,
    allocation_for_cpu_core_hours,
    allocation_for_gpu_hours,
    for_resource_scope_flag
  ) |>
  dplyr::arrange(
    year,
    dplyr::desc(allocation_share_for)
  ) |>
  print(n = 20, width = Inf)
```

```
# A tibble: 133 x 7
    year research_field_reported        allocation_share_for allocation_for_ksu
   <dbl> <chr>                                        <dbl>              <dbl>
 1  2014 Earth & Environmental Sciences               0.42              35072.
 2  2014 Physical Sciences                            0.22              18371.
 3  2014 Chemical Sciences                            0.17              14196.
 4  2014 Engineering & Technology                     0.11               9186.
 5  2014 Biological Sciences                          0.04               3340.
 6  2014 Other / Unreported                           0.03               2505.
 7  2014 Mathematical Sciences                        0.01                835.
 8  2015 Earth & Environmental Sciences               0.39              34605.
```

```
 9  2015 Physical Sciences                              0.23              20408.
10  2015 Chemical Sciences                              0.21              18634.
11  2015 Engineering & Technology                       0.13              11535.
12  2015 Biological Sciences                            0.04               3549.
13  2016 Earth & Environmental Sciences                 0.302             29328.
14  2016 Mathematical Sciences                          0.235             22821.
15  2016 Chemical Sciences                              0.197             19131.
16  2016 Engineering & Technology                       0.165             16023.
17  2016 Biological Sciences                            0.039              3787.
18  2016 Physical Sciences                              0.032              3108.
19  2016 Other / Unreported                             0.03               2913.
20  2017 Engineering                                    0.305             30540
   allocation_for_cpu_core_hours allocation_for_gpu_hours
                           <dbl>                   <dbl>
 1                      34721379                    9742.
 2                      18187389                    5103.
 3                      14053892.                   3943.
 4                       9093694.                   2552.
 5                       3306798                     928.
 6                       2480098.                    696.
 7                        826700.                    232.
 8                      34189829.                  11535.
 9                      20163232.                   6803.
10                      18409908.                   6211.
11                      11396610.                   3845.
12                       3506649.                   1183.
13                      28887907.                  12220.
14                      22479000.                   9509.
15                      18844098.                   7971.
16                      15783128.                   6676.
17                       3730557.                   1578.
18                       3060970.                   1295.
19                       2869660.                   1214.
20                      29990280                   15270
   for_resource_scope_flag
   <chr>
 1 Annual resource allocation distributed using FoR shares
 2 Annual resource allocation distributed using FoR shares
 3 Annual resource allocation distributed using FoR shares
 4 Annual resource allocation distributed using FoR shares
 5 Annual resource allocation distributed using FoR shares
 6 Annual resource allocation distributed using FoR shares
 7 Annual resource allocation distributed using FoR shares
 8 Annual resource allocation distributed using FoR shares
 9 Annual resource allocation distributed using FoR shares
10 Annual resource allocation distributed using FoR shares
```

```
11 Annual resource allocation distributed using FoR shares
12 Annual resource allocation distributed using FoR shares
13 Annual resource allocation distributed using FoR shares
14 Annual resource allocation distributed using FoR shares
15 Annual resource allocation distributed using FoR shares
16 Annual resource allocation distributed using FoR shares
17 Annual resource allocation distributed using FoR shares
18 Annual resource allocation distributed using FoR shares
19 Annual resource allocation distributed using FoR shares
20 Pawsey-wide FoR scope; not directly comparable to NCMAS
# i 113 more rows
```

```r
# Validate coverage and completeness
stopifnot(
  nrow(for_converted) == 133,
  identical(
    sort(unique(for_converted$year)),
    as.numeric(2014:2026)
  ),
  all(!is.na(for_converted$research_field_reported)),
  all(!is.na(for_converted$allocation_share_for)),
  all(!is.na(for_converted$allocation_for_ksu)),
  all(!is.na(for_converted$allocation_for_cpu_core_hours)),
  all(!is.na(for_converted$allocation_for_gpu_hours))
)

# Check annual allocation shares
for_share_summary <- for_converted |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    n_fields = dplyr::n(),
    total_share = sum(allocation_share_for),
    .groups = "drop"
  )

print(for_share_summary, n = Inf)
```

```
# A tibble: 13 x 3
    year n_fields total_share
   <dbl>    <int>       <dbl>
 1  2014        7           1
 2  2015        5           1
 3  2016        7           1
 4  2017       17           1
 5  2018        8           1
 6  2019        9           1
```

```
 7  2020        8       1
 8  2021        9       1
 9  2022       13       1.00
10  2023       12       1.000
11  2024       17       1
12  2025       12       1.000
13  2026        9       1.000
```

```
stopifnot(
  all(abs(for_share_summary$total_share - 1) < 1e-6)
)
```

8.4.2 Research Field Allocation Composition

The original research-field labels are reviewed before harmonisation. Reported categories may differ in terminology, granularity and classification structure across years.

```
for_category_review <- for_converted |>
  dplyr::group_by(research_field_reported) |>
  dplyr::summarise(
    n_years = dplyr::n_distinct(year),
    first_year = min(year),
    last_year = max(year),
    total_allocation_share = sum(
      allocation_share_for,
      na.rm = TRUE
    ),
    .groups = "drop"
  ) |>
  dplyr::arrange(research_field_reported)

print(for_category_review, n = Inf, width = Inf)
```

```
# A tibble: 51 x 5
   research_field_reported                          n_years first_year last_year
   <chr>                                              <int>      <dbl>     <dbl>
 1 01 Mathematical Sciences                               3       2022      2024
 2 02 Physical Sciences                                   3       2022      2024
 3 03 Chemical Sciences                                   3       2022      2024
 4 04 Earth Sciences                                      3       2022      2024
 5 05 Environmental Sciences                              3       2022      2024
 6 06 Biological Sciences                                 3       2022      2024
 7 07 Agricultural and Veterinary Sciences                3       2022      2024
 8 08 Information and Computing Sciences                  3       2022      2024
 9 09 Engineering                                         3       2022      2024
10 10 Technology                                          3       2022      2024
```

11	11 Medical and Health Sciences	3	2022	2024
12	14 Economics	1	2024	2024
13	16 Studies in Human Society	1	2024	2024
14	17 Psychology and Cognitive Sciences	3	2022	2024
15	20 Language, Communication and Culture	1	2024	2024
16	21 History and Archaeology	1	2024	2024
17	22 Philosophy and Religious Studies	1	2022	2022
18	30 Agricultural, veterinary and food sciences	1	2025	2025
19	31 Biological sciences	2	2025	2026
20	32 Biomedical and clinical sciences	2	2025	2026
21	33 Built environment and design	1	2025	2025
22	34 Chemical sciences	2	2025	2026
23	37 Earth sciences	2	2025	2026
24	38 Economics	1	2025	2025
25	40 Engineering	2	2025	2026
26	41 Environmental sciences	2	2025	2026
27	46 Information and computing sciences	2	2025	2026
28	49 Mathematical sciences	2	2025	2026
29	51 Physical sciences	2	2025	2026
30	Agricultural And Veterinary Sciences	1	2017	2017
31	Biological Sciences	8	2014	2021
32	Built Environment and Design	1	2017	2017
33	Chemical Sciences	8	2014	2021
34	Earth & Environmental Sciences	7	2014	2021
35	Earth Sciences (exc. Geosciences)	1	2017	2017
36	Engineering	1	2017	2017
37	Engineering & Technology	7	2014	2021
38	Environmental Sciences	1	2017	2017
39	Geosciences	1	2017	2017
40	History and Archaeology	1	2017	2017
41	Information & Computing Sciences	5	2017	2021
42	Mathematical Sciences	4	2014	2018
43	Medical & Health Sciences	5	2017	2021
44	Other / Unclassified	3	2019	2021
45	Other / Unreported	5	2014	2021
46	Pawsey Operations	1	2017	2017
47	Physical Sciences	7	2014	2021
48	Physical Sciences (exc. radio astronomy)	1	2017	2017
49	Radioastronomy	1	2017	2017
50	Technology	1	2017	2017
51	Unclassified / unallocated FoR share	1	2024	2024

	total_allocation_share
	<dbl>
1	0.0296
2	0.804
3	0.512

4	0.384
5	0.0164
6	0.107
7	0.00555
8	0.0235
9	1.02
10	0.0783
11	0.0112
12	0.00123
13	0.000185
14	0.00526
15	0.000246
16	0.000246
17	0.000215
18	0.00203
19	0.142
20	0.0197
21	0.00118
22	0.359
23	0.263
24	0.00388
25	0.747
26	0.00936
27	0.0615
28	0.0110
29	0.379
30	0.0045
31	0.390
32	0
33	1.53
34	2.20
35	0.0155
36	0.305
37	1.12
38	0.0078
39	0.158
40	0.0031
41	0.0326
42	0.257
43	0.139
44	0.438
45	0.0604
46	0.0047
47	1.08
48	0.101
49	0.146

```
50              0.0096
51              0.00123
```

8.4.3 Harmonise Research Fields for Longitudinal Comparison

For the 2014–2026 comparison, original reported labels are consolidated into nine broad analytical groups.

Physical Sciences and Chemical Sciences remain separate. Earth and Environmental Sciences are combined where source reporting differs in granularity. Engineering and Technology are similarly grouped for longitudinal visualisation.

The mappings are analytical decisions, not official one-to-one classification conversions.

```
# Create broad analytical groups from original reported labels
for_visual_data <- for_converted |>
  dplyr::mutate(
    reported_clean = research_field_reported |>
      stringr::str_remove("^\\s*\\d{2}\\s+") |>
      stringr::str_to_lower() |>
      stringr::str_replace_all("&", " and ") |>
      stringr::str_squish(),

    for_visual_group = dplyr::case_when(

      # Physical Sciences
      stringr::str_detect(
        reported_clean,
        "physical sciences|radioastronomy|radio astronomy"
      ) ~ "Physical Sciences",

      # Chemical Sciences - kept separate
      stringr::str_detect(
        reported_clean,
        "chemical sciences"
      ) ~ "Chemical Sciences",

      # Earth and Environmental Sciences
      stringr::str_detect(
        reported_clean,
        "earth|environmental|geoscien"
      ) ~ "Earth & Environmental Sciences",

      # Engineering and Technology
      stringr::str_detect(
```

```
        reported_clean,
        "engineering|technology"
      ) ~ "Engineering & Technology",

      # Biological and Agricultural Sciences
      stringr::str_detect(
        reported_clean,
        "biological|agricultural|veterinary|food sciences"
      ) ~ "Biological & Agricultural Sciences",

      # Medical and Health Sciences
      stringr::str_detect(
        reported_clean,
        "medical|health|biomedical|clinical"
      ) ~ "Medical & Health Sciences",

      # Mathematics and Computing
      stringr::str_detect(
        reported_clean,
        "mathematical|computing|information"
      ) ~ "Mathematics & Computing",

      # Social Sciences, Humanities and Built Environment
      stringr::str_detect(
        reported_clean,
        paste0(
          "psychology|economics|history|archaeology|",
          "human society|language|communication|culture|",
          "philosophy|religious|built environment|design"
        )
      ) ~ "Social Sciences, Humanities & Built Environment",

      # Unclassified and operational categories
      stringr::str_detect(
        reported_clean,
        "other|unreported|unclassified|pawsey operations"
      ) ~ "Unclassified & Operational",

      TRUE ~ NA_character_
    )
  )

# Review original labels and assigned groups
for_mapping_review <- for_visual_data |>
  dplyr::distinct(
```

```
    research_field_reported,
    for_visual_group
  ) |>
  dplyr::arrange(
    for_visual_group,
    research_field_reported
  )

print(for_mapping_review, n = Inf, width = Inf)
```

```
# A tibble: 51 x 2
   research_field_reported
   <chr>
 1 06 Biological Sciences
 2 07 Agricultural and Veterinary Sciences
 3 30 Agricultural, veterinary and food sciences
 4 31 Biological sciences
 5 Agricultural And Veterinary Sciences
 6 Biological Sciences
 7 03 Chemical Sciences
 8 34 Chemical sciences
 9 Chemical Sciences
10 04 Earth Sciences
11 05 Environmental Sciences
12 37 Earth sciences
13 41 Environmental sciences
14 Earth & Environmental Sciences
15 Earth Sciences (exc. Geosciences)
16 Environmental Sciences
17 Geosciences
18 09 Engineering
19 10 Technology
20 40 Engineering
21 Engineering
22 Engineering & Technology
23 Technology
24 01 Mathematical Sciences
25 08 Information and Computing Sciences
26 46 Information and computing sciences
27 49 Mathematical sciences
28 Information & Computing Sciences
29 Mathematical Sciences
30 11 Medical and Health Sciences
31 32 Biomedical and clinical sciences
32 Medical & Health Sciences
```

```
33 02 Physical Sciences
34 51 Physical sciences
35 Physical Sciences
36 Physical Sciences (exc. radio astronomy)
37 Radioastronomy
38 14 Economics
39 16 Studies in Human Society
40 17 Psychology and Cognitive Sciences
41 20 Language, Communication and Culture
42 21 History and Archaeology
43 22 Philosophy and Religious Studies
44 33 Built environment and design
45 38 Economics
46 Built Environment and Design
47 History and Archaeology
48 Other / Unclassified
49 Other / Unreported
50 Pawsey Operations
51 Unclassified / unallocated FoR share
   for_visual_group
   <chr>
 1 Biological & Agricultural Sciences
 2 Biological & Agricultural Sciences
 3 Biological & Agricultural Sciences
 4 Biological & Agricultural Sciences
 5 Biological & Agricultural Sciences
 6 Biological & Agricultural Sciences
 7 Chemical Sciences
 8 Chemical Sciences
 9 Chemical Sciences
10 Earth & Environmental Sciences
11 Earth & Environmental Sciences
12 Earth & Environmental Sciences
13 Earth & Environmental Sciences
14 Earth & Environmental Sciences
15 Earth & Environmental Sciences
16 Earth & Environmental Sciences
17 Earth & Environmental Sciences
18 Engineering & Technology
19 Engineering & Technology
20 Engineering & Technology
21 Engineering & Technology
22 Engineering & Technology
23 Engineering & Technology
24 Mathematics & Computing
25 Mathematics & Computing
```

```
26 Mathematics & Computing
27 Mathematics & Computing
28 Mathematics & Computing
29 Mathematics & Computing
30 Medical & Health Sciences
31 Medical & Health Sciences
32 Medical & Health Sciences
33 Physical Sciences
34 Physical Sciences
35 Physical Sciences
36 Physical Sciences
37 Physical Sciences
38 Social Sciences, Humanities & Built Environment
39 Social Sciences, Humanities & Built Environment
40 Social Sciences, Humanities & Built Environment
41 Social Sciences, Humanities & Built Environment
42 Social Sciences, Humanities & Built Environment
43 Social Sciences, Humanities & Built Environment
44 Social Sciences, Humanities & Built Environment
45 Social Sciences, Humanities & Built Environment
46 Social Sciences, Humanities & Built Environment
47 Social Sciences, Humanities & Built Environment
48 Unclassified & Operational
49 Unclassified & Operational
50 Unclassified & Operational
51 Unclassified & Operational
```

```r
# Check for unmapped categories
unmapped_for <- for_visual_data |>
  dplyr::filter(is.na(for_visual_group)) |>
  dplyr::distinct(research_field_reported)

print(unmapped_for, n = Inf)
```

```
# A tibble: 0 x 1
# i 1 variable: research_field_reported <chr>
```

```r
stopifnot(nrow(unmapped_for) == 0)

# Confirm nine broad groups
stopifnot(
  dplyr::n_distinct(for_visual_data$for_visual_group) == 9
)
```

8.4.4 Annual Research Field Allocation Composition, 2014–2026

```
# Aggregate allocation shares by year and analytical group
for_plot_data <- for_visual_data |>
  dplyr::group_by(year, for_visual_group) |>
  dplyr::summarise(
    allocation_share = sum(allocation_share_for),
    .groups = "drop"
  )

# Validate annual allocation shares
for_plot_validation <- for_plot_data |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    total_share = sum(allocation_share),
    .groups = "drop"
  )

stopifnot(
  all(abs(for_plot_validation$total_share - 1) < 1e-6)
)

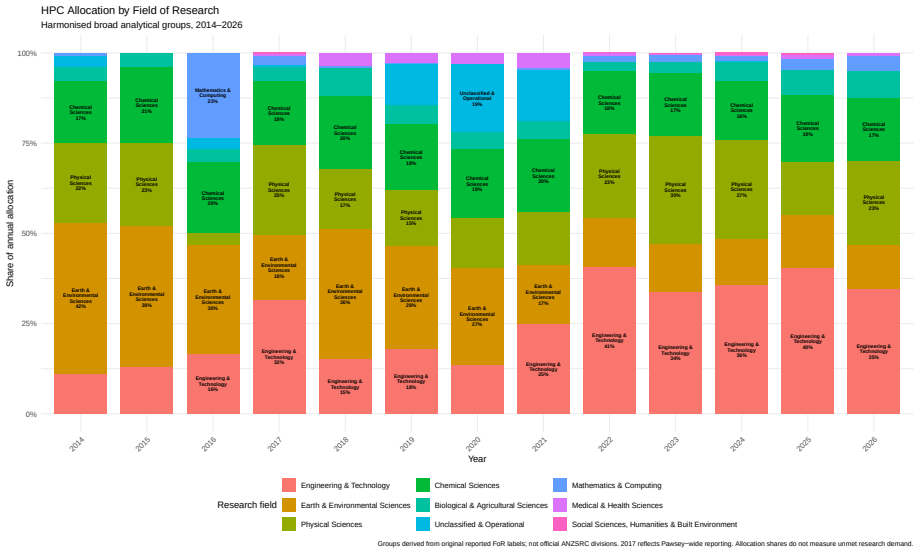
# Rank groups by average share across all 13 years
for_group_order <- for_plot_data |>
  tidyr::complete(
    year = 2014:2026,
    for_visual_group,
    fill = list(allocation_share = 0)
  ) |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    average_share = mean(allocation_share),
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share)) |>
  dplyr::pull(for_visual_group)

for_plot_data <- for_plot_data |>
  dplyr::mutate(
    for_visual_group = factor(
      for_visual_group,
      levels = for_group_order
    )
  )
```

```r
# Save harmonised FoR shares for longitudinal evaluation
saveRDS(
  for_plot_data,
  "assets/data/data converted/for_2014_2026_harmonised_shares.rds"
)

# Plot
ggplot(
  for_plot_data,
  aes(
    x = factor(year),
    y = allocation_share,
    fill = for_visual_group
  )
) +
  geom_col(
    width = 0.8,
    position = position_stack(reverse = TRUE)
  ) +
  geom_text(
    aes(
      label = ifelse(
        allocation_share >= 0.15,
        paste0(
          stringr::str_wrap(
            as.character(for_visual_group),
            width = 15
          ),
          "\n",
          scales::percent(
            allocation_share,
            accuracy = 1
          )
        ),
        ""
      )
    ),
    position = position_stack(
      vjust = 0.5,
      reverse = TRUE
    ),
    size = 2.1,
    colour = "black",
    fontface = "bold",
    lineheight = 0.9
```

```
  ) +
  scale_y_continuous(
    labels = scales::label_percent(),
    breaks = seq(0, 1, 0.25)
  ) +
  coord_cartesian(ylim = c(0, 1)) +
  labs(
    title = "HPC Allocation by Field of Research",
    subtitle = "Harmonised broad analytical groups, 2014-2026",
    x = "Year",
    y = "Share of annual allocation",
    fill = "Research field",
    caption = paste(
      "Groups derived from original reported FoR labels;",
      "not official ANZSRC divisions.",
      "2017 reflects Pawsey-wide reporting.",
      "Allocation shares do not measure unmet research demand."
    )
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(
      angle = 45,
      hjust = 1
    ),
    legend.position = "bottom",
    legend.text = element_text(size = 9),
    plot.caption = element_text(size = 8)
  ) +
  guides(
    fill = guide_legend(ncol = 3)
  )
```



For longitudinal comparison, original reported FoR categories are harmonised into nine broad analytical groups. Physical Sciences and Chemical Sciences are retained separately to avoid obscuring their individual allocation patterns.

These groupings reduce differences in naming and reporting granularity, but they do not eliminate all classification or scope inconsistencies. In particular, 2017 uses Pawsey-wide reporting, and historical Technology categories are grouped with Engineering for analytical visualisation without implying a verified official FoR 2020 alignment.

The resulting chart describes changes in the composition of allocated HPC resources, not directly observed researcher demand.

8.4.5 Field of Research Reporting Periods

The 2014–2020 and 2021–2026 periods are examined separately to assess differences in research-field allocation composition.

The earlier period uses FoR 2008 division codes, while the later period uses the existing FoR 2020 mappings. These are analytical reporting periods and do not imply that all source datasets adopted a new classification system in the same year.

Physical Sciences and Chemical Sciences are retained separately in both periods. Combined categories are used where necessary to improve visual consistency, without modifying the original classifications.

Allocation shares represent the distribution of allocated HPC resources, not directly observed researcher demand. The 2017 FoR data have a different Pawsey-wide reporting scope and should be interpreted cautiously.

8.4.5.1 Earlier Reporting Period: 2014–2020

```
for_2008_detail <- for_converted |>
  dplyr::filter(dplyr::between(year, 2014, 2020)) |>
  dplyr::mutate(
    for_standard_2008 = dplyr::case_when(
      for_2008_division_code == "01" ~ "Mathematical Sciences",
      for_2008_division_code == "02" ~ "Physical Sciences",
      for_2008_division_code == "03" ~ "Chemical Sciences",
      for_2008_division_code %in% c(
        "04", "05", "04 + 05"
      ) ~ "Earth & Environmental Sciences",
      for_2008_division_code == "06" ~ "Biological Sciences",
      for_2008_division_code == "07" ~ "Agricultural & Veterinary Sciences",
      for_2008_division_code == "08" ~ "Information & Computing Sciences",
      for_2008_division_code %in% c(
        "09", "10", "09 + 10"
      ) ~ "Engineering & Technology",
      for_2008_division_code == "11" ~ "Medical & Health Sciences",
      for_2008_division_code == "12" ~ "Built Environment & Design",
      for_2008_division_code == "21" ~ "History & Archaeology",
      is.na(for_2008_division_code) ~ "Other / Unclassified",
      TRUE ~ NA_character_
    )
  )

# Check unmatched classifications
unmapped_2008 <- for_2008_detail |>
  dplyr::filter(is.na(for_standard_2008)) |>
  dplyr::distinct(
    year,
    research_field_reported,
    for_2008_division_code
  )

print(unmapped_2008, n = Inf)
```

```
# A tibble: 0 x 3
# i 3 variables: year <dbl>, research_field_reported <chr>,
#   for_2008_division_code <chr>
```

```
stopifnot(nrow(unmapped_2008) == 0)

# Aggregate allocation shares
for_2008_plot <- for_2008_detail |>
  dplyr::group_by(year, for_standard_2008) |>
```

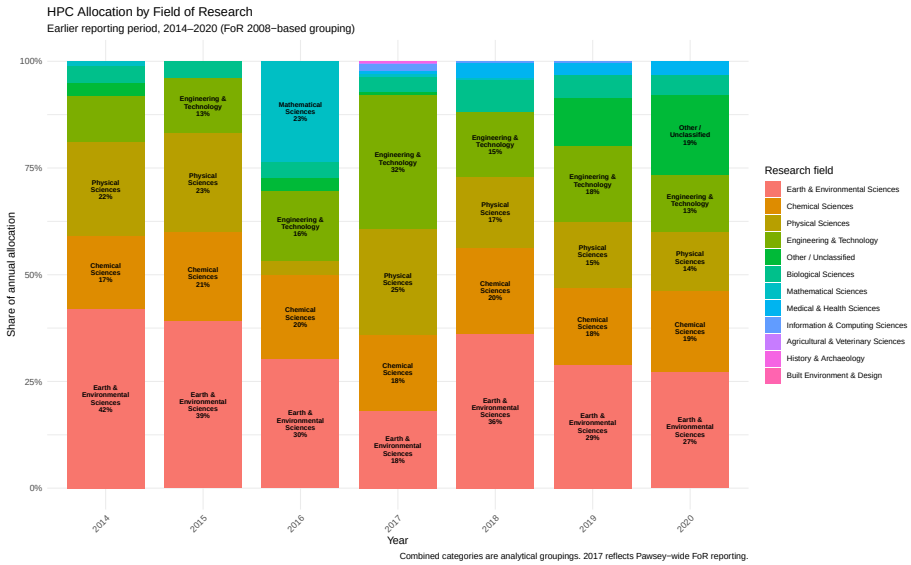
```
  dplyr::summarise(
    allocation_share = sum(allocation_share_for),
    .groups = "drop"
  )
```

```
# Order categories by average allocation share
for_2008_order <- for_2008_plot |>
  dplyr::group_by(for_standard_2008) |>
  dplyr::summarise(
    average_share = sum(allocation_share) / 7,
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share)) |>
  dplyr::pull(for_standard_2008)

for_2008_plot <- for_2008_plot |>
  dplyr::mutate(
    for_standard_2008 = factor(
      for_standard_2008,
      levels = for_2008_order
    )
  )

ggplot(
  for_2008_plot,
  aes(
    x = factor(year),
    y = allocation_share,
    fill = for_standard_2008
  )
) +
  geom_col(
    width = 0.8,
    position = position_stack(reverse = TRUE)
  ) +
  geom_text(
    aes(
      label = ifelse(
        allocation_share >= 0.12,
        paste0(
          stringr::str_wrap(
            as.character(for_standard_2008),
            width = 15
          ),
          "\n",
```

```
          scales::percent(
            allocation_share,
            accuracy = 1
          )
        ),
        ""
      )
    ),
    position = position_stack(
      vjust = 0.5,
      reverse = TRUE
    ),
    size = 2.4,
    colour = "black",
    fontface = "bold",
    lineheight = 0.9
  ) +
  scale_y_continuous(
    labels = scales::label_percent(),
    breaks = seq(0, 1, 0.25)
  ) +
  coord_cartesian(ylim = c(0, 1)) +
  labs(
    title = "HPC Allocation by Field of Research",
    subtitle = "Earlier reporting period, 2014-2020 (FoR 2008-based grouping)",
    x = "Year",
    y = "Share of annual allocation",
    fill = "Research field",
    caption = paste(
      "Combined categories are analytical groupings.",
      "2017 reflects Pawsey-wide FoR reporting."
    )
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1),
    legend.position = "right",
    legend.text = element_text(size = 8)
  )
```



8.4.5.2 Later Reporting Period: 2021–2026

```
for_2020_detail <- for_converted |>
  dplyr::filter(dplyr::between(year, 2021, 2026)) |>
  dplyr::mutate(
    for_2020_visual = dplyr::case_when(

      for_standard_2020 %in% c(
        "Engineering",
        "Engineering + Technology",
        "Technology - no direct FoR 2020 alignment"
      ) ~ "Engineering & Technology",

      for_standard_2020 == "Physical sciences" ~
        "Physical Sciences",

      for_standard_2020 == "Chemical sciences" ~
        "Chemical Sciences",

      for_standard_2020 %in% c(
        "Earth sciences",
        "Environmental sciences",
        "Earth sciences + Environmental sciences"
      ) ~ "Earth & Environmental Sciences",

      for_standard_2020 %in% c(
```

```
        "Biological sciences",
        "Agricultural, veterinary and food sciences"
      ) ~ "Biological & Agricultural Sciences",

      for_standard_2020 %in% c(
        "Biomedical and clinical sciences",
        "Health sciences",
        "Biomedical and clinical sciences + Health sciences"
      ) ~ "Medical & Health Sciences",

      for_standard_2020 %in% c(
        "Mathematical sciences",
        "Information and computing sciences"
      ) ~ "Mathematics & Computing",

      for_standard_2020 %in% c(
        "Psychology",
        "Economics",
        "History, heritage and archaeology",
        "Language, communication and culture",
        "Philosophy and religious studies",
        "Human society",
        "Built environment and design"
      ) ~ "Social Sciences, Humanities & Built Environment",

      for_standard_2020 == "Other / Unclassified" ~
        "Other / Unclassified",

      TRUE ~ NA_character_
    )
  )

# Check unmatched classifications
unmapped_2020 <- for_2020_detail |>
  dplyr::filter(is.na(for_2020_visual)) |>
  dplyr::distinct(
    year,
    research_field_reported,
    for_standard_2020
  )

print(unmapped_2020, n = Inf)
```

```
# A tibble: 0 x 3
# i 3 variables: year <dbl>, research_field_reported <chr>,
```

```
#   for_standard_2020 <chr>
```

```
stopifnot(nrow(unmapped_2020) == 0)

# Aggregate allocation shares
for_2020_plot <- for_2020_detail |>
  dplyr::group_by(year, for_2020_visual) |>
  dplyr::summarise(
    allocation_share = sum(allocation_share_for),
    .groups = "drop"
  )

# Save harmonised FoR 2020 shares for recent-period evaluation
saveRDS(
  for_2020_plot,
  "assets/data/data converted/for_2020_harmonised_shares.rds"
)
```

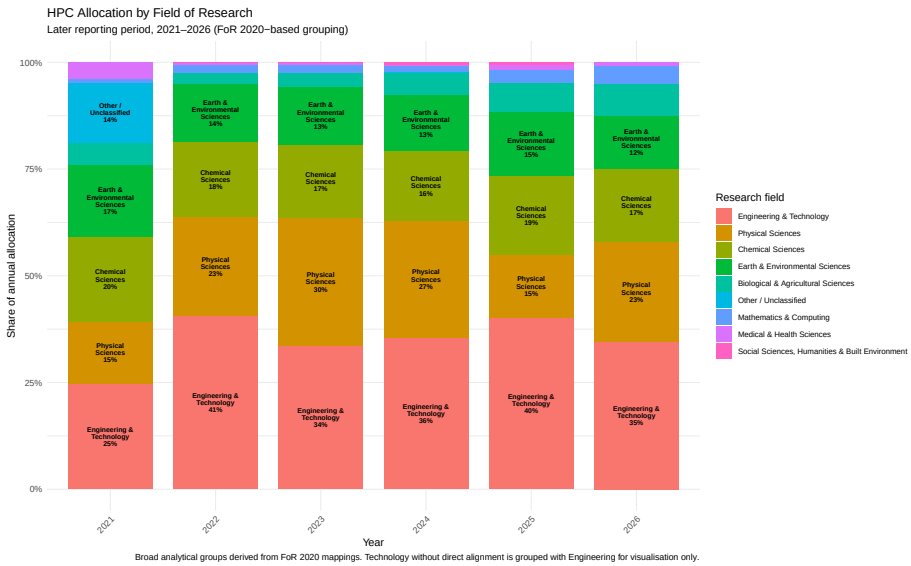
```
# Order categories by average allocation share
for_2020_order <- for_2020_plot |>
  dplyr::group_by(for_2020_visual) |>
  dplyr::summarise(
    average_share = sum(allocation_share) / 6,
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share)) |>
  dplyr::pull(for_2020_visual)

for_2020_plot <- for_2020_plot |>
  dplyr::mutate(
    for_2020_visual = factor(
      for_2020_visual,
      levels = for_2020_order
    )
  )

ggplot(
  for_2020_plot,
  aes(
    x = factor(year),
    y = allocation_share,
    fill = for_2020_visual
  )
) +
  geom_col(
    width = 0.8,
```

```
  position = position_stack(reverse = TRUE)
) +
geom_text(
  aes(
    label = ifelse(
      allocation_share >= 0.12,
      paste0(
        stringr::str_wrap(
          as.character(for_2020_visual),
          width = 15
        ),
        "\n",
        scales::percent(
          allocation_share,
          accuracy = 1
        )
      ),
      ""
    )
  ),
  position = position_stack(
    vjust = 0.5,
    reverse = TRUE
  ),
  size = 2.4,
  colour = "black",
  fontface = "bold",
  lineheight = 0.9
) +
scale_y_continuous(
  labels = scales::label_percent(),
  breaks = seq(0, 1, 0.25)
) +
coord_cartesian(ylim = c(0, 1)) +
labs(
  title = "HPC Allocation by Field of Research",
  subtitle = "Later reporting period, 2021-2026 (FoR 2020-based grouping)",
  x = "Year",
  y = "Share of annual allocation",
  fill = "Research field",
  caption = paste(
    "Broad analytical groups derived from FoR 2020 mappings.",
    "Technology without direct alignment is grouped with Engineering",
    "for visualisation only."
  )
```

```
  ) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1),
    legend.position = "right",
    legend.text = element_text(size = 8)
  )
```



8.4.5.3 Validation of Reporting-Period Allocations

```
for_period_validation <- dplyr::bind_rows(
  for_2008_plot |>
    dplyr::transmute(
      year,
      allocation_share,
      period = "2014-2020"
    ),
  for_2020_plot |>
    dplyr::transmute(
      year,
      allocation_share,
      period = "2021-2026"
    )
) |>
  dplyr::group_by(period, year) |>
  dplyr::summarise(
```

```
    total_share = sum(allocation_share),
    .groups = "drop"
  )

print(for_period_validation, n = Inf)
```

```
# A tibble: 13 x 3
   period     year total_share
   <chr>     <dbl>       <dbl>
 1 2014-2020  2014       1
 2 2014-2020  2015       1
 3 2014-2020  2016       1
 4 2014-2020  2017       1
 5 2014-2020  2018       1
 6 2014-2020  2019       1
 7 2014-2020  2020       1
 8 2021-2026  2021       1
 9 2021-2026  2022       1.00
10 2021-2026  2023       1.000
11 2021-2026  2024       1
12 2021-2026  2025       1.000
13 2021-2026  2026       1.000
```

```
stopifnot(
  nrow(for_period_validation) == 13,
  all(abs(for_period_validation$total_share - 1) < 1e-6)
)
```

The two reporting-period charts provide complementary views of historical FoR allocation composition. The earlier chart uses FoR 2008 division-based groupings, while the later chart uses existing FoR 2020 mappings.

These analytical categories are not necessarily equivalent to individual official ANZSRC divisions. Physical Sciences and Chemical Sciences are kept separate in both periods to improve interpretability.

The 2017 Pawsey-wide observation remains subject to a reporting-scope limitation. Changes in allocation composition should therefore not be interpreted as changes in researcher demand without further evidence.

8.4.6 Field of Research: Initial Findings

This section summarises the historical distribution of HPC allocations across research disciplines. The analysis focuses on allocation shares, rather than treating these shares as direct measures of researcher demand.

8.4.6.1 Dominant Research Fields, 2014–2026

```
# Calculate average annual allocation shares across all 13 years
for_overall_findings <- for_plot_data |>
  dplyr::mutate(
    for_visual_group = as.character(for_visual_group)
  ) |>
  tidyr::complete(
    year = 2014:2026,
    for_visual_group,
    fill = list(allocation_share = 0)
  ) |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    average_share = mean(allocation_share),
    min_share = min(allocation_share),
    max_share = max(allocation_share),
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share))

for_overall_findings |>
  dplyr::mutate(
    dplyr::across(
      c(average_share, min_share, max_share),
      ~ scales::percent(.x, accuracy = 0.1)
    )
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 9 x 4
  for_visual_group                                average_share min_share
  <chr>                                           <chr>         <chr>
1 Engineering & Technology                        25.2%         11.0%
2 Earth & Environmental Sciences                  23.5%         12.3%
3 Physical Sciences                               19.3%         3.2%
4 Chemical Sciences                               18.5%         16.3%
5 Biological & Agricultural Sciences              5.0%          2.8%
6 Unclassified & Operational                      3.9%          0.0%
7 Mathematics & Computing                         3.2%          0.0%
8 Medical & Health Sciences                       1.3%          0.0%
9 Social Sciences, Humanities & Built Environment 0.1%          0.0%
  max_share
  <chr>
1 40.7%
2 42.0%
```

```
3 29.8%
4 21.0%
5 7.8%
6 18.6%
7 23.5%
8 4.0%
9 0.6%
```

The table ranks broad research-field groups by their average annual allocation share across 2014–2026. These are unweighted averages of annual percentages, not shares of cumulative allocated service units. Consequently, they describe the typical annual composition rather than each discipline's contribution to total historical allocation volume.

8.4.6.2 Changes in Research Field Composition

```
# Compare the first and final years
for_composition_change <- for_plot_data |>
  dplyr::mutate(
    for_visual_group = as.character(for_visual_group)
  ) |>
  dplyr::filter(year %in% c(2014, 2026)) |>
  tidyr::complete(
    year = c(2014, 2026),
    for_visual_group,
    fill = list(allocation_share = 0)
  ) |>
  tidyr::pivot_wider(
    names_from = year,
    values_from = allocation_share,
    names_prefix = "share_"
  ) |>
  dplyr::mutate(
    change_percentage_points =
      (share_2026 - share_2014) * 100
  ) |>
  dplyr::arrange(
    dplyr::desc(change_percentage_points)
  )

for_composition_change |>
  dplyr::mutate(
    share_2014 = scales::percent(share_2014, accuracy = 0.1),
    share_2026 = scales::percent(share_2026, accuracy = 0.1),
    change_percentage_points = round(
      change_percentage_points, 1
```

```
    )
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 8 x 4
  for_visual_group                    share_2014 share_2026
  <chr>                               <chr>      <chr>
1 Engineering & Technology            11.0%      34.5%
2 Biological & Agricultural Sciences  4.0%       7.5%
3 Mathematics & Computing             1.0%       4.3%
4 Physical Sciences                   22.0%      23.2%
5 Medical & Health Sciences           0.0%       0.7%
6 Chemical Sciences                   17.0%      17.4%
7 Unclassified & Operational          3.0%       0.0%
8 Earth & Environmental Sciences      42.0%      12.3%
  change_percentage_points
                     <dbl>
1                     23.5
2                      3.5
3                      3.3
4                      1.2
5                      0.7
6                      0.4
7                     -3
8                    -29.7
```

Changes are expressed in percentage points. A positive change indicates that a research-field group accounted for a larger proportion of annual HPC allocations in 2026 than in 2014. It does not necessarily indicate an increase in absolute allocated resources or unmet researcher demand.

8.4.6.3 Comparison of Reporting Periods

```
# Compare dominant categories within each reporting period
for_period_findings <- dplyr::bind_rows(

  for_2008_plot |>
    dplyr::transmute(
      period = "2014-2020",
      year,
      research_field = as.character(for_standard_2008),
      allocation_share
    ),

  for_2020_plot |>
    dplyr::transmute(
```

```
      period = "2021-2026",
      year,
      research_field = as.character(for_2020_visual),
      allocation_share
    )

) |>
  dplyr::group_by(period, research_field) |>
  dplyr::summarise(
    total_share = sum(allocation_share),
    .groups = "drop"
  ) |>
  dplyr::mutate(
    average_annual_share = dplyr::if_else(
      period == "2014-2020",
      total_share / 7,
      total_share / 6
    )
  ) |>
  dplyr::group_by(period) |>
  dplyr::slice_max(
    order_by = average_annual_share,
    n = 5,
    with_ties = FALSE
  ) |>
  dplyr::arrange(
    period,
    dplyr::desc(average_annual_share),
    .by_group = TRUE
  ) |>
  dplyr::ungroup()

for_period_findings |>
  dplyr::select(
    period,
    research_field,
    average_annual_share
  ) |>
  dplyr::mutate(
    average_annual_share = scales::percent(
      average_annual_share,
      accuracy = 0.1
    )
  ) |>
  print(n = Inf, width = Inf)
```

```
# A tibble: 10 x 3
   period    research_field                     average_annual_share
   <chr>     <chr>                              <chr>
 1 2014-2020 Earth & Environmental Sciences     31.6%
 2 2014-2020 Chemical Sciences                  19.0%
 3 2014-2020 Physical Sciences                  16.9%
 4 2014-2020 Engineering & Technology           16.9%
 5 2014-2020 Other / Unclassified               5.2%
 6 2021-2026 Engineering & Technology           34.9%
 7 2021-2026 Physical Sciences                  22.1%
 8 2021-2026 Chemical Sciences                  17.8%
 9 2021-2026 Earth & Environmental Sciences     14.0%
10 2021-2026 Biological & Agricultural Sciences 5.2%
```

The reporting-period comparison identifies the five largest analytical research-field categories within each period. Because the earlier period uses FoR 2008-based groupings and the later period uses FoR 2020-based groupings, their rankings are descriptive rather than strictly equivalent classification comparisons.

The 2017 Pawsey-wide observation also introduces a reporting-scope difference within the earlier period.

8.4.6.4 Implications for HPC Demand Forecasting

The historical FoR analysis provides evidence about the disciplinary composition of allocated HPC resources. It may help identify research areas relevant to future capacity planning and provide context for forecast scenarios.

However, three limitations must be considered:

1. Allocation shares represent the distribution of available resources, not directly observed unmet researcher demand.
2. Changes in research-field classification and reporting scope limit direct comparability across all years.
3. FoR-specific CPU and GPU resource measures are proportionally allocated estimates, rather than directly observed resource consumption by discipline.

Accordingly, the FoR results should primarily inform the interpretation of aggregate demand forecasts and the development of disciplinary scenarios. They should not be used as independent discipline-level demand forecasts without additional evidence.

Chapter 9

Forecasting Models

9.1 Purpose

This chapter aims to develop transparent, reproducible, and statistically defensible forecasting models to estimate Australian national research supercomputing demand through 2030.

Building on the harmonised historical dataset and initial findings presented in Chapters 03 - 06, the analysis evaluates alternative approaches for modelling annual researcher demand, expressed in CPU-equivalent core-hours.

The modelling framework has four objectives:

1. Establish a suitable historical demand series for forecasting, accounting for data quality, reporting differences, and conversion assumptions.
2. Develop and compare candidate forecasting models, prioritising simple and interpretable approaches such as linear and log-linear trends.
3. Examine the sensitivity of projected demand to the substantial increase observed in 2026, distinguishing potential structural changes from temporary fluctuations.
4. Generate alternative demand projections for 2027 - 2030 to support national HPC infrastructure planning and policy discussions.

Following expert statistical feedback, particular emphasis is placed on model simplicity, transparent assumptions, and the limitations of extrapolating trends from a relatively short and heterogeneous historical dataset.

The forecasts are intended to provide evidence-informed planning scenarios rather than precise predictions of future researcher demand.

9.2 Forecast Target and Data Suitability

9.2.1 Forecast Target

The forecasting analysis considers three related but distinct measures of Australian national research supercomputing demand:

1. Overall HPC demand, expressed in thousands of service units (kSU).
2. CPU demand, expressed in CPU core-hours.
3. GPU demand, expressed in GPU-hours.

These measures represent harmonised estimates of computing resources requested by researchers, rather than allocated resources or available computing capacity. CPU and GPU demand are modelled separately because their resource units and conversion assumptions are not directly interchangeable.

The overall kSU series provides a reference for aggregate demand, while the CPU and GPU series support resource-specific forecasts relevant to future infrastructure planning.

9.2.2 Historical Data Suitability

The historical demand series provides a consistent analytical unit for forecasting, but it should not be interpreted as a fully homogeneous set of directly observed demand measurements.

Depending on the year and reporting source, historical estimates include directly reported demand, derived estimates, and proxy-based measures. Differences in facility coverage, reporting definitions, and CPU/GPU conversion assumptions introduce uncertainty into comparisons across years.

The relatively small number of annual observations also limits the complexity of forecasting models that can be estimated reliably. Consequently, simple trend-based models will be prioritised, while more flexible time-series approaches will be used as benchmarks.

The substantial increase in reported demand in 2026 presents an additional uncertainty. Available historical observations do not establish whether this increase represents a sustained shift in researcher demand, a temporary fluctuation, or a difference in measurement or reporting scope.

Accordingly, the modelling analysis will consider both the complete historical series (2014–2026) and a sensitivity specification excluding 2026 (2014–2025). The latter is not treated as a more accurate historical dataset, but as a means of assessing how strongly the latest observation influences projected demand.

9.2.3 Forecasting Horizon

The forecasting horizon covers four annual periods:

- 2027
- 2028
- 2029
- 2030

Forecasts will be reported in the original units of each target: kSU for overall HPC demand, CPU core-hours for CPU demand, and GPU-hours for GPU demand.

The GPU analysis will explicitly distinguish historical scenario-based estimates from later reported or derived values, rather than treating all historical GPU observations as equally reliable.

The resulting forecasts are conditional projections based on historical patterns and modelling assumptions. They should not be interpreted as precise estimates of future HPC requirements, particularly given uncertainty surrounding emerging GPU-intensive workloads, changes in research activity, and future infrastructure investment.

9.3 Data Preparation and Validation

The forecasting analysis uses the converted annual demand and allocation dataset produced in Chapter 05. The dataset is loaded directly from its saved RDS file to ensure that this chapter can be executed independently of the preceding analysis.

The preparation process retains the historical demand observations without modifying their original values. Validation checks confirm the expected annual coverage, uniqueness of observations, and availability of positive demand values.

```
# Load the converted annual dataset
annual_converted <- readRDS(
  "assets/data/data converted/annual_demand_allocation_converted.rds"
)

# Retain the three forecasting targets
forecast_data <- annual_converted |>
  dplyr::transmute(
    year = as.integer(year),
    demand_ksu = as.numeric(demand_ksu),
    demand_cpu_core_hours = as.numeric(demand_cpu_core_hours),
    demand_gpu_hours = as.numeric(demand_gpu_hours)
  ) |>
  dplyr::arrange(year)

# Validate annual coverage and values
stopifnot(
```

```
  identical(forecast_data$year, 2014:2026),
  !anyDuplicated(forecast_data$year),
  all(is.finite(as.matrix(forecast_data[, -1]))),
  all(as.matrix(forecast_data[, -1]) >= 0),
  all(forecast_data$demand_ksu > 0),
  all(forecast_data$demand_cpu_core_hours > 0)
)

# Confirm historical coverage
cat(
  "Historical period:", min(forecast_data$year),
  "-", max(forecast_data$year), "\n",
  "Observations:", nrow(forecast_data), "\n"
)
```

```
Historical period: 2014 - 2026
 Observations: 13
```

9.4 Candidate Forecasting Models

9.4.1 Overall HPC Demand: Linear and Log-Linear Trends

Two initial trend models are fitted to the harmonised annual HPC demand series for 2014–2026.

The linear trend models demand directly in kSU, assuming a constant absolute change over time. The log-linear trend models the natural logarithm of demand, allowing for approximately constant proportional growth.

Both specifications are retained for subsequent evaluation. Neither is assumed to be the preferred model before out-of-sample validation.

```
# Use a centred time index for numerical stability
overall_model_data <- forecast_data |>
  dplyr::mutate(time = year - 2014)

# Linear trend: demand in kSU
overall_linear <- lm(
  demand_ksu ~ time,
  data = overall_model_data
)

# Log-linear trend: log(demand in kSU)
overall_loglinear <- lm(
  log(demand_ksu) ~ time,
  data = overall_model_data
)
```

```r
# Inspect fitted coefficients
summary(overall_linear)
```

```
Call:
lm(formula = demand_ksu ~ time, data = overall_model_data)

Residuals:
    Min      1Q  Median      3Q     Max
-274921 -105309  -35637  129560  387964

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   -89052      98432  -0.905    0.385
time          161992      13920  11.637 1.59e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 187800 on 11 degrees of freedom
Multiple R-squared:  0.9249,    Adjusted R-squared:  0.918
F-statistic: 135.4 on 1 and 11 DF,  p-value: 1.594e-07
```

```r
summary(overall_loglinear)
```

```
Call:
lm(formula = log(demand_ksu) ~ time, data = overall_model_data)

Residuals:
     Min       1Q   Median       3Q      Max
-0.44334 -0.12749 -0.03195  0.14904  0.40933

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 11.94897    0.13106   91.17  < 2e-16 ***
time         0.23348    0.01854   12.60 7.05e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2501 on 11 degrees of freedom
Multiple R-squared:  0.9352,    Adjusted R-squared:  0.9293
F-statistic: 158.7 on 1 and 11 DF,  p-value: 7.054e-08
```

The linear model estimates an annual absolute change in HPC demand, whereas the log-linear model estimates an approximately constant annual percentage growth rate.

The log-linear model's fitted values must be back-transformed to kSU for comparison with observed demand. Simple exponentiation produces median predictions under a log-normal error assumption; mean predictions require appropriate bias adjustment.

Model performance and forecast accuracy will be evaluated in Chapter 08 using consistent original-scale error measures.

9.4.2 Initial Model Findings

The linear trend model estimates an average annual increase of approximately 161,992 kSU over 2014–2026. In contrast, the log-linear model estimates an annual compounded growth rate of approximately 26.3%.

Both models indicate an increasing historical demand trend. However, their assumptions differ substantially: the linear specification implies constant absolute growth, while the log-linear specification implies exponential growth.

Although both models exhibit relatively high in-sample R-squared values, these statistics are calculated on different response scales and cannot be used directly to select the preferred model.

Given the limited historical sample, heterogeneous demand estimates, and uncertainty surrounding the 2026 observation, neither model is selected at this stage. Their predictive performance will be assessed using consistent original-scale forecast errors in Chapter 08.

9.5 Sensitivity Analysis: Influence of 2026

The substantial increase in reported HPC demand in 2026 introduces uncertainty into the estimation of future demand trends.

To assess the influence of this observation, the linear and log-linear trend models are refitted using historical observations from 2014 - 2025. These results are compared with the corresponding models fitted to the complete 2014–2026 dataset.

This analysis does not assume that the 2026 observation is incorrect or should be excluded from the final forecasting model. Instead, it evaluates how sensitive the estimated trend is to the inclusion of the latest observation.

9.5.1 Trend Sensitivity Comparison

```
# Historical data excluding 2026
overall_pre2026_data <- overall_model_data |>
  dplyr::filter(year <= 2025)

# Refit the same model specifications
```

```r
overall_linear_pre2026 <- lm(
  demand_ksu ~ time,
  data = overall_pre2026_data
)

overall_loglinear_pre2026 <- lm(
  log(demand_ksu) ~ time,
  data = overall_pre2026_data
)

# Compare trend coefficients
trend_sensitivity <- data.frame(
  model = c("Linear trend", "Log-linear trend"),
  full_2014_2026 = c(
    coef(overall_linear)[["time"]],
    coef(overall_loglinear)[["time"]]
  ),
  excluding_2026 = c(
    coef(overall_linear_pre2026)[["time"]],
    coef(overall_loglinear_pre2026)[["time"]]
  )
)

trend_sensitivity$percentage_change <- with(
  trend_sensitivity,
  100 * (full_2014_2026 - excluding_2026) /
    abs(excluding_2026)
)

knitr::kable(
  trend_sensitivity,
  digits = 3,
  col.names = c(
    "Model",
    "Trend with 2026",
    "Trend without 2026",
    "Change (%)"
  ),
  caption = "Sensitivity of estimated HPC demand trends to the inclusion of 2026"
)
```

Table 9.1: Sensitivity of estimated HPC demand trends to the inclusion of 2026

Model	Trend with 2026	Trend without 2026	Change (%)
Linear trend	161992.127	144357.416	12.216
Log-linear trend	0.233	0.239	-2.422

The comparison evaluates changes in the estimated annual trend coefficients. For the linear model, the coefficient represents the annual absolute change in demand measured in kSU. For the log-linear model, the coefficient represents the annual change in log demand, which can be converted into an implied compounded growth rate.

Changes in these coefficients indicate how strongly the latest observation influences the fitted historical trend. They do not, on their own, establish whether the 2026 increase represents a structural change in underlying researcher demand.

9.5.2 Sensitivity Analysis Findings

Including the 2026 observation increases the estimated linear trend from approximately 144,357 to 161,992 kSU per year, representing a 12.2% increase in the annual absolute growth coefficient.

In contrast, the log-linear trend coefficient decreases slightly when 2026 is included, from approximately 0.239 to 0.233, equivalent to a reduction of approximately 2.4% in the estimated log-scale trend coefficient.

These results indicate that the latest observation has a greater relative influence on the estimated absolute growth rate than on the proportional growth trend.

Nevertheless, the relatively small change in the log-linear coefficient does not necessarily imply that its future forecasts are insensitive to 2026, since changes in the fitted intercept may also influence projected demand levels.

The analysis therefore supports retaining both model specifications for further evaluation. It does not establish whether the 2026 increase represents a temporary fluctuation or a sustained change in researcher demand.

9.6 CPU Demand Forecasting Models

CPU demand is modelled separately from overall HPC demand to examine the historical growth of CPU-related computing requirements.

The analysis uses the harmonised annual CPU demand estimates from 2014–2026, expressed in millions of CPU core-hours for readability. These values include directly reported and derived estimates, and therefore remain subject to the conversion assumptions documented in Chapter 05.

As with overall HPC demand, two simple trend specifications are initially considered: linear and log-linear regression.

9.6.1 CPU Linear and Log-Linear Trends

```
# Prepare CPU demand in millions of core-hours
cpu_model_data <- forecast_data |>
  dplyr::transmute(
    year,
    time = year - 2014,
    cpu_million_hours = demand_cpu_core_hours / 1e6
  )

# Linear trend
cpu_linear <- lm(
  cpu_million_hours ~ time,
  data = cpu_model_data
)

# Log-linear trend
cpu_loglinear <- lm(
  log(cpu_million_hours) ~ time,
  data = cpu_model_data
)

# Inspect fitted coefficients
summary(cpu_linear)
```

```
Call:
lm(formula = cpu_million_hours ~ time, data = cpu_model_data)

Residuals:
    Min      1Q  Median      3Q     Max
-264.98  -27.39  -14.20   95.63  189.16

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    73.81      73.88   0.999    0.339
time           81.01      10.45   7.753 8.79e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 141 on 11 degrees of freedom
Multiple R-squared:  0.8453,    Adjusted R-squared:  0.8312
```

```
F-statistic: 60.11 on 1 and 11 DF,  p-value: 8.792e-06
```

```
summary(cpu_loglinear)
```

```

Call:
lm(formula = log(cpu_million_hours) ~ time, data = cpu_model_data)

Residuals:
     Min       1Q   Median       3Q      Max
-0.42810 -0.19869 -0.03323  0.03456  0.52239

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.09400    0.15779  32.284 3.00e-12 ***
time         0.17010    0.02231   7.623 1.03e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.301 on 11 degrees of freedom
Multiple R-squared:  0.8408,    Adjusted R-squared:  0.8264
F-statistic: 58.11 on 1 and 11 DF,  p-value: 1.03e-05
```

The linear model estimates the annual absolute change in CPU demand, measured in millions of core-hours. The log-linear model estimates the proportional growth trend.

Both specifications are retained for subsequent evaluation. Their in-sample R-squared values are not directly comparable because the models use different response scales.

The analysis does not assume that all historical CPU demand estimates are equally reliable. Differences in reporting methods, resource allocation shares, and conversion assumptions remain relevant when interpreting the fitted trends.

9.6.2 CPU Trend Findings

The linear trend model estimates an average annual increase in CPU demand of approximately 81.01 million core-hours over 2014 - 2026.

The log-linear model estimates a compounded annual growth rate of approximately 18.5%, indicating substantial proportional growth in the harmonised historical CPU demand series.

Both models identify an increasing trend, although the estimated growth patterns differ. The linear model assumes constant absolute growth, whereas the log-linear model assumes approximately constant proportional growth.

The models' in-sample R-squared values cannot be directly compared because

their dependent variables are expressed on different scales. Forecast accuracy will therefore be evaluated using consistent original-scale error measures in Chapter 08.

The findings should also be interpreted cautiously because historical CPU demand includes derived estimates based on allocation shares and conversion assumptions. Consequently, the fitted trends may reflect changes in measurement and reporting practices as well as changes in underlying researcher demand.

Neither specification is selected as the preferred CPU forecasting model at this stage.

9.6.3 CPU Demand Sensitivity Analysis

The influence of the 2026 observation is assessed by comparing CPU demand trends fitted over 2014–2026 with equivalent models fitted over 2014 - 2025.

This sensitivity analysis examines whether the latest observation materially changes the estimated absolute and proportional growth rates. Excluding 2026 is a modelling sensitivity specification, not a correction to the historical dataset.

```
# Prepare historical CPU demand excluding 2026
cpu_pre2026_data <- cpu_model_data |>
  dplyr::filter(year <= 2025)

# Refit linear and log-linear models
cpu_linear_pre2026 <- lm(
  cpu_million_hours ~ time,
  data = cpu_pre2026_data
)

cpu_loglinear_pre2026 <- lm(
  log(cpu_million_hours) ~ time,
  data = cpu_pre2026_data
)

# Compare estimated trend coefficients
cpu_trend_sensitivity <- data.frame(
  model = c("Linear trend", "Log-linear trend"),
  full_2014_2026 = c(
    coef(cpu_linear)[["time"]],
    coef(cpu_loglinear)[["time"]]
  ),
  excluding_2026 = c(
    coef(cpu_linear_pre2026)[["time"]],
    coef(cpu_loglinear_pre2026)[["time"]]
  )
)
```

```r
# Calculate percentage change in trend coefficients
cpu_trend_sensitivity$percentage_change <- with(
  cpu_trend_sensitivity,
  100 * (full_2014_2026 - excluding_2026) /
    abs(excluding_2026)
)

knitr::kable(
  cpu_trend_sensitivity,
  digits = 3,
  col.names = c(
    "Model",
    "Trend with 2026",
    "Trend without 2026",
    "Change (%)"
  ),
  caption = "Sensitivity of CPU demand trends to the inclusion of 2026"
)
```

Table 9.2: Sensitivity of CPU demand trends to the inclusion of 2026

Model	Trend with 2026	Trend without 2026	Change (%)
Linear trend	81.005	72.407	11.875
Log-linear trend	0.170	0.171	-0.436

The linear trend coefficient represents the estimated annual absolute change in CPU demand, measured in millions of core-hours.

The log-linear coefficient represents annual change in log demand and can be converted into an implied compounded growth rate.

Differences between the fitted coefficients indicate the sensitivity of estimated historical CPU demand trends to the latest observation. They do not establish whether changes in 2026 are temporary or structural, nor do they determine which model has superior

9.6.4 CPU Sensitivity Findings

Including the 2026 observation increases the estimated linear CPU demand trend from approximately 72.41 to 81.01 million core-hours per year, representing an increase of 11.9% in the annual absolute growth coefficient.

In contrast, the log-linear trend coefficient changes only slightly, decreasing by approximately 0.4% when 2026 is included.

These results suggest that the 2026 observation has a greater relative influence on the estimated absolute CPU demand growth rate than on the proportional growth trend.

However, the limited change in the log-linear coefficient does not necessarily imply that projected CPU demand is insensitive to 2026, as changes in the fitted intercept may also affect future forecasts.

The sensitivity analysis does not establish whether the 2026 increase represents a sustained change in underlying researcher demand or reflects temporary conditions, reporting differences, or conversion assumptions.

Both linear and log-linear CPU models are therefore retained for subsequent forecast evaluation in Chapter 08.

9.7 GPU Demand Forecasting Models

GPU demand is analysed separately from overall HPC and CPU demand because GPU-hours represent a distinct resource measure with different conversion assumptions and reporting characteristics.

The harmonised GPU demand series covers 2014 - 2026. However, the historical observations are not equally suitable for forecasting.

For 2014 - 2021, GPU demand values were constructed using hypothetical conversion scenarios rather than directly reported GPU demand. These estimates are useful for exploring historical resource requirements but should not be treated as independently observed GPU demand.

For 2022 - 2026, the dataset contains a combination of reported and derived GPU demand estimates. Consequently, the later observations provide a more relevant basis for preliminary GPU forecasting, although they remain subject to reporting and conversion uncertainty.

Given these limitations, the initial GPU modelling analysis will focus on 2022–2026. Any fitted trends will be interpreted as exploratory scenarios rather than reliable estimates of long-term GPU demand growth.

9.7.1 GPU Data Provenance and Modelling Period

```
# Retain GPU demand and its methodological metadata
gpu_model_data <- annual_converted |>
  dplyr::transmute(
    year = as.integer(year),
    gpu_million_hours = as.numeric(demand_gpu_hours) / 1e6,
    gpu_demand_method = gpu_demand_method,
    gpu_demand_confidence = gpu_demand_confidence
  ) |>
```

```
  dplyr::arrange(year)

# Define the later historical period for exploratory modelling
gpu_recent_data <- gpu_model_data |>
  dplyr::filter(year >= 2022, year <= 2026) |>
  dplyr::mutate(time = year - 2022)

# Validate recent GPU demand values
stopifnot(
  identical(gpu_recent_data$year, 2022:2026),
  all(is.finite(gpu_recent_data$gpu_million_hours)),
  all(gpu_recent_data$gpu_million_hours > 0)
)

# Review the values and provenance used for modelling
print(gpu_recent_data, n = Inf)
```

```
# A tibble: 5 x 5
   year gpu_million_hours gpu_demand_method          gpu_demand_confidence  time
  <int>             <dbl> <chr>                      <chr>                 <dbl>
1  2022              2.66 Estimated using GPU alloc~ Moderate-
low: alloca~     0
2  2023              1.35 Reported Setonix GPU dema~ Higher: reported GPU~     1
3  2024              2.30 Estimated using GPU alloc~ Moderate-
low: alloca~     2
4  2025              2.39 Estimated using GPU alloc~ Moderate-
low: alloca~     3
5  2026              3.00 Reported Setonix GPU dema~ Higher: reported GPU~     4
```

The 2022–2026 modelling period contains only five annual observations. This substantially limits the reliability of estimated growth trends and makes complex time-series models inappropriate.

The provenance information is retained to distinguish reported GPU demand from estimates derived using allocation shares or conversion assumptions.

The resulting models will be used to explore plausible growth trajectories and assess sensitivity to assumptions, rather than to claim a statistically established long-term GPU demand trend.

9.7.2 GPU Linear and Log-Linear Trends

Two exploratory trend models are fitted to GPU demand estimates for 2022–2026. The linear specification assumes a constant annual absolute change in GPU-hours, while the log-linear specification assumes approximately constant proportional growth.

Given the short modelling period and differences in data provenance, these models are intended to illustrate possible growth trajectories rather than establish reliable long-term GPU demand forecasts.

```
# Linear trend: GPU demand in millions of GPU-hours
gpu_linear <- lm(
  gpu_million_hours ~ time,
  data = gpu_recent_data
)

# Log-linear trend: natural logarithm of GPU demand
gpu_loglinear <- lm(
  log(gpu_million_hours) ~ time,
  data = gpu_recent_data
)

# Inspect fitted coefficients
summary(gpu_linear)
```

```
Call:
lm(formula = gpu_million_hours ~ time, data = gpu_recent_data)

Residuals:
       1        2        3        4        5
 0.66789 -0.82096 -0.04181 -0.12506  0.31994

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   1.9954     0.4980   4.007   0.0279 *
time          0.1720     0.2033   0.846   0.4597
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6429 on 3 degrees of freedom
Multiple R-squared:  0.1926,    Adjusted R-squared:  -0.0765
F-statistic: 0.7158 on 1 and 3 DF,  p-value: 0.4597
```

```
summary(gpu_loglinear)
```

```
Call:
lm(formula = log(gpu_million_hours) ~ time, data = gpu_recent_data)

Residuals:
       1        2        3        4        5
 0.32642 -0.43696  0.01619 -0.02718  0.12153
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.65314    0.25030   2.609   0.0797 .
time         0.08126    0.10218   0.795   0.4846
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3231 on 3 degrees of freedom
Multiple R-squared:  0.1741,    Adjusted R-squared:  -0.1012
F-statistic: 0.6324 on 1 and 3 DF,  p-value: 0.4846
```

The linear model estimates the annual absolute change in GPU demand, measured in millions of GPU-hours.

The log-linear model estimates the proportional annual growth trend, which can be expressed as a compounded percentage rate using the exponential transformation of its time coefficient.

The two specifications are not directly comparable using their in-sample R-squared values because they use different response scales.

With only five annual estimates, model coefficients are highly sensitive to individual observations, conversion assumptions, and changes in reporting coverage. Any projections derived from these models will therefore be presented as exploratory scenarios with substantial uncertainty.

9.7.3 GPU Trend Findings

The linear trend model estimates an annual increase in GPU demand of approximately 0.172 million GPU-hours over 2022–2026.

The log-linear model estimates an implied compounded annual growth rate of approximately 8.5%. However, neither trend coefficient is statistically significant at the 5% level.

Both models exhibit relatively low in-sample explanatory power, with R-squared values of approximately 0.193 and 0.174, respectively. These values are calculated on different response scales and should not be directly compared to select a model.

The historical GPU demand estimates fluctuate substantially across the five-year period. Furthermore, the series combines reported Setonix-based demand with estimates derived from GPU allocation shares and working conversion factors.

Consequently, the fitted trends provide limited evidence of a stable historical GPU demand growth pattern. The absence of a statistically significant trend should not be interpreted as evidence that future GPU demand will remain unchanged.

Given the small sample size and differences in data provenance, both models are retained only as exploratory forecasting scenarios. Further sensitivity analysis and explicit treatment of conversion uncertainty are necessary before drawing conclusions about GPU demand through 2030.

9.7.4 GPU Demand Sensitivity Analysis

The sensitivity of GPU demand trends to the 2026 observation is examined by comparing models fitted over 2022–2026 with equivalent models fitted over 2022 - 2025.

This analysis is particularly important because the recent GPU demand series contains only five annual estimates, combining reported Setonix-based demand with values derived using GPU allocation shares and conversion assumptions.

Excluding 2026 reduces the modelling sample to four observations. The resulting comparison is therefore exploratory and should not be interpreted as evidence of a structural change or as a reliable basis for selecting the preferred forecasting model.

```
# Prepare GPU demand estimates excluding 2026
gpu_pre2026_data <- gpu_recent_data |>
  dplyr::filter(year <= 2025)

# Refit linear and log-linear models
gpu_linear_pre2026 <- lm(
  gpu_million_hours ~ time,
  data = gpu_pre2026_data
)

gpu_loglinear_pre2026 <- lm(
  log(gpu_million_hours) ~ time,
  data = gpu_pre2026_data
)

# Compare estimated trend coefficients
gpu_trend_sensitivity <- data.frame(
  model = c("Linear trend", "Log-linear trend"),
  full_2022_2026 = c(
    coef(gpu_linear)[["time"]],
    coef(gpu_loglinear)[["time"]]
  ),
  excluding_2026 = c(
    coef(gpu_linear_pre2026)[["time"]],
    coef(gpu_loglinear_pre2026)[["time"]]
  )
)
```

```r
# Calculate percentage change in trend coefficients
gpu_trend_sensitivity$percentage_change <- with(
  gpu_trend_sensitivity,
  100 * (full_2022_2026 - excluding_2026) /
    abs(excluding_2026)
)

knitr::kable(
  gpu_trend_sensitivity,
  digits = 3,
  col.names = c(
    "Model",
    "Trend with 2026",
    "Trend without 2026",
    "Change (%)"
  ),
  caption = "Sensitivity of GPU demand trends to the inclusion of 2026"
)
```

Table 9.3: Sensitivity of GPU demand trends to the inclusion of 2026

Model	Trend with 2026	Trend without 2026	Change (%)
Linear trend	0.172	0.012	1330.816
Log-linear trend	0.081	0.020	296.466

The linear trend coefficient represents the estimated annual absolute change in GPU demand, measured in millions of GPU-hours.

The log-linear coefficient represents annual change in log GPU demand and can be converted into an implied compounded growth rate.

Because the modelling period is short, the estimated coefficients may change substantially when individual observations are included or excluded. Percentage changes should also be interpreted carefully when the reference trend coefficient is close to zero.

This sensitivity analysis evaluates model dependence on the latest observation but does not establish whether the 2026 increase reflects sustained demand growth, temporary variation, or changes in reporting and conversion assumptions.

9.7.5 GPU Sensitivity Findings

Including the 2026 observation increases the estimated linear GPU demand trend from approximately 0.012 to 0.172 million GPU-hours per year, representing a relative increase of approximately 1,330.8% in the trend coefficient.

The log-linear trend coefficient also increases substantially, from approximately 0.020 to 0.081, corresponding to a relative increase of approximately 296.5%.

These large percentage changes indicate that both GPU trend specifications are highly sensitive to the inclusion of the 2026 observation. However, the percentage changes are amplified by the relatively small trend coefficients estimated from the 2022 - 2025 sample.

Unlike the overall HPC and CPU demand models, where the log-linear trend coefficients were comparatively stable, both GPU specifications exhibit substantial sensitivity to the latest observation.

These findings must be interpreted cautiously because the full GPU modelling sample contains only five annual estimates, while the reduced sample contains four. Furthermore, the historical series combines reported Setonix-based demand with allocation-share estimates and working conversion factors.

The results do not establish whether the increase in 2026 represents sustained growth in GPU requirements or reflects temporary variation and differences in reporting methods.

Consequently, neither trend specification provides a sufficiently robust basis for selecting a definitive GPU forecast. Both are retained as exploratory scenarios, with additional consideration of historical conversion uncertainty required before projecting demand through 2030.

9.8 Historical GPU Conversion Scenario Assessment

The historical GPU conversion scenarios developed in Chapter 05 provide an indication of how assumptions about service-unit conversion affect estimated GPU resource requirements.

For 2014 - 2021, the low, baseline, and high GPU-hour series are hypothetical conversion scenarios rather than independently observed GPU demand. They should therefore not be treated as additional historical observations or automatically included in the recent-period GPU forecasting models.

The baseline scenario uses an assumed conversion of 36 service units per GPU-hour. Alternative conversion assumptions generate different GPU-hour estimates from the same underlying service-unit quantities.

These scenarios represent measurement and conversion uncertainty, rather than evidence of alternative historical demand growth patterns.

9.8.1 Historical Scenario Consistency

```
# Load historical GPU conversion scenarios
historical_gpu_scenarios <- readRDS(
  "assets/data/data converted/historical_gpu_scenarios.rds"
)

# Validate scenario coverage and ordering
stopifnot(
  isTRUE(all.equal(
    sort(as.numeric(historical_gpu_scenarios$year)),
    as.numeric(2014:2021)
  )),
  all(
    historical_gpu_scenarios$demand_gpu_hours_low <=
      historical_gpu_scenarios$demand_gpu_hours_baseline
  ),
  all(
    historical_gpu_scenarios$demand_gpu_hours_baseline <=
      historical_gpu_scenarios$demand_gpu_hours_high
  )
)

# Compare scenario values with the converted baseline
gpu_baseline_check <- historical_gpu_scenarios |>
  dplyr::select(
    year,
    demand_gpu_hours_baseline
  ) |>
  dplyr::left_join(
    forecast_data |>
      dplyr::select(year, demand_gpu_hours),
    by = "year"
  )

stopifnot(
  isTRUE(all.equal(
    gpu_baseline_check$demand_gpu_hours_baseline,
    gpu_baseline_check$demand_gpu_hours,
    tolerance = 1e-6
  ))
)

cat(
  "Historical GPU scenario period: 2014-2021\n",
```

```
  "Scenario observations:", nrow(historical_gpu_scenarios), "\n",
  "Baseline consistency: confirmed\n"
)
```

```
Historical GPU scenario period: 2014-2021
 Scenario observations: 8
 Baseline consistency: confirmed
```

This validation confirms that the historical GPU conversion scenarios cover 2014 - 2021, preserve the expected low-to-high ordering, and reproduce the baseline GPU demand estimates contained in the converted annual dataset.

The scenarios will be retained as evidence of historical conversion uncertainty. They will not be treated as observed GPU demand when fitting the 2022 - 2026 exploratory trend models.

Any future low, baseline, and high GPU demand projections must distinguish uncertainty in the historical conversion assumptions from uncertainty in future demand growth.

9.9 Demand Projections for 2027–2030

9.9.1 Overall HPC Demand Projections

Four candidate projection trajectories are generated for overall HPC demand using the previously fitted linear and log-linear trend models.

The projections include full-history specifications fitted over 2014 - 2026 and sensitivity specifications fitted over 2014 - 2025.

This comparison illustrates the effects of model functional form and the inclusion of the 2026 observation on projected demand through 2030.

```
# Forecast horizon
future_years <- data.frame(
  year = 2027:2030,
  time = 2027:2030 - 2014
)

# Generate projections in kSU
overall_projections <- future_years |>
  dplyr::transmute(
    year,
    linear_full = as.numeric(
      predict(overall_linear, newdata = future_years)
    ),
    loglinear_full = exp(as.numeric(
      predict(overall_loglinear, newdata = future_years)
```

```
    )),
    linear_pre2026 = as.numeric(
      predict(overall_linear_pre2026, newdata = future_years)
    ),
    loglinear_pre2026 = exp(as.numeric(
      predict(overall_loglinear_pre2026, newdata = future_years)
    ))
  )

knitr::kable(
  overall_projections,
  digits = 0,
  format.args = list(big.mark = ","),
  col.names = c(
    "Year",
    "Linear (Full)",
    "Log-linear (Full)",
    "Linear (Excl. 2026)",
    "Log-linear (Excl. 2026)"
  ),
  caption = "Candidate overall HPC demand projections, 2027-2030 (kSU)"
)
```

Table 9.4: Candidate overall HPC demand projections, 2027–2030 (kSU)

Year	Linear (Full)	Log-linear (Full)	Linear (Excl. 2026)	Log-linear (Excl. 2026)
2,027	2,016,845	3,217,801	1,852,255	3,396,628
2,028	2,178,838	4,064,043	1,996,612	4,314,830
2,029	2,340,830	5,132,835	2,140,970	5,481,247
2,030	2,502,822	6,482,707	2,285,327	6,962,979

The linear models project demand based on constant annual absolute growth, while the log-linear models project demand based on approximately constant proportional growth.

The log-linear projections are obtained by exponentiating predicted log demand. These represent back-transformed central trajectories, rather than bias-adjusted conditional mean forecasts.

The differences between trajectories reflect model specification and sensitivity to the 2026 observation. They should not be interpreted as formal prediction intervals or statistically established uncertainty bounds.

Final model selection and forecast accuracy assessment will be conducted in Chapter 08.

9.9.2 Overall HPC Projection Findings

The candidate models produce substantially different projections of overall HPC demand through 2030.

Under the linear trend specifications, projected demand in 2030 ranges from approximately 2.29 million to 2.50 million kSU. The log-linear specifications produce considerably higher projections, ranging from approximately 6.48 million to 6.96 million kSU.

The differences reflect the models' underlying growth assumptions. Linear models extrapolate constant annual absolute increases, whereas log-linear models extrapolate compounded proportional growth.

Excluding the 2026 observation reduces the 2030 linear projection but increases the corresponding log-linear projection. This demonstrates that the effect of the latest observation depends on the functional form of the forecasting model.

The wide range of projected outcomes highlights the importance of model selection and forecast evaluation, particularly given the short and heterogeneous historical dataset.

These trajectories are preliminary model-based projections, not formal prediction intervals or validated low, baseline, and high demand scenarios. Their relative predictive accuracy will be evaluated in Chapter 08 before selecting a preferred forecasting approach.

9.9.3 CPU Demand Projections

Four candidate CPU demand trajectories are generated for 2027 - 2030 using the previously fitted linear and log-linear trend models.

The full-history models use observations from 2014 - 2026, while the sensitivity specifications exclude the 2026 observation.

CPU demand is expressed in millions of core-hours, consistent with the resource-specific historical analysis in Chapter 06.

These projections examine how future CPU requirements differ under constant absolute growth and compounded proportional growth assumptions.

```
# Use the existing 2027-2030 forecast horizon
cpu_future_years <- data.frame(
  year = 2027:2030,
  time = (2027:2030) - 2014
)

# Generate candidate CPU demand projections
cpu_projections <- cpu_future_years |>
  dplyr::transmute(
    year,
```

```
    linear_full = as.numeric(
      predict(cpu_linear, newdata = cpu_future_years)
    ),
    loglinear_full = exp(as.numeric(
      predict(cpu_loglinear, newdata = cpu_future_years)
    )),
    linear_pre2026 = as.numeric(
      predict(cpu_linear_pre2026, newdata = cpu_future_years)
    ),
    loglinear_pre2026 = exp(as.numeric(
      predict(cpu_loglinear_pre2026, newdata = cpu_future_years)
    ))
  )

knitr::kable(
  cpu_projections,
  digits = 2,
  format.args = list(big.mark = ","),
  col.names = c(
    "Year",
    "Linear (Full)",
    "Log-linear (Full)",
    "Linear (Excl. 2026)",
    "Log-linear (Excl. 2026)"
  ),
  caption = paste(
    "Candidate CPU demand projections,",
    "2027-2030 (million core-hours)"
  )
)
```

Table 9.5: Candidate CPU demand projections, 2027–2030 (million core-hours)

Year	Linear (Full)	Log-linear (Full)	Linear (Excl. 2026)	Log-linear (Excl. 2026)
2,027	1,126.88	1,488.26	1,046.63	1,498.65
2,028	1,207.89	1,764.23	1,119.04	1,777.87
2,029	1,288.89	2,091.36	1,191.45	2,109.11
2,030	1,369.90	2,479.16	1,263.85	2,502.07

The linear specifications extrapolate constant annual absolute increases in CPU demand, whereas the log-linear specifications extrapolate compounded proportional growth.

The log-linear projections are back-transformed from predicted log demand using

exponentiation. They represent central trajectories on the original scale rather than bias-adjusted conditional mean forecasts.

Differences between the projected trajectories reflect model specification and sensitivity to the 2026 observation. They are not formal prediction intervals or validated low, baseline, and high demand scenarios.

The predictive performance of these models will be assessed in Chapter 08 before selecting a preferred CPU forecasting approach.

9.9.4 CPU Projection Findings

The candidate CPU demand models produce substantially different projections for 2027 - 2030, reflecting their alternative growth assumptions.

Under the linear specifications, projected CPU demand in 2030 ranges from approximately 1,263.85 to 1,369.90 million core-hours. The log-linear specifications produce higher projections, ranging from approximately 2,479.16 to 2,502.07 million core-hours.

These differences arise primarily from the models' functional forms. Linear models assume constant annual absolute increases in CPU demand, while log-linear models extrapolate compounded proportional growth.

Excluding the 2026 observation reduces the 2030 linear projection by approximately 7.7%, whereas the corresponding log-linear projection increases by approximately 0.9%.

The results indicate that projected CPU demand is more sensitive to the choice of growth specification than to the inclusion of 2026 within the models examined.

However, these trajectories remain preliminary and are conditional on historical reporting methods, conversion assumptions, and model specifications. They should not be interpreted as validated forecasts or formal uncertainty bounds.

The preferred CPU forecasting approach will be determined through model evaluation and predictive accuracy assessment in Chapter 08.

9.9.5 GPU Demand Projections

Four exploratory GPU demand trajectories are generated for 2027 - 2030 using the previously fitted linear and log-linear trend models.

The full-period specifications use GPU demand estimates from 2022 - 2026, while the sensitivity specifications use 2022 - 2025, excluding the 2026 observation.

GPU demand is expressed in millions of GPU-hours, consistent with the resource-specific analysis in Chapter 06.

Unlike the overall HPC and CPU projections, these trajectories are based on a very short and heterogeneous historical series. The estimates combine reported

Setonix-based GPU demand with values derived from allocation shares and working conversion factors.

Consequently, the projections are exploratory model trajectories rather than validated forecasts of future GPU requirements.

```
# Prepare the 2027-2030 forecast horizon
gpu_future_years <- data.frame(
  year = 2027:2030,
  time = (2027:2030) - 2022
)

# Generate exploratory GPU demand projections
gpu_projections <- gpu_future_years |>
  dplyr::transmute(
    year,
    linear_full = as.numeric(
      predict(gpu_linear, newdata = gpu_future_years)
    ),
    loglinear_full = exp(as.numeric(
      predict(gpu_loglinear, newdata = gpu_future_years)
    )),
    linear_pre2026 = as.numeric(
      predict(gpu_linear_pre2026, newdata = gpu_future_years)
    ),
    loglinear_pre2026 = exp(as.numeric(
      predict(gpu_loglinear_pre2026, newdata = gpu_future_years)
    ))
  )

knitr::kable(
  gpu_projections,
  digits = 3,
  col.names = c(
    "Year",
    "Linear (Full)",
    "Log-linear (Full)",
    "Linear (Excl. 2026)",
    "Log-linear (Excl. 2026)"
  ),
  caption = paste(
    "Exploratory GPU demand projections,",
    "2027-2030 (million GPU-hours)"
  )
)
```

Table 9.6: Exploratory GPU demand projections, 2027–2030 (million GPU-hours)

Year	Linear (Full)	Log-linear (Full)	Linear (Excl. 2026)	Log-linear (Excl. 2026)
2027	2.855	2.885	2.215	2.262
2028	3.027	3.129	2.227	2.309
2029	3.199	3.394	2.240	2.357
2030	3.371	3.681	2.252	2.406

The linear specifications extrapolate constant annual absolute changes in GPU demand, while the log-linear specifications extrapolate compounded proportional growth.

Exponentiating the log-linear predictions produces central trajectories on the original scale, not bias-adjusted conditional mean forecasts.

The historical GPU conversion scenarios for 2014–2021 are not included as additional observations in these models. Those scenarios represent conversion uncertainty rather than directly observed GPU demand.

Given the limited number of recent observations, differences in data provenance, and strong sensitivity to the inclusion of 2026, these projections should not be interpreted as reliable estimates of future GPU demand or as formal uncertainty intervals.

Their suitability for forecasting will be considered further in Chapter 08.

9.9.6 GPU Projection Findings

The exploratory GPU demand models produce different projections for 2027 - 2030, reflecting both their growth assumptions and sensitivity to the inclusion of 2026.

Under the linear specifications, projected GPU demand in 2030 ranges from approximately 2.252 to 3.371 million GPU-hours. The log-linear specifications produce projections ranging from approximately 2.406 to 3.681 million GPU-hours.

Excluding 2026 substantially reduces the projected GPU demand under both model specifications. For the linear model, the 2030 projection increases by approximately 49.7% when 2026 is included.

These results indicate that GPU projections are particularly sensitive to the historical modelling period. This contrasts with the CPU projections, where differences between linear and log-linear growth assumptions were more pronounced.

However, the GPU modelling sample contains only five annual estimates, with a mixture of reported Setonix-based demand and allocation-share-derived values. The models also exhibit weak evidence of a stable historical growth trend.

Consequently, the projected trajectories should be interpreted as exploratory illustrations of possible GPU demand growth, rather than reliable forecasts or formal uncertainty bounds.

The historical 2014–2021 GPU conversion scenarios remain separate from these future growth projections, as they represent uncertainty in conversion assumptions rather than independently observed historical GPU demand.

Further model evaluation and explicit scenario development will be required before selecting a preferred GPU demand projection for infrastructure planning.

9.10 Field of Research Demand Modelling Framework

9.10.1 Purpose and Modelling Approach

Field of Research (FoR) demand is considered as an additional dimension of national HPC planning, complementing the overall HPC, CPU, and GPU demand projections developed in this chapter.

Unlike the aggregate and resource-specific demand measures, the harmonised FoR dataset primarily describes the distribution of allocated computing resources across research fields.

Consequently, FoR-specific demand cannot be treated as independently observed researcher demand.

A top-down allocation-share approach is therefore proposed to estimate indicative FoR-level demand through 2030.

Under this approach, the selected overall HPC demand forecast will be distributed across harmonised FoR groups using explicit assumptions about their future allocation shares.

The modelling relationship is:

$$\widehat{D}_{f,t} = \widehat{D}_t \times s_{f,t}$$

where:

- $\widehat{D}_{f,t}$ represents indicative demand for FoR group f in year t, expressed in kSU.
- $\widehat{D}_t$ represents the selected overall HPC demand forecast in year t, expressed in kSU.
- $s_{f,t}$ represents the assumed allocation share of FoR group f in year t.

The assumed shares will be constrained to be non-negative and sum to one within each forecast year, ensuring that the indicative FoR demand estimates reconcile with the overall HPC demand projection.

9.10.2 Historical FoR Data Suitability

The historical FoR analysis in Chapter 06 covers 2014 - 2026 and includes changes in research classification standards, reporting coverage, and analytical grouping.

The earlier period (2014 - 2020) is primarily organised using FoR 2008 classifications, while the later period (2021 - 2026) uses FoR 2020 classifications.

Although Chapter 06 harmonised these classifications into broader analytical groups for visualisation, the resulting categories are not perfectly equivalent across periods.

Furthermore, allocation shares reflect the distribution of computing resources awarded to research fields, rather than the full amount of computing resources requested by researchers.

For these reasons, historical FoR allocation shares will be treated as indicators of research activity composition rather than direct measures of FoR demand.

9.10.3 FoR Allocation-Share Scenarios

Three approaches will be considered for developing future FoR allocation-share assumptions:

1. Recent composition: hold the latest available harmonised FoR allocation shares constant through 2030.
2. Recent-period average: use average allocation shares from the FoR 2020 reporting period (2021 - 2026).
3. Composition-change sensitivity: examine how alternative share assumptions affect indicative demand across broad research fields.

These approaches are proposed modelling assumptions, not statistically validated predictions of future research-field composition.

The suitability of the historical shares and their variation across recent years will be assessed in Chapter 08.

Final assumptions will be selected in Chapter 09, and indicative FoR demand scenarios for 2027 - 2030 will be generated in Chapter 10.

9.10.4 Interpretation and Limitations

The resulting FoR projections will represent indicative distributions of overall forecast HPC demand under specified allocation-share assumptions.

They should not be interpreted as independently forecast researcher requests by FoR, nor as evidence that historical allocation shares fully capture unmet computing demand.

Differences in FoR classification, facility coverage, and allocation practices may affect the comparability of historical shares.

The top-down approach preserves consistency with the selected overall HPC demand projection while making the underlying assumptions transparent.

These limitations will be retained when interpreting FoR-specific planning implications.

9.11 Modelling Summary and Limitations

9.11.1 Summary of Candidate Demand Projections

This chapter developed linear and log-linear forecasting models for three quantitative measures of Australian national research supercomputing demand: overall HPC demand in kSU, CPU demand in core-hours, and GPU demand in GPU-hours.

In addition, a top-down allocation-share framework was established to develop indicative Field of Research (FoR) demand scenarios. Unlike the three quantitative forecasting targets, FoR-specific demand is not independently modelled from observed researcher requests.

The models were fitted using harmonised historical estimates, with alternative specifications excluding 2026 to assess the influence of the latest observation.

The resulting candidate projections for 2030 are summarised below.

Demand measure	Linear models	Log-linear models	Unit
Overall HPC	2.29–2.50 million	6.48–6.96 million	kSU
CPU	1,263.85–1,369.90 million	2,479.16–2,502.07 million	CPU core-hours
GPU	2.252–3.371 million	2.406–3.681 million	GPU-hours

These ranges reflect the alternative full-period and 2026-excluded model specifications. They are not statistical prediction intervals or validated low, baseline, and high demand scenarios.

9.11.2 Key Modelling Findings

The overall HPC and CPU analyses identify substantial historical demand growth. However, projected demand through 2030 differs considerably between linear and

log-linear specifications, particularly because compounded proportional growth produces increasingly divergent trajectories.

For overall HPC and CPU demand, excluding 2026 increases or decreases projected demand depending on the model specification. The estimated log-linear growth coefficients are comparatively stable, although changes in fitted intercepts can still affect future projections.

GPU demand exhibits substantially greater uncertainty. The exploratory models use only five annual estimates from 2022 - 2026, combining reported Setonix-based demand with allocation-share-derived estimates.

Both GPU trend specifications are sensitive to the inclusion of 2026, and neither provides strong statistical evidence of a stable historical growth pattern. Consequently, GPU projections should be interpreted as exploratory trajectories rather than reliable long-term forecasts.

For Field of Research, the proposed modelling approach distributes the selected overall HPC demand forecast across harmonised research fields using explicit allocation-share assumptions. This approach preserves consistency with aggregate demand projections while recognising that historical FoR allocation shares do not directly measure researcher demand. The resulting FoR projections will be developed as indicative planning scenarios rather than independently estimated demand forecasts.

9.11.3 Limitations and Implications for Forecast Evaluation

Several limitations affect the interpretation of the candidate projections.

First, the historical demand series combines directly reported, derived, and proxy-based estimates, with differences in reporting coverage and conversion assumptions.

Second, the overall HPC and CPU models are fitted to only 13 annual observations, limiting the reliability of complex forecasting approaches and long-horizon extrapolation.

Third, the historical GPU conversion scenarios for 2014 - 2021 represent hypothetical estimates rather than independently observed demand. They are retained as evidence of conversion uncertainty but are not included in the recent-period GPU trend models.

Fourth, the substantial increase in reported demand in 2026 has not been established as a sustained structural change. Its influence is therefore examined through alternative model specifications rather than assumed to persist.

Fifth, the FoR modelling framework relies on historical allocation shares rather than directly observed research-field demand. Changes in FoR classification standards, reporting coverage, and allocation practices limit historical comparability.

Consequently, FoR-level projections will depend on explicit assumptions about future research-field composition.

Finally, the candidate models have not yet been selected using out-of-sample predictive accuracy. In-sample fit statistics alone are insufficient for identifying the most defensible forecasting approach.

Chapter 08 will evaluate the overall HPC, CPU, and GPU candidate models using time-series cross-validation, consistent forecast error measures, and residual diagnostics where appropriate.

It will also examine the stability and comparability of historical FoR allocation shares to inform the development of FoR-specific planning scenarios.

The resulting evidence will support model selection in Chapter 09 and the development of national HPC demand scenarios through 2030 in Chapter 10.

Chapter 10

Model Evaluation

10.1 Purpose

This chapter evaluates the predictive performance and statistical suitability of the candidate forecasting models developed in Chapter 07.

The evaluation considers three distinct measures of Australian national research supercomputing demand:

1. Overall HPC demand, expressed in thousands of service units (kSU).
2. CPU demand, expressed in millions of CPU core-hours.
3. GPU demand, expressed in millions of GPU-hours.

The objective is to identify forecasting approaches that provide a defensible basis for estimating researcher demand through 2030, while recognising limitations in historical data quality, comparability, and sample size.

Model selection will prioritise predictive accuracy, simplicity, interpretability, and robustness to influential observations, particularly the increase in reported demand in 2026.

10.2 Evaluation Framework

10.2.1 Candidate Models and Benchmarks

The evaluation considers four forecasting approaches:

1. Naïve benchmark: forecasts future demand using the most recent available observation.
2. Drift benchmark: extrapolates the average annual change between the first and most recent observations in the training sample.
3. Linear trend: estimates a constant annual absolute change in demand.

4. Log-linear trend: estimates approximately constant proportional growth using a regression model fitted to log-transformed demand.

Naïve and drift forecasts provide simple reference points against which the regression-based models can be evaluated.

The linear and log-linear specifications follow the modelling framework established in Chapter 07, including sensitivity analyses excluding the 2026 observation.

More complex models will not be prioritised unless their additional complexity can be justified by the available data and demonstrated predictive performance.

10.2.2 Time-Series Cross-Validation

Forecast accuracy will be assessed using rolling-origin time-series cross-validation.

At each evaluation origin, models will be estimated using only observations available up to that year. Forecasts will then be compared with subsequent historical observations.

This approach preserves chronological ordering and avoids using future observations to estimate models evaluated on earlier periods.

One-step-ahead evaluation will provide the primary comparison because the historical series are relatively short.

Where sufficient observations are available, additional forecast horizons may be examined to assess the stability of longer-term extrapolations.

The evaluation will distinguish between the longer overall HPC and CPU series covering 2014–2026 and the shorter GPU modelling series covering 2022 - 2026.

For GPU demand, the small number of recent observations severely limits the number of meaningful cross-validation origins. Any accuracy comparisons will therefore be considered exploratory.

10.2.3 Forecast Accuracy Measures

Forecast errors will be evaluated on the original scale of each demand measure using:

1. Root Mean Squared Error (RMSE), which gives greater weight to larger forecast errors.
2. Mean Absolute Error (MAE), which measures the average absolute forecast error.

Accuracy measures will be calculated separately for overall HPC, CPU, and GPU demand because their units are different.

For log-linear models, predictions will be back-transformed to the original demand scale before calculating forecast errors. The evaluation will distinguish

exponentiated log-scale predictions from bias-adjusted mean forecasts where relevant.

Models will be compared using common evaluation origins within each demand measure to ensure a fair comparison.

10.2.4 Model Selection Principles

The preferred forecasting approach will not be selected solely on the basis of the lowest numerical error.

Model selection will also consider:

- Stability across evaluation periods and forecast origins.
- Sensitivity to the inclusion of the 2026 observation.
- Residual behaviour and evidence of systematic model errors.
- Interpretability and consistency with historical data provenance.
- The risks of extrapolating growth patterns through 2030.

Particular caution will be applied to GPU demand because its historical modelling period is short and includes estimates derived from different reporting and conversion methods.

The evaluation results presented in this chapter will provide the statistical evidence for model selection in Chapter 09 and the development of demand scenarios in Chapter 10.

10.3 Rolling-Origin Cross-Validation

10.3.1 Overall HPC Demand: One-Step-Ahead Evaluation

The initial evaluation focuses on overall HPC demand, measured in thousands of service units (kSU).

Rolling-origin cross-validation is applied using an initial training period of 2014 - 2020. At each subsequent origin, the training sample expands by one year, and a one-step-ahead forecast is generated for the following year.

This produces six evaluation years, from 2021 to 2026.

Four forecasting approaches are evaluated using identical training periods and forecast origins:

1. Naïve benchmark.
2. Drift benchmark.
3. Linear trend.
4. Log-linear trend.

Each model is re-estimated using only the observations available at the corresponding forecast origin. This avoids information leakage from future observations.

```
# Load harmonised annual demand data
annual_converted <- readRDS(
  "assets/data/data converted/annual_demand_allocation_converted.rds"
)

# Prepare overall HPC demand series
overall_cv_data <- annual_converted |>
  dplyr::transmute(
    year = as.integer(year),
    demand_ksu = as.numeric(demand_ksu)
  ) |>
  dplyr::arrange(year)

# Validate historical series
stopifnot(
  identical(overall_cv_data$year, 2014:2026),
  all(is.finite(overall_cv_data$demand_ksu)),
  all(overall_cv_data$demand_ksu > 0)
)

# Define common one-step-ahead evaluation years
cv_test_years <- 2021:2026

# Generate forecasts using expanding training windows
overall_cv_predictions <- dplyr::bind_rows(
  lapply(cv_test_years, function(test_year) {

    training <- overall_cv_data |>
      dplyr::filter(year < test_year) |>
      dplyr::mutate(time = year - 2014)

    actual <- overall_cv_data$demand_ksu[
      overall_cv_data$year == test_year
    ]

    future <- data.frame(time = test_year - 2014)

    # Naïve: last observed demand
    naive_forecast <- tail(training$demand_ksu, 1)

    # Drift: average change from first to last observation
    drift_forecast <- tail(training$demand_ksu, 1) +
```

```
      (
        tail(training$demand_ksu, 1) -
          training$demand_ksu[1]
      ) / (nrow(training) - 1)

    # Linear trend fitted only to training data
    linear_fit <- lm(
      demand_ksu ~ time,
      data = training
    )

    linear_forecast <- as.numeric(
      predict(linear_fit, newdata = future)
    )

    # Log-linear trend fitted only to training data
    loglinear_fit <- lm(
      log(demand_ksu) ~ time,
      data = training
    )

    loglinear_forecast <- exp(as.numeric(
      predict(loglinear_fit, newdata = future)
    ))

    data.frame(
      year = test_year,
      actual = actual,
      naive = naive_forecast,
      drift = drift_forecast,
      linear = linear_forecast,
      loglinear = loglinear_forecast
    )
  })
)

# Display the predictions for inspection
knitr::kable(
  overall_cv_predictions,
  digits = 0,
  format.args = list(big.mark = ","),
  col.names = c(
    "Year",
    "Actual (kSU)",
    "Naïve",
```

```
    "Drift",
    "Linear",
    "Log-linear"
  ),
  caption = paste(
    "Rolling-origin one-step-ahead predictions",
    "for overall HPC demand, 2021-2026 (kSU)"
  )
)
```

Table 10.1: Rolling-origin one-step-ahead predictions for overall HPC demand, 2021–2026 (kSU)

Year	Actual (kSU)	Naïve	Drift	Linear	Log-linear
2,021	1,193,811	607,980	687,417	743,407	935,745
2,022	1,162,230	1,193,811	1,345,590	1,063,196	1,397,147
2,023	1,333,240	1,162,230	1,291,089	1,239,332	1,736,678
2,024	1,425,560	1,333,240	1,466,783	1,415,618	2,081,974
2,025	1,564,643	1,425,560	1,554,980	1,563,079	2,382,752
2,026	2,242,817	1,564,643	1,694,942	1,707,897	2,673,820

All four approaches produce forecasts for the same six test years, allowing their prediction errors to be compared consistently.

The log-linear forecasts are exponentiated predictions from the fitted log-demand models. They are evaluated on the original kSU scale and represent back-transformed central predictions, rather than bias-adjusted conditional means.

The resulting predictions will be used to calculate RMSE and MAE in the next stage of the evaluation.

10.3.2 Overall HPC Forecast Accuracy

Forecast accuracy is evaluated using root mean squared error (RMSE) and mean absolute error (MAE) across the six one-step-ahead predictions for 2021 - 2026.

Both measures are calculated on the original kSU scale. Lower values indicate better historical predictive accuracy over the common evaluation period.

```
# Calculate forecast accuracy for each candidate model
overall_cv_accuracy <- dplyr::bind_rows(
  lapply(
    c("naive", "drift", "linear", "loglinear"),
    function(model_name) {

      errors <- overall_cv_predictions[[model_name]] -
```

```
        overall_cv_predictions$actual

      data.frame(
        model = model_name,
        RMSE = sqrt(mean(errors^2)),
        MAE = mean(abs(errors))
      )
    }
  )
) |>
  dplyr::arrange(RMSE)

knitr::kable(
  overall_cv_accuracy,
  digits = 0,
  format.args = list(big.mark = ","),
  col.names = c("Model", "RMSE (kSU)", "MAE (kSU)"),
  caption = paste(
    "One-step-ahead cross-validation accuracy",
    "for overall HPC demand, 2021-2026"
  )
)
```

Table 10.2: One-step-ahead cross-validation accuracy for overall HPC demand, 2021–2026

Model	RMSE (kSU)	MAE (kSU)
linear	290,898	198,295
drift	314,588	221,778
naive	378,864	283,000
loglinear	511,614	466,991

The accuracy comparison uses identical test years for all four models, ensuring that differences in RMSE and MAE are not caused by different evaluation periods.

RMSE places greater weight on large forecast errors, whereas MAE provides a measure of typical absolute forecast error.

These metrics provide evidence for subsequent model selection in Chapter 09, alongside considerations of model simplicity, data provenance, and sensitivity to influential observations.

10.3.3 Overall HPC Accuracy Findings

Rolling-origin one-step-ahead cross-validation over 2021–2026 indicates that the linear trend model achieves the lowest forecast errors among the four evaluated approaches.

The linear model records an RMSE of approximately 290,898 kSU and an MAE of 198,295 kSU. The drift benchmark ranks second, with an RMSE of 314,588 kSU and an MAE of 221,778 kSU.

The naïve benchmark produces an RMSE of 378,864 kSU, while the log-linear model records the highest RMSE of 511,614 kSU and MAE of 466,991 kSU.

These findings indicate that the linear trend provides the strongest one-step-ahead predictive performance over the selected evaluation period.

Although the log-linear model exhibited a relatively high in-sample R-squared value in Chapter 07, its weaker cross-validation accuracy demonstrates that in-sample fit alone is insufficient for selecting a forecasting model.

The evaluation is nevertheless based on only six forecast errors and includes the substantial demand increase in 2026. Consequently, the model rankings may be sensitive to the selected evaluation period and individual observations.

Furthermore, one-step-ahead accuracy does not necessarily translate into superior multi-year forecasting performance through 2030.

10.3.4 CPU Demand: One-Step-Ahead Evaluation

CPU demand is evaluated using expanding-window rolling-origin cross-validation, consistent with the methodology applied to overall HPC demand.

The initial training period covers 2014–2020, with one-step-ahead forecasts generated for each year from 2021 to 2026.

Four approaches are evaluated: naïve, drift, linear trend, and log-linear trend. Each model uses only observations available before its corresponding evaluation year.

CPU demand is expressed in millions of core-hours, consistent with the modelling approach established in Chapter 07.

```
# Prepare CPU demand in millions of core-hours
cpu_cv_data <- annual_converted |>
  dplyr::transmute(
    year = as.integer(year),
    cpu_million_hours = as.numeric(demand_cpu_core_hours) / 1e6
  ) |>
  dplyr::arrange(year)

# Validate the CPU demand series
```

```r
stopifnot(
  identical(cpu_cv_data$year, 2014:2026),
  all(is.finite(cpu_cv_data$cpu_million_hours)),
  all(cpu_cv_data$cpu_million_hours > 0)
)

# Use the same evaluation years as overall HPC
cpu_cv_predictions <- dplyr::bind_rows(
  lapply(2021:2026, function(test_year) {

    training <- cpu_cv_data |>
      dplyr::filter(year < test_year) |>
      dplyr::mutate(time = year - 2014)

    actual <- cpu_cv_data$cpu_million_hours[
      cpu_cv_data$year == test_year
    ]

    future <- data.frame(time = test_year - 2014)

    # Naïve benchmark
    naive_forecast <- tail(training$cpu_million_hours, 1)

    # Drift benchmark
    drift_forecast <- tail(training$cpu_million_hours, 1) +
      (
        tail(training$cpu_million_hours, 1) -
          training$cpu_million_hours[1]
      ) / (nrow(training) - 1)

    # Linear trend
    linear_fit <- lm(
      cpu_million_hours ~ time,
      data = training
    )

    linear_forecast <- as.numeric(
      predict(linear_fit, newdata = future)
    )

    # Log-linear trend
    loglinear_fit <- lm(
      log(cpu_million_hours) ~ time,
      data = training
    )
```

```
    loglinear_forecast <- exp(as.numeric(
      predict(loglinear_fit, newdata = future)
    ))

    data.frame(
      year = test_year,
      actual = actual,
      naive = naive_forecast,
      drift = drift_forecast,
      linear = linear_forecast,
      loglinear = loglinear_forecast
    )
  })
)

knitr::kable(
  cpu_cv_predictions,
  digits = 2,
  format.args = list(big.mark = ","),
  col.names = c(
    "Year",
    "Actual (million core-hours)",
    "Naïve",
    "Drift",
    "Linear",
    "Log-linear"
  ),
  caption = paste(
    "Rolling-origin one-step-ahead CPU demand predictions,",
    "2021-2026 (million core-hours)"
  )
)
```

Table 10.3: Rolling-origin one-step-ahead CPU demand predictions, 2021–2026 (million core-hours)

Year	Actual (million core-hours)	Naïve	Drift	Linear	Log-linear
2,021	736.48	294.87	322.34	554.60	609.86
2,022	521.21	736.48	823.11	705.46	819.78
2,023	775.47	521.21	570.11	698.65	832.94
2,024	863.24	775.47	847.18	792.17	975.91
2,025	942.05	863.24	936.56	884.99	1,120.88
2,026	1,235.04	942.05	1,015.87	974.23	1,263.29

The evaluation uses identical forecast origins across all four models, allowing a consistent comparison of their predictive performance.

The log-linear forecasts are back-transformed to the original CPU demand scale before evaluation. These exponentiated values represent central predictions rather than bias-adjusted conditional mean forecasts.

The resulting predictions will be used to calculate RMSE and MAE in the next stage.

10.3.5 CPU Forecast Accuracy

Forecast accuracy is assessed using root mean squared error (RMSE) and mean absolute error (MAE), calculated from the six rolling-origin one-step-ahead predictions for 2021 - 2026.

Both measures are expressed in millions of CPU core-hours.

```
# Calculate CPU forecast accuracy
cpu_cv_accuracy <- dplyr::bind_rows(
  lapply(c("naive", "drift", "linear", "loglinear"), function(model) {

    errors <- cpu_cv_predictions[[model]] -
      cpu_cv_predictions$actual

    data.frame(
      model = model,
      RMSE = sqrt(mean(errors^2)),
      MAE = mean(abs(errors))
    )
  })
) |>
  dplyr::arrange(RMSE)

knitr::kable(
  cpu_cv_accuracy,
  digits = 2,
  col.names = c(
    "Model",
    "RMSE (million core-hours)",
    "MAE (million core-hours)"
  ),
  caption = "CPU demand forecast accuracy, 2021-2026"
)
```

Table 10.4: CPU demand forecast accuracy, 2021–2026

Model	RMSE (million core-hours)	MAE (million core-hours)
linear	157.72	138.65
loglinear	160.18	133.74
drift	242.61	193.69
naive	260.05	228.45

Lower RMSE and MAE values indicate better predictive accuracy. RMSE places greater weight on large forecast errors, while MAE represents the average absolute deviation between predicted and actual CPU demand.

10.3.6 CPU Accuracy Findings

Rolling-origin one-step-ahead cross-validation over 2021 - 2026 indicates that the linear and log-linear trend models provide the strongest predictive performance among the four evaluated CPU forecasting approaches.

The linear model achieves the lowest RMSE of 157.72 million CPU core-hours, followed closely by the log-linear model at 160.18 million CPU core-hours.

However, the log-linear model achieves the lowest MAE of 133.74 million CPU core-hours, compared with 138.65 million CPU core-hours for the linear model.

Both trend models outperform the drift and naïve benchmarks under both accuracy measures.

The difference in model rankings indicates that no single CPU trend specification dominates across both evaluation criteria. The linear model performs slightly better under RMSE, which places greater weight on larger errors, while the log-linear model performs slightly better under MAE.

These findings also demonstrate the importance of out-of-sample evaluation rather than relying solely on in-sample model fit.

Nevertheless, the evaluation is based on only six one-step-ahead forecast errors. Model rankings may be sensitive to individual observations, particularly the substantial increase in CPU demand observed in 2026.

Furthermore, one-step-ahead predictive accuracy does not necessarily establish superior performance for multi-year projections through 2030.

Both linear and log-linear trend models are therefore retained as competitive candidates for CPU demand forecasting. Final model selection will be addressed in Chapter 09.

10.3.7 GPU Demand: One-Step-Ahead Evaluation

GPU demand is evaluated using expanding-window rolling-origin cross-validation, consistent with the approach applied to overall HPC and CPU demand.

However, the GPU evaluation uses only the 2022–2026 period, as earlier GPU demand estimates were constructed using hypothetical conversion scenarios rather than independently reported GPU demand.

The initial training period covers 2022 - 2024, with one-step-ahead forecasts generated for 2025 and 2026.

Four approaches are evaluated: naïve, drift, linear trend, and log-linear trend.

Each model is estimated using only observations available before the corresponding evaluation year.

GPU demand is expressed in millions of GPU-hours.

```
# Prepare the recent GPU demand series
gpu_cv_data <- annual_converted |>
  dplyr::transmute(
    year = as.integer(year),
    gpu_million_hours = as.numeric(demand_gpu_hours) / 1e6
  ) |>
  dplyr::filter(year >= 2022, year <= 2026) |>
  dplyr::arrange(year)

# Validate recent GPU data
stopifnot(
  identical(gpu_cv_data$year, 2022:2026),
  all(is.finite(gpu_cv_data$gpu_million_hours)),
  all(gpu_cv_data$gpu_million_hours > 0)
)

# Generate one-step-ahead predictions
gpu_cv_predictions <- dplyr::bind_rows(
  lapply(2025:2026, function(test_year) {

    training <- gpu_cv_data |>
      dplyr::filter(year < test_year) |>
      dplyr::mutate(time = year - 2022)

    actual <- gpu_cv_data$gpu_million_hours[
      gpu_cv_data$year == test_year
    ]

    future <- data.frame(time = test_year - 2022)
```

```
    # Naïve benchmark
    naive_forecast <- tail(training$gpu_million_hours, 1)

    # Drift benchmark
    drift_forecast <- tail(training$gpu_million_hours, 1) +
      (
        tail(training$gpu_million_hours, 1) -
          training$gpu_million_hours[1]
      ) / (nrow(training) - 1)

    # Linear trend
    linear_fit <- lm(
      gpu_million_hours ~ time,
      data = training
    )

    linear_forecast <- as.numeric(
      predict(linear_fit, newdata = future)
    )

    # Log-linear trend
    loglinear_fit <- lm(
      log(gpu_million_hours) ~ time,
      data = training
    )

    loglinear_forecast <- exp(as.numeric(
      predict(loglinear_fit, newdata = future)
    ))

    data.frame(
      year = test_year,
      actual = actual,
      naive = naive_forecast,
      drift = drift_forecast,
      linear = linear_forecast,
      loglinear = loglinear_forecast
    )
  })
)

knitr::kable(
  gpu_cv_predictions,
  digits = 3,
  col.names = c(
```

```
    "Year",
    "Actual (million GPU-hours)",
    "Naïve",
    "Drift",
    "Linear",
    "Log-linear"
  ),
  caption = paste(
    "Rolling-origin one-step-ahead GPU demand predictions,",
    "2025-2026 (million GPU-hours)"
  )
)
```

Table 10.5: Rolling-origin one-step-ahead GPU demand predictions, 2025–2026 (million GPU-hours)

Year	Actual (million GPU-hours)	Naïve	Drift	Linear	Log-linear
2025	2.386	2.298	2.115	1.737	1.742
2026	3.003	2.386	2.294	2.203	2.216

The evaluation uses the same forecast origins for all fouapproaches, allowing a consistent comparison of their one-step-ahead predictions.

However, the results are based on only two evaluation observations. Furthermore, the historical GPU series combines reported and allocation-share-derived estimates, introducing additional measurement uncertainty.

The resulting predictions will therefore be interpreted as exploratory evidence rather than a reliable basis for selecting a preferred GPU forecasting model.

10.3.8 GPU Forecast Accuracy

Forecast accuracy is assessed using root mean squared error (RMSE) and mean absolute error (MAE), calculated from the two rolling-origin one-step-ahead predictions for 2025 - 2026.

Both measures are expressed in millions of GPU-hours.

```
# Calculate GPU forecast accuracy
gpu_cv_accuracy <- dplyr::bind_rows(
  lapply(c("naive", "drift", "linear", "loglinear"), function(model) {

    errors <- gpu_cv_predictions[[model]] -
      gpu_cv_predictions$actual

    data.frame(
```

```
      model = model,
      RMSE = sqrt(mean(errors^2)),
      MAE = mean(abs(errors))
    )
  })
) |>
  dplyr::arrange(RMSE)

knitr::kable(
  gpu_cv_accuracy,
  digits = 3,
  col.names = c(
    "Model",
    "RMSE (million GPU-hours)",
    "MAE (million GPU-hours)"
  ),
  caption = "GPU demand forecast accuracy, 2025-2026"
)
```

Table 10.6: GPU demand forecast accuracy, 2025–2026

Model	RMSE (million GPU-hours)	MAE (million GPU-hours)
naive	0.441	0.353
drift	0.537	0.490
loglinear	0.719	0.715
linear	0.729	0.725

Lower RMSE and MAE values indicate better predictive accuracy. However, with only two evaluation observations, these metrics provide limited evidence of how the models would perform under different historical conditions or over longer forecast horizons.

10.3.9 GPU Accuracy Findings

Rolling-origin one-step-ahead cross-validation over 2025 - 2026 indicates that the naïve benchmark achieves the lowest forecast errors among the four evaluated GPU forecasting approaches.

The naïve model records an RMSE of approximately 0.441 million GPU-hours and an MAE of 0.353 million GPU-hours.

The drift benchmark ranks second, with an RMSE of 0.537 million GPU-hours and an MAE of 0.490 million GPU-hours.

The log-linear and linear trend models produce higher forecast errors, with RMSE values of approximately 0.719 and 0.729 million GPU-hours, respectively.

These results suggest that simple persistence-based forecasting performs better than trend extrapolation over the selected GPU evaluation period.

However, the evaluation contains only two one-step-ahead forecast errors, substantially fewer than the six observations used for overall HPC and CPU evaluation.

Furthermore, the GPU demand series combines reported and allocation-share-derived estimates, introducing additional measurement uncertainty.

The findings therefore provide insufficient evidence to establish a stable or generalisable ranking of GPU forecasting models.

In particular, the naïve model's superior performance over 2025 - 2026 does not establish that future GPU demand will remain constant or that a naïve forecast will be appropriate for the multi-year horizon through 2030.

The naïve benchmark is therefore identified as the strongest candidate under the current evaluation criteria, rather than selected as the final GPU forecasting model.

Final model selection and consideration of alternative GPU demand scenarios will be addressed in Chapter 09.

10.4 Field of Research Allocation-Share Evaluation

10.4.1 Purpose and Evaluation Approach

The Field of Research (FoR) evaluation examines the historical distribution and stability of HPC resource allocations across research fields.

Unlike the overall HPC, CPU, and GPU evaluations, this analysis does not assess independently observed FoR-specific researcher demand forecasts.

Instead, it evaluates whether historical allocation shares provide a sufficiently consistent basis for developing indicative FoR demand scenarios through 2030.

The analysis focuses on 2021–2026, corresponding to the FoR 2020 classification period used in Chapter 06.

Three considerations guide the evaluation:

1. Historical stability of FoR allocation shares.
2. Variation between the latest-year composition and the recent-period average.
3. Sensitivity to classification, reporting coverage, and allocation-based measurement assumptions.

The results will inform the selection of FoR share assumptions in Chapter 09 and their application to overall HPC demand forecasts in Chapter 10.

10.4.2 Harmonised FoR Data Preparation

The FoR evaluation reuses the harmonised allocation-share datasets developed in Chapter 06, ensuring consistency between historical analysis and model evaluation.

Two reporting periods are considered:

- The full historical period (2014 - 2026), using broad harmonised analytical groups to examine long-term allocation composition.
- The recent period (2021–2026), using FoR 2020-based analytical groups to evaluate recent allocation-share stability.

The full-period analysis provides historical context, while the recent-period analysis informs the assumptions used for indicative FoR demand scenarios.

```
# Load the harmonised datasets saved in Chapter 06
for_longterm <- readRDS(
  "assets/data/data converted/for_2014_2026_harmonised_shares.rds"
)

for_recent <- readRDS(
  "assets/data/data converted/for_2020_harmonised_shares.rds"
)

# Inspect structure and coverage
cat("Long-term FoR data:\n")
```

```
Long-term FoR data:
```

```
print(names(for_longterm))
```

```
[1] "year"             "for_visual_group" "allocation_share"
```

```
print(range(for_longterm$year))
```

```
[1] 2014 2026
```

```
cat("Rows:", nrow(for_longterm), "\n\n")
```

```
Rows: 99
```

```
cat("Recent-period FoR data:\n")
```

```
Recent-period FoR data:
```

```
print(names(for_recent))
```

```
[1] "year"             "for_2020_visual"  "allocation_share"
```

```
print(range(for_recent$year))
```

```
[1] 2021 2026
```

```
cat("Rows:", nrow(for_recent), "\n")
```

```
Rows: 48
```

These datasets preserve the research-field grouping established in Chapter 06. No additional classification mapping is applied during model evaluation.

10.4.3 Historical FoR Allocation-Share Stability

Historical FoR allocation shares are examined over 2021 - 2026 to assess changes in the distribution of HPC resources across research fields.

The analysis uses the harmonised analytical groups established in Chapter 06. These groups provide a consistent reporting structure for comparison but should not be interpreted as perfectly equivalent official FoR classifications.

For each research group, the evaluation calculates:

- Average allocation share over 2021 - 2026.
- Latest allocation share in 2026.
- Minimum and maximum annual allocation shares.
- Difference between the 2026 share and the historical average.

These descriptive measures assess the stability of historical resource allocation composition. They do not directly measure FoR-specific researcher demand.

```
# Validate long-term allocation shares
for_longterm_totals <- for_longterm |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    total_share = sum(allocation_share),
    .groups = "drop"
  )

# Validate recent-period allocation shares
for_recent_totals <- for_recent |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    total_share = sum(allocation_share),
    .groups = "drop"
  )

knitr::kable(
  for_longterm_totals,
  digits = 4,
  caption = "Annual FoR allocation-share totals, 2014-2026"
)
```

Table 10.7: Annual FoR allocation-share totals, 2014–2026

year	total_share
2014	1
2015	1
2016	1
2017	1
2018	1
2019	1
2020	1
2021	1
2022	1
2023	1
2024	1
2025	1
2026	1

```
knitr::kable(
  for_recent_totals,
  digits = 4,
  caption = "Annual FoR allocation-share totals, 2021-2026"
)
```

Table 10.8: Annual FoR allocation-share totals, 2021–2026

year	total_share
2021	1
2022	1
2023	1
2024	1
2025	1
2026	1

The difference between the latest-year allocation share and the recent-period average is expressed in percentage points.

Large differences may indicate changing research-field composition, reporting effects, or variation in allocation practices. They should not automatically be interpreted as changes in underlying researcher demand.

```
# Summarise recent FoR allocation-share stability (2021-2026)
for_share_stability <- for_recent |>
  dplyr::group_by(for_2020_visual) |>
  dplyr::summarise(
```

```
    years_observed = dplyr::n_distinct(year),
    average_share = mean(allocation_share),
    latest_share = allocation_share[year == 2026][1],
    minimum_share = min(allocation_share),
    maximum_share = max(allocation_share),
    .groups = "drop"
  ) |>
  dplyr::mutate(
    change_percentage_points =
      (latest_share - average_share) * 100
  ) |>
  dplyr::arrange(dplyr::desc(average_share))

knitr::kable(
  for_share_stability |>
    dplyr::mutate(
      dplyr::across(
        c(average_share, latest_share, minimum_share, maximum_share),
        ~ .x * 100
      )
    ),
  digits = 2,
  col.names = c(
    "FoR Group",
    "Years",
    "Average (%)",
    "2026 (%)",
    "Minimum (%)",
    "Maximum (%)",
    "2026 vs Average (pp)"
  ),
  caption = "FoR allocation-share stability, 2021-2026"
)
```

Table 10.9: FoR allocation-share stability, 2021–2026

FoR Group	Years	Average (%)	2026 (%)	Minimum (%)	Maximum (%)	2026 vs Average (pp)
Engineering & Technology	6	34.88	34.54	24.73	40.65	-0.34
Physical Sciences	6	22.13	23.23	14.51	29.83	1.10
Chemical Sciences	6	17.84	17.38	16.31	19.99	-0.46
Earth & Environmental Sciences	6	14.00	12.35	12.35	16.67	-1.65

FoR Group	Years	Average (%)	2026 (%)	Minimum (%)	Maximum (%)	2026 vs Average (pp)
Other / Unclassified	2	7.07	NA	0.12	14.02	NA
Biological & Agricultural Sciences	6	5.17	7.45	2.80	7.45	2.29
Mathematics & Computing	6	2.23	4.33	0.83	4.33	2.09
Medical & Health Sciences	6	1.18	0.71	0.19	3.97	-0.46
Social Sciences, Humanities & Built Environment	4	0.31	NA	0.04	0.58	NA

Engineering & Technology accounted for the largest average allocation share during 2021–2026 (34.88%), followed by Physical Sciences (22.13%) and Chemical Sciences (17.84%). Compared with their recent-period averages, the 2026 shares were higher for Biological & Agricultural Sciences (+2.29 percentage points) and Mathematics & Computing (+2.09 percentage points). Two categories were not represented in the 2026 summary and therefore have no calculated latest-year difference. These results describe changes in allocation composition rather than directly observed FoR-specific researcher demand.

```
for_recent |>
  dplyr::group_by(for_2020_visual) |>
  dplyr::summarise(
    years_present = paste(sort(unique(year)), collapse = ", "),
    years_observed = dplyr::n_distinct(year),
    .groups = "drop"
  ) |>
  knitr::kable(
    caption = "FoR category coverage, 2021-2026"
  )
```

Table 10.10: FoR category coverage, 2021–2026

for_2020_visual	years_present	years_observed
Biological & Agricultural Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Chemical Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Earth & Environmental Sciences	2021, 2022, 2023, 2024, 2025, 2026	6

for_2020_visual	years_present	years_observed
Engineering & Technology	2021, 2022, 2023, 2024, 2025, 2026	6
Mathematics & Computing	2021, 2022, 2023, 2024, 2025, 2026	6
Medical & Health Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Other / Unclassified	2021, 2024	2
Physical Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Social Sciences, Humanities & Built Environment	2022, 2023, 2024, 2025	4

Seven harmonised FoR groups have allocation-share observations for all six years from 2021 to 2026. Two additional groups have incomplete coverage: Other / Unclassified appears in 2021 and 2024, while Social Sciences, Humanities & Built Environment appears from 2022 to 2025.

Consequently, the latest-year comparisons are restricted to groups with observed 2026 allocation shares. The incomplete categories are retained for transparency, but their absence in particular years is not automatically interpreted as zero allocation or zero researcher demand.

```
print(for_longterm_totals, n = Inf)
```

```
# A tibble: 13 x 2
    year total_share
   <dbl>       <dbl>
 1  2014       1
 2  2015       1
 3  2016       1
 4  2017       1
 5  2018       1
 6  2019       1
 7  2020       1
 8  2021       1
 9  2022       1.00
10  2023       1.000
11  2024       1
12  2025       1.000
13  2026       1.000
```

10.4.4 Long-Term FoR Allocation-Share Variation, 2014–2026

The harmonised historical dataset is examined to assess variation in research-field allocation composition across 2014–2026.

Unlike the recent-period evaluation, this analysis covers both FoR reporting standards. The harmonised analytical groups support broad historical comparison, but classification and reporting-coverage differences remain important limitations.

```
# Summarise long-term allocation-share variation
for_longterm_stability <- for_longterm |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    years_observed = dplyr::n_distinct(year),
    average_share = mean(allocation_share),
    minimum_share = min(allocation_share),
    maximum_share = max(allocation_share),
    share_range_pp = (maximum_share - minimum_share) * 100,
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share))

knitr::kable(
  for_longterm_stability |>
    dplyr::mutate(
      dplyr::across(
        c(average_share, minimum_share, maximum_share),
        ~ .x * 100
      )
    ),
  digits = 2,
  col.names = c(
    "FoR Group",
    "Years",
    "Average (%)",
    "Minimum (%)",
    "Maximum (%)",
    "Range (pp)"
  ),
  caption = "Long-term FoR allocation-share variation, 2014-2026"
)
```

Table 10.11: Long-term FoR allocation-share variation, 2014–2026

FoR Group	Years	Average (%)	Minimum (%)	Maximum (%)	Range (pp)
Engineering & Technology	13	25.20	11.00	40.65	29.65
Earth & Environmental Sciences	13	23.49	12.35	42.00	29.65
Physical Sciences	13	19.33	3.20	29.83	26.63
Chemical Sciences	13	18.48	16.31	21.00	4.69
Unclassified & Operational	7	7.21	0.12	18.63	18.51
Biological & Agricultural Sciences	13	5.01	2.80	7.82	5.02
Mathematics & Computing	12	3.46	0.08	23.50	23.42
Medical & Health Sciences	10	1.69	0.19	3.97	3.78
Social Sciences, Humanities & Built Environment	5	0.31	0.04	0.58	0.54

The range is measured in percentage points and describes the difference between the highest and lowest observed annual allocation shares within each harmonised group.

These measures provide descriptive historical context rather than a direct assessment of FoR-specific researcher demand.

The long-term evaluation shows substantial variation in allocation composition across several harmonised research-field groups.

Engineering & Technology recorded the largest average allocation share (25.20%), followed by Earth & Environmental Sciences (23.49%), Physical Sciences (19.33%), and Chemical Sciences (18.48%).

Engineering & Technology and Earth & Environmental Sciences each exhibited a range of 29.65 percentage points, while Mathematics & Computing varied by 23.42 percentage points. In contrast, Chemical Sciences showed a narrower range of 4.69 percentage points.

These historical differences should be interpreted cautiously. The series spans changes in FoR classification standards and reporting coverage, while several categories have incomplete annual observations. Consequently, the full historical period provides contextual evidence of allocation-composition variation rather than a directly comparable basis for extrapolating FoR-specific demand.

The recent FoR 2020 reporting period (2021–2026) therefore provides the more relevant basis for evaluating allocation-share assumptions for subsequent scenario development.

10.4.5 Earlier FoR Allocation-Share Variation, 2014–2020

The earlier reporting period is evaluated separately to examine historical HPC allocation composition under the FoR 2008 classification framework.

The analysis uses the harmonised analytical groups developed in Chapter 06. For each group, the average, minimum, maximum, and range of annual allocation shares are calculated over 2014–2020. These measures provide descriptive evidence of allocation-share variation within the earlier reporting period.

```
for_earlier <- for_longterm |>
  dplyr::filter(dplyr::between(year, 2014, 2020))

for_earlier_stability <- for_earlier |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    years_observed = dplyr::n_distinct(year),
    average_share = mean(allocation_share),
    minimum_share = min(allocation_share),
    maximum_share = max(allocation_share),
    share_range_pp = (maximum_share - minimum_share) * 100,
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(average_share))

knitr::kable(
  for_earlier_stability |>
    dplyr::mutate(
      dplyr::across(
        c(average_share, minimum_share, maximum_share),
        ~ .x * 100
      )
    ),
  digits = 2,
  col.names = c(
    "FoR Group",
    "Years",
    "Average (%)",
    "Minimum (%)",
    "Maximum (%)",
    "Range (pp)"
  ),
  caption = "FoR allocation-share variation, 2014-2020"
)
```

Table 10.12: FoR allocation-share variation, 2014–2020

FoR Group	Years	Average (%)	Minimum (%)	Maximum (%)	Range (pp)
Earth & Environmental Sciences	7	31.62	18.09	42.00	23.91
Chemical Sciences	7	19.02	17.00	21.00	4.00
Physical Sciences	7	16.92	3.20	24.77	21.57
Engineering & Technology	7	16.90	11.00	31.50	20.50
Unclassified & Operational	5	7.26	0.49	18.63	18.14
Biological & Agricultural Sciences	7	4.88	3.90	7.82	3.92
Mathematics & Computing	6	4.69	0.08	23.50	23.42
Medical & Health Sciences	4	2.47	0.48	3.51	3.03
Social Sciences, Humanities & Built Environment	1	0.31	0.31	0.31	0.00

The analysis retains groups with incomplete annual coverage rather than assuming that missing records represent zero allocation. Historical variation may also reflect differences in reporting scope and allocation practices.

The earlier reporting period (2014–2020) was characterised by the largest average allocation share in Earth & Environmental Sciences (31.62%), followed by Chemical Sciences (19.02%), Physical Sciences (16.92%), and Engineering & Technology (16.90%).

Allocation-share variation differed across research fields. Earth & Environmental Sciences exhibited a range of 23.91 percentage points, while Physical Sciences and Engineering & Technology varied by 21.57 and 20.50 percentage points, respectively. Chemical Sciences showed a narrower range of 4.00 percentage points across all seven years.

Several groups had incomplete annual coverage, including Mathematics & Computing (six years), Medical & Health Sciences (four years), and Social Sciences, Humanities & Built Environment (one year). Consequently, ranges calculated from limited observations should not be interpreted as evidence of stable allocation composition.

These results describe historical resource allocation patterns under the earlier FoR reporting framework, rather than directly observed research-field demand.

10.4.6 Comparison of FoR Reporting Periods, 2014–2020 and 2021–2026

The earlier and recent reporting periods are compared using the harmonised analytical groups established in Chapter 06.

The comparison examines differences in average annual allocation shares between 2014–2020 and 2021–2026. Using a common analytical grouping supports broad longitudinal comparison, although it does not eliminate differences between FoR classification standards or reporting coverage.

```
# Calculate average allocation shares within each reporting period
for_period_summary <- for_longterm |>
  dplyr::mutate(
    reporting_period = dplyr::if_else(
      year <= 2020,
      "2014-2020",
      "2021-2026"
    )
  ) |>
  dplyr::group_by(for_visual_group, reporting_period) |>
  dplyr::summarise(
    years_observed = dplyr::n_distinct(year),
    average_share = mean(allocation_share),
    .groups = "drop"
  )

# Compare period averages for each harmonised FoR group
for_period_comparison <- for_period_summary |>
  tidyr::pivot_wider(
    names_from = reporting_period,
    values_from = c(average_share, years_observed),
    names_glue = "{.value}_{reporting_period}"
  ) |>
  dplyr::mutate(
    change_pp =
      (`average_share_2021-2026` -
       `average_share_2014-2020`) * 100
  ) |>
  dplyr::arrange(dplyr::desc(
    `average_share_2021-2026`
  ))

# Display percentage shares and percentage-point differences
knitr::kable(
  for_period_comparison |>
    dplyr::transmute(
      `FoR Group` = for_visual_group,
      `Earlier Years` = `years_observed_2014-2020`,
      `Recent Years` = `years_observed_2021-2026`,
      `2014-2020 Average (%)` =
        `average_share_2014-2020` * 100,
```

```
      `2021-2026 Average (%)` =
        `average_share_2021-2026` * 100,
      `Difference (pp)` = change_pp
    ),
  digits = 2,
  caption = "Comparison of harmonised FoR allocation shares across reporting periods"
)
```

Table 10.13: Comparison of harmonised FoR allocation shares across reporting periods

FoR Group	Earlier Years	Recent Years	2014–2020 Average (%)	2021–2026 Average (%)	Difference (pp)
Engineering & Technology	7	6	16.90	34.88	17.98
Physical Sciences	7	6	16.92	22.13	5.21
Chemical Sciences	7	6	19.02	17.84	-1.18
Earth & Environmental Sciences	7	6	31.62	14.00	-17.61
Unclassified & Operational	5	2	7.26	7.07	-0.19
Biological & Agricultural Sciences	7	6	4.88	5.17	0.29
Mathematics & Computing	6	6	4.69	2.23	-2.46
Medical & Health Sciences	4	6	2.47	1.18	-1.30
Social Sciences, Humanities & Built Environment	1	4	0.31	0.31	0.00

A positive percentage-point difference indicates a higher average recorded allocation share in 2021–2026 than in 2014–2020; a negative difference indicates a lower average share. These differences should not automatically be interpreted as changes in researcher demand, particularly where classification, category coverage, or reporting scope differs between periods.

The comparison indicates substantial differences in recorded allocation composition between the two FoR reporting periods.

Engineering & Technology increased from an average allocation share of 16.90% in 2014–2020 to 34.88% in 2021–2026, a difference of +17.98 percentage points. In contrast, Earth & Environmental Sciences decreased from 31.62% to 14.00%, a difference of −17.61 percentage points. Physical Sciences also recorded a higher

average share in the recent period (+5.21 percentage points), while Mathematics & Computing recorded a lower average (−2.46 percentage points).

These differences should not be interpreted as directly observed changes in research-field demand. The comparison spans different FoR classification standards and potentially different reporting scopes. Although the analytical grouping supports broad longitudinal comparison, it does not establish complete equivalence between research fields across periods.

Furthermore, period averages are calculated across available annual observations, and some categories have incomplete coverage. For these groups, differences may partly reflect the number and timing of recorded observations rather than sustained changes in allocation composition.

Overall, the results demonstrate the importance of evaluating the two reporting periods separately. The 2014–2020 data provide historical context, while the more recent 2021–2026 allocation composition provides a more relevant reference for assessing potential FoR share assumptions for the 2030 scenarios.

10.4.7 FoR Evaluation Findings and Limitations

The Field of Research (FoR) evaluation provides evidence on the historical composition and variability of HPC resource allocations across research fields. Unlike the overall HPC, CPU, and GPU evaluations, which assess forecasting performance against historical resource measures, the FoR evaluation is descriptive. Its purpose is to assess the stability and comparability of allocation shares rather than to establish predictive accuracy for research-field demand.

Historical allocation-share patterns

The 2021 - 2026 evaluation shows that Engineering & Technology accounted for the largest average allocation share (34.88%), followed by Physical Sciences (22.13%), Chemical Sciences (17.84%), and Earth & Environmental Sciences (14.00%). These categories were observed in all six years, although their shares varied over time.

The longer 2014 - 2026 evaluation indicates that allocation composition was not constant across the historical period. Engineering & Technology and Earth & Environmental Sciences each exhibited a range of 29.65 percentage points, while Physical Sciences exhibited a range of 26.63 percentage points. Chemical Sciences was comparatively stable, with a range of 4.69 percentage points.

The separate 2014 - 2020 evaluation provides additional context. Earth & Environmental Sciences accounted for the largest average allocation share (31.62%), followed by Chemical Sciences (19.02%), Physical Sciences (16.92%), and Engineering & Technology (16.90%). Several categories exhibited substantial variation during this earlier period.

Differences between reporting periods

Comparing the two periods using harmonised analytical categories reveals substantial differences in average allocation composition. Engineering & Technology increased from 16.90% in 2014 - 2020 to 34.88% in 2021 - 2026, a difference of +17.98 percentage points. Conversely, Earth & Environmental Sciences decreased from 31.62% to 14.00%, a difference of -17.61 percentage points. Physical Sciences recorded a higher average share (+5.21 percentage points), while Mathematics & Computing recorded a lower average (-2.46 percentage points).

These differences indicate changes in the recorded distribution of HPC allocations across research fields. They do not, however, establish equivalent changes in underlying researcher demand.

Data and interpretation limitations

Several limitations affect the interpretation of these findings:

1. **Allocation versus demand:** The FoR dataset represents allocated HPC resources, not directly observed research-field requests or unconstrained computational needs. Allocation shares may therefore reflect resource availability, allocation decisions, and reporting practices as well as researcher activity.

2. **Classification comparability:** The historical analysis spans FoR 2008 and FoR 2020 reporting frameworks. Although the categories were harmonised for analytical purposes in Chapter 06, they are not necessarily equivalent to official classifications across all years.

3. **Incomplete category coverage:** Not all research-field groups are observed in every year. Missing observations were not automatically interpreted as zero allocations. Consequently, averages and ranges for categories with limited coverage are less informative than those for consistently reported categories.

4. **Reporting scope:** Historical differences in facility coverage, resource accounting, and reporting conventions may influence the recorded allocation shares. Apparent changes in research-field composition cannot be attributed exclusively to changes in research priorities.

5. **Period-average interpretation:** The comparison uses unweighted averages of available annual allocation shares. These describe the typical recorded composition within each period but do not represent shares of the total HPC resources allocated across all years combined.

6. **Limited predictive validation:** The available FoR allocation-share series does not provide a consistent, directly observed research-field demand target suitable for evaluating field-specific demand forecasts using the same rolling-origin procedure applied to the overall HPC, CPU, and GPU series.

Implications for model selection and scenarios

The FoR evaluation supports the use of historical allocation composition as a

reference for developing research-field scenarios, but not as a direct forecast of field-specific researcher demand.

The more recent 2021 - 2026 shares provide a relevant reference for a baseline allocation-composition assumption, subject to their observed variability and reporting limitations. The earlier period and the full historical series provide additional context for examining alternative compositions and assessing sensitivity to changes in research-field allocation patterns.

Accordingly, Chapter 09 will distinguish models selected through quantitative forecast-error evaluation from FoR assumptions supported by descriptive historical evidence. Any subsequent application of FoR shares to projected overall HPC demand will be presented as a scenario-based allocation of projected demand, rather than an independently validated forecast of demand within each research field.

Chapter 11

Model Selection

11.1 Purpose and Approach

This chapter selects the forecasting models and analytical assumptions that will inform projections of Australian national research high-performance computing (HPC) demand to 2030.

The selection process builds on the candidate models developed in Chapter 07 and the historical forecast evaluation conducted in Chapter 08. Overall HPC demand, CPU demand, and GPU demand are considered separately because they have different measurement units, historical coverage, and forecasting uncertainties.

Model selection is guided by predictive performance, statistical adequacy, interpretability, and the reliability of the underlying data. Particular attention is given to the substantial increase in recorded overall HPC demand in 2026 and whether this observation should influence the central forecast or be examined through sensitivity scenarios.

Field of Research (FoR) allocation shares are assessed separately. These shares describe the historical distribution of allocated resources rather than directly observed research-field demand. They are therefore used to inform scenario assumptions, not to select independently validated field-specific demand models.

11.2 Model Selection Criteria

The following criteria are used to assess the candidate forecasting models.

1. **Out-of-sample forecast accuracy:** Rolling-origin cross-validation results from Chapter 08 provide the principal quantitative evidence. Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) are evaluated

within each resource series, rather than compared directly across different measurement units.

2. **Consistency across evaluation metrics:** Models are assessed using both RMSE and MAE. RMSE places greater weight on large forecast errors, while MAE describes the average absolute forecast error. Differences between their rankings are considered when interpreting model performance.
3. **Statistical adequacy:** Residual behaviour and evidence of systematic forecast errors are considered where supported by the available data. Diagnostic results are interpreted cautiously because the annual series contain relatively few observations.
4. **Interpretability and plausibility:** Selected models should provide transparent assumptions about future demand growth and produce forecasts that can be explained in the context of national research infrastructure planning.
5. **Historical data quality and coverage:** Model complexity is assessed relative to the length, consistency, and provenance of each resource series. Short or partly estimated historical series do not necessarily support complex forecasting models.
6. **Sensitivity to recent observations:** The influence of the 2026 demand increase is considered when interpreting fitted trends and selecting models. The increase is not automatically assumed to represent a permanent acceleration in demand.
7. **Suitability for projections to 2030:** The selected models must support interpretable forecasts over the 2027–2030 horizon. Forecast uncertainty and alternative growth assumptions will be examined in Chapter 10.

No single criterion determines model selection automatically. Out-of-sample accuracy provides the primary quantitative comparison, while the remaining criteria support a transparent assessment of model suitability for policy-oriented forecasting.

11.3 Overall HPC Demand Model Selection

11.3.1 Comparison of Candidate Models

Four candidate models were evaluated for overall HPC demand: linear trend, log-linear trend, drift, and naïve forecasting. Their predictive performance was assessed in Chapter 08 using expanding-window, one-step-ahead rolling-origin cross-validation.

The initial training period covered 2014 - 2020, followed by six evaluation years from 2021 to 2026. Each model was re-estimated using the observations available

before the corresponding forecast year, avoiding the use of future observations in model fitting.

Model	RMSE (kSU)	MAE (kSU)	RMSE Rank
Linear trend	290,898	198,295	1
Drift	314,588	221,778	2
Naïve	378,864	283,000	3
Log-linear trend	511,614	466,991	4

The linear trend model achieved the lowest RMSE and MAE, indicating the strongest historical one-step-ahead predictive performance among the evaluated candidates. Its RMSE was approximately 7.5% lower than the drift model and 23.2% lower than the naïve benchmark.

The agreement between RMSE and MAE rankings provides additional support for the linear trend model. In contrast, the log-linear trend model produced substantially larger forecast errors, suggesting that its proportional-growth assumption was less effective over the historical evaluation period.

These findings support the linear trend model as the leading candidate for the central overall HPC demand forecast. However, one-step-ahead performance does not guarantee equivalent accuracy over the longer 2027 - 2030 forecasting horizon. The final selection must also consider the sensitivity of fitted trends to recent observations and the interpretation of historical demand measurements.

11.3.2 Sensitivity to the 2026 Observation

The 2026 observation warrants particular attention when selecting a model for overall HPC demand. Recorded demand increased from 1,564,643 kSU in 2025 to 2,242,817 kSU in 2026, representing an increase of approximately 43.3%.

Over the same period, recorded allocations decreased slightly from 876,200 kSU to 872,288 kSU. Consequently, the ratio of allocated resources to recorded demand declined from approximately 56.0% to 38.9%. This divergence highlights the importance of distinguishing requested HPC resources from the resources ultimately allocated.

The increase in recorded demand may reflect changes in computational requirements, the composition of applications, or other factors affecting the reporting of requested resources. However, the available analysis does not establish the relative contribution of these explanations. The increase should therefore not be attributed to a specific technological or institutional development without supporting evidence.

11.3.2.1 Implications for Trend Estimation

The candidate models developed in Chapter 07 include trend specifications estimated using the available historical observations, alongside sensitivity analysis examining the influence of the latest observation.

Including 2026 allows the fitted models to incorporate the most recent reported demand. However, a substantial increase near the end of a short annual series can influence estimated trend coefficients and subsequent forecasts.

Excluding 2026 provides a sensitivity comparison showing how projected demand changes when the latest observation does not influence model estimation. This comparison is not intended to imply that the 2026 observation is incorrect or should be discarded.

The two estimation periods therefore serve different analytical purposes:

- **Including 2026:** Represents a model fitted to the complete available historical series.
- **Excluding 2026:** Provides a sensitivity benchmark for assessing the influence of the latest demand increase.

The preferred model should not be determined solely by whether its forecast is higher or lower. Instead, the comparison should consider historical predictive performance, the plausibility of extrapolated growth, and the uncertainty associated with the most recent observation.

11.3.2.2 Interpretation for Model Selection

The rolling-origin evaluation in Chapter 08 identified the linear trend as the strongest candidate for overall HPC demand based on both RMSE and MAE.

Nevertheless, this evaluation primarily assesses one-step-ahead forecast performance. It does not independently establish the reliability of forecasts over the full 2027–2030 horizon, particularly following a substantial increase in the final observed year.

Accordingly, the linear trend remains the leading candidate for the central overall HPC demand forecast, while the sensitivity of projections to the inclusion of 2026 must be considered before finalising the forecasting assumptions.

The resulting differences in projected demand will inform the uncertainty and alternative-growth scenarios developed in Chapter 10.

11.3.3 Overall HPC Model Selection Decision

Based on the rolling-origin evaluation and sensitivity analysis, the linear trend model fitted to the complete 2014–2026 historical series is selected as the central statistical forecasting specification for overall HPC demand.

This selection is supported by three considerations.

First, the linear trend achieved the lowest out-of-sample RMSE (290,898 kSU) and MAE (198,295 kSU) among the four candidate models evaluated in Chapter 08. It therefore provides the strongest historical one-step-ahead predictive performance within the tested model set.

Second, the linear specification assumes a constant absolute annual increase rather than sustained compounded percentage growth. This produces a more conservative and interpretable extrapolation than the log-linear model, which generated substantially larger forecast errors and much higher long-horizon projections.

Third, the model remains sensitive to the inclusion of the latest observation, but this sensitivity can be quantified and communicated transparently.

The full-history linear model estimates an annual increase of approximately 161,992 kSU and projects overall HPC demand of approximately 2.50 million kSU in 2030.

When the model is fitted using observations from 2014–2025, the estimated annual increase decreases to approximately 144,357 kSU, producing a 2030 projection of approximately 2.29 million kSU.

The difference between these projections demonstrates the influence of the 2026 observation on the estimated trend. It does not represent a formal prediction interval or establish a statistically validated lower bound for future demand.

The following specifications are therefore retained for subsequent analysis:

- **Central statistical projection:** Linear trend fitted to the complete 2014–2026 historical series.
- **Historical sensitivity comparison:** Linear trend fitted to 2014–2025, excluding the latest observation.
- **Alternative growth assumption:** Log-linear trend, retained as an exploratory illustration of compounded growth rather than a preferred forecasting model.

The central projection should be interpreted as a model-based extrapolation of recorded and harmonised HPC demand, not a definitive estimate of unconstrained future researcher requirements.

The selected model and sensitivity comparisons will inform the development of explicit demand scenarios in Chapter 10. These scenarios will distinguish modelling assumptions from statistical forecast uncertainty and acknowledge limitations in historical data comparability.

11.4 CPU Demand Model Selection

11.4.1 Comparison of Candidate Models

CPU demand was evaluated separately from overall HPC demand to account for differences in resource measurement, historical coverage, and computational requirements.

Four candidate forecasting models were compared: linear trend, log-linear trend, drift, and naïve forecasting.

As described in Chapter 08, model performance was assessed using expanding-window, one-step-ahead rolling-origin cross-validation, with an initial training period of 2014–2020 and six evaluation years from 2021 to 2026.

CPU demand is expressed in millions of core-hours.

Model	RMSE (million core-hours)	MAE (million core-hours)	RMSE Rank
Linear trend	157.72	138.65	1
Log-linear trend	160.18	133.74	2
Drift	242.61	193.69	3
Naïve	260.05	228.45	4

The linear trend model achieved the lowest RMSE of 157.72 million core-hours, approximately 1.5% lower than the log-linear trend model.

However, the log-linear model achieved the lowest MAE of 133.74 million core-hours, compared with 138.65 million core-hours for the linear model.

Unlike the overall HPC evaluation, the two accuracy measures therefore produce different rankings for the leading CPU models.

This distinction is important because RMSE places greater weight on larger forecast errors, whereas MAE measures the average magnitude of errors without additional quadratic weighting.

The relatively small difference in RMSE suggests that the linear and log-linear models have comparable historical one-step-ahead forecasting performance. Consequently, selecting the preferred CPU model requires consideration of the implied growth assumptions and their consequences for longer-term projections.

11.4.2 Assessment of Growth Assumptions

The linear trend model assumes a constant absolute annual increase in CPU demand, while the log-linear trend model assumes approximately constant proportional growth.

Chapter 07 demonstrated that these assumptions generate substantially different projections for 2030.

The linear specifications project CPU demand between approximately 1,263.85 and 1,369.90 million core-hours in 2030, depending on whether the 2026 observation is included.

By comparison, the log-linear specifications project approximately 2,479.16 to 2,502.07 million core-hours.

The considerably higher log-linear projections arise from the compounding of proportional growth over the forecast horizon.

Although the log-linear model performs slightly better on MAE, the historical evaluation does not provide sufficiently strong evidence that compounded growth will persist through 2030.

The linear trend is therefore more suitable as the leading candidate for a transparent central statistical projection, while the log-linear specification remains useful for examining alternative growth assumptions.

These projections are model-based trajectories rather than statistical prediction intervals or independently validated demand scenarios.

11.4.3 CPU Model Selection Decision

The linear trend model fitted to the complete 2014–2026 historical series is selected as the central statistical forecasting specification for CPU demand.

This decision reflects the combined assessment of historical forecast accuracy, model interpretability, and the implications of extrapolating alternative growth assumptions through 2030.

The linear model achieved the lowest rolling-origin RMSE of 157.72 million core-hours, compared with 160.18 million core-hours for the log-linear model. However, the log-linear model achieved a slightly lower MAE of 133.74 million core-hours, compared with 138.65 million core-hours for the linear model.

These differences are relatively small and do not establish a statistically significant advantage for either model. Both specifications therefore remain plausible candidates based on their historical one-step-ahead forecasting performance.

The principal distinction lies in their growth assumptions. The linear model extrapolates an approximately constant absolute increase in CPU demand, whereas the log-linear model extrapolates compounded proportional growth.

Chapter 07 showed that these assumptions produce substantially different 2030 projections. The linear specifications generate approximately 1.26–1.37 billion core-hours, while the log-linear specifications generate approximately 2.48–2.50 billion core-hours.

Given the short annual series, differences in historical measurement and conversion assumptions, and uncertainty regarding the persistence of recent demand growth, the evidence is insufficient to establish compounded growth as the most appropriate central forecasting assumption.

The linear model is therefore preferred for the central statistical projection because it provides competitive historical forecast accuracy while retaining a simpler and more transparent growth assumption.

The following specifications are retained:

- **Central statistical projection:** Linear trend fitted to the complete 2014–2026 historical series.
- **Historical sensitivity comparison:** Linear trend fitted to 2014–2025, excluding the latest observation.
- **Alternative growth assumption:** Log-linear trend, retained to examine the implications of sustained proportional growth.

The 2026-excluded linear specification provides a sensitivity comparison rather than an independently validated lower bound. Similarly, the log-linear specification represents an alternative growth trajectory, not a statistically established upper bound.

The selected CPU model should be interpreted as a projection of the harmonised historical CPU demand measure, subject to the provenance, conversion, and reporting limitations documented in earlier chapters.

Chapter 10 will use these specifications to examine alternative CPU demand trajectories through 2030 without treating differences between models as formal prediction intervals.

11.5 GPU Demand Model Selection

11.5.1 Comparison of Candidate Models

GPU demand was evaluated separately from overall HPC and CPU demand because of its shorter historical coverage, differences in data provenance, and greater uncertainty surrounding the measurement of computational requirements.

The GPU forecasting models developed in Chapter 07 use five annual demand estimates covering 2022–2026. These estimates combine reported Setonix-based demand information with values derived using allocation shares and working conversion assumptions.

Four candidate models were assessed in Chapter 08: naïve forecasting, drift, linear trend, and log-linear trend.

An expanding-window, one-step-ahead rolling-origin evaluation was conducted using the available recent-period observations. Unlike the overall HPC and CPU evaluations, the GPU comparison contains only two test years, 2025 and 2026.

Model	RMSE (million GPU-hours)	MAE (million GPU-hours)	RMSE Rank
Naïve	0.441	0.353	1
Drift	0.537	0.490	2
Log-linear trend	0.719	0.715	3
Linear trend	0.729	0.725	4

The naïve model achieved the lowest RMSE and MAE among the evaluated candidates. It therefore provided the strongest historical one-step-ahead predictive performance within this limited evaluation.

However, the comparison is based on only two forecast errors per model. Consequently, the resulting accuracy rankings may be strongly influenced by individual observations and should not be interpreted as robust evidence of long-term predictive superiority.

The naïve model assumes that future GPU demand remains at the most recently observed level. Although this assumption performed comparatively well during the evaluation period, it does not establish that GPU demand will remain constant through 2030.

Similarly, the higher forecast errors recorded by the linear and log-linear models do not establish that future GPU demand growth is absent. They indicate that these trend specifications were less accurate over the two available historical test years.

The results therefore support the naïve model as a short-term statistical benchmark, rather than a definitive long-term forecasting model.

11.5.2 Sensitivity to the 2026 Observation

Chapter 07 examined the sensitivity of GPU trend estimates to the inclusion of the 2026 observation.

For the linear trend specification, the estimated annual increase changed from approximately 0.012 million GPU-hours when fitted to 2022–2025 to approximately 0.172 million GPU-hours when 2026 was included.

The log-linear trend coefficient also increased, from approximately 0.020 to 0.081.

These changes demonstrate that the estimated GPU growth trajectory is highly sensitive to the latest observation. However, the large relative percentage changes partly reflect the small trend coefficients estimated from the shorter 2022–2025 sample.

The resulting 2030 projections also differ substantially between specifications.

The linear trend models produce projected GPU demand ranging from approximately 2.252 to 3.371 million GPU-hours, while the log-linear models produce

projections ranging from approximately 2.406 to 3.681 million GPU-hours.

These values represent alternative model-based trajectories, not formal prediction intervals or statistically validated lower and upper demand bounds.

The short modelling period and mixture of reported and estimated demand values limit confidence in extrapolating any single trend through 2030.

Furthermore, the historical GPU conversion scenarios developed in Chapter 05 represent uncertainty associated with conversion assumptions. They should not be interpreted as additional observed GPU demand data or as independent evidence of future growth.

Consequently, the GPU model selection process must distinguish between short-term predictive performance, sensitivity to individual observations, and the broader uncertainty surrounding future GPU requirements.

11.5.3 GPU Model Selection Decision

The naïve forecasting model is retained as the preferred statistical benchmark for GPU demand because it achieved the lowest RMSE (0.441 million GPU-hours) and MAE (0.353 million GPU-hours) in the rolling-origin evaluation conducted in Chapter 08.

However, no single GPU forecasting model is selected as a sufficiently validated central projection through 2030.

This decision reflects three considerations.

First, the evaluation contains only two one-step-ahead test observations, covering 2025 and 2026. Although the naïve model performed best within this comparison, the limited number of forecast errors does not provide robust evidence that it will remain the most accurate model over a four-year forecasting horizon.

Second, the GPU modelling series contains only five annual estimates from 2022–2026. These estimates combine reported Setonix-based demand information with values derived using allocation shares and working conversion assumptions. Differences in data provenance limit the comparability of observations and the reliability of estimated growth trends.

Third, the linear and log-linear specifications are highly sensitive to the inclusion of 2026. Chapter 07 demonstrated substantial changes in their estimated trend coefficients and resulting 2030 projections when the latest observation was included or excluded.

The naïve model is therefore retained as a reference trajectory representing unchanged demand relative to the latest observation. It should not be interpreted as evidence that future GPU requirements will remain constant.

The linear and log-linear trend models are retained as exploratory growth trajectories. These specifications illustrate how projected demand changes

under alternative assumptions about absolute and proportional growth, but their historical performance and sensitivity do not justify selecting either as a definitive long-term forecast.

The following approach is adopted for subsequent scenario analysis:

- **Statistical reference benchmark:** Naïve model, representing unchanged GPU demand from the latest observation.
- **Exploratory absolute-growth trajectory:** Linear trend fitted to the available 2022–2026 estimates.
- **Exploratory proportional-growth trajectory:** Log-linear trend fitted to the available 2022–2026 estimates.
- **Historical sensitivity comparisons:** Alternative trend specifications excluding 2026, used to assess dependence on the latest observation.
- **Conversion uncertainty:** Historical GPU conversion scenarios retained separately from future growth assumptions.

These trajectories are not assigned probabilities or interpreted as statistically validated lower, central, and upper demand bounds.

Consequently, the GPU analysis provides an exploratory assessment of possible future resource requirements rather than a definitive forecast of national research GPU demand.

Chapter 10 will present these alternatives transparently, distinguishing historical measurement uncertainty from assumptions about future demand growth.

11.6 Field of Research Scenario Assumptions

11.6.1 Purpose and Analytical Approach

Field of Research (FoR) is incorporated into the forecasting framework to examine how projected national HPC demand could be distributed across research disciplines under alternative allocation-composition assumptions.

Unlike overall HPC, CPU, and GPU demand, the harmonised FoR dataset primarily represents historically allocated resources rather than directly observed requests for computational resources within individual research fields.

Consequently, FoR-specific demand cannot be independently forecast and validated using the same approach applied to the aggregate resource series.

Instead, a top-down scenario framework is adopted. Under this approach, projected overall HPC demand is distributed across harmonised FoR categories using explicit assumptions about future allocation shares.

These estimates represent hypothetical distributions of projected demand based on historical allocation patterns, rather than independently validated forecasts of researcher demand within each field.

11.6.2 Historical Evidence for Scenario Assumptions

Chapter 08 evaluated harmonised FoR allocation shares across 2014–2026, including separate assessments of the earlier 2014–2020 and recent 2021–2026 periods.

During 2021–2026, Engineering & Technology accounted for the largest average allocation share (34.88%), followed by Physical Sciences (22.13%), Chemical Sciences (17.84%), and Earth & Environmental Sciences (14.00%).

The longer historical comparison identified substantial differences in recorded allocation composition between reporting periods.

Engineering & Technology increased from an average allocation share of 16.90% in 2014–2020 to 34.88% in 2021–2026, representing a difference of +17.98 percentage points.

Conversely, Earth & Environmental Sciences decreased from 31.62% to 14.00%, a difference of -17.61 percentage points.

Physical Sciences increased by 5.21 percentage points, while Mathematics & Computing decreased by 2.46 percentage points.

These changes demonstrate that the recorded distribution of HPC allocations has varied over time. However, they cannot be interpreted exclusively as changes in underlying researcher demand because the historical series spans different FoR classification standards, reporting coverage, and allocation practices.

The observed differences therefore provide evidence for exploring alternative allocation-composition assumptions, but not for extrapolating discipline-specific demand growth rates directly.

11.6.3 Selection of FoR Scenario Assumptions

Three allocation-composition approaches are proposed for scenario development in Chapter 10.

Recent-period composition

The first approach uses the harmonised allocation-share patterns observed during 2021–2026 as a reference for a contemporary allocation-composition scenario.

This approach reflects the more recent reporting framework and provides a transparent historical basis for distributing projected overall HPC demand across research fields.

However, the observed annual shares are not constant, and some categories are missing from particular years. The implementation must therefore define how category coverage, missing observations, and share normalisation are handled.

Latest-year composition

The second approach uses the 2026 allocation-share distribution as an alternative reference.

This approach illustrates the implications of maintaining the latest observed composition throughout the projection horizon.

However, a single year's allocation shares may reflect temporary conditions or reporting-specific effects. The 2026 composition should therefore not automatically be considered more representative of future researcher demand than the recent-period average.

Historical-composition sensitivity

The third approach uses the earlier 2014–2020 allocation composition to examine how projected demand distributions change under a different historical allocation pattern.

This approach is intended as a sensitivity comparison rather than a direct prediction that earlier allocation patterns will return.

Differences in FoR classification frameworks, reporting scope, and category coverage must be considered before constructing comparable scenario shares.

11.6.4 Interpretation and Limitations

The proposed FoR scenarios are conditional on two distinct assumptions:

1. The selected overall HPC model provides an indicative trajectory of total national research HPC demand.
2. Historical FoR allocation shares provide a reasonable reference for constructing hypothetical distributions of that projected demand.

Neither assumption establishes that the resulting FoR-level estimates represent directly observed or independently validated future researcher demand.

The allocation shares must also form a complete and internally consistent distribution for each scenario. Missing categories should not automatically be treated as zero allocations, and any normalisation or category exclusions must be documented.

Furthermore, changes in allocation composition may reflect resource constraints, institutional priorities, facility coverage, or reporting conventions rather than changes in unconstrained computational requirements.

For these reasons, FoR scenarios will be presented as allocation-share-based distributions of projected overall demand, not as independent field-specific demand forecasts.

Chapter 10 will implement these assumptions and examine how alternative historical compositions affect the indicative distribution of HPC requirements through 2030.

11.7 Selected Models and Limitations

11.7.1 Summary of Model Selection Decisions

The forecasting models and analytical assumptions selected in this chapter are summarised below.

Demand measure	Selected approach	Supporting evidence	Role in Chapter 10
Overall HPC	Full-history linear trend	Lowest RMSE and MAE across six evaluation years	Central statistical projection
CPU	Full-history linear trend	Lowest RMSE; competitive MAE across six evaluation years	Central statistical projection
GPU	Naïve benchmark with exploratory trend models	Lowest naïve RMSE and MAE, but only two evaluation years	Reference benchmark and exploratory growth trajectories
FoR	Historical allocation-share assumptions	Descriptive stability and period-comparison analysis	Conditional distributions of projected overall HPC demand

The selected approaches reflect differences in data availability, predictive evidence, and the interpretation of each demand measure.

For overall HPC, the linear trend provides the strongest historical one-step-ahead forecasting performance among the evaluated models. Its full-history specification is selected as the central statistical projection, while the 2026-excluded specification is retained for sensitivity analysis.

For CPU demand, the linear trend is also selected as the central statistical projection. Although the log-linear model achieved a slightly lower MAE, the linear model achieved the lowest RMSE and avoids assuming that historical proportional growth will necessarily continue through 2030.

For GPU demand, the naïve model is retained as the preferred statistical benchmark. However, the short modelling period, limited evaluation sample, and sensitivity of trend estimates to 2026 prevent the selection of a sufficiently validated central long-term forecast. Linear and log-linear trajectories are retained for exploratory scenario analysis.

For FoR, no independent demand forecasting model is selected. Instead, historical allocation-share patterns are used to construct conditional distributions of projected overall HPC demand across research fields.

11.7.2 Forecasting Uncertainty and Data Limitations

Several limitations apply across the selected forecasting approaches.

Historical measurement and comparability

The harmonised historical series combine directly reported observations, derived estimates, and proxy-based values. Differences in resource accounting, reporting coverage, and conversion assumptions may affect comparability across years.

Consequently, the resulting forecasts should be interpreted as projections of the constructed historical demand measures rather than definitive estimates of unconstrained national researcher requirements.

Limited annual observations

The overall HPC and CPU forecasting series contain 13 annual observations covering 2014–2026.

Although rolling-origin cross-validation provides useful out-of-sample evidence, the six available evaluation years limit the precision with which differences in predictive performance can be established.

The GPU series contains only five recent annual estimates and two evaluation years, resulting in substantially greater model-selection uncertainty.

Sensitivity to recent demand changes

The substantial increase in recorded overall HPC demand in 2026 affects estimated historical trends and projected future demand.

The influence of this observation is examined using alternative specifications that exclude 2026. These comparisons assess sensitivity but do not establish whether the observed increase represents a temporary change, a sustained acceleration, or differences in reporting and measurement.

Long-horizon extrapolation

The rolling-origin evaluations primarily assess one-step-ahead predictive performance.

Forecasts extending to 2030 require additional assumptions about the persistence of historical trends. Strong performance over a one-year horizon does not guarantee equivalent performance over four years.

Resource-specific interpretation

Overall HPC demand is expressed in kSU, CPU demand in core-hours, and GPU demand in GPU-hours.

These measures describe different aspects of computing requirements and should not be directly added or compared without an explicitly justified conversion framework.

FoR allocation-share interpretation

FoR allocation shares represent the recorded distribution of allocated resources rather than directly observed field-specific demand.

The resulting scenarios depend on assumptions about how historical allocation composition may relate to the future distribution of computational requirements. They are not independently validated FoR demand forecasts.

11.7.3 Implications for Scenario Development

The model-selection findings establish the analytical basis for Chapter 10.

The full-history linear trend models will provide central statistical projections for overall HPC and CPU demand, with 2026-excluded specifications used to examine sensitivity to the latest observation.

Alternative log-linear trajectories will illustrate the implications of compounded proportional growth. These alternatives will not automatically be interpreted as statistically established upper demand bounds.

For GPU demand, the naïve benchmark and exploratory trend specifications will be presented separately, reflecting the limited evidence available for selecting a definitive long-term forecasting model.

For FoR, alternative historical allocation compositions will be applied to projected overall HPC demand to examine conditional distributions across research fields.

Scenario development will distinguish three sources of uncertainty:

1. **Historical measurement uncertainty:** Differences in data provenance, reporting scope, and conversion assumptions.
2. **Model uncertainty:** Differences between forecasting specifications and their underlying growth assumptions.
3. **Scenario uncertainty:** Alternative assumptions about future demand growth and the distribution of computational requirements.

These sources of uncertainty will not be treated as interchangeable.

The scenarios will therefore provide transparent, assumption-dependent projections for national HPC infrastructure planning rather than probabilistic estimates of all possible future outcomes.

11.8 Chapter Summary

This chapter selected forecasting approaches for overall HPC, CPU, and GPU demand using the predictive evaluation results from Chapter 08 and the modelling

evidence developed in Chapter 07.

Linear trend models were selected as the central statistical specifications for overall HPC and CPU demand because they provide competitive historical forecast accuracy and interpretable absolute-growth assumptions.

For GPU demand, the naïve model was retained as a statistical benchmark, while alternative trend models were retained for exploratory analysis because of limited historical evidence.

FoR allocation-share assumptions were selected as a framework for distributing projected overall HPC demand across research fields, subject to classification and reporting limitations.

These decisions provide the foundation for the demand scenarios developed in Chapter 10, where alternative growth assumptions, historical sensitivity, and resource-specific uncertainties will be examined through 2030.

Chapter 12

Demand Scenarios: 2027–2030

12.1 Purpose and Scope

This chapter develops indicative scenarios for Australian national research high-performance computing (HPC) demand through 2030, drawing on the forecasting models evaluated in Chapter 08 and selected in Chapter 09.

The analysis considers four dimensions:

1. Overall HPC demand, expressed in kSU.
2. CPU demand, expressed in million CPU core-hours.
3. GPU demand, expressed in million GPU-hours.
4. Field of Research (FoR) allocation-share-based distributions of projected overall HPC demand.

The projection period covers 2027–2030, with 2026 representing the latest available historical observation.

The objective is to provide transparent, reproducible estimates of potential future computational requirements that can inform infrastructure planning and policy discussions.

The scenarios are conditional on the available historical data, model specifications, and stated assumptions. They should not be interpreted as definitive predictions of unconstrained national HPC demand.

12.2 Scenario Framework

The scenario framework distinguishes central statistical projections, historical sensitivity comparisons, and exploratory growth trajectories.

Demand measure	Reference approach	Alternative assumptions
Overall HPC	Full-history linear trend	Linear trend excluding 2026; log-linear growth
CPU	Full-history linear trend	Linear trend excluding 2026; log-linear growth
GPU	Naïve statistical benchmark	Linear and log-linear exploratory trajectories; sensitivity to 2026
FoR	Recent-period allocation composition	Latest-year and earlier-period allocation compositions

For overall HPC and CPU demand, the full-history linear trend models provide the central statistical projections because they demonstrated the strongest or most competitive historical forecasting performance.

The 2026-excluded linear specifications are retained to examine the sensitivity of projected demand to the latest observation.

Log-linear models are retained to illustrate the implications of compounded proportional growth. These projections are alternative modelling assumptions rather than statistically validated upper bounds.

For GPU demand, the naïve model serves as the statistical reference benchmark. Linear and log-linear projections are presented as exploratory trajectories because the available historical series is short and the rolling-origin evaluation contains only two test years.

FoR scenarios use historical allocation compositions to distribute projected overall HPC demand across harmonised research categories. These distributions are conditional allocation-share scenarios, not independently validated FoR-specific demand forecasts.

12.3 Interpretation of Scenario Uncertainty

Three sources of uncertainty are considered throughout the chapter:

- **Measurement uncertainty:** Differences in historical reporting, resource accounting, and conversion assumptions.
- **Model uncertainty:** Differences in forecasting specifications and their underlying growth assumptions.

- **Scenario uncertainty:** Differences in assumed future growth and allocation composition.

The scenarios are not assigned probabilities, and the differences between model projections should not be interpreted as formal prediction intervals.

All numerical results will retain their original resource units. Overall HPC kSU, CPU core-hours, and GPU-hours will be analysed separately rather than combined without a justified conversion framework.

The analysis prioritises transparency, reproducibility, and clear communication of uncertainty over unnecessary model complexity.

12.4 Overall HPC Demand Scenarios

12.4.1 Forecast Construction

Overall HPC demand projections are generated for 2027–2030 using the linear and log-linear models developed in Chapter 07 and evaluated in Chapter 08.

The full-history linear trend is adopted as the central statistical projection, consistent with the model-selection decision in Chapter 09.

A linear trend excluding 2026 is retained to examine sensitivity to the latest historical observation. Full-history and 2026-excluded log-linear specifications are also presented to illustrate alternative compounded-growth assumptions.

All projections are expressed in kSU and represent model-based estimates of the constructed historical demand measure, rather than forecasts of allocated computing capacity.

12.4.2 Annual Demand Projections

```
# Reuse the fitted models from Chapter 07.
# This chunk requires those model objects to have been created
# earlier in the current Quarto rendering session.

required_models <- c(
  "overall_linear",
  "overall_linear_pre2026",
  "overall_loglinear",
  "overall_loglinear_pre2026"
)

missing_models <- required_models[
  !vapply(required_models, exists, logical(1), inherits = TRUE)
]
```

```
if (length(missing_models) > 0) {
  stop(
    "Missing Chapter 07 model objects: ",
    paste(missing_models, collapse = ", "),
    ". Fit these models from the original historical data ",
    "before generating Chapter 10 projections."
  )
}

hpc_future <- data.frame(
  year = 2027:2030,
  time = (2027:2030) - 2014
)

hpc_scenarios <- hpc_future |>
  dplyr::transmute(
    year,
    central_linear = as.numeric(
      predict(overall_linear, newdata = hpc_future)
    ),
    linear_excl_2026 = as.numeric(
      predict(overall_linear_pre2026, newdata = hpc_future)
    ),
    loglinear_full = exp(as.numeric(
      predict(overall_loglinear, newdata = hpc_future)
    )),
    loglinear_excl_2026 = exp(as.numeric(
      predict(overall_loglinear_pre2026, newdata = hpc_future)
    ))
  )

stopifnot(
  nrow(hpc_scenarios) == 4,
  all(is.finite(as.matrix(hpc_scenarios[, -1]))),
  all(as.matrix(hpc_scenarios[, -1]) > 0)
)


hpc_scenarios |>
  dplyr::mutate(
    year = as.character(year)
  ) |>
  knitr::kable(
    digits = 0,
    format.args = list(big.mark = ","),
```

```
    col.names = c(
      "Year",
      "Central linear",
      "Linear (excl. 2026)",
      "Log-linear (full)",
      "Log-linear (excl. 2026)"
    ),
    caption = "Overall HPC demand scenarios, 2027-2030 (kSU)"
  )
```

Table 12.1: Overall HPC demand scenarios, 2027–2030 (kSU)

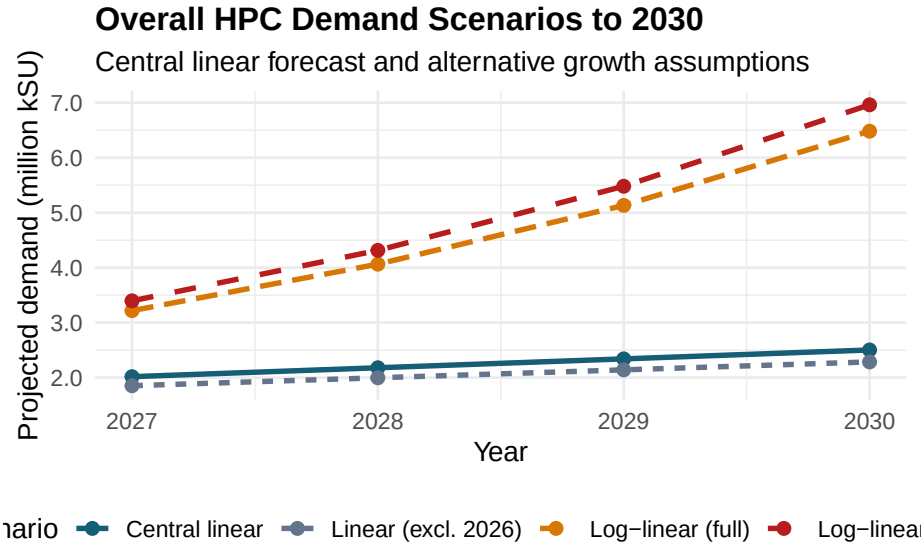
Year	Central linear	Linear (excl. 2026)	Log-linear (full)	Log-linear (excl. 2026)
2027	2,016,845	1,852,255	3,217,801	3,396,628
2028	2,178,838	1,996,612	4,064,043	4,314,830
2029	2,340,830	2,140,970	5,132,835	5,481,247
2030	2,502,822	2,285,327	6,482,707	6,962,979

```
#| label: fig-overall-hpc-scenarios
#| fig-cap: "Projected overall HPC demand under alternative model specifications, 2027-2030.
#| fig-width: 10
#| fig-height: 6

# Convert scenario data into long format
hpc_plot_data <- hpc_scenarios |>
  tidyr::pivot_longer(
    cols = -year,
    names_to = "scenario",
    values_to = "demand_ksu"
  ) |>
  dplyr::mutate(
    scenario = factor(
      scenario,
      levels = c(
        "central_linear",
        "linear_excl_2026",
        "loglinear_full",
        "loglinear_excl_2026"
      ),
      labels = c(
        "Central linear",
        "Linear (excl. 2026)",
        "Log-linear (full)",
```

```
        "Log-linear (excl. 2026)"
      )
    )
  )

ggplot2::ggplot(
  hpc_plot_data,
  ggplot2::aes(
    x = year,
    y = demand_ksu / 1e6,
    colour = scenario,
    linetype = scenario,
    group = scenario
  )
) +
  ggplot2::geom_line(linewidth = 1.1) +
  ggplot2::geom_point(size = 2.3) +
  ggplot2::scale_x_continuous(
    breaks = 2027:2030
  ) +
  ggplot2::scale_y_continuous(
    labels = scales::label_number(accuracy = 0.1)
  ) +
  ggplot2::scale_colour_manual(
    values = c(
      "Central linear" = "#155E75",
      "Linear (excl. 2026)" = "#64748B",
      "Log-linear (full)" = "#D97706",
      "Log-linear (excl. 2026)" = "#B91C1C"
    )
  ) +
  ggplot2::labs(
    title = "Overall HPC Demand Scenarios to 2030",
    subtitle = "Central linear forecast and alternative growth assumptions",
    x = "Year",
    y = "Projected demand (million kSU)",
    colour = "Scenario",
    linetype = "Scenario"
  ) +
  ggplot2::theme_minimal(base_size = 12) +
  ggplot2::theme(
    legend.position = "bottom",
    plot.title = ggplot2::element_text(face = "bold")
  )
```



12.4.3 Statistical Prediction Intervals for the Central Forecast

To complement the alternative growth scenarios, 80% and 95% model-based prediction intervals are calculated around the selected full-history linear trend.

These intervals describe uncertainty in future annual observations conditional on the fitted linear regression model and its error assumptions.

They are distinct from the sensitivity comparisons and alternative growth trajectories presented above.

```
# Calculate 80% and 95% prediction intervals
# for the selected overall HPC linear model

hpc_pi_80 <- predict(
  overall_linear,
  newdata = hpc_future,
  interval = "prediction",
  level = 0.80
)

hpc_pi_95 <- predict(
  overall_linear,
  newdata = hpc_future,
  interval = "prediction",
  level = 0.95
```

```
)

# Combine the results into one table
hpc_prediction_intervals <- data.frame(
  year = hpc_future$year,
  central_forecast = hpc_pi_95[, "fit"],
  lower_80 = hpc_pi_80[, "lwr"],
  upper_80 = hpc_pi_80[, "upr"],
  lower_95 = hpc_pi_95[, "lwr"],
  upper_95 = hpc_pi_95[, "upr"]
)

# Confirm intervals are internally consistent
stopifnot(
  all(hpc_prediction_intervals$lower_95 <=
      hpc_prediction_intervals$lower_80),
  all(hpc_prediction_intervals$lower_80 <=
      hpc_prediction_intervals$central_forecast),
  all(hpc_prediction_intervals$central_forecast <=
      hpc_prediction_intervals$upper_80),
  all(hpc_prediction_intervals$upper_80 <=
      hpc_prediction_intervals$upper_95)
)

hpc_prediction_intervals |>
  dplyr::mutate(year = as.character(year)) |>
  knitr::kable(
    digits = 0,
    format.args = list(big.mark = ","),
    col.names = c(
      "Year",
      "Central forecast",
      "Lower 80%",
      "Upper 80%",
      "Lower 95%",
      "Upper 95%"
    ),
    caption = paste(
      "Model-based prediction intervals for",
      "overall HPC demand, 2027-2030 (kSU)"
    )
  )
```

Table 12.2: Model-based prediction intervals for overall HPC demand, 2027–2030 (kSU)

Year	Central forecast	Lower 80%	Upper 80%	Lower 95%	Upper 95%
2027	2,016,845	1,719,770	2,313,921	1,537,276	2,496,415
2028	2,178,838	1,872,803	2,484,872	1,684,806	2,672,870
2029	2,340,830	2,024,948	2,656,711	1,830,902	2,850,757
2030	2,502,822	2,176,287	2,829,357	1,975,696	3,029,948

```
ggplot2::ggplot(
  hpc_prediction_intervals,
  ggplot2::aes(x = year)
) +
  # 95% prediction interval (wider band)
  ggplot2::geom_ribbon(
    ggplot2::aes(
      ymin = lower_95 / 1e6,
      ymax = upper_95 / 1e6,
      fill = "95% prediction interval"
    ),
    alpha = 0.35
  ) +

  # 80% prediction interval (narrower band)
  ggplot2::geom_ribbon(
    ggplot2::aes(
      ymin = lower_80 / 1e6,
      ymax = upper_80 / 1e6,
      fill = "80% prediction interval"
    ),
    alpha = 0.30
  ) +

  # Central linear forecast
  ggplot2::geom_line(
    ggplot2::aes(
      y = central_forecast / 1e6
    ),
    colour = "#155E75",
    linewidth = 1.2
  ) +

  ggplot2::geom_point(
    ggplot2::aes(
```

```
    y = central_forecast / 1e6
  ),
  colour = "#155E75",
  size = 2.5
) +

# X-axis
ggplot2::scale_x_continuous(
  breaks = 2027:2030
) +

# Y-axis
ggplot2::scale_y_continuous(
  labels = scales::label_number(
    accuracy = 0.1
  )
) +

# Prediction interval legend
ggplot2::scale_fill_manual(
  name = "Uncertainty bands",
  values = c(
    "80% prediction interval" = "#0891B2",
    "95% prediction interval" = "#94A3B8"
  ),
  breaks = c(
    "80% prediction interval",
    "95% prediction interval"
  )
) +

# Chart labels
ggplot2::labs(
  title = "Overall HPC Forecast Uncertainty to 2030",
  subtitle = paste(
    "Central linear projection with",
    "model-based prediction intervals"
  ),
  x = "Year",
  y = "Projected demand (million kSU)",
  caption = paste(
    "Bands are conditional on linear regression error assumptions;",
    "measurement and model-selection uncertainty are excluded."
  )
) +
```

```
# Styling
ggplot2::theme_minimal(
  base_size = 12
) +

ggplot2::theme(
  plot.title = ggplot2::element_text(
    face = "bold"
  ),
  legend.position = "bottom",
  legend.title = ggplot2::element_text(
    face = "bold"
  ),
  plot.caption = ggplot2::element_text(
    size = 9,
    colour = "grey40"
  )
)
```

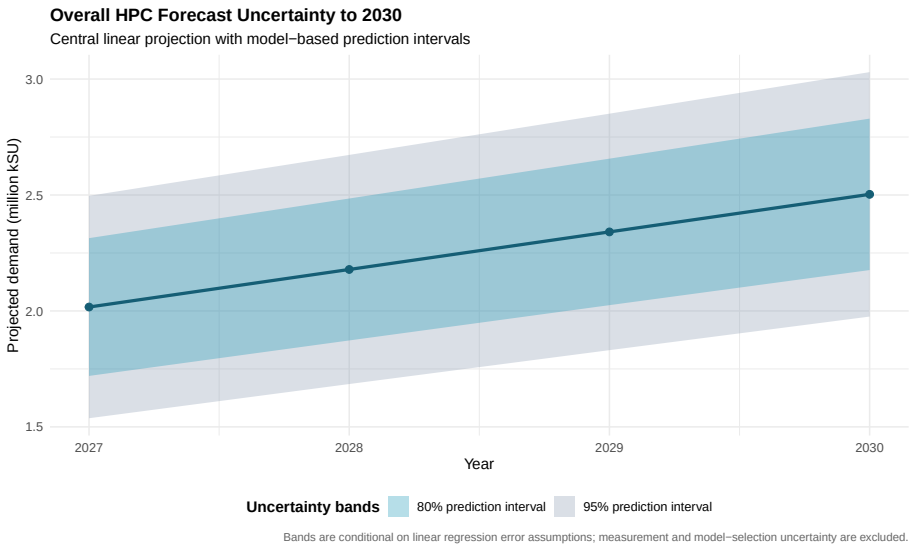


Figure 12.1: Central overall HPC demand projection with 80% and 95% model-based prediction intervals, 2027–2030.

12.4.4 Scenario Results and Interpretation

The overall HPC demand scenarios indicate continued growth in national research computing requirements through 2030, although the projected magnitude

depends substantially on the selected model specification.

Central linear projection

The selected full-history linear model projects overall HPC demand increasing from approximately 2.017 million kSU in 2027 to 2.503 million kSU in 2030.

This corresponds to an estimated annual absolute increase of approximately 161,992 kSU, consistent with the fitted historical trend.

The central projection is preferred because the linear model achieved the lowest rolling-origin RMSE and MAE among the evaluated overall HPC models in Chapter 08.

However, this result reflects historical one-step-ahead predictive performance and does not guarantee equivalent accuracy over the full projection horizon.

Sensitivity to the 2026 observation

Excluding the 2026 observation reduces the projected 2030 demand from 2,502,822 kSU to 2,285,327 kSU.

The difference is approximately 217,495 kSU, equivalent to 8.7% of the central projection.

This demonstrates that the latest observation materially influences the estimated linear growth trajectory.

The sensitivity comparison does not establish whether the 2026 increase represents a sustained change in researcher demand, a temporary development, or differences in historical reporting.

Alternative compounded-growth trajectories

The full-history log-linear model projects approximately 6.483 million kSU in 2030, while the model excluding 2026 projects approximately 6.963 million kSU.

These values are substantially higher than the central linear projection because the log-linear specifications extrapolate compounded proportional growth rather than a constant annual absolute increase.

However, Chapter 08 found that the log-linear model had substantially higher rolling-origin forecast errors than the linear model.

The log-linear projections are therefore retained as exploratory alternative growth trajectories rather than equally supported central forecasts or statistically validated upper demand bounds.

Notably, excluding 2026 increases the log-linear projection, despite reducing the corresponding linear projection. This illustrates how sensitivity to historical observations depends on the model specification.

Implications for infrastructure planning

The central statistical projection suggests that national research HPC requirements could reach approximately 2.50 million kSU by 2030 if the fitted historical linear trend continues.

The sensitivity analysis indicates that the projected magnitude depends partly on the interpretation of the latest historical observation.

The alternative log-linear trajectories demonstrate how substantially greater requirements could emerge under continued compounded growth, although the available historical evaluation provides weaker support for these specifications.

These results highlight the importance of maintaining flexibility in infrastructure planning and distinguishing projected researcher requirements from actual allocated computing capacity.

The scenarios should be interpreted as conditional, model-based estimates rather than definitive predictions of future national HPC demand.

12.5 CPU Demand Scenarios

12.5.1 Forecast Construction

CPU demand projections are developed for 2027–2030 using the linear and log-linear trend models fitted to the harmonised 2014–2026 CPU demand series in Chapter 07.

The full-history linear trend is selected as the central statistical projection, consistent with the model-selection findings in Chapter 09.

A linear specification excluding 2026 is retained to examine sensitivity to the latest observation. Full-history and 2026-excluded log-linear models are also considered as alternative compounded-growth trajectories.

CPU demand is expressed in **millions of CPU core-hours**, consistent with the historical modelling framework. These estimates are analysed separately from overall HPC demand in kSU and GPU demand in GPU-hours.

12.5.2 Annual CPU Demand Projections

```
# Generate the 2027-2030 forecast horizon
cpu_future <- data.frame(
  year = 2027:2030,
  time = (2027:2030) - 2014
)

# Generate CPU demand projections
# Units: million CPU core-hours
cpu_scenarios <- cpu_future |>
```

```
  dplyr::transmute(
    year,
    central_linear = as.numeric(
      predict(cpu_linear, newdata = cpu_future)
    ),
    linear_excl_2026 = as.numeric(
      predict(cpu_linear_pre2026, newdata = cpu_future)
    ),
    loglinear_full = exp(as.numeric(
      predict(cpu_loglinear, newdata = cpu_future)
    )),
    loglinear_excl_2026 = exp(as.numeric(
      predict(cpu_loglinear_pre2026, newdata = cpu_future)
    ))
  )

# Validate the generated projections
stopifnot(
  nrow(cpu_scenarios) == 4,
  identical(cpu_scenarios$year, 2027:2030),
  all(is.finite(as.matrix(cpu_scenarios[, -1]))),
  all(as.matrix(cpu_scenarios[, -1]) > 0)
)

# Display the annual scenario table
cpu_scenarios |>
  dplyr::mutate(
    year = as.character(year)
  ) |>
  knitr::kable(
    digits = 2,
    format.args = list(big.mark = ","),
    col.names = c(
      "Year",
      "Central linear",
      "Linear (excl. 2026)",
      "Log-linear (full)",
      "Log-linear (excl. 2026)"
    ),
    caption = paste(
      "CPU demand scenarios, 2027–2030",
      "(million CPU core-hours)"
    )
  )
```

Table 12.3: CPU demand scenarios, 2027–2030 (million CPU core-hours)

Year	Central linear	Linear (excl. 2026)	Log-linear (full)	Log-linear (excl. 2026)
2027	1,126.88	1,046.63	1,488.26	1,498.65
2028	1,207.89	1,119.04	1,764.23	1,777.87
2029	1,288.89	1,191.45	2,091.36	2,109.11
2030	1,369.90	1,263.85	2,479.16	2,502.07

```
# Convert CPU scenarios to long format
cpu_plot_data <- cpu_scenarios |>
  tidyr::pivot_longer(
    cols = -year,
    names_to = "scenario",
    values_to = "demand_million_core_hours"
  ) |>
  dplyr::mutate(
    scenario = factor(
      scenario,
      levels = c(
        "central_linear",
        "linear_excl_2026",
        "loglinear_full",
        "loglinear_excl_2026"
      ),
      labels = c(
        "Central linear",
        "Linear (excl. 2026)",
        "Log-linear (full)",
        "Log-linear (excl. 2026)"
      )
    )
  )

ggplot2::ggplot(
  cpu_plot_data,
  ggplot2::aes(
    x = year,
    y = demand_million_core_hours,
    colour = scenario,
    linetype = scenario,
    group = scenario
  )
) +
```

```
  ggplot2::geom_line(linewidth = 1.1) +
  ggplot2::geom_point(size = 2.3) +
  ggplot2::scale_x_continuous(
    breaks = 2027:2030
  ) +
  ggplot2::scale_y_continuous(
    labels = scales::label_comma(accuracy = 1)
  ) +
  ggplot2::scale_colour_manual(
    values = c(
      "Central linear" = "#155E75",
      "Linear (excl. 2026)" = "#64748B",
      "Log-linear (full)" = "#D97706",
      "Log-linear (excl. 2026)" = "#B91C1C"
    )
  ) +
  ggplot2::labs(
    title = "CPU Demand Scenarios to 2030",
    subtitle = "Central linear forecast and alternative growth assumptions",
    x = "Year",
    y = "Projected demand (million CPU core-hours)",
    colour = "Scenario",
    linetype = "Scenario"
  ) +
  ggplot2::theme_minimal(base_size = 12) +
  ggplot2::theme(
    legend.position = "bottom",
    plot.title = ggplot2::element_text(face = "bold")
  )
```

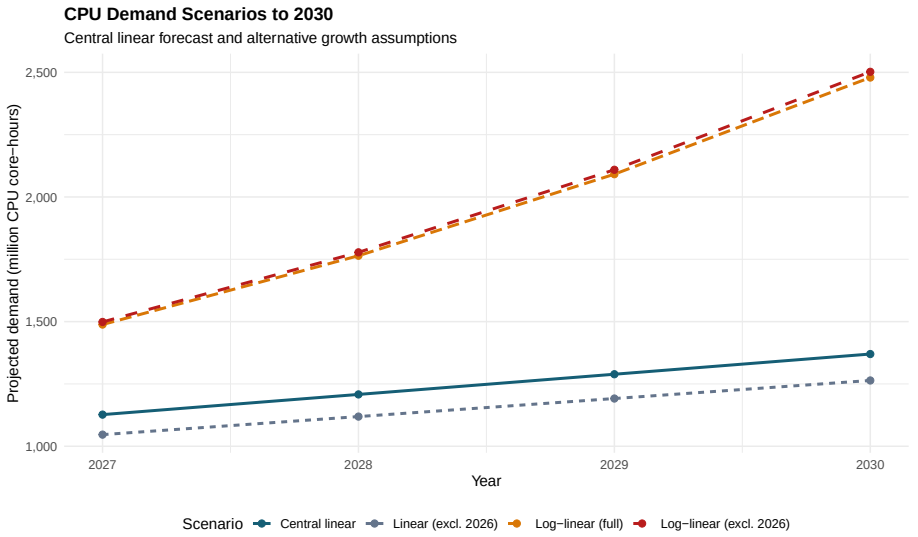


Figure 12.2: Projected CPU demand under alternative model specifications, 2027–2030.

12.5.3 Statistical Prediction Intervals for the Central CPU Forecast

To complement the alternative CPU growth scenarios, 80% and 95% model-based prediction intervals are calculated around the selected full-history linear trend.

These intervals describe uncertainty in future annual CPU demand observations conditional on the fitted linear regression model and its error assumptions.

They are distinct from the alternative growth trajectories and sensitivity comparisons presented above.

```
# Calculate prediction intervals for the central CPU model
# Units: million CPU core-hours

cpu_pi_80 <- predict(
  cpu_linear,
  newdata = cpu_future,
  interval = "prediction",
  level = 0.80
)

cpu_pi_95 <- predict(
  cpu_linear,
  newdata = cpu_future,
```

```
  interval = "prediction",
  level = 0.95
)

# Combine forecasts and interval bounds
cpu_prediction_intervals <- data.frame(
  year = cpu_future$year,
  central_forecast = cpu_pi_95[, "fit"],
  lower_80 = cpu_pi_80[, "lwr"],
  upper_80 = cpu_pi_80[, "upr"],
  lower_95 = cpu_pi_95[, "lwr"],
  upper_95 = cpu_pi_95[, "upr"]
)

# Check interval ordering and consistency
stopifnot(
  nrow(cpu_prediction_intervals) == 4,
  all(is.finite(as.matrix(cpu_prediction_intervals[, -1]))),
  all(
    cpu_prediction_intervals$lower_95 <=
      cpu_prediction_intervals$lower_80
  ),
  all(
    cpu_prediction_intervals$lower_80 <=
      cpu_prediction_intervals$central_forecast
  ),
  all(
    cpu_prediction_intervals$central_forecast <=
      cpu_prediction_intervals$upper_80
  ),
  all(
    cpu_prediction_intervals$upper_80 <=
      cpu_prediction_intervals$upper_95
  ),
  isTRUE(all.equal(
    cpu_prediction_intervals$central_forecast,
    cpu_scenarios$central_linear,
    tolerance = 1e-8
  ))
)

# Display prediction interval table
cpu_prediction_intervals |>
  dplyr::mutate(
    year = as.character(year)
```

```
  ) |>
  knitr::kable(
    digits = 2,
    format.args = list(big.mark = ","),
    col.names = c(
      "Year",
      "Central forecast",
      "Lower 80%",
      "Upper 80%",
      "Lower 95%",
      "Upper 95%"
    ),
    caption = paste(
      "Model-based prediction intervals for CPU demand,",
      "2027-2030 (million CPU core-hours)"
    )
  )
```

Table 12.4: Model-based prediction intervals for CPU demand, 2027–2030 (million CPU core-hours)

Year	Central forecast	Lower 80%	Upper 80%	Lower 95%	Upper 95%
2027	1,126.88	903.90	1,349.86	766.92	1,486.84
2028	1,207.89	978.18	1,437.59	837.07	1,578.70
2029	1,288.89	1,051.80	1,525.99	906.15	1,671.64
2030	1,369.90	1,124.80	1,614.99	974.24	1,765.55

```
ggplot2::ggplot(
  cpu_prediction_intervals,
  ggplot2::aes(x = year)
) +
  # 95% prediction interval (wider band)
  ggplot2::geom_ribbon(
    ggplot2::aes(
      ymin = lower_95,
      ymax = upper_95,
      fill = "95% prediction interval"
    ),
    alpha = 0.35
  ) +

  # 80% prediction interval (narrower band)
  ggplot2::geom_ribbon(
    ggplot2::aes(
```

```
    ymin = lower_80,
    ymax = upper_80,
    fill = "80% prediction interval"
  ),
  alpha = 0.30
) +

# Central linear forecast
ggplot2::geom_line(
  ggplot2::aes(y = central_forecast),
  colour = "#155E75",
  linewidth = 1.2
) +

ggplot2::geom_point(
  ggplot2::aes(y = central_forecast),
  colour = "#155E75",
  size = 2.5
) +

# Axes
ggplot2::scale_x_continuous(
  breaks = 2027:2030
) +

ggplot2::scale_y_continuous(
  labels = scales::label_comma(accuracy = 1)
) +

# Prediction interval legend
ggplot2::scale_fill_manual(
  name = "Uncertainty bands",
  values = c(
    "80% prediction interval" = "#0891B2",
    "95% prediction interval" = "#94A3B8"
  ),
  breaks = c(
    "80% prediction interval",
    "95% prediction interval"
  )
) +

# Labels
ggplot2::labs(
  title = "CPU Forecast Uncertainty to 2030",
```

```
    subtitle = paste(
      "Central linear projection with",
      "model-based prediction intervals"
    ),
    x = "Year",
    y = "Projected demand (million CPU core-hours)",
    caption = paste(
      "Bands are conditional on linear regression error assumptions;",
      "measurement and model-selection uncertainty are excluded."
    )
  ) +

  # Styling
  ggplot2::theme_minimal(base_size = 12) +

  ggplot2::theme(
    plot.title = ggplot2::element_text(
      face = "bold"
    ),
    legend.position = "bottom",
    legend.title = ggplot2::element_text(
      face = "bold"
    ),
    plot.caption = ggplot2::element_text(
      size = 9,
      colour = "grey40"
    )
  )
```

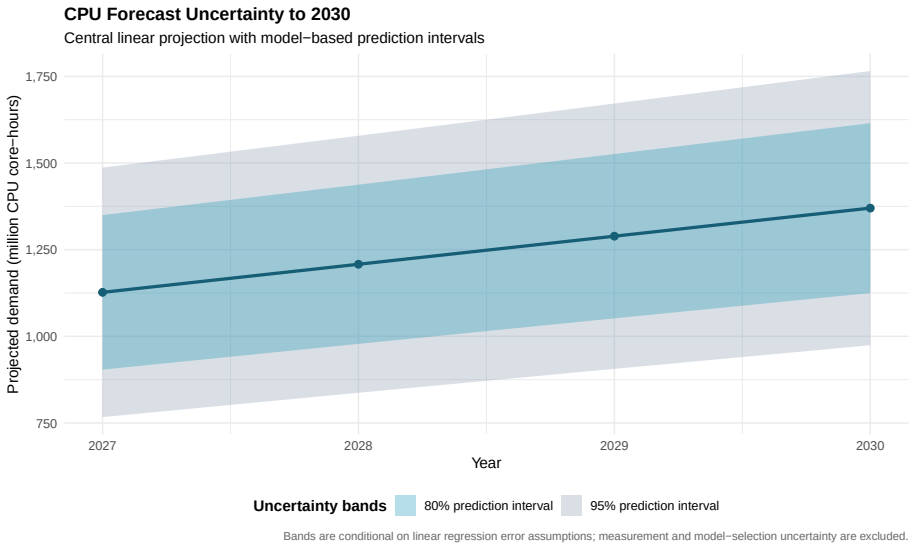
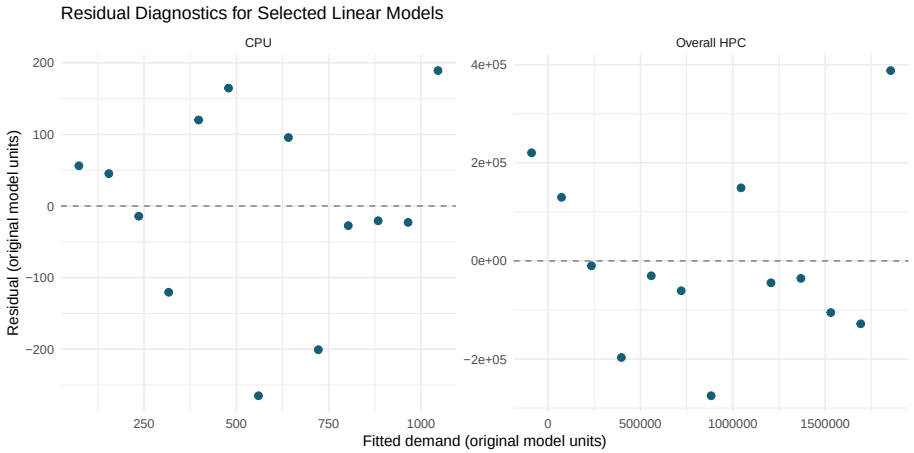


Figure 12.3: Central CPU demand projection with 80% and 95% model-based prediction intervals, 2027–2030.

```
# Compare residual diagnostics for selected linear models
diagnostic_data <- dplyr::bind_rows(
  data.frame(
    model = "Overall HPC",
    fitted = fitted(overall_linear),
    residual = residuals(overall_linear)
  ),
  data.frame(
    model = "CPU",
    fitted = fitted(cpu_linear),
    residual = residuals(cpu_linear)
  )
)

ggplot2::ggplot(
  diagnostic_data,
  ggplot2::aes(x = fitted, y = residual)
) +
  ggplot2::geom_hline(
    yintercept = 0,
    linetype = "dashed",
    colour = "grey50"
  ) +
  ggplot2::geom_point(
```

```
    colour = "#155E75",
    size = 2.5
  ) +
  ggplot2::facet_wrap(
    ~ model,
    scales = "free"
  ) +
  ggplot2::labs(
    title = "Residual Diagnostics for Selected Linear Models",
    x = "Fitted demand (original model units)",
    y = "Residual (original model units)"
  ) +
  ggplot2::theme_minimal(base_size = 12)
```




```
# Ljung-Box tests for residual autocorrelation
diagnostic_tests <- data.frame(
  model = c("Overall HPC", "CPU"),
  ljung_box_p_value = c(
    Box.test(
      residuals(overall_linear),
      lag = 2,
      type = "Ljung-Box",
      fitdf = 0
    )$p.value,
    Box.test(
      residuals(cpu_linear),
      lag = 2,
      type = "Ljung-Box",
      fitdf = 0
    )$p.value
```

```
  )
)

knitr::kable(
  diagnostic_tests,
  digits = 4,
  caption = "Exploratory Ljung-Box residual autocorrelation tests"
)
```

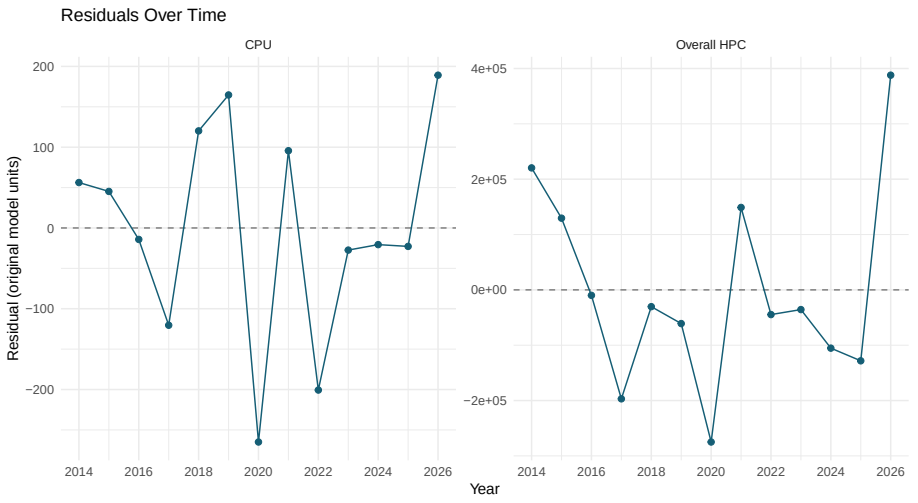
Table 12.5: Exploratory Ljung-Box residual autocorrelation tests

model	ljung_box_p_value
Overall HPC	0.8775
CPU	0.3608

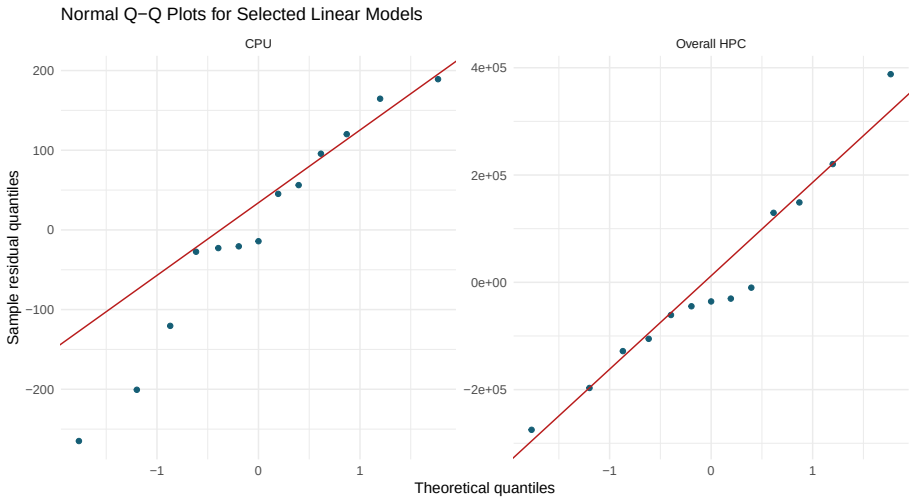
```
# Construct diagnostics using each model's fitted observation years
additional_diagnostics <- dplyr::bind_rows(
  data.frame(
    model = "Overall HPC",
    year = 2014:2026,
    residual = as.numeric(residuals(overall_linear))
  ),
  data.frame(
    model = "CPU",
    year = 2014:2026,
    residual = as.numeric(residuals(cpu_linear))
  )
)

# Residuals over time
ggplot2::ggplot(
  additional_diagnostics,
  ggplot2::aes(x = year, y = residual)
) +
  ggplot2::geom_hline(
    yintercept = 0,
    linetype = "dashed",
    colour = "grey50"
  ) +
  ggplot2::geom_line(colour = "#155E75") +
  ggplot2::geom_point(colour = "#155E75", size = 2) +
  ggplot2::facet_wrap(~ model, scales = "free_y") +
  ggplot2::scale_x_continuous(
    breaks = seq(2014, 2026, by = 2)
```

```
  ) +
  ggplot2::labs(
    title = "Residuals Over Time",
    x = "Year",
    y = "Residual (original model units)"
  ) +
  ggplot2::theme_minimal(base_size = 12)
```




```
# Normal Q-Q plots
ggplot2::ggplot(
  additional_diagnostics,
  ggplot2::aes(sample = residual)
) +
  ggplot2::stat_qq(colour = "#155E75") +
  ggplot2::stat_qq_line(colour = "#B91C1C") +
  ggplot2::facet_wrap(~ model, scales = "free") +
  ggplot2::labs(
    title = "Normal Q-Q Plots for Selected Linear Models",
    x = "Theoretical quantiles",
    y = "Sample residual quantiles"
  ) +
  ggplot2::theme_minimal(base_size = 12)
```



12.5.4 Assessment of Prediction Interval Assumptions

The 80% and 95% prediction intervals presented for overall HPC and CPU demand are derived from the selected full-history linear regression models using conventional prediction interval calculations.

Their interpretation depends on the adequacy of the fitted models and their underlying error assumptions, including independence, approximately constant variance, and normality.

Residual diagnostics were examined for both selected models using residual-versus-fitted plots, residuals over time, normal Q–Q plots, and exploratory Ljung–Box tests.

The Ljung–Box tests at lag 2 returned p-values of 0.8775 for overall HPC and 0.3608 for CPU. Neither test detected statistically significant residual autocorrelation at the 5% significance level.

However, the limited sample of 13 annual observations reduces the power of these tests. Non-significant results should therefore not be interpreted as confirmation of residual independence.

For overall HPC, the residuals fluctuate around zero, with a relatively large positive residual in 2026. The normal Q–Q plot shows approximate alignment with the reference line but some departures at the distribution tails.

For CPU, the residuals also fluctuate around zero, although several years exhibit substantial positive or negative deviations. The Q–Q plot shows departures from the normal reference line, particularly in the lower tail.

The residual plots do not establish a definitive pattern of changing variance, but the available evidence is insufficient to confirm homoscedasticity or normality.

Accordingly, the reported 80% and 95% prediction intervals are interpreted as **conditional, model-based uncertainty estimates** rather than empirically validated coverage bounds.

These intervals reflect uncertainty under the selected linear regression specifications. They do not incorporate uncertainty arising from historical measurement differences, alternative model selection, future structural changes, or shifts in research computing policy and infrastructure.

The alternative log-linear trajectories and 2026 sensitivity analyses should therefore be considered alongside the prediction intervals when assessing possible future computational requirements.

The prediction intervals provide an additional perspective on statistical uncertainty, but they do not represent comprehensive lower and upper limits for national HPC or CPU demand through 2030.

12.6 GPU Demand Scenarios

12.6.1 Forecast Construction and Limitations

GPU demand scenarios are developed for 2027–2030 using the recent 2022–2026 historical series, expressed in million GPU-hours.

Unlike the overall HPC and CPU analyses, the GPU series contains only five annual estimates, combining reported Setonix-based demand with estimates derived from GPU allocation shares and working conversion assumptions.

Chapter 08 found that the naïve benchmark achieved the lowest rolling-origin RMSE and MAE among the evaluated GPU models. However, the evaluation covered only two forecast years, providing insufficient evidence to establish a definitive long-term forecasting model.

Accordingly, the naïve approach is retained as a constant-demand statistical benchmark. Linear and log-linear specifications, fitted both with and without the 2026 observation, are retained as exploratory growth trajectories.

The historical GPU conversion scenarios for 2014–2021 are not included as additional observations in these forecasting models because they represent hypothetical conversion assumptions rather than independently observed GPU demand.

The resulting projections should be interpreted as conditional scenarios rather than validated long-term forecasts.

12.6.2 Annual GPU Demand Projections

```
# Forecast horizon: 2027-2030
# Chapter 07 GPU models use time = year - 2022
```

```r
gpu_future <- data.frame(
  year = 2027:2030,
  time = (2027:2030) - 2022
)

# Retrieve the final observed/estimated GPU demand value
# from the data used to fit the full-history linear model.
gpu_training <- model.frame(gpu_linear)

gpu_training_years <- 2022 + gpu_training$time

stopifnot(
  nrow(gpu_training) == 5,
  identical(as.integer(gpu_training_years), 2022:2026),
  all(is.finite(gpu_training$gpu_million_hours))
)

gpu_last_value <- gpu_training$gpu_million_hours[
  which(gpu_training_years == 2026)
]

# Generate five conditional GPU demand scenarios
# Units: million GPU-hours

gpu_scenarios <- gpu_future |>
  dplyr::transmute(
    year,

    naive_benchmark = rep(
      gpu_last_value,
      length(year)
    ),

    linear_full = as.numeric(
      predict(gpu_linear, newdata = gpu_future)
    ),

    linear_excl_2026 = as.numeric(
      predict(gpu_linear_pre2026, newdata = gpu_future)
    ),

    loglinear_full = exp(as.numeric(
      predict(gpu_loglinear, newdata = gpu_future)
    )),
```

```
    loglinear_excl_2026 = exp(as.numeric(
      predict(gpu_loglinear_pre2026, newdata = gpu_future)
    ))
  )

# Validate projections
stopifnot(
  nrow(gpu_scenarios) == 4,
  identical(gpu_scenarios$year, 2027:2030),
  all(is.finite(as.matrix(gpu_scenarios[, -1]))),
  all(as.matrix(gpu_scenarios[, -1]) > 0)
)

# Display GPU scenario table
gpu_scenarios |>
  dplyr::mutate(
    year = as.character(year)
  ) |>
  knitr::kable(
    digits = 3,
    format.args = list(big.mark = ","),
    col.names = c(
      "Year",
      "Naïve benchmark",
      "Linear (full)",
      "Linear (excl. 2026)",
      "Log-linear (full)",
      "Log-linear (excl. 2026)"
    ),
    caption = paste(
      "Exploratory GPU demand scenarios,",
      "2027–2030 (million GPU-hours)"
    )
  )
```

Table 12.6: Exploratory GPU demand scenarios, 2027–2030 (million GPU-hours)

Year	Naïve bench-mark	Linear (full)	Linear (excl. 2026)	Log-linear (full)	Log-linear (excl. 2026)
2027	3.003	2.855	2.215	2.885	2.262
2028	3.003	3.027	2.227	3.129	2.309
2029	3.003	3.199	2.240	3.394	2.357
2030	3.003	3.371	2.252	3.681	2.406

```
# Convert scenario projections into long format
gpu_plot_data <- gpu_scenarios |>
  tidyr::pivot_longer(
    cols = -year,
    names_to = "scenario",
    values_to = "demand_million_gpu_hours"
  ) |>
  dplyr::mutate(
    scenario = factor(
      scenario,
      levels = c(
        "naive_benchmark",
        "linear_full",
        "linear_excl_2026",
        "loglinear_full",
        "loglinear_excl_2026"
      ),
      labels = c(
        "Naïve benchmark",
        "Linear (full)",
        "Linear (excl. 2026)",
        "Log-linear (full)",
        "Log-linear (excl. 2026)"
      )
    )
  )

ggplot2::ggplot(
  gpu_plot_data,
  ggplot2::aes(
    x = year,
    y = demand_million_gpu_hours,
    colour = scenario,
    linetype = scenario,
    group = scenario
  )
) +
  ggplot2::geom_line(linewidth = 1.1) +
  ggplot2::geom_point(size = 2.4) +

  ggplot2::scale_x_continuous(
    breaks = 2027:2030
  ) +

  ggplot2::scale_y_continuous(
```

```
  labels = scales::label_number(accuracy = 0.1)
) +

ggplot2::scale_colour_manual(
  values = c(
    "Naïve benchmark" = "#155E75",
    "Linear (full)" = "#2563EB",
    "Linear (excl. 2026)" = "#64748B",
    "Log-linear (full)" = "#D97706",
    "Log-linear (excl. 2026)" = "#B91C1C"
  )
) +

ggplot2::labs(
  title = "Exploratory GPU Demand Scenarios to 2030",
  subtitle = "Naïve benchmark and alternative growth assumptions",
  x = "Year",
  y = "Projected demand (million GPU-hours)",
  colour = "Scenario",
  linetype = "Scenario",
  caption = paste(
    "Based on five annual GPU estimates (2022-2026);",
    "trajectories are exploratory, not statistical prediction bounds."
  )
) +

ggplot2::theme_minimal(base_size = 12) +

ggplot2::theme(
  plot.title = ggplot2::element_text(face = "bold"),
  legend.position = "bottom",
  legend.title = ggplot2::element_text(face = "bold"),
  legend.box = "vertical",
  plot.caption = ggplot2::element_text(
    size = 9,
    colour = "grey40"
  )
) +

ggplot2::guides(
  colour = ggplot2::guide_legend(nrow = 2),
  linetype = ggplot2::guide_legend(nrow = 2)
)
```

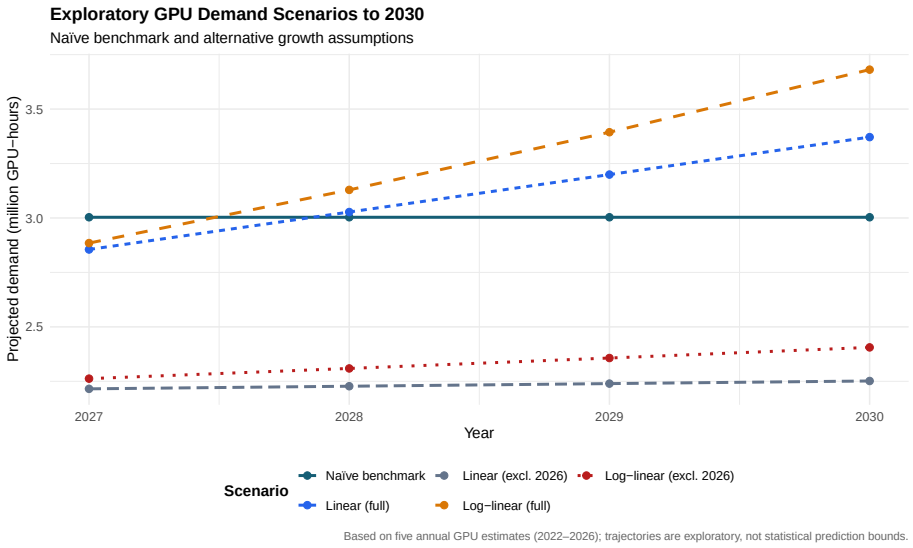


Figure 12.4: Exploratory GPU demand scenarios, 2027–2030. The naïve benchmark assumes constant demand; trend trajectories reflect alternative growth assumptions.

12.6.3 GPU Scenario Results and Interpretation

The GPU demand scenarios produce substantially different trajectories for 2027–2030, reflecting both alternative growth assumptions and the influence of the 2026 observation.

The naïve benchmark maintains the 2026 GPU demand estimate of approximately 3.003 million GPU-hours throughout the forecast horizon. Although this benchmark achieved the lowest rolling-origin RMSE and MAE in Chapter 08, its evaluation was based on only two forecast years. Its relatively favourable short-term performance therefore does not establish that GPU demand will remain constant through 2030.

Under the full-history linear specification, projected GPU demand increases from approximately 2.855 million GPU-hours in 2027 to 3.371 million GPU-hours in 2030. The corresponding full-history log-linear scenario reaches approximately 3.681 million GPU-hours by 2030, reflecting the assumption of compounded proportional growth.

Excluding 2026 produces considerably lower projections. The 2030 linear projection falls to approximately 2.252 million GPU-hours, while the corresponding log-linear projection reaches approximately 2.406 million GPU-hours.

The difference between the full-history and 2026-excluded linear projections is approximately 1.119 million GPU-hours in 2030, equivalent to an increase of

approximately 49.7% relative to the 2026-excluded projection.

Across the four exploratory trend specifications, projected GPU demand in 2030 ranges from approximately **2.252 to 3.681 million GPU-hours**. Including the naïve benchmark does not change these minimum and maximum values.

This range should not be interpreted as a statistical prediction interval or as validated lower and upper limits for future GPU demand. It represents the spread of projections generated by the selected modelling assumptions.

The substantial differences between scenarios are particularly important given the limitations of the historical GPU series. Only five annual estimates are available for 2022–2026, combining reported Setonix-based demand with estimates derived from allocation shares and working conversion assumptions.

Consequently, the available evidence is insufficient to establish a definitive central GPU demand forecast. The naïve benchmark provides a constant-demand reference, while the linear and log-linear specifications illustrate alternative growth trajectories.

Unlike the overall HPC and CPU analyses, formal 80% and 95% prediction intervals are not reported for GPU demand. With such a short and heterogeneous historical series, conventional regression intervals would be highly sensitive to modelling assumptions and difficult to interpret as reliable coverage bounds.

For national research computing planning, these results indicate that GPU demand should be considered through multiple conditional scenarios rather than a single point estimate. Improved GPU-specific demand reporting, consistent measurement units, and additional annual observations would strengthen future forecasting and infrastructure planning.

The GPU scenarios should therefore be interpreted as **exploratory planning inputs**, not independently validated forecasts of national GPU requirements.

12.7 Field of Research Allocation-Share Scenarios

12.7.1 Purpose and Methodology

The Field of Research (FoR) scenarios examine how projected overall HPC demand for 2027–2030 could be distributed across research fields under alternative historical allocation-share assumptions.

These scenarios use the harmonised FoR allocation shares developed in Chapter 06 and the central overall HPC linear projection developed earlier in this chapter.

The analysis adopts a top-down approach:

[Illustrative FoR demand*{f,t} = Projected overall HPC demand*{t} ×Assumed allocation share_{f}
]

where (f) denotes the research-field group and (t) denotes the forecast year.

The resulting values are expressed in kSU. They represent hypothetical distributions of projected overall demand, **not independently estimated or observed FoR-level researcher demand**.

Two historical reference periods are considered:

- **Recent-period composition (2021–2026):** Uses the harmonised FoR 2020 classification groups to reflect more recent allocation patterns.
- **Longer-term composition (2014–2026):** Uses the harmonised analytical FoR groups to examine historical allocation patterns across the wider period.

The two classification structures are not assumed to be directly interchangeable. Harmonisation supports broad analytical comparisons but does not establish complete equivalence between official FoR classification versions.

Annual allocation shares are examined before constructing scenario assumptions. Category coverage, reporting scope, and changes in classification are retained as methodological limitations.

In particular, the 2017 FoR observation requires caution because its reporting scope differs from other years.

The scenarios hold assumed research-field shares constant over 2027–2030. This is a transparent planning assumption rather than evidence that future disciplinary demand composition will remain unchanged.

12.7.2 Validate Historical Allocation Shares

```
# Load harmonised FoR allocation-share datasets from Chapter 06

for_recent <- readRDS(
  "assets/data/data converted/for_2020_harmonised_shares.rds"
)

for_longterm <- readRDS(
  "assets/data/data converted/for_2014_2026_harmonised_shares.rds"
)

# Confirm expected columns
stopifnot(
  all(c("year", "for_2020_visual", "allocation_share") %in%
        names(for_recent)),
  all(c("year", "for_visual_group", "allocation_share") %in%
        names(for_longterm))
)
```

```r
# Validate historical coverage
# Compare year values without requiring identical numeric types

stopifnot(
  isTRUE(all.equal(
    sort(unique(as.numeric(for_recent$year))),
    as.numeric(2021:2026)
  )),
  isTRUE(all.equal(
    sort(unique(as.numeric(for_longterm$year))),
    as.numeric(2014:2026)
  )),
  all(is.finite(for_recent$allocation_share)),
  all(is.finite(for_longterm$allocation_share)),
  all(for_recent$allocation_share >= 0),
  all(for_longterm$allocation_share >= 0)
)


# Check annual allocation-share totals
for_share_checks <- dplyr::bind_rows(
  for_recent |>
    dplyr::group_by(year) |>
    dplyr::summarise(
      total_share = sum(allocation_share),
      .groups = "drop"
    ) |>
    dplyr::mutate(period = "Recent: 2021-2026"),

  for_longterm |>
    dplyr::group_by(year) |>
    dplyr::summarise(
      total_share = sum(allocation_share),
      .groups = "drop"
    ) |>
    dplyr::mutate(period = "Long-term: 2014-2026")
)

stopifnot(
  all(abs(for_share_checks$total_share - 1) < 1e-6)
)

knitr::kable(
  for_share_checks |>
    dplyr::select(period, year, total_share) |>
```

```
    dplyr::mutate(year = as.character(year)),
  digits = 4,
  col.names = c(
    "Historical period",
    "Year",
    "Total allocation share"
  ),
  caption = "Validation of annual harmonised FoR allocation shares"
)
```

Table 12.7: Validation of annual harmonised FoR allocation shares

Historical period	Year	Total allocation share
Recent: 2021–2026	2021	1
Recent: 2021–2026	2022	1
Recent: 2021–2026	2023	1
Recent: 2021–2026	2024	1
Recent: 2021–2026	2025	1
Recent: 2021–2026	2026	1
Long-term: 2014–2026	2014	1
Long-term: 2014–2026	2015	1
Long-term: 2014–2026	2016	1
Long-term: 2014–2026	2017	1
Long-term: 2014–2026	2018	1
Long-term: 2014–2026	2019	1
Long-term: 2014–2026	2020	1
Long-term: 2014–2026	2021	1
Long-term: 2014–2026	2022	1
Long-term: 2014–2026	2023	1
Long-term: 2014–2026	2024	1
Long-term: 2014–2026	2025	1
Long-term: 2014–2026	2026	1

12.7.3 Historical FoR Share Assumptions

For each historical reference period, an unweighted mean of the annual allocation shares is calculated for each harmonised research-field group.

These averages describe the historical distribution of recorded allocations. They are not estimates of independently observed research demand.

The recent-period and longer-term shares are calculated separately because their harmonised category structures and underlying reporting histories differ.

```r
# Recent-period mean shares: 2021-2026
recent_for_shares <- for_recent |>
  dplyr::group_by(for_2020_visual) |>
  dplyr::summarise(
    mean_share = mean(allocation_share),
    years_reported = dplyr::n_distinct(year),
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(mean_share))

# Longer-term mean shares: 2014-2026
longterm_for_shares <- for_longterm |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    mean_share = mean(allocation_share),
    years_reported = dplyr::n_distinct(year),
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(mean_share))

# Check the total of the unweighted mean shares
share_summary <- data.frame(
  reference_period = c(
    "Recent: 2021-2026",
    "Long-term: 2014-2026"
  ),
  total_mean_share = c(
    sum(recent_for_shares$mean_share),
    sum(longterm_for_shares$mean_share)
  ),
  number_of_groups = c(
    nrow(recent_for_shares),
    nrow(longterm_for_shares)
  )
)

knitr::kable(
  share_summary,
  digits = 4,
  col.names = c(
    "Reference period",
    "Sum of mean shares",
    "Number of FoR groups"
  ),
  caption = "Summary of historical FoR allocation-share assumptions"
```

```
)
```

Table 12.8: Summary of historical FoR allocation-share assumptions

Reference period	Sum of mean shares	Number of FoR groups
Recent: 2021–2026	1.0482	9
Long-term: 2014–2026	1.0417	9

```
# Display recent-period group assumptions
recent_for_shares |>
  dplyr::mutate(
    mean_share_percent = 100 * mean_share
  ) |>
  dplyr::select(
    for_2020_visual,
    years_reported,
    mean_share_percent
  ) |>
  knitr::kable(
    digits = 2,
    col.names = c(
      "Recent FoR group",
      "Years represented",
      "Mean annual allocation share (%)"
    ),
    caption = "Recent-period FoR allocation-share assumptions, 2021-2026"
  )
```

Table 12.9: Recent-period FoR allocation-share assumptions, 2021–2026

Recent FoR group	Years represented	Mean annual allocation share (%)
Engineering & Technology	6	34.88
Physical Sciences	6	22.13
Chemical Sciences	6	17.84
Earth & Environmental Sciences	6	14.00
Other / Unclassified	2	7.07
Biological & Agricultural Sciences	6	5.17
Mathematics & Computing	6	2.23
Medical & Health Sciences	6	1.18
Social Sciences, Humanities & Built Environment	4	0.31

12.7.4 Earlier-Period FoR Share Assumptions, 2014–2020

The earlier-period scenario uses the 2014–2020 subset of the harmonised historical FoR allocation-share dataset. This provides a reference for comparing allocation composition before and after 2021.

The earlier and recent periods are non-overlapping, although differences in classification, facility coverage, and reporting scope limit direct comparability.

```
# Earlier historical period: 2014-2020

for_earlier <- for_longterm |>
  dplyr::filter(year >= 2014, year <= 2020)

# Calculate observed-year mean allocation shares
earlier_for_shares <- for_earlier |>
  dplyr::group_by(for_visual_group) |>
  dplyr::summarise(
    mean_share = mean(allocation_share),
    years_reported = dplyr::n_distinct(year),
    .groups = "drop"
  ) |>
  dplyr::arrange(dplyr::desc(mean_share))

# Normalise to create conditional scenario weights
earlier_for_weights <- earlier_for_shares |>
  dplyr::mutate(
    scenario_share = mean_share / sum(mean_share)
  )

# Validate
stopifnot(
  identical(sort(unique(as.integer(for_earlier$year))), 2014:2020),
  abs(sum(earlier_for_weights$scenario_share) - 1) < 1e-10
)

# Display earlier-period assumptions
earlier_for_weights |>
  dplyr::mutate(
    observed_mean_pct = 100 * mean_share,
    scenario_share_pct = 100 * scenario_share
  ) |>
  dplyr::select(
    for_visual_group,
    years_reported,
    observed_mean_pct,
    scenario_share_pct
```

```
  ) |>
  knitr::kable(
    digits = 2,
    col.names = c(
      "FoR group",
      "Years represented",
      "Observed-year mean (%)",
      "Normalised scenario share (%)"
    ),
    caption = "Earlier-period FoR scenario weights, 2014-2020"
  )
```

Table 12.10: Earlier-period FoR scenario weights, 2014–2020

FoR group	Years represented	Observed-year mean (%)	Normalised scenario share (%)
Earth & Environmental Sciences	7	31.62	30.38
Chemical Sciences	7	19.02	18.28
Physical Sciences	7	16.92	16.26
Engineering & Technology	7	16.90	16.24
Unclassified & Operational	5	7.26	6.98
Biological & Agricultural Sciences	7	4.88	4.69
Mathematics & Computing	6	4.69	4.51
Medical & Health Sciences	4	2.47	2.38
Social Sciences, Humanities & Built Environment	1	0.31	0.30

```
# Identify which years contain each FoR group

for_category_coverage <- for_recent |>
  dplyr::group_by(for_2020_visual) |>
  dplyr::summarise(
    years_present = paste(sort(unique(year)), collapse = ", "),
    years_reported = dplyr::n_distinct(year),
    .groups = "drop"
  ) |>
  dplyr::arrange(years_reported, for_2020_visual)

knitr::kable(
  for_category_coverage,
  col.names = c(
```

```
    "FoR group",
    "Years present",
    "Number of years"
  ),
  caption = "Coverage of recent FoR groups, 2021-2026"
)
```

Table 12.11: Coverage of recent FoR groups, 2021–2026

FoR group	Years present	Number of years
Other / Unclassified	2021, 2024	2
Social Sciences, Humanities & Built Environment	2022, 2023, 2024, 2025	4
Biological & Agricultural Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Chemical Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Earth & Environmental Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Engineering & Technology	2021, 2022, 2023, 2024, 2025, 2026	6
Mathematics & Computing	2021, 2022, 2023, 2024, 2025, 2026	6
Medical & Health Sciences	2021, 2022, 2023, 2024, 2025, 2026	6
Physical Sciences	2021, 2022, 2023, 2024, 2025, 2026	6

‘

```
knitr::kable(
  share_summary,
  digits = 6,
  col.names = c(
    "Reference period",
    "Sum of mean shares",
    "Number of FoR groups"
  ),
  caption = "Completeness of historical FoR share assumptions"
)
```

Table 12.12: Completeness of historical FoR share assumptions

Reference period	Sum of mean shares	Number of FoR groups
Recent: 2021–2026	1.048180	9
Long-term: 2014–2026	1.041746	9

```
# Construct conditional scenario weights by normalising
# observed-year mean allocation shares.
#
# IMPORTANT:
# Missing category-years are not treated as observed zeros.
# Normalisation is a scenario assumption, not data recovery.

recent_for_weights <- recent_for_shares |>
  dplyr::mutate(
    scenario_share = mean_share / sum(mean_share)
  )

longterm_for_weights <- longterm_for_shares |>
  dplyr::mutate(
    scenario_share = mean_share / sum(mean_share)
  )

stopifnot(
  abs(sum(recent_for_weights$scenario_share) - 1) < 1e-10,
  abs(sum(longterm_for_weights$scenario_share) - 1) < 1e-10
)

# Show recent-period scenario assumptions
recent_for_weights |>
  dplyr::mutate(
    observed_mean_pct = 100 * mean_share,
    scenario_share_pct = 100 * scenario_share
  ) |>
  dplyr::select(
    for_2020_visual,
    years_reported,
    observed_mean_pct,
    scenario_share_pct
  ) |>
  knitr::kable(
    digits = 2,
    col.names = c(
      "FoR group",
```

```
      "Years represented",
      "Observed-year mean (%)",
      "Normalised scenario share (%)"
    ),
    caption = paste(
      "Recent FoR scenario weights, 2021-2026:",
      "normalised observed-year means"
    )
  )
```

Table 12.13: Recent FoR scenario weights, 2021–2026: normalised observed-year means

FoR group	Years represented	Observed-year mean (%)	Normalised scenario share (%)
Engineering & Technology	6	34.88	33.28
Physical Sciences	6	22.13	21.11
Chemical Sciences	6	17.84	17.02
Earth & Environmental Sciences	6	14.00	13.36
Other / Unclassified	2	7.07	6.75
Biological & Agricultural Sciences	6	5.17	4.93
Mathematics & Computing	6	2.23	2.13
Medical & Health Sciences	6	1.18	1.12
Social Sciences, Humanities & Built Environment	4	0.31	0.30

12.7.5 Recent-Period FoR Allocation Scenario, 2027–2030

The recent-period scenario applies the normalised 2021–2026 FoR allocation-share assumptions to the central overall HPC demand projection for 2027–2030.

The calculation assumes that the historical allocation composition remains constant over the forecast horizon. It is intended to illustrate a possible distribution of projected national HPC demand rather than provide independently validated demand forecasts for individual research fields.

```
# Apply recent-period FoR weights to projected overall HPC demand
# Units: kSU

for_recent_projections <- tidyr::crossing(
  year = hpc_scenarios$year,
  for_2020_visual = recent_for_weights$for_2020_visual
```

```
) |>
  dplyr::left_join(
    hpc_scenarios |>
      dplyr::select(
        year,
        overall_hpc_ksu = central_linear
      ),
    by = "year"
  ) |>
  dplyr::left_join(
    recent_for_weights |>
      dplyr::select(
        for_2020_visual,
        scenario_share
      ),
    by = "for_2020_visual"
  ) |>
  dplyr::mutate(
    illustrative_for_ksu =
      overall_hpc_ksu * scenario_share
  )

# Confirm no missing values and correct annual totals
stopifnot(
  all(is.finite(for_recent_projections$illustrative_for_ksu)),
  all(for_recent_projections$illustrative_for_ksu >= 0)
)

for_recent_total_check <- for_recent_projections |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    allocated_ksu = sum(illustrative_for_ksu),
    projected_overall_ksu = dplyr::first(overall_hpc_ksu),
    difference_ksu = allocated_ksu - projected_overall_ksu,
    .groups = "drop"
  )

stopifnot(
  all(abs(for_recent_total_check$difference_ksu) < 1e-6)
)

# Display the 2030 illustrative allocation by FoR group
for_recent_projections |>
  dplyr::filter(year == 2030) |>
  dplyr::arrange(dplyr::desc(illustrative_for_ksu)) |>
```

```
  dplyr::mutate(
    scenario_share_pct = 100 * scenario_share
  ) |>
  dplyr::select(
    for_2020_visual,
    scenario_share_pct,
    illustrative_for_ksu
  ) |>
  knitr::kable(
    digits = 2,
    format.args = list(big.mark = ","),
    col.names = c(
      "FoR group",
      "Scenario share (%)",
      "Illustrative 2030 allocation (kSU)"
    ),
    caption = paste(
      "Illustrative 2030 FoR distribution of projected",
      "overall HPC demand: recent-period share assumptions"
    )
  )
```

Table 12.14: Illustrative 2030 FoR distribution of projected overall HPC demand: recent-period share assumptions

FoR group	Scenario share (%)	Illustrative 2030 allocation (kSU)
Engineering & Technology	33.28	832,907.52
Physical Sciences	21.11	528,465.99
Chemical Sciences	17.02	425,998.36
Earth & Environmental Sciences	13.36	334,352.53
Other / Unclassified	6.75	168,852.53
Biological & Agricultural Sciences	4.93	123,417.70
Mathematics & Computing	2.13	53,319.84
Medical & Health Sciences	1.12	28,081.41
Social Sciences, Humanities & Built Environment	0.30	7,426.00

12.7.6 Longer-Term FoR Allocation Scenario, 2027–2030

The longer-term scenario applies normalised historical FoR allocation-share assumptions from 2014–2026 to the central overall HPC demand projection.

This provides an alternative distribution based on a broader historical reference period. The resulting projections are illustrative and should not be interpreted

as independently validated research-field demand forecasts.

Differences between the recent-period and longer-term scenarios may reflect changes in research allocation patterns, FoR classification structures, facility coverage, and reporting practices. They should not automatically be interpreted as changes in underlying researcher demand.

```
# Apply longer-term FoR weights to projected overall HPC demand
# Units: kSU

for_longterm_projections <- tidyr::crossing(
  year = hpc_scenarios$year,
  for_visual_group = longterm_for_weights$for_visual_group
) |>
  dplyr::left_join(
    hpc_scenarios |>
      dplyr::select(
        year,
        overall_hpc_ksu = central_linear
      ),
    by = "year"
  ) |>
  dplyr::left_join(
    longterm_for_weights |>
      dplyr::select(
        for_visual_group,
        scenario_share
      ),
    by = "for_visual_group"
  ) |>
  dplyr::mutate(
    illustrative_for_ksu =
      overall_hpc_ksu * scenario_share
  )

# Validate projections
stopifnot(
  all(is.finite(for_longterm_projections$illustrative_for_ksu)),
  all(for_longterm_projections$illustrative_for_ksu >= 0)
)

# Confirm allocations sum to projected overall HPC demand
for_longterm_total_check <- for_longterm_projections |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    allocated_ksu = sum(illustrative_for_ksu),
```

```
    projected_overall_ksu = dplyr::first(overall_hpc_ksu),
    difference_ksu = allocated_ksu - projected_overall_ksu,
    .groups = "drop"
  )

stopifnot(
  all(abs(for_longterm_total_check$difference_ksu) < 1e-6)
)

# Display the 2030 illustrative distribution
for_longterm_projections |>
  dplyr::filter(year == 2030) |>
  dplyr::arrange(dplyr::desc(illustrative_for_ksu)) |>
  dplyr::mutate(
    scenario_share_pct = 100 * scenario_share
  ) |>
  dplyr::select(
    for_visual_group,
    scenario_share_pct,
    illustrative_for_ksu
  ) |>
  knitr::kable(
    digits = 2,
    format.args = list(big.mark = ","),
    col.names = c(
      "FoR group",
      "Scenario share (%)",
      "Illustrative 2030 allocation (kSU)"
    ),
    caption = paste(
      "Illustrative 2030 FoR distribution of projected",
      "overall HPC demand: longer-term share assumptions"
    )
  )
```

Table 12.15: Illustrative 2030 FoR distribution of projected overall HPC demand: longer-term share assumptions

FoR group	Scenario share (%)	Illustrative 2030 allocation (kSU)
Engineering & Technology	24.19	605,404.67
Earth & Environmental Sciences	22.55	564,290.59
Physical Sciences	18.55	464,339.49
Chemical Sciences	17.74	443,903.19
Unclassified & Operational	6.92	173,129.72

FoR group	Scenario share (%)	Illustrative 2030 allocation (kSU)
Biological & Agricultural Sciences	4.81	120,389.40
Mathematics & Computing	3.32	83,183.85
Medical & Health Sciences	1.63	40,713.90
Social Sciences, Humanities & Built Environment	0.30	7,467.06

12.7.7 Earlier-Period FoR Allocation Scenario, 2027–2030

This scenario applies the normalised 2014–2020 allocation-share assumptions to the central overall HPC demand projection. It provides an earlier historical reference against which the recent-period and longer-term scenarios can be compared.

```
# Apply 2014-2020 FoR weights to the 2027-2030
# central overall HPC demand forecast.
# Units: kSU

for_earlier_projections <- tidyr::crossing(
  year = hpc_scenarios$year,
  for_visual_group = earlier_for_weights$for_visual_group
) |>
  dplyr::left_join(
    hpc_scenarios |>
      dplyr::select(
        year,
        overall_hpc_ksu = central_linear
      ),
    by = "year"
  ) |>
  dplyr::left_join(
    earlier_for_weights |>
      dplyr::select(
        for_visual_group,
        scenario_share
      ),
    by = "for_visual_group"
  ) |>
  dplyr::mutate(
    illustrative_for_ksu = overall_hpc_ksu * scenario_share
  )

# Validate that all projected values are finite and non-negative
stopifnot(
```

```
  all(is.finite(for_earlier_projections$illustrative_for_ksu)),
  all(for_earlier_projections$illustrative_for_ksu >= 0)
)

# Check that the FoR totals equal the overall HPC forecast
for_earlier_total_check <- for_earlier_projections |>
  dplyr::group_by(year) |>
  dplyr::summarise(
    allocated_ksu = sum(illustrative_for_ksu),
    projected_overall_ksu = dplyr::first(overall_hpc_ksu),
    .groups = "drop"
  )

stopifnot(
  all(abs(
    for_earlier_total_check$allocated_ksu -
    for_earlier_total_check$projected_overall_ksu
  ) < 1e-6)
)

# Display 2030 illustrative distribution
for_earlier_projections |>
  dplyr::filter(year == 2030) |>
  dplyr::arrange(dplyr::desc(illustrative_for_ksu)) |>
  dplyr::mutate(
    scenario_share_pct = 100 * scenario_share
  ) |>
  dplyr::select(
    for_visual_group,
    scenario_share_pct,
    illustrative_for_ksu
  ) |>
  knitr::kable(
    digits = 2,
    format.args = list(big.mark = ","),
    col.names = c(
      "FoR group",
      "Scenario share (%)",
      "Illustrative 2030 allocation (kSU)"
    ),
    caption = paste(
      "Illustrative 2030 FoR distribution:",
      "earlier-period 2014-2020 assumptions"
    )
  )
```

Table 12.16: Illustrative 2030 FoR distribution: earlier-period 2014–2020 assumptions

FoR group	Scenario share (%)	Illustrative 2030 allocation (kSU)
Earth & Environmental Sciences	30.38	760,374.43
Chemical Sciences	18.28	457,454.62
Physical Sciences	16.26	406,985.16
Engineering & Technology	16.24	406,401.10
Unclassified & Operational	6.98	174,598.90
Biological & Agricultural Sciences	4.69	117,258.18
Mathematics & Computing	4.51	112,831.93
Medical & Health Sciences	2.38	59,462.23
Social Sciences, Humanities & Built Environment	0.30	7,455.32

```
for_comparison <- dplyr::bind_rows(
  earlier_for_weights |>
    dplyr::transmute(
      for_group = as.character(for_visual_group),
      period = "Earlier (2014-2020)",
      share_pct = 100 * scenario_share
    ),
  recent_for_weights |>
    dplyr::transmute(
      for_group = as.character(for_2020_visual),
      period = "Recent (2021-2026)",
      share_pct = 100 * scenario_share
    ),
  longterm_for_weights |>
    dplyr::transmute(
      for_group = as.character(for_visual_group),
      period = "Longer-term (2014-2026)",
      share_pct = 100 * scenario_share
    )
) |>
  dplyr::mutate(
    period = factor(
      period,
      levels = c(
        "Earlier (2014-2020)",
        "Recent (2021-2026)",
        "Longer-term (2014-2026)"
      )
```

```
    )
  )

# Order research fields by their maximum scenario share
group_order <- for_comparison |>
  dplyr::group_by(for_group) |>
  dplyr::summarise(
    max_share = max(share_pct),
    .groups = "drop"
  ) |>
  dplyr::arrange(max_share) |>
  dplyr::pull(for_group)

for_comparison <- for_comparison |>
  dplyr::mutate(
    for_group = factor(for_group, levels = group_order)
  )

ggplot2::ggplot(
  for_comparison,
  ggplot2::aes(
    x = share_pct,
    y = for_group,
    fill = period
  )
) +
  ggplot2::geom_col(
    position = ggplot2::position_dodge(
      width = 0.8,
      preserve = "single"
    ),
    width = 0.7
  ) +
  ggplot2::scale_fill_manual(
    values = c(
      "Earlier (2014-2020)" = "#64748B",
      "Recent (2021-2026)" = "#155E75",
      "Longer-term (2014-2026)" = "#D97706"
    )
  ) +
  ggplot2::scale_x_continuous(
    labels = function(x) paste0(x, "%"),
    expand = ggplot2::expansion(mult = c(0, 0.05))
  ) +
  ggplot2::labs(
```

```
    title = "FoR Allocation-Share Scenarios for 2030",
    subtitle = "Three historical reference-period assumptions",
    x = "Assumed share of projected overall HPC demand",
    y = NULL,
    fill = "Reference period",
    caption = paste(
      "Shares are normalised observed-year means;",
      "classification, coverage and missing-category limitations apply.",
      "Unclassified categories are not assumed equivalent."
    )
  ) +
  ggplot2::theme_minimal(base_size = 12) +
  ggplot2::theme(
    plot.title = ggplot2::element_text(face = "bold"),
    legend.position = "bottom",
    plot.caption = ggplot2::element_text(
      size = 8,
      colour = "grey40"
    )
  )
```

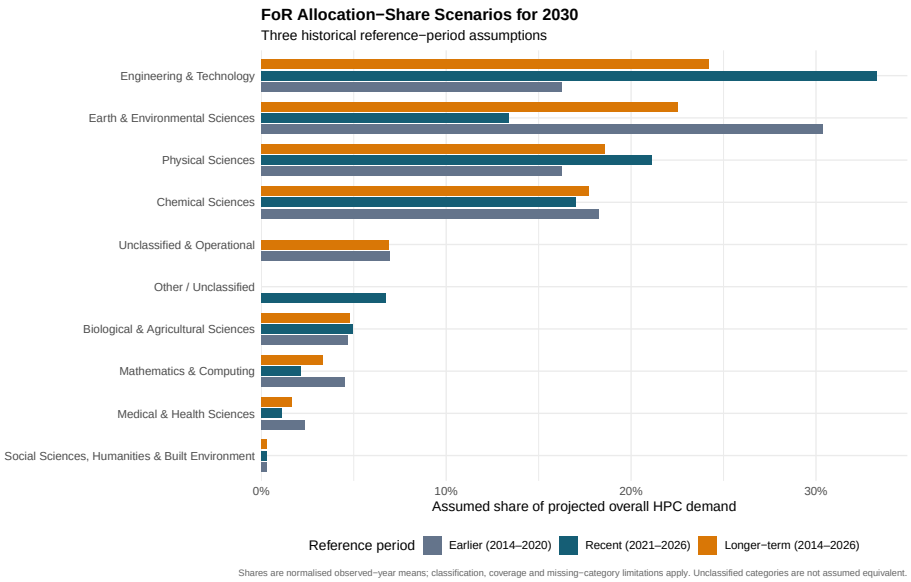


Figure 12.5: Comparison of normalised FoR allocation-share assumptions across three historical reference periods.

12.7.8 Comparison and Interpretation of FoR Scenarios

The three FoR allocation-share scenarios illustrate how the distribution of projected overall HPC demand varies under different historical reference-period assumptions.

All scenarios apply their respective normalised historical allocation-share weights to the same central overall HPC demand projection. Consequently, differences between the scenarios reflect the assumed disciplinary composition rather than differences in the total national demand forecast.

Under the earlier-period assumptions (2014–2020), Earth & Environmental Sciences receives the largest illustrative share at 30.38%, followed by Chemical Sciences (18.28%), Physical Sciences (16.26%), and Engineering & Technology (16.24%).

Under the recent-period assumptions (2021–2026), Engineering & Technology becomes the largest group at 33.28%, followed by Physical Sciences (21.11%), Chemical Sciences (17.02%), and Earth & Environmental Sciences (13.36%).

The shift in normalised scenario weights is substantial for the two largest changing groups. Engineering & Technology increases by approximately 17.04 percentage points between the earlier and recent reference periods, while Earth & Environmental Sciences decreases by approximately 17.02 percentage points.

The longer-term scenario (2014–2026) provides a third reference. Engineering & Technology accounts for 24.19%, Earth & Environmental Sciences for 22.55%, Physical Sciences for 18.55%, and Chemical Sciences for 17.74%.

These differences illustrate the sensitivity of FoR-level planning allocations to the selected historical reference period.

However, they should not automatically be interpreted as changes in underlying researcher demand. The historical data span different FoR classification structures, and differences in facility coverage, reporting scope, category completeness, and allocation practices may contribute to the observed variation.

In addition, several categories are not reported consistently across all years. The scenario weights were constructed by normalising observed-year mean allocation shares to sum to 100%, without treating unreported category-years as confirmed zero allocations.

The resulting FoR projections therefore represent conditional distributions of projected overall HPC demand, not independently validated demand forecasts for individual research fields.

The earlier-period, recent-period, and longer-term scenarios provide complementary planning perspectives. Their differences indicate that infrastructure planning should consider uncertainty in disciplinary composition rather than relying exclusively on a single historical allocation pattern.

The analysis does not establish which historical composition will best represent future demand. Further work would require consistently classified FoR-level demand requests, comparable facility reporting, and improved coverage of currently incomplete categories.

12.8 Overall Scenario Synthesis and Planning Implications

The scenario analysis combines projections of overall national HPC demand, CPU demand, exploratory GPU demand, and illustrative Field of Research (FoR) distributions through 2030. Together, these provide complementary perspectives on potential future demand for Australian national research computing infrastructure.

12.8.1 Central Forecast and Alternative Growth Assumptions

The selected linear model projects overall HPC demand to reach approximately **2.50 million kSU by 2030**, compared with approximately **2.29 million kSU** when the 2026 observation is excluded.

The log-linear alternatives produce substantially higher 2030 projections, ranging from approximately **6.48 to 6.96 million kSU**. However, their weaker historical cross-validation performance means they should be treated as alternative growth scenarios rather than equally supported central forecasts.

For CPU demand, the central linear projection reaches approximately **1,369.90 million CPU core-hours in 2030**. The 2026-excluded linear model projects approximately **1,263.85 million CPU core-hours**, while the log-linear alternatives project approximately **2,479–2,502 million CPU core-hours**.

The difference between linear and log-linear projections demonstrates the importance of the assumed growth structure when forecasting several years beyond a relatively short historical series.

12.8.2 Forecast Uncertainty and Model Adequacy

Model-based 80% and 95% prediction intervals were calculated for the selected overall HPC and CPU linear models.

For 2030, the overall HPC central forecast of approximately **2.50 million kSU** has a model-based 95% prediction interval of approximately **1.98–3.03 million kSU**.

The corresponding CPU central forecast of approximately **1,369.90 million CPU core-hours** has a model-based 95% prediction interval of approximately **974.24–1,765.55 million CPU core-hours**.

These intervals describe uncertainty conditional on the fitted linear regression models and their error assumptions. They do not fully account for measurement uncertainty, changes in reporting scope, model selection, future infrastructure developments, or structural changes in researcher demand.

Residual plots, normal Q-Q plots, and exploratory Ljung–Box tests provide additional model diagnostics. However, the small annual samples limit the strength of statistical conclusions. The absence of statistically significant residual autocorrelation should not be interpreted as proof that all model assumptions hold.

12.8.3 GPU Demand and Emerging Infrastructure Requirements

GPU demand is presented through exploratory scenarios rather than a definitive long-term forecast.

The naïve benchmark remains constant at approximately **3.003 million GPU-hours** through 2030. Alternative linear and log-linear specifications produce 2030 projections ranging from approximately **2.25 to 3.68 million GPU-hours**.

These trajectories are sensitive to the inclusion of 2026 and rely on only five annual GPU estimates for 2022–2026. The underlying series also combines reported and inferred demand measures.

Consequently, the GPU scenarios should be interpreted as alternative planning assumptions, not reliable statistical bounds on future GPU requirements.

12.8.4 Disciplinary Composition and FoR Scenarios

Three historical FoR allocation-share reference periods were considered: longer-term (2014–2026), earlier (2014–2020), and recent (2021–2026).

Each reference-period composition was applied to the same central overall HPC demand projection. This isolates the effect of alternative historical allocation-share assumptions on the illustrative distribution of projected demand across research fields.

The earlier-period scenario assigns the largest share to Earth & Environmental Sciences, while the recent-period scenario assigns the largest share to Engineering & Technology. The longer-term scenario provides an additional historical reference.

These differences demonstrate sensitivity to the selected reference period. They cannot be attributed exclusively to changes in researcher demand because FoR classification, facility coverage, reporting scope, and incomplete category observations may also affect the results.

The FoR projections are therefore illustrative top-down distributions, not independently estimated forecasts of demand by research field.

12.8.5 Implications for National HPC Planning

The analysis suggests that national HPC planning should consider a range of plausible demand trajectories rather than rely on a single extrapolation.

The central linear projections provide transparent reference estimates supported by the available historical cross-validation results. Alternative growth specifications and 2026-excluded sensitivity models illustrate how substantially projections can change under different assumptions.

CPU and GPU requirements should be assessed separately because they represent different computational resources and are not directly interchangeable.

Similarly, FoR allocation-share scenarios can support discussions about potential disciplinary composition, but they should not be used as definitive evidence of future demand by research field.

A key limitation is that recorded requests, allocations, and estimated resource requirements do not always measure the same concept. In particular, allocated resources may be constrained by available capacity and should not automatically be interpreted as the full level of researcher demand.

For infrastructure investment decisions, these projections are most appropriately used alongside facility-level capacity assessments, evidence of unmet researcher demand, expected changes in computing technologies, and direct engagement with research communities.

12.8.6 Chapter Conclusion

The forecasting exercise provides a reproducible, evidence-based set of scenarios for Australian national research HPC demand through 2030.

The central results are useful as planning reference points, while the sensitivity analyses, prediction intervals, GPU scenarios, and FoR distributions make important sources of uncertainty explicit.

The analysis does not establish a single certain future demand level. Instead, it provides a transparent framework for evaluating alternative assumptions and identifying where improved data collection would most strengthen future national HPC planning.

Chapter 13

References and AI Acknowledgement

13.1 References

Australian Academy of Science. 2024. "The Future Computing Needs of the Australian Science Sector." https://www.science.org.au/.

Australian Academy of Science. 2026. *Internship Project Brief: Forecasting Demand for National Research Supercomputing.*

National Computational Infrastructure. 2014a. *Annual Report 2013/14.* https://nci.org.au/sites/default/files/annual-reports/2019-07/About%20Us_Annual%20Reports_2013-14.pdf.

National Computational Infrastructure. 2014b. "NCMAS Information for Applicants." https://nci.org.au/news-events/news/ncmas-information-applicants.

National Computational Infrastructure. 2015. *Annual Report 2014/15.* https://nci.org.au/sites/default/files/annual-reports/2019-07/About%20Us_Annual%20Reports_2014-15.pdf.

National Computational Infrastructure. 2016. *Annual Report 2015/16.* https://nci.org.au/sites/default/files/annual-reports/2019-07/About%20Us_Annual%20Reports_2015-16.pdf.

National Computational Infrastructure. 2017. *Annual Report 2016/17.* https://nci.org.au/sites/default/files/annual-reports/2019-07/About%20Us_Annual%20Reports_2016-17.pdf.

National Computational Infrastructure. 2018a. *Annual Report 2017/18*. https://nci.org.au/sites/default/files/annual-reports/2019-07/About%20Us_Annual%20Reports_2017-18.pdf.

National Computational Infrastructure. 2018b. "NCMAS Allocations Now Available Online." https://nci.org.au/news-events/news/ncmas-allocations-now-available-online.

National Computational Infrastructure. 2019. *Annual Report 2018/19*. https://nci.org.au/sites/default/files/annual-reports/2019-11/NCI%20Report%202019%20web.pdf.

National Computational Infrastructure. 2020a. *Annual Report 2019/20*. https://nci.org.au/sites/default/files/annual-reports/2021-01/NCI%20Annual%20report%202019-20%20web_final_opt2.pdf.

National Computational Infrastructure. 2020b. *Compute Allocations 2019/20*. https://nci.org.au/sites/default/files/annual-reports/2021-01/Allocations%202019-20%20for%20website%20FINAL%20with%20pagenb.pdf.

National Computational Infrastructure. 2020c. *NCMAS 2020 Information for Applicants*. https://nci.org.au/sites/default/files/documents/2019-10/NCMAS-2020-InformationForApplicants.pdf.

National Computational Infrastructure. 2021. *Compute Allocations 2020–2021*. https://nci.org.au/sites/default/files/annual-reports/2021-12/NCI%20Compute%20Allocations%202020-2021_0.pdf.

National Computational Infrastructure. 2022. *NCMAS 2023 Information for Applicants*. https://my.nci.org.au/mancini-assets/media/mas/ncmas/2023/NCMAS2023-NCMAS_2023_Info4Applicants_v5-20220830T164026AEST.pdf.

National Computational Infrastructure. 2023. *NCMAS 2024 Information for Applicants*. https://my.nci.org.au/mancini-assets/media/mas/ncmas/2024/NCMAS2024-NCMAS_2024_-_Information_for_Applicants_20231109-20231109T131740AEDT.pdf.

National Computational Infrastructure. 2025. *NCMAS 2026 Information for Applicants*. https://my.nci.org.au/mancini-assets/media/mas/ncmas/2026/NCMAS2026-Information_for_Applicants-20251023T122535AEDT.pdf.

National Computational Infrastructure. n.d. "Who We Are." Accessed October 4, 2026. https://nci.org.au/.

National Computational Merit Allocation Scheme. 2021. "NCMAS 2021 Out-

comes." https://my.nci.org.au/mancini/ncmas/2021/outcomes.

National Computational Merit Allocation Scheme. 2022. "NCMAS 2022 Outcomes." https://my.nci.org.au/mancini/ncmas/2022/outcomes.

National Computational Merit Allocation Scheme. 2023. "NCMAS 2023 Outcomes." https://my.nci.org.au/mancini/ncmas/2023/outcomes.

National Computational Merit Allocation Scheme. 2024. "NCMAS 2024 Outcomes." https://my.nci.org.au/mancini/ncmas/2024/outcomes.

National Computational Merit Allocation Scheme. 2025. "NCMAS 2025 Outcomes." https://my.nci.org.au/mancini/ncmas/2025/outcomes.

National Computational Merit Allocation Scheme. 2026. "NCMAS 2026 Outcomes." https://my.nci.org.au/mancini/ncmas/2026/outcomes.

Pawsey Supercomputing Research Centre. 2015. *Annual Report 2014/15*. https://pawsey.org.au/wp-content/uploads/2018/10/ANNUALREPORT14-15v11web.pdf.

Pawsey Supercomputing Research Centre. 2017. *Annual Report 2016/17*. https://pawsey.org.au/wp-content/uploads/2018/10/Pawsey-Annual-Report-2016-2017.pdf.

Pawsey Supercomputing Research Centre. 2019. "Applications Open for 2019 Access to Pawsey Supercomputers." https://pawsey.org.au/applications-open-for-2019-access-to-pawsey-supercomputers/.

Pawsey Supercomputing Research Centre. n.d. "About Pawsey." Accessed October 4, 2026. https://pawsey.org.au/.

13.2 AI Acknowledgement

Generative AI tools, including ChatGPT, were used to support aspects of this project such as brainstorming, reviewing code, improving written expression, and assisting with the organisation of the analysis and documentation.

All data preparation, analytical decisions, model development, interpretation of results, and final conclusions were reviewed and verified by the author.

AI-generated suggestions were not treated as primary sources. Information and assumptions used in the analysis were checked against the original sources where appropriate.